

Delay accounting

Tasks encounter delays in execution when they wait for some kernel resource to become available e.g. a runnable task may wait for a free CPU to run on.

The per-task delay accounting functionality measures the delays experienced by a task while

- a) waiting for a CPU (while being runnable)
- b) completion of synchronous block I/O initiated by the task
- c) swapping in pages
- d) memory reclaim

and makes these statistics available to userspace through the taskstats interface.

Such delays provide feedback for setting a task's cpu priority, io priority and rss limit values appropriately. Long delays for important tasks could be a trigger for raising its corresponding priority.

The functionality, through its use of the taskstats interface, also provides delay statistics aggregated for all tasks (or threads) belonging to a thread group (corresponding to a traditional Unix process). This is a commonly needed aggregation that is more efficiently done by the kernel.

Userspace utilities, particularly resource management applications, can also aggregate delay statistics into arbitrary groups. To enable this, delay statistics of a task are available both during its lifetime as well as on its exit, ensuring continuous and complete monitoring can be done.

Interface

Delay accounting uses the taskstats interface which is described in detail in a separate document in this directory. Taskstats returns a generic data structure to userspace corresponding to per-pid and per-tgid statistics. The delay accounting functionality populates specific fields of this structure. See

include/linux/taskstats.h

for a description of the fields pertaining to delay accounting.

It will generally be in the form of counters returning the cumulative delay seen for cpu, sync block I/O, swapin, memory reclaim etc.

Taking the difference of two successive readings of a given counter (say `cpu_delay_total`) for a task will give the delay experienced by the task waiting for the corresponding resource in that interval.

When a task exits, records containing the per-task statistics are sent to userspace without requiring a command. If it is the last exiting task of a thread group, the per-tgid statistics are also sent. More details are given in the taskstats interface description.

The getdelays.c userspace utility in this directory allows simple commands to be run and the corresponding delay statistics to be displayed. It also serves as an example of using the taskstats interface.

Usage

Compile the kernel with
 CONFIG_TASK_DELAY_ACCT=y
 CONFIG_TASKSTATS=y

Delay accounting is enabled by default at boot up.
 To disable, add
 nodelayacct
 to the kernel boot options. The rest of the instructions below assume this has not been done.

After the system has booted up, use a utility similar to getdelays.c to access the delays seen by a given task or a task group (tgid). The utility also allows a given command to be executed and the corresponding delays to be seen.

General format of the getdelays command

getdelays [-t tgid] [-p pid] [-c cmd...]

Get delays, since system boot, for pid 10
 # ./getdelays -p 10
 (output similar to next case)

Get sum of delays, since system boot, for all pids with tgid 5
 # ./getdelays -t 5

CPU	count	real total	virtual total	delay total
	7876	92005750	100000000	24001500
IO	count	delay total		
	0	0		
SWAP	count	delay total		
	0	0		
RECLAIM	count	delay total		
	0	0		

Get delays seen in executing a given simple command
 # ./getdelays -c ls /

bin	data1	data3	data5	dev	home	media	opt	root	srv	sys	usr
boot	data2	data4	data6	etc	lib	mnt	proc	sbin	subdomain	tmp	var

CPU	count	real total	virtual total	delay total
	6	4000250	4000000	0
IO	count	delay total		

delay-accounting.txt.txt

	0	0
SWAP	count	delay total
	0	0
RECLAIM	count	delay total
	0	0