

```

/* getdelays.c
 *
 * Utility to get per-pid and per-tgid delay accounting statistics
 * Also illustrates usage of the taskstats interface
 *
 * Copyright (C) Shailabh Nagar, IBM Corp. 2005
 * Copyright (C) Balbir Singh, IBM Corp. 2006
 * Copyright (c) Jay Lan, SGI. 2006
 *
 * Compile with
 * gcc -I/usr/src/linux/include getdelays.c -o getdelays
 */

#include <stdio.h>
#include <stdlib.h>
#include <errno.h>
#include <unistd.h>
#include <poll.h>
#include <string.h>
#include <fcntl.h>
#include <sys/types.h>
#include <sys/stat.h>
#include <sys/socket.h>
#include <signal.h>

#include <linux/genetlink.h>
#include <linux/taskstats.h>
#include <linux/cgroupstats.h>

/*
 * Generic macros for dealing with netlink sockets. Might be duplicated
 * elsewhere. It is recommended that commercial grade applications use
 * libnl or libnetlink and use the interfaces provided by the library
 */
#define GENLMSG_DATA(glh) ((void *) (NLMSG_DATA(glh) + GENL_HDRLEN))
#define GENLMSG_PAYLOAD(glh) (NLMSG_PAYLOAD(glh, 0) - GENL_HDRLEN)
#define NLA_DATA(na) ((void *) ((char *) (na) + NLA_HDRLEN))
#define NLA_PAYLOAD(len) (len - NLA_HDRLEN)

#define err(code, fmt, arg...) \
do { \
    fprintf(stderr, fmt, ##arg); \
    exit(code); \
} while (0)

int done;
int rcvbufsz;
char name[100];
int dbg;
int print_delays;
int print_io_accounting;
int print_task_context_switch_counts;
__u64 stime, utime;

#define PRINTF(fmt, arg...) { \
    if (dbg) { \

```

```

                                getdelays.c.txt
                                printf(fmt, ##arg);
                                \
                                \
        }
    }

/* Maximum size of response requested or message sent */
#define MAX_MSG_SIZE    1024
/* Maximum number of cpus expected to be specified in a cpumask */
#define MAX_CPUS        32

struct msgtemplate {
    struct nlmsgghdr n;
    struct genlmsgghdr g;
    char buf[MAX_MSG_SIZE];
};

char cpumask[100+6*MAX_CPUS];

static void usage(void)
{
    fprintf(stderr, "getdelays [-dilv] [-w logfile] [-r bufsize] "
        "[-m cpumask] [-t tgid] [-p pid]\n");
    fprintf(stderr, "  -d: print delayacct stats\n");
    fprintf(stderr, "  -i: print IO accounting (works only with -p)\n");
    fprintf(stderr, "  -l: listen forever\n");
    fprintf(stderr, "  -v: debug on\n");
    fprintf(stderr, "  -C: container path\n");
}

/*
 * Create a raw netlink socket and bind
 */
static int create_nl_socket(int protocol)
{
    int fd;
    struct sockaddr_nl local;

    fd = socket(AF_NETLINK, SOCK_RAW, protocol);
    if (fd < 0)
        return -1;

    if (rcvbufsz)
        if (setsockopt(fd, SOL_SOCKET, SO_RCVBUF,
            &rcvbufsz, sizeof(rcvbufsz)) < 0) {
            fprintf(stderr, "Unable to set socket rcv buf size "
                "to %d\n",
                rcvbufsz);
            return -1;
        }

    memset(&local, 0, sizeof(local));
    local.nl_family = AF_NETLINK;

    if (bind(fd, (struct sockaddr *) &local, sizeof(local)) < 0)
        goto error;

    return fd;
}

```

```

error:
    close(fd);
    return -1;
}

static int send_cmd(int sd, __u16 nlmsg_type, __u32 nlmsg_pid,
                    __u8 genl_cmd, __u16 nla_type,
                    void *nla_data, int nla_len)
{
    struct nlattr *na;
    struct sockaddr_nl nladdr;
    int r, buflen;
    char *buf;

    struct msgtemplate msg;

    msg.n.nlmsg_len = NLMSG_LENGTH(GENL_HDRLEN);
    msg.n.nlmsg_type = nlmsg_type;
    msg.n.nlmsg_flags = NLM_F_REQUEST;
    msg.n.nlmsg_seq = 0;
    msg.n.nlmsg_pid = nlmsg_pid;
    msg.g.cmd = genl_cmd;
    msg.g.version = 0x1;
    na = (struct nlattr *) GENLMSG_DATA(&msg);
    na->nla_type = nla_type;
    na->nla_len = nla_len + 1 + NLA_HDRLEN;
    memcpy(NLA_DATA(na), nla_data, nla_len);
    msg.n.nlmsg_len += NLMSG_ALIGN(na->nla_len);

    buf = (char *) &msg;
    buflen = msg.n.nlmsg_len;
    memset(&nladdr, 0, sizeof(nladdr));
    nladdr.nl_family = AF_NETLINK;
    while ((r = sendto(sd, buf, buflen, 0, (struct sockaddr *) &nladdr,
                      sizeof(nladdr))) < buflen) {
        if (r > 0) {
            buf += r;
            buflen -= r;
        } else if (errno != EAGAIN)
            return -1;
    }
    return 0;
}

/*
 * Probe the controller in genetlink to find the family id
 * for the TASKSTATS family
 */
static int get_family_id(int sd)
{
    struct {
        struct nlmsgghdr n;
        struct genlmsgghdr g;
        char buf[256];
    }

```

getdelays.c.txt

```
} ans;

int id = 0, rc;
struct nlattr *na;
int rep_len;

strcpy(name, TASKSTATS_GENL_NAME);
rc = send_cmd(sd, GENL_ID_CTRL, getpid(), CTRL_CMD_GETFAMILY,
              CTRL_ATTR_FAMILY_NAME, (void *)name,
              strlen(TASKSTATS_GENL_NAME)+1);

rep_len = recv(sd, &ans, sizeof(ans), 0);
if (ans.n.nlmsg_type == NLMSG_ERROR ||
    (rep_len < 0) || !NLMSG_OK((&ans.n), rep_len))
    return 0;

na = (struct nlattr *) GENLMSG_DATA(&ans);
na = (struct nlattr *) ((char *) na + NLA_ALIGN(na->nla_len));
if (na->nla_type == CTRL_ATTR_FAMILY_ID) {
    id = *(__u16 *) NLA_DATA(na);
}
return id;
}

static void print_delayacct(struct taskstats *t)
{
    printf("\n\nCPU    %15s%15s%15s%15s\n"
           "      %15llu%15llu%15llu%15llu\n"
           "IO      %15s%15s\n"
           "      %15llu%15llu\n"
           "SWAP    %15s%15s\n"
           "      %15llu%15llu\n"
           "RECLAIM %12s%15s\n"
           "      %15llu%15llu\n",
           "count", "real total", "virtual total", "delay total",
           (unsigned long long)t->cpu_count,
           (unsigned long long)t->cpu_run_real_total,
           (unsigned long long)t->cpu_run_virtual_total,
           (unsigned long long)t->cpu_delay_total,
           "count", "delay total",
           (unsigned long long)t->blkio_count,
           (unsigned long long)t->blkio_delay_total,
           "count", "delay total",
           (unsigned long long)t->swpin_count,
           (unsigned long long)t->swpin_delay_total,
           "count", "delay total",
           (unsigned long long)t->freepages_count,
           (unsigned long long)t->freepages_delay_total);
}

static void task_context_switch_counts(struct taskstats *t)
{
    printf("\n\nTask    %15s%15s\n"
           "      %15llu%15llu\n",
           "voluntary", "nonvoluntary",
           (unsigned long long)t->nvcsw, (unsigned long long)t->nivcsw);
}
```

```

}

static void print_cgroupstats(struct cgroupstats *c)
{
    printf("sleeping %llu, blocked %llu, running %llu, stopped %llu, "
           "uninterruptible %llu\n", (unsigned long long)c->nr_sleeping,
           (unsigned long long)c->nr_io_wait,
           (unsigned long long)c->nr_running,
           (unsigned long long)c->nr_stopped,
           (unsigned long long)c->nr_uninterruptible);
}

static void print_ioacct(struct taskstats *t)
{
    printf("%s: read=%llu, write=%llu, cancelled_write=%llu\n",
           t->ac_comm,
           (unsigned long long)t->read_bytes,
           (unsigned long long)t->write_bytes,
           (unsigned long long)t->cancelled_write_bytes);
}

int main(int argc, char *argv[])
{
    int c, rc, rep_len, aggr_len, len2;
    int cmd_type = TASKSTATS_CMD_ATTR_UNSPEC;
    __u16 id;
    __u32 mypid;

    struct nlattr *na;
    int nl_sd = -1;
    int len = 0;
    pid_t tid = 0;
    pid_t rtid = 0;

    int fd = 0;
    int count = 0;
    int write_file = 0;
    int maskset = 0;
    char *logfile = NULL;
    int loop = 0;
    int containerset = 0;
    char containerpath[1024];
    int cfd = 0;

    struct msgtemplate msg;

    while (1) {
        c = getopt(argc, argv, "qdiw:r:m:t:p:vlC:");
        if (c < 0)
            break;

        switch (c) {
            case 'd':
                printf("print delayacct stats ON\n");
                print_delays = 1;

```

getdelays.c.txt

```
        break;
case 'i':
    printf("printing IO accounting\n");
    print_io_accounting = 1;
    break;
case 'q':
    printf("printing task/process context switch rates\n");
    print_task_context_switch_counts = 1;
    break;
case 'C':
    containerset = 1;
    strncpy(containerpath, optarg, strlen(optarg) + 1);
    break;
case 'w':
    logfile = strdup(optarg);
    printf("write to file %s\n", logfile);
    write_file = 1;
    break;
case 'r':
    rcvbufsz = atoi(optarg);
    printf("receive buf size %d\n", rcvbufsz);
    if (rcvbufsz < 0)
        err(1, "Invalid rcv buf size\n");
    break;
case 'm':
    strncpy(cpumask, optarg, sizeof(cpumask));
    maskset = 1;
    printf("cpumask %s maskset %d\n", cpumask, maskset);
    break;
case 't':
    tid = atoi(optarg);
    if (!tid)
        err(1, "Invalid tgid\n");
    cmd_type = TASKSTATS_CMD_ATTR_TGID;
    break;
case 'p':
    tid = atoi(optarg);
    if (!tid)
        err(1, "Invalid pid\n");
    cmd_type = TASKSTATS_CMD_ATTR_PID;
    break;
case 'v':
    printf("debug on\n");
    dbg = 1;
    break;
case 'l':
    printf("listen forever\n");
    loop = 1;
    break;
default:
    usage();
    exit(-1);
}

if (write_file) {
```

```

                                getdelays.c.txt
fd = open(logfile, O_WRONLY | O_CREAT | O_TRUNC,
           S_IRUSR | S_IWUSR | S_IRGRP | S_IROTH);
if (fd == -1) {
    perror("Cannot open output file\n");
    exit(1);
}

if ((nl_sd = create_nl_socket(NETLINK_GENERIC)) < 0)
    err(1, "error creating Netlink socket\n");

mypid = getpid();
id = get_family_id(nl_sd);
if (!id) {
    fprintf(stderr, "Error getting family id, errno %d\n", errno);
    goto err;
}
PRINTF("family id %d\n", id);

if (maskset) {
    rc = send_cmd(nl_sd, id, mypid, TASKSTATS_CMD_GET,
                  TASKSTATS_CMD_ATTR_REGISTER_CPUMASK,
                  &cpumask, strlen(cpumask) + 1);
    PRINTF("Sent register cpumask, retval %d\n", rc);
    if (rc < 0) {
        fprintf(stderr, "error sending register cpumask\n");
        goto err;
    }
}

if (tid && containerset) {
    fprintf(stderr, "Select either -t or -C, not both\n");
    goto err;
}

if (tid) {
    rc = send_cmd(nl_sd, id, mypid, TASKSTATS_CMD_GET,
                  cmd_type, &tid, sizeof(__u32));
    PRINTF("Sent pid/tgid, retval %d\n", rc);
    if (rc < 0) {
        fprintf(stderr, "error sending tid/tgid cmd\n");
        goto done;
    }
}

if (containerset) {
    cfd = open(containerpath, O_RDONLY);
    if (cfd < 0) {
        perror("error opening container file");
        goto err;
    }
    rc = send_cmd(nl_sd, id, mypid, CGROUPSTATS_CMD_GET,
                  CGROUPSTATS_CMD_ATTR_FD, &cfd, sizeof(__u32));
    if (rc < 0) {
        perror("error sending cgroupstats command");
    }
}

```

```

                                getdelays.c.txt
                                goto err;
                                }
                                }
if (!maskset && !tid && !containerset) {
    usage();
    goto err;
}

do {
    int i;

    rep_len = recv(nl_sd, &msg, sizeof(msg), 0);
    PRINTF("received %d bytes\n", rep_len);

    if (rep_len < 0) {
        fprintf(stderr, "nonfatal reply error: errno %d\n",
            errno);
        continue;
    }
    if (msg.n.nlmsg_type == NLMSG_ERROR ||
        !NLMSG_OK((&msg.n), rep_len)) {
        struct nlmsgerr *err = NLMSG_DATA(&msg);
        fprintf(stderr, "fatal reply error,  errno %d\n",
            err->error);
        goto done;
    }

    PRINTF("nlmsg_hdr size=%zu, nlmsg_len=%d, rep_len=%d\n",
        sizeof(struct nlmsg_hdr), msg.n.nlmsg_len, rep_len);

    rep_len = GENLMSG_PAYLOAD(&msg.n);

    na = (struct nlattr *) GENLMSG_DATA(&msg);
    len = 0;
    i = 0;
    while (len < rep_len) {
        len += NLA_ALIGN(na->nla_len);
        switch (na->nla_type) {
            case TASKSTATS_TYPE_AGGR_TGID:
                /* Fall through */
            case TASKSTATS_TYPE_AGGR_PID:
                aggr_len = NLA_PAYLOAD(na->nla_len);
                len2 = 0;
                /* For nested attributes, na follows */
                na = (struct nlattr *) NLA_DATA(na);
                done = 0;
                while (len2 < aggr_len) {
                    switch (na->nla_type) {
                        case TASKSTATS_TYPE_PID:
                            rtid = *(int *) NLA_DATA(na);
                            if (print_delays)
                                printf("PID\t%d\n",
rtid);
                            break;
                        case TASKSTATS_TYPE_TGID:

```



```

getdelays.c.txt
rtid = *(int *) NLA_DATA(na);
if (print_delays)
    printf("TGID\t%d\n",
rtid);

break;
case TASKSTATS_TYPE_STATS:
    count++;
    if (print_delays)
        print_delayacct((struct
taskstats *) NLA_DATA(na));

    if (print_io_accounting)
        print_ioacct((struct
taskstats *) NLA_DATA(na));

    if
(print_task_context_switch_counts)
task_context_switch_counts((struct taskstats *) NLA_DATA(na));
    if (fd) {
        if (write(fd,
NLA_DATA(na), na->nla_len) < 0) {
            err(1, "write
error\n");
        }
    }
    if (!loop)
        goto done;
    break;
default:
    fprintf(stderr, "Unknown nested"
        " nla_type %d\n",
        na->nla_type);
    break;
}
len2 += NLA_ALIGN(na->nla_len);
na = (struct nlattr *) ((char *) na +
len2);
}
break;

case CGROUPSTATS_TYPE_CGROUP_STATS:
    print_cgroupstats(NLA_DATA(na));
    break;
default:
    fprintf(stderr, "Unknown nla_type %d\n",
        na->nla_type);
    break;
}
na = (struct nlattr *) (GENLMSG_DATA(&msg) + len);
}
} while (loop);
done:
if (maskset) {
    rc = send_cmd(nl_sd, id, mypid, TASKSTATS_CMD_GET,
        TASKSTATS_CMD_ATTR_DEREGISTER_CPUMASK,
        &cpumask, strlen(cpumask) + 1);
    printf("Sent deregister mask, retval %d\n", rc);

```

```

                                getdelays.c.txt
        if (rc < 0)
            err(rc, "error sending deregister cpumask\n");
    }
err:
    close(nl_sd);
    if (fd)
        close(fd);
    if (cfd)
        close(cfd);
    return 0;
}

```