## The struct taskstats

This document contains an explanation of the struct taskstats fields.

There are three different groups of fields in the struct taskstats:

1) Common and basic accounting fields
 If CONFIG\_TASKSTATS is set, the taskstats interface is enabled and
 the common fields and basic accounting fields are collected for
 delivery at do exit() of a task.

2) Delay accounting fields

These fields are placed between /\* Delay accounting fields start \*/ and

/\* Delay accounting fields end \*/

Their values are collected if CONFIG\_TASK\_DELAY\_ACCT is set.

3) Extended accounting fields

These fields are placed between
/\* Extended accounting fields start \*/
and
/\* Extended accounting fields end \*/

Their values are collected if CONFIG\_TASK\_XACCT is set.

- 4) Per-task and per-thread context switch count statistics
- 5) Time accounting for SMT machines
- 6) Extended delay accounting fields for memory reclaim

Future extension should add fields to the end of the taskstats struct, and should not change the relative position of each field within the struct.

```
struct taskstats {
```

```
1) Common and basic accounting fields:
       /* The version number of this struct. This field is always set to
        * TAKSTATS_VERSION, which is defined in linux/taskstats.h>.
        * Each time the struct is changed, the value should be incremented.
        */
        u16
               version;
       /* The exit code of a task. */
                                        /* Exit status */
               ac exitcode;
        u32
        /* The accounting flags of a task as defined in linux/acct.h>
        * Defined values are AFORK, ASU, ACOMPAT, ACORE, and AXSIG.
        */
       __u8
                                        /* Record flags */
               ac flag;
       /* The value of task_nice() of a task. */
                                        /* task nice */
        u8
               ac_nice;
        /* The name of the command that started this task. */
               ac_comm[TS_COMM_LEN]; /* Command name */
```

第1页

## taskstats-struct.txt.txt

```
/* The scheduling discipline as set in task->policy field. */
        __ u8
                ac sched;
                                        /* Scheduling discipline */
                ac pad[3];
         u8
                                        /* User ID */
         u32
                ac uid;
        _u32
                                        /* Group ID */
                ac_gid;
                                        /* Process ID */
        u32
                ac pid;
                                        /* Parent process ID */
        u32
                ac ppid;
        /* The time when a task begins, in [secs] since 1970. */
        u32
                ac btime;
                                        /* Begin time [sec since 1970] */
        /* The elapsed time of a task, in [usec]. */
        __u64
                ac etime;
                                        /* Elapsed time [usec] */
        /* The user CPU time of a task, in [usec]. */
                ac utime;
                                        /* User CPU time [usec] */
        u64
        /* The system CPU time of a task, in [usec]. */
                                        /* System CPU time [usec] */
                ac stime;
        /* The minor page fault count of a task, as set in task->min_flt. */
                                        /* Minor Page Fault Count */
        u64
                ac minflt;
        /* The major page fault count of a task, as set in task->maj flt. */
                                        /* Major Page Fault Count */
        u64
                ac majflt;
2) Delay accounting fields:
        /* Delay accounting fields start
         * All values, until the comment "Delay accounting fields end" are
         * available only if delay accounting is enabled, even though the last
         * few fields are not delays
         * xxx count is the number of delay values recorded
         * xxx delay total is the corresponding cumulative delay in nanoseconds
         * xxx_delay_total wraps around to zero on overflow
         * xxx_count incremented regardless of overflow
         */
        /* Delay waiting for cpu, while runnable
         * count, delay total NOT updated atomically
         */
         _u64
                cpu count;
                cpu delay total;
        u64
        /* Following four fields atomically updated using task->delays->lock */
        /* Delay waiting for synchronous block I/O to complete
         * does not account for delays in I/O submission
         */
         u64
                blkio count;
                blkio delay total;
        u64
                                     第 2 页
```

## taskstats-struct.txt.txt

```
/* Delay waiting for page fault I/O (swap in only) */
        __u64
                swapin count;
        u64
                swapin delay total;
        /* cpu "wall-clock" running time
         st On some architectures, value will adjust for cpu time stolen
         * from the kernel in involuntary waits due to virtualization.
         * Value is cumulative, in nanoseconds, without a corresponding count
         * and wraps around to zero silently on overflow
         */
        u64
                cpu run real total;
        /* cpu "virtual" running time
         * Uses time intervals seen by the kernel i.e. no adjustment
         * for kernel's involuntary waits due to virtualization.
         * Value is cumulative, in nanoseconds, without a corresponding count
         * and wraps around to zero silently on overflow
               cpu run virtual total;
         u64
        /* Delay accounting fields end */
        /* version 1 ends here */
3) Extended accounting fields
        /* Extended accounting fields start */
        /* Accumulated RSS usage in duration of a task, in MBytes-usecs.
         * The current rss usage is added to this counter every time
         * a tick is charged to a task's system time. So, at the end we
         * will have memory usage multiplied by system time. Thus an
         * average usage per system time unit can be calculated.
         */
                                        /* accumulated RSS usage in MB-usec */
        u64
                coremem;
        /* Accumulated virtual memory usage in duration of a task.
         * Same as acct rss mem1 above except that we keep track of VM usage.
         */
        u64
                virtmem:
                                        /* accumulated VM usage in MB-usec */
        /* High watermark of RSS usage in duration of a task, in KBytes. */
                                        /* High-watermark of RSS usage */
        u64
               hiwater rss;
        /* High watermark of VM usage in duration of a task, in KBytes. */
        __u64
                hiwater_vm;
                                        /* High-water virtual memory usage */
        /* The following four fields are I/O statistics of a task. */
        u64
                read char;
                                        /* bytes read */
                                        /* bytes written */
                write char:
         u64
        __u64
                read_syscalls:
                                        /* read syscalls */
                write_syscalls;
                                        /* write syscalls */
         u64
        /* Extended accounting fields end */
4) Per-task and per-thread statistics
                                        /* Context voluntary switch counter */
        u64 nvcsw:
                                     第 3 页
```

## taskstats-struct.txt.txt \_\_u64 nivcsw; /\* Context involuntary switch counter \*/ 5) Time accounting for SMT machines /\* utime scaled on frequency etc \*/ \_\_u64 ac\_utimescaled; \_\_u64 ac stimescaled; /\* stime scaled on frequency etc \*/ cpu\_scaled\_run\_real\_total; /\* scaled cpu\_run\_real\_total \*/ \_\_u64 6) Extended delay accounting fields for memory reclaim /\* Delay waiting for memory reclaim \*/ \_\_u64 freepages\_count; freepages\_delay\_total; \_\_u64 }