

memory.txt.txt
Kernel Memory Layout on ARM Linux

Russell King <rmk@arm.linux.org.uk>
November 17, 2005 (2.6.15)

This document describes the virtual memory layout which the Linux kernel uses for ARM processors. It indicates which regions are free for platforms to use, and which are used by generic code.

The ARM CPU is capable of addressing a maximum of 4GB virtual memory space, and this must be shared between user space processes, the kernel, and hardware devices.

As the ARM architecture matures, it becomes necessary to reserve certain regions of VM space for use for new facilities; therefore this document may reserve more VM space over time.

| Start | End | Use |
|---------------|---------------|---|
| ffff8000 | fffffff | copy_user_page / clear_user_page use. For SAllxx and Xscale, this is used to setup a minicache mapping. |
| ffff4000 | fffffff | cache aliasing on ARMv6 and later CPUs. |
| ffff1000 | ffff7fff | Reserved. Platforms must not use this address range. |
| ffff0000 | ffff0fff | CPU vector page. The CPU vectors are mapped here if the CPU supports vector relocation (control register V bit.) |
| fffe0000 | fffeffff | XScale cache flush area. This is used in proc-xscale.S to flush the whole data cache. Free for other usage on non-XScale. |
| fff00000 | fffdffff | Fixmap mapping region. Addresses provided by fix_to_virt() will be located here. |
| ffc00000 | ffeffff | DMA memory mapping region. Memory returned by the dma_alloc_xxx functions will be dynamically mapped here. |
| ff000000 | ffbffff | Reserved for future expansion of DMA mapping region. |
| VMALLOC_END | feffff | Free for platform use, recommended. VMALLOC_END must be aligned to a 2MB boundary. |
| VMALLOC_START | VMALLOC_END-1 | vmalloc() / ioremap() space. Memory returned by vmalloc/ioremap will be dynamically placed in this region. VMALLOC_START may be based upon the value of the high_memory variable. |

memory.txt.txt

| | | |
|---------------|---------------|---|
| PAGE_OFFSET | high_memory-1 | Kernel direct-mapped RAM region. This maps the platforms RAM, and typically maps all platform RAM in a 1:1 relationship. |
| PKMAP_BASE | PAGE_OFFSET-1 | Permanent kernel mappings One way of mapping HIGHMEM pages into kernel space. |
| MODULES_VADDR | MODULES_END-1 | Kernel module space Kernel modules inserted via insmod are placed here using dynamic mappings. |
| 00001000 | TASK_SIZE-1 | User space mappings Per-thread mappings are placed here via the mmap() system call. |
| 00000000 | 00000fff | CPU vector page / null pointer trap CPUs which do not support vector remapping place their vector page here. NULL pointer dereferences by both the kernel and user space are also caught via this mapping. |

Please note that mappings which collide with the above areas may result in a non-bootable kernel, or may cause the kernel to (eventually) panic at run time.

Since future CPUs may impact the kernel mapping layout, user programs must not access any memory which is not mapped inside their 0x0001000 to TASK_SIZE address range. If they wish to access these areas, they must set up their own mappings using open() and mmap().