

atomic_ops.txt
Semantics and Behavior of Atomic and
Bitmask Operations

David S. Miller

This document is intended to serve as a guide to Linux port maintainers on how to implement atomic counter, bitops, and spinlock interfaces properly.

The `atomic_t` type should be defined as a signed integer. Also, it should be made opaque such that any kind of cast to a normal C integer type will fail. Something like the following should suffice:

```
typedef struct { volatile int counter; } atomic_t;
```

Historically, `counter` has been declared `volatile`. This is now discouraged. See `Documentation/volatile-considered-harmful.txt` for the complete rationale.

`local_t` is very similar to `atomic_t`. If the counter is per CPU and only updated by one CPU, `local_t` is probably more appropriate. Please see `Documentation/local_ops.txt` for the semantics of `local_t`.

The first operations to implement for `atomic_t`'s are the initializers and plain reads.

```
#define ATOMIC_INIT(i)          { (i) }  
#define atomic_set(v, i)      ((v)->counter = (i))
```

The first macro is used in definitions, such as:

```
static atomic_t my_counter = ATOMIC_INIT(1);
```

The initializer is atomic in that the return values of the atomic operations are guaranteed to be correct reflecting the initialized value if the initializer is used before runtime. If the initializer is used at runtime, a proper implicit or explicit read memory barrier is needed before reading the value with `atomic_read` from another thread.

The second interface can be used at runtime, as in:

```
struct foo { atomic_t counter; };  
...  
  
struct foo *k;  
  
k = kmalloc(sizeof(*k), GFP_KERNEL);  
if (!k)  
    return -ENOMEM;  
atomic_set(&k->counter, 0);
```

The setting is atomic in that the return values of the atomic operations by all threads are guaranteed to be correct reflecting either the value that has been set with this operation or set with another operation. A proper implicit or explicit memory barrier is needed before the value set with the operation is guaranteed to be readable with `atomic_read` from another thread.

Next, we have:

```
#define atomic_read(v) ((v)->counter)
```

which simply reads the counter value currently visible to the calling thread. The read is atomic in that the return value is guaranteed to be one of the values initialized or modified with the interface operations if a proper implicit or explicit memory barrier is used after possible runtime initialization by any other thread and the value is modified only with the interface operations. `atomic_read` does not guarantee that the runtime initialization by any other thread is visible yet, so the user of the interface must take care of that with a proper implicit or explicit memory barrier.

*** WARNING: `atomic_read()` and `atomic_set()` DO NOT IMPLY BARRIERS! ***

Some architectures may choose to use the `volatile` keyword, barriers, or inline assembly to guarantee some degree of immediacy for `atomic_read()` and `atomic_set()`. This is not uniformly guaranteed, and may change in the future, so all users of `atomic_t` should treat `atomic_read()` and `atomic_set()` as simple C statements that may be reordered or optimized away entirely by the compiler or processor, and explicitly invoke the appropriate compiler and/or memory barrier for each use case. Failure to do so will result in code that may suddenly break when used with different architectures or compiler optimizations, or even changes in unrelated code which changes how the compiler optimizes the section accessing `atomic_t` variables.

*** YOU HAVE BEEN WARNED! ***

Now, we move onto the atomic operation interfaces typically implemented with the help of assembly code.

```
void atomic_add(int i, atomic_t *v);
void atomic_sub(int i, atomic_t *v);
void atomic_inc(atomic_t *v);
void atomic_dec(atomic_t *v);
```

These four routines add and subtract integral values to/from the given `atomic_t` value. The first two routines pass explicit integers by which to make the adjustment, whereas the latter two use an implicit adjustment value of "1".

One very important aspect of these two routines is that they DO NOT require any explicit memory barriers. They need only perform the `atomic_t` counter update in an SMP safe manner.

Next, we have:

```
int atomic_inc_return(atomic_t *v);
int atomic_dec_return(atomic_t *v);
```

These routines add 1 and subtract 1, respectively, from the given `atomic_t` and return the new counter value after the operation is performed.

atomic_ops.txt

Unlike the above routines, it is required that explicit memory barriers are performed before and after the operation. It must be done such that all memory operations before and after the atomic operation calls are strongly ordered with respect to the atomic operation itself.

For example, it should behave as if a `smp_mb()` call existed both before and after the atomic operation.

If the atomic instructions used in an implementation provide explicit memory barrier semantics which satisfy the above requirements, that is fine as well.

Let's move on:

```
int atomic_add_return(int i, atomic_t *v);
int atomic_sub_return(int i, atomic_t *v);
```

These behave just like `atomic_{inc,dec}_return()` except that an explicit counter adjustment is given instead of the implicit "1". This means that like `atomic_{inc,dec}_return()`, the memory barrier semantics are required.

Next:

```
int atomic_inc_and_test(atomic_t *v);
int atomic_dec_and_test(atomic_t *v);
```

These two routines increment and decrement by 1, respectively, the given atomic counter. They return a boolean indicating whether the resulting counter value was zero or not.

It requires explicit memory barrier semantics around the operation as above.

```
int atomic_sub_and_test(int i, atomic_t *v);
```

This is identical to `atomic_dec_and_test()` except that an explicit decrement is given instead of the implicit "1". It requires explicit memory barrier semantics around the operation.

```
int atomic_add_negative(int i, atomic_t *v);
```

The given increment is added to the given atomic counter value. A boolean is return which indicates whether the resulting counter value is negative. It requires explicit memory barrier semantics around the operation.

Then:

```
int atomic_xchg(atomic_t *v, int new);
```

This performs an atomic exchange operation on the atomic variable `v`, setting the given new value. It returns the old value that the atomic variable `v` had just before the operation.

```

                                atomic_ops.txt
int atomic_cmpxchg(atomic_t *v, int old, int new);

```

This performs an atomic compare exchange operation on the atomic value `v`, with the given `old` and `new` values. Like all `atomic_xxx` operations, `atomic_cmpxchg` will only satisfy its atomicity semantics as long as all other accesses of `*v` are performed through `atomic_xxx` operations.

`atomic_cmpxchg` requires explicit memory barriers around the operation.

The semantics for `atomic_cmpxchg` are the same as those defined for 'cas' below.

Finally:

```
int atomic_add_unless(atomic_t *v, int a, int u);
```

If the atomic value `v` is not equal to `u`, this function adds `a` to `v`, and returns non zero. If `v` is equal to `u` then it returns zero. This is done as an atomic operation.

`atomic_add_unless` requires explicit memory barriers around the operation unless it fails (returns 0).

`atomic_inc_not_zero`, equivalent to `atomic_add_unless(v, 1, 0)`

If a caller requires memory barrier semantics around an `atomic_t` operation which does not return a value, a set of interfaces are defined which accomplish this:

```

void smp_mb__before_atomic_dec(void);
void smp_mb__after_atomic_dec(void);
void smp_mb__before_atomic_inc(void);
void smp_mb__after_atomic_inc(void);

```

For example, `smp_mb__before_atomic_dec()` can be used like so:

```

obj->dead = 1;
smp_mb__before_atomic_dec();
atomic_dec(&obj->ref_count);

```

It makes sure that all memory operations preceding the `atomic_dec()` call are strongly ordered with respect to the atomic counter operation. In the above example, it guarantees that the assignment of "1" to `obj->dead` will be globally visible to other cpus before the atomic counter decrement.

Without the explicit `smp_mb__before_atomic_dec()` call, the implementation could legally allow the atomic counter update visible to other cpus before the "`obj->dead = 1;`" assignment.

The other three interfaces listed are used to provide explicit ordering with respect to memory operations after an `atomic_dec()` call (`smp_mb__after_atomic_dec()`) and around `atomic_inc()` calls (`smp_mb__{before,after}_atomic_inc()`).

A missing memory barrier in the cases where they are required by the `atomic_t` implementation above can have disastrous results. Here is an example, which follows a pattern occurring frequently in the Linux kernel. It is the use of atomic counters to implement reference counting, and it works such that once the counter falls to zero it can be guaranteed that no other entity can be accessing the object:

```
static void obj_list_add(struct obj *obj, struct list_head *head)
{
    obj->active = 1;
    list_add(&obj->list, head);
}
```

```
static void obj_list_del(struct obj *obj)
{
    list_del(&obj->list);
    obj->active = 0;
}
```

```
static void obj_destroy(struct obj *obj)
{
    BUG_ON(obj->active);
    kfree(obj);
}
```

```
struct obj *obj_list_peek(struct list_head *head)
{
    if (!list_empty(head)) {
        struct obj *obj;

        obj = list_entry(head->next, struct obj, list);
        atomic_inc(&obj->refcnt);
        return obj;
    }
    return NULL;
}
```

```
void obj_poke(void)
{
    struct obj *obj;

    spin_lock(&global_list_lock);
    obj = obj_list_peek(&global_list);
    spin_unlock(&global_list_lock);

    if (obj) {
        obj->ops->poke(obj);
        if (atomic_dec_and_test(&obj->refcnt))
            obj_destroy(obj);
    }
}
```

```
void obj_timeout(struct obj *obj)
{
    spin_lock(&global_list_lock);
    obj_list_del(obj);
}
```

```

                                atomic_ops.txt
spin_unlock(&global_list_lock);

if (atomic_dec_and_test(&obj->refcnt))
    obj_destroy(obj);
}

```

(This is a simplification of the ARP queue management in the generic neighbour discover code of the networking. Olaf Kirch found a bug wrt. memory barriers in `kfree_skb()` that exposed the `atomic_t` memory barrier requirements quite clearly.)

Given the above scheme, it must be the case that the `obj->active` update done by the `obj` list deletion be visible to other processors before the atomic counter decrement is performed.

Otherwise, the counter could fall to zero, yet `obj->active` would still be set, thus triggering the assertion in `obj_destroy()`. The error sequence looks like this:

<pre> cpu 0 obj_poke() obj = obj_list_peek(); ... gains ref to obj, refcnt=2 atomic_dec_and_test() ... refcount drops to 0 ... obj_destroy() BUG() triggers since obj->active still seen as one </pre>	<pre> cpu 1 obj_timeout() obj_list_del(obj); obj->active = 0 visibility delayed ... atomic_dec_and_test() ... refcnt drops to 1 ... obj->active update visibility occurs </pre>
--	---

With the memory barrier semantics required of the `atomic_t` operations which return values, the above sequence of memory visibility can never happen. Specifically, in the above case the `atomic_dec_and_test()` counter decrement would not become globally visible until the `obj->active` update does.

As a historical note, 32-bit Sparc used to only allow usage of 24-bits of its `atomic_t` type. This was because it used 8 bits as a spinlock for SMP safety. Sparc32 lacked a "compare and swap" type instruction. However, 32-bit Sparc has since been moved over to a "hash table of spinlocks" scheme, that allows the full 32-bit counter to be realized. Essentially, an array of spinlocks are indexed into based upon the address of the `atomic_t` being operated on, and that lock protects the atomic operation. Parisc uses the same scheme.

Another note is that the `atomic_t` operations returning values are extremely slow on an old 386.

We will now cover the atomic bitmask operations. You will find that their SMP and memory barrier semantics are similar in shape and scope

to the atomic_t ops above.

Native atomic bit operations are defined to operate on objects aligned to the size of an "unsigned long" C data type, and are least of that size. The endianness of the bits within each "unsigned long" are the native endianness of the cpu.

```
void set_bit(unsigned long nr, volatile unsigned long *addr);
void clear_bit(unsigned long nr, volatile unsigned long *addr);
void change_bit(unsigned long nr, volatile unsigned long *addr);
```

These routines set, clear, and change, respectively, the bit number indicated by "nr" on the bit mask pointed to by "ADDR".

They must execute atomically, yet there are no implicit memory barrier semantics required of these interfaces.

```
int test_and_set_bit(unsigned long nr, volatile unsigned long *addr);
int test_and_clear_bit(unsigned long nr, volatile unsigned long *addr);
int test_and_change_bit(unsigned long nr, volatile unsigned long *addr);
```

Like the above, except that these routines return a boolean which indicates whether the changed bit was set BEFORE the atomic bit operation.

WARNING! It is incredibly important that the value be a boolean, ie. "0" or "1". Do not try to be fancy and save a few instructions by declaring the above to return "long" and just returning something like "old_val & mask" because that will not work.

For one thing, this return value gets truncated to int in many code paths using these interfaces, so on 64-bit if the bit is set in the upper 32-bits then testers will never see that.

One great example of where this problem crops up are the thread_info flag operations. Routines such as test_and_set_ti_thread_flag() chop the return value into an int. There are other places where things like this occur as well.

These routines, like the atomic_t counter operations returning values, require explicit memory barrier semantics around their execution. All memory operations before the atomic bit operation call must be made visible globally before the atomic bit operation is made visible. Likewise, the atomic bit operation must be visible globally before any subsequent memory operation is made visible. For example:

```
obj->dead = 1;
if (test_and_set_bit(0, &obj->flags))
    /* ... */;
obj->killed = 1;
```

The implementation of test_and_set_bit() must guarantee that "obj->dead = 1;" is visible to cpus before the atomic memory operation done by test_and_set_bit() becomes visible. Likewise, the atomic memory operation done by test_and_set_bit() must become visible before "obj->killed = 1;" is visible.

Finally there is the basic operation:

```
int test_bit(unsigned long nr, __const__ volatile unsigned long *addr);
```

Which returns a boolean indicating if bit "nr" is set in the bitmask pointed to by "addr".

If explicit memory barriers are required around `clear_bit()` (which does not return a value, and thus does not need to provide memory barrier semantics), two interfaces are provided:

```
void smp_mb__before_clear_bit(void);
void smp_mb__after_clear_bit(void);
```

They are used as follows, and are akin to their `atomic_t` operation brothers:

```
/* All memory operations before this call will
 * be globally visible before the clear_bit().
 */
smp_mb__before_clear_bit();
clear_bit( ... );

/* The clear_bit() will be visible before all
 * subsequent memory operations.
 */
smp_mb__after_clear_bit();
```

There are two special bitops with lock barrier semantics (acquire/release, same as spinlocks). These operate in the same way as their `non_lock/unlock` postfix variants, except that they are to provide acquire/release semantics, respectively. This means they can be used for `bit_spin_trylock` and `bit_spin_unlock` type operations without specifying any more barriers.

```
int test_and_set_bit_lock(unsigned long nr, unsigned long *addr);
void clear_bit_unlock(unsigned long nr, unsigned long *addr);
void __clear_bit_unlock(unsigned long nr, unsigned long *addr);
```

The `__clear_bit_unlock` version is non-atomic, however it still implements unlock barrier semantics. This can be useful if the lock itself is protecting the other bits in the word.

Finally, there are non-atomic versions of the bitmask operations provided. They are used in contexts where some other higher-level SMP locking scheme is being used to protect the bitmask, and thus less expensive non-atomic operations may be used in the implementation. They have names similar to the above bitmask operation interfaces, except that two underscores are prefixed to the interface name.

```
void __set_bit(unsigned long nr, volatile unsigned long *addr);
void __clear_bit(unsigned long nr, volatile unsigned long *addr);
void __change_bit(unsigned long nr, volatile unsigned long *addr);
int __test_and_set_bit(unsigned long nr, volatile unsigned long *addr);
int __test_and_clear_bit(unsigned long nr, volatile unsigned long
*addr);
```


atomic_ops.txt

```
int __test_and_change_bit(unsigned long nr, volatile unsigned long
*addr);
```

These non-atomic variants also do not require any special memory barrier semantics.

The routines `xchg()` and `cmpxchg()` need the same exact memory barriers as the atomic and bit operations returning values.

Spinlocks and rwlocks have memory barrier expectations as well. The rule to follow is simple:

- 1) When acquiring a lock, the implementation must make it globally visible before any subsequent memory operation.
- 2) When releasing a lock, the implementation must make it such that all previous memory operations are globally visible before the lock release.

Which finally brings us to `_atomic_dec_and_lock()`. There is an architecture-neutral version implemented in `lib/dec_and_lock.c`, but most platforms will wish to optimize this in assembler.

```
int _atomic_dec_and_lock(atomic_t *atomic, spinlock_t *lock);
```

Atomically decrement the given counter, and if will drop to zero atomically acquire the given spinlock and perform the decrement of the counter to zero. If it does not drop to zero, do nothing with the spinlock.

It is actually pretty simple to get the memory barrier correct. Simply satisfy the spinlock grab requirements, which is make sure the spinlock operation is globally visible before any subsequent memory operation.

We can demonstrate this operation more clearly if we define an abstract atomic operation:

```
long cas(long *mem, long old, long new);
```

"cas" stands for "compare and swap". It atomically:

- 1) Compares "old" with the value currently at "mem".
- 2) If they are equal, "new" is written to "mem".
- 3) Regardless, the current value at "mem" is returned.

As an example usage, here is what an atomic counter update might look like:

```
void example_atomic_inc(long *counter)
{
    long old, new, ret;

    while (1) {
        old = *counter;
        new = old + 1;
```

atomic_ops.txt

```
        ret = cas(counter, old, new);
        if (ret == old)
            break;
    }
}
```

Let's use `cas()` in order to build a pseudo-C `atomic_dec_and_lock()`:

```
int _atomic_dec_and_lock(atomic_t *atomic, spinlock_t *lock)
{
    long old, new, ret;
    int went_to_zero;

    went_to_zero = 0;
    while (1) {
        old = atomic_read.atomic;
        new = old - 1;
        if (new == 0) {
            went_to_zero = 1;
            spin_lock(lock);
        }
        ret = cas.atomic, old, new);
        if (ret == old)
            break;
        if (went_to_zero) {
            spin_unlock(lock);
            went_to_zero = 0;
        }
    }

    return went_to_zero;
}
```

Now, as far as memory barriers go, as long as `spin_lock()` strictly orders all subsequent memory operations (including the `cas()`) with respect to itself, things will be fine.

Said another way, `_atomic_dec_and_lock()` must guarantee that a counter dropping to zero is never made visible before the spinlock being acquired.

Note that this also means that for the case where the counter is not dropping to zero, there are no memory ordering requirements.