

## cputopology.txt

Export CPU topology info via sysfs. Items (attributes) are similar to /proc/cpuinfo.

1) /sys/devices/system/cpu/cpuX/topology/physical\_package\_id:

physical package id of cpuX. Typically corresponds to a physical socket number, but the actual value is architecture and platform dependent.

2) /sys/devices/system/cpu/cpuX/topology/core\_id:

the CPU core ID of cpuX. Typically it is the hardware platform's identifier (rather than the kernel's). The actual value is architecture and platform dependent.

3) /sys/devices/system/cpu/cpuX/topology/thread\_siblings:

internal kernel map of cpuX's hardware threads within the same core as cpuX

4) /sys/devices/system/cpu/cpuX/topology/core\_siblings:

internal kernel map of cpuX's hardware threads within the same physical\_package\_id.

To implement it in an architecture-neutral way, a new source file, drivers/base/topology.c, is to export the 4 attributes.

For an architecture to support this feature, it must define some of these macros in include/asm-XXX/topology.h:

```
#define topology_physical_package_id(cpu)
#define topology_core_id(cpu)
#define topology_thread_cpumask(cpu)
#define topology_core_cpumask(cpu)
```

The type of \*\*\_id is int.

The type of siblings is (const) struct cpumask \*.

To be consistent on all architectures, include/linux/topology.h provides default definitions for any of the above macros that are not defined by include/asm-XXX/topology.h:

- 1) physical\_package\_id: -1
- 2) core\_id: 0
- 3) thread\_siblings: just the given CPU
- 4) core\_siblings: just the given CPU

Additionally, CPU topology information is provided under /sys/devices/system/cpu and includes these files. The internal source for the output is in brackets ("[]").

kernel\_max: the maximum CPU index allowed by the kernel configuration.  
[NR\_CPUS-1]

offline: CPUs that are not online because they have been  
HOTPLUGGED off (see cpu-hotplug.txt) or exceed the limit

```

                                cputopology.txt
of CPUs allowed by the kernel configuration (kernel_max
above). [ $\sim$ cpu_online_mask + cpus  $\geq$  NR_CPUS]

online:    CPUs that are online and being scheduled [cpu_online_mask]

possible:  CPUs that have been allocated resources and can be
            brought online if they are present. [cpu_possible_mask]

present:   CPUs that have been identified as being present in the
            system. [cpu_present_mask]

```

The format for the above output is compatible with `cpulist_parse()` [see `<linux/cpumask.h>`]. Some examples follow.

In this example, there are 64 CPUs in the system but cpus 32-63 exceed the kernel max which is limited to 0..31 by the NR\_CPUS config option being 32. Note also that CPUs 2 and 4-31 are not online but could be brought online as they are both present and possible.

```

kernel_max: 31
  offline: 2, 4-31, 32-63
    online: 0-1, 3
  possible: 0-31
    present: 0-31

```

In this example, the NR\_CPUS config option is 128, but the kernel was started with `possible_cpus=144`. There are 4 CPUs in the system and `cpu2` was manually taken offline (and is the only CPU that can be brought online.)

```

kernel_max: 127
  offline: 2, 4-127, 128-143
    online: 0-1, 3
  possible: 0-127
    present: 0-3

```

See `cpu-hotplug.txt` for the `possible_cpus=NUM` kernel start parameter as well as more information on the various cpumasks.