

DMA-API.txt  
Dynamic DMA mapping using the generic device  
=====

James E. J. Bottomley <James.Bottomley@HansenPartnership.com>

This document describes the DMA API. For a more gentle introduction of the API (and actual examples) see Documentation/DMA-API-HOWTO.txt.

This API is split into two pieces. Part I describes the API. Part II describes the extensions to the API for supporting non-consistent memory machines. Unless you know that your driver absolutely has to support non-consistent platforms (this is usually only legacy platforms) you should only use the API described in part I.

Part I - dma\_ API  
-----

To get the dma\_ API, you must #include <linux/dma-mapping.h>

Part Ia - Using large dma-coherent buffers  
-----

```
void *
dma_alloc_coherent(struct device *dev, size_t size,
                  dma_addr_t *dma_handle, gfp_t flag)
```

Consistent memory is memory for which a write by either the device or the processor can immediately be read by the processor or device without having to worry about caching effects. (You may however need to make sure to flush the processor's write buffers before telling devices to read that memory.)

This routine allocates a region of <size> bytes of consistent memory. It also returns a <dma\_handle> which may be cast to an unsigned integer the same width as the bus and used as the physical address base of the region.

Returns: a pointer to the allocated region (in the processor's virtual address space) or NULL if the allocation failed.

Note: consistent memory can be expensive on some platforms, and the minimum allocation length may be as big as a page, so you should consolidate your requests for consistent memory as much as possible. The simplest way to do that is to use the dma\_pool calls (see below).

The flag parameter (dma\_alloc\_coherent only) allows the caller to specify the GFP\_ flags (see kmalloc) for the allocation (the implementation may choose to ignore flags that affect the location of the returned memory, like GFP\_DMA).

```
void
dma_free_coherent(struct device *dev, size_t size, void *cpu_addr,
                  dma_addr_t dma_handle)
```

## DMA-API.txt

Free the region of consistent memory you previously allocated. `dev`, `size` and `dma_handle` must all be the same as those passed into the `consistent allocate`. `cpu_addr` must be the virtual address returned by the `consistent allocate`.

Note that unlike their sibling allocation calls, these routines may only be called with IRQs enabled.

### Part Ib - Using small dma-coherent buffers

---

To get this part of the `dma_ API`, you must `#include <linux/dmapool.h>`

Many drivers need lots of small dma-coherent memory regions for DMA descriptors or I/O buffers. Rather than allocating in units of a page or more using `dma_alloc_coherent()`, you can use DMA pools. These work much like a `struct kmem_cache`, except that they use the dma-coherent allocator, not `__get_free_pages()`. Also, they understand common hardware constraints for alignment, like queue heads needing to be aligned on N-byte boundaries.

```
struct dma_pool *  
dma_pool_create(const char *name, struct device *dev,  
                size_t size, size_t align, size_t alloc);
```

The `pool create()` routines initialize a pool of dma-coherent buffers for use with a given device. It must be called in a context which can sleep.

The "name" is for diagnostics (like a `struct kmem_cache` name); `dev` and `size` are like what you'd pass to `dma_alloc_coherent()`. The device's hardware alignment requirement for this type of data is "align" (which is expressed in bytes, and must be a power of two). If your device has no boundary crossing restrictions, pass 0 for `alloc`; passing 4096 says memory allocated from this pool must not cross 4KByte boundaries.

```
void *dma_pool_alloc(struct dma_pool *pool, gfp_t gfp_flags,  
                    dma_addr_t *dma_handle);
```

This allocates memory from the pool; the returned memory will meet the size and alignment requirements specified at creation time. Pass `GFP_ATOMIC` to prevent blocking, or if it's permitted (not in `interrupt`, not holding SMP locks), pass `GFP_KERNEL` to allow blocking. Like `dma_alloc_coherent()`, this returns two values: an address usable by the cpu, and the dma address usable by the pool's device.

```
void dma_pool_free(struct dma_pool *pool, void *vaddr,  
                  dma_addr_t addr);
```

This puts memory back into the pool. The pool is what was passed to the pool allocation routine; the cpu (`vaddr`) and dma addresses are what were returned when that routine allocated the memory being freed.

```
void dma_pool_destroy(struct dma_pool *pool);
```

The pool destroy() routines free the resources of the pool. They must be called in a context which can sleep. Make sure you've freed all allocated memory back to the pool before you destroy it.

#### Part Ic - DMA addressing limitations

---

```
int
dma_supported(struct device *dev, u64 mask)
```

Checks to see if the device can support DMA to the memory described by mask.

Returns: 1 if it can and 0 if it can't.

Notes: This routine merely tests to see if the mask is possible. It won't change the current mask settings. It is more intended as an internal API for use by the platform than an external API for use by driver writers.

```
int
dma_set_mask(struct device *dev, u64 mask)
```

Checks to see if the mask is possible and updates the device parameters if it is.

Returns: 0 if successful and a negative error if not.

```
int
dma_set_coherent_mask(struct device *dev, u64 mask)
```

Checks to see if the mask is possible and updates the device parameters if it is.

Returns: 0 if successful and a negative error if not.

```
u64
dma_get_required_mask(struct device *dev)
```

This API returns the mask that the platform requires to operate efficiently. Usually this means the returned mask is the minimum required to cover all of memory. Examining the required mask gives drivers with variable descriptor sizes the opportunity to use smaller descriptors as necessary.

Requesting the required mask does not alter the current mask. If you wish to take advantage of it, you should issue a dma\_set\_mask() call to set the mask to the value returned.

#### Part Id - Streaming DMA mappings

```
-----  
dma_addr_t  
dma_map_single(struct device *dev, void *cpu_addr, size_t size,  
               enum dma_data_direction direction)
```

Maps a piece of processor virtual memory so it can be accessed by the device and returns the physical handle of the memory.

The direction for both api's may be converted freely by casting. However the dma\_ API uses a strongly typed enumerator for its direction:

DMA_NONE	no direction (used for debugging)
DMA_TO_DEVICE	data is going from the memory to the device
DMA_FROM_DEVICE	data is coming from the device to the memory
DMA_BIDIRECTIONAL	direction isn't known

Notes: Not all memory regions in a machine can be mapped by this API. Further, regions that appear to be physically contiguous in kernel virtual space may not be contiguous as physical memory. Since this API does not provide any scatter/gather capability, it will fail if the user tries to map a non-physically contiguous piece of memory. For this reason, it is recommended that memory mapped by this API be obtained only from sources which guarantee it to be physically contiguous (like kmalloc).

Further, the physical address of the memory must be within the dma\_mask of the device (the dma\_mask represents a bit mask of the addressable region for the device. I.e., if the physical address of the memory anded with the dma\_mask is still equal to the physical address, then the device can perform DMA to the memory). In order to ensure that the memory allocated by kmalloc is within the dma\_mask, the driver may specify various platform-dependent flags to restrict the physical memory range of the allocation (e.g. on x86, GFP\_DMA guarantees to be within the first 16Mb of available physical memory, as required by ISA devices).

Note also that the above constraints on physical contiguity and dma\_mask may not apply if the platform has an IOMMU (a device which supplies a physical to virtual mapping between the I/O memory bus and the device). However, to be portable, device driver writers may *\*not\** assume that such an IOMMU exists.

Warnings: Memory coherency operates at a granularity called the cache line width. In order for memory mapped by this API to operate correctly, the mapped region must begin exactly on a cache line boundary and end exactly on one (to prevent two separately mapped regions from sharing a single cache line). Since the cache line size may not be known at compile time, the API will not enforce this requirement. Therefore, it is recommended that driver writers who don't take special care to determine the cache line size at run time only map virtual regions that begin and end on page boundaries (which are guaranteed also to be cache line boundaries).

DMA\_TO\_DEVICE synchronisation must be done after the last modification

of the memory region by the software and before it is handed off to the driver. Once this primitive is used, memory covered by this primitive should be treated as read-only by the device. If the device may write to it at any point, it should be DMA\_BIDIRECTIONAL (see below).

DMA\_FROM\_DEVICE synchronisation must be done before the driver accesses data that may be changed by the device. This memory should be treated as read-only by the driver. If the driver needs to write to it at any point, it should be DMA\_BIDIRECTIONAL (see below).

DMA\_BIDIRECTIONAL requires special handling: it means that the driver isn't sure if the memory was modified before being handed off to the device and also isn't sure if the device will also modify it. Thus, you must always sync bidirectional memory twice: once before the memory is handed off to the device (to make sure all memory changes are flushed from the processor) and once before the data may be accessed after being used by the device (to make sure any processor cache lines are updated with data that the device may have changed).

```
void
dma_unmap_single(struct device *dev, dma_addr_t dma_addr, size_t size,
                 enum dma_data_direction direction)
```

Unmaps the region previously mapped. All the parameters passed in must be identical to those passed in (and returned) by the mapping API.

```
dma_addr_t
dma_map_page(struct device *dev, struct page *page,
             unsigned long offset, size_t size,
             enum dma_data_direction direction)

void
dma_unmap_page(struct device *dev, dma_addr_t dma_address, size_t size,
               enum dma_data_direction direction)
```

API for mapping and unmapping for pages. All the notes and warnings for the other mapping APIs apply here. Also, although the <offset> and <size> parameters are provided to do partial page mapping, it is recommended that you never use these unless you really know what the cache width is.

```
int
dma_mapping_error(struct device *dev, dma_addr_t dma_addr)
```

In some circumstances dma\_map\_single and dma\_map\_page will fail to create a mapping. A driver can check for these errors by testing the returned dma address with dma\_mapping\_error(). A non-zero return value means the mapping could not be created and the driver should take appropriate action (e.g. reduce current DMA mapping usage or delay and try again later).

```
int
dma_map_sg(struct device *dev, struct scatterlist *sg,
           int nents, enum dma_data_direction direction)
```

Returns: the number of physical segments mapped (this may be shorter

than <nents> passed in if some elements of the scatter/gather list are physically or virtually adjacent and an IOMMU maps them with a single entry).

Please note that the sg cannot be mapped again if it has been mapped once. The mapping process is allowed to destroy information in the sg.

As with the other mapping interfaces, dma\_map\_sg can fail. When it does, 0 is returned and a driver must take appropriate action. It is critical that the driver do something, in the case of a block driver aborting the request or even oopsing is better than doing nothing and corrupting the filesystem.

With scatterlists, you use the resulting mapping like this:

```
int i, count = dma_map_sg(dev, sglist, nents, direction);
struct scatterlist *sg;

for_each_sg(sglist, sg, count, i) {
    hw_address[i] = sg_dma_address(sg);
    hw_len[i] = sg_dma_len(sg);
}
```

where nents is the number of entries in the sglist.

The implementation is free to merge several consecutive sglist entries into one (e.g. with an IOMMU, or if several pages just happen to be physically contiguous) and returns the actual number of sg entries it mapped them to. On failure 0, is returned.

Then you should loop count times (note: this can be less than nents times) and use sg\_dma\_address() and sg\_dma\_len() macros where you previously accessed sg->address and sg->length as shown above.

```
void
dma_unmap_sg(struct device *dev, struct scatterlist *sg,
             int nhwentries, enum dma_data_direction direction)
```

Unmap the previously mapped scatter/gather list. All the parameters must be the same as those and passed in to the scatter/gather mapping API.

Note: <nents> must be the number you passed in, *\*not\** the number of physical entries returned.

```
void
dma_sync_single_for_cpu(struct device *dev, dma_addr_t dma_handle, size_t size,
                       enum dma_data_direction direction)

void
dma_sync_single_for_device(struct device *dev, dma_addr_t dma_handle, size_t
size,
                           enum dma_data_direction direction)

void
dma_sync_sg_for_cpu(struct device *dev, struct scatterlist *sg, int nelems,
                    enum dma_data_direction direction)

void
```

## DMA-API.txt

```
dma_sync_sg_for_device(struct device *dev, struct scatterlist *sg, int nelems,  
                        enum dma_data_direction direction)
```

Synchronise a single contiguous or scatter/gather mapping for the cpu and device. With the sync\_sg API, all the parameters must be the same as those passed into the single mapping API. With the sync\_single API, you can use dma\_handle and size parameters that aren't identical to those passed into the single mapping API to do a partial sync.

Notes: You must do this:

- Before reading values that have been written by DMA from the device (use the DMA\_FROM\_DEVICE direction)
- After writing values that will be written to the device using DMA (use the DMA\_TO\_DEVICE) direction
- before \*and\* after handing memory to the device if the memory is DMA\_BIDIRECTIONAL

See also dma\_map\_single().

```
dma_addr_t  
dma_map_single_attrs(struct device *dev, void *cpu_addr, size_t size,  
                     enum dma_data_direction dir,  
                     struct dma_attrs *attrs)
```

```
void  
dma_unmap_single_attrs(struct device *dev, dma_addr_t dma_addr,  
                       size_t size, enum dma_data_direction dir,  
                       struct dma_attrs *attrs)
```

```
int  
dma_map_sg_attrs(struct device *dev, struct scatterlist *sgl,  
                 int nents, enum dma_data_direction dir,  
                 struct dma_attrs *attrs)
```

```
void  
dma_unmap_sg_attrs(struct device *dev, struct scatterlist *sgl,  
                  int nents, enum dma_data_direction dir,  
                  struct dma_attrs *attrs)
```

The four functions above are just like the counterpart functions without the \_attrs suffixes, except that they pass an optional struct dma\_attrs\*.

struct dma\_attrs encapsulates a set of "dma attributes". For the definition of struct dma\_attrs see linux/dma-attrs.h.

The interpretation of dma attributes is architecture-specific, and each attribute should be documented in Documentation/DMA-attributes.txt.

If struct dma\_attrs\* is NULL, the semantics of each of these functions is identical to those of the corresponding function without the \_attrs suffix. As a result dma\_map\_single\_attrs() can generally replace dma\_map\_single(), etc.

As an example of the use of the \*\_attrs functions, here's how

## DMA-API.txt

you could pass an attribute DMA\_ATTR\_FOO when mapping memory for DMA:

```
#include <linux/dma-attrs.h>
/* DMA_ATTR_FOO should be defined in linux/dma-attrs.h and
 * documented in Documentation/DMA-attributes.txt */
...

DEFINE_DMA_ATTRS(attrs);
dma_set_attr(DMA_ATTR_FOO, &attrs);
....
n = dma_map_sg_attrs(dev, sg, nents, DMA_TO_DEVICE, &attr);
....
```

Architectures that care about DMA\_ATTR\_FOO would check for its presence in their implementations of the mapping and unmapping routines, e.g.:

```
void whizco_dma_map_sg_attrs(struct device *dev, dma_addr_t dma_addr,
                             size_t size, enum dma_data_direction dir,
                             struct dma_attrs *attrs)
{
    ....
    int foo = dma_get_attr(DMA_ATTR_FOO, attrs);
    ....
    if (foo)
        /* twizzle the frobnozzle */
    ....
}
```

## Part II - Advanced dma\_ usage

---

Warning: These pieces of the DMA API should not be used in the majority of cases, since they cater for unlikely corner cases that don't belong in usual drivers.

If you don't understand how cache line coherency works between a processor and an I/O device, you should not be using this part of the API at all.

```
void *
dma_alloc_noncoherent(struct device *dev, size_t size,
                      dma_addr_t *dma_handle, gfp_t flag)
```

Identical to dma\_alloc\_coherent() except that the platform will choose to return either consistent or non-consistent memory as it sees fit. By using this API, you are guaranteeing to the platform that you have all the correct and necessary sync points for this memory in the driver should it choose to return non-consistent memory.

Note: where the platform can return consistent memory, it will guarantee that the sync points become nops.

Warning: Handling non-consistent memory is a real pain. You should only ever use this API if you positively know your driver will be



required to work on one of the rare (usually non-PCI) architectures that simply cannot make consistent memory.

```
void
dma_free_noncoherent(struct device *dev, size_t size, void *cpu_addr,
                    dma_addr_t dma_handle)
```

Free memory allocated by the nonconsistent API. All parameters must be identical to those passed in (and returned by `dma_alloc_noncoherent()`).

```
int
dma_is_consistent(struct device *dev, dma_addr_t dma_handle)
```

Returns true if the device `dev` is performing consistent DMA on the memory area pointed to by the `dma_handle`.

```
int
dma_get_cache_alignment(void)
```

Returns the processor cache alignment. This is the absolute minimum alignment *and* width that you must observe when either mapping memory or doing partial flushes.

Notes: This API may return a number *larger* than the actual cache line, but it will guarantee that one or more cache lines fit exactly into the width returned by this call. It will also always be a power of two for easy alignment.

```
void
dma_cache_sync(struct device *dev, void *vaddr, size_t size,
               enum dma_data_direction direction)
```

Do a partial sync of memory that was allocated by `dma_alloc_noncoherent()`, starting at virtual address `vaddr` and continuing on for `size`. Again, you *must* observe the cache line boundaries when doing this.

```
int
dma_declare_coherent_memory(struct device *dev, dma_addr_t bus_addr,
                           dma_addr_t device_addr, size_t size, int
                           flags)
```

Declare region of memory to be handed out by `dma_alloc_coherent` when it's asked for coherent memory for this device.

`bus_addr` is the physical address to which the memory is currently assigned in the bus responding region (this will be used by the platform to perform the mapping).

`device_addr` is the physical address the device needs to be programmed with actually to address this memory (this will be handed out as the `dma_addr_t` in `dma_alloc_coherent()`).

`size` is the size of the area (must be multiples of `PAGE_SIZE`).

flags can be or'd together and are:

`DMA_MEMORY_MAP` - request that the memory returned from `dma_alloc_coherent()` be directly writable.

`DMA_MEMORY_IO` - request that the memory returned from `dma_alloc_coherent()` be addressable using `read/write/memcpy_toio` etc.

One or both of these flags must be present.

`DMA_MEMORY_INCLUDES_CHILDREN` - make the declared memory be allocated by `dma_alloc_coherent` of any child devices of this one (for memory residing on a bridge).

`DMA_MEMORY_EXCLUSIVE` - only allocate memory from the declared regions. Do not allow `dma_alloc_coherent()` to fall back to system memory when it's out of memory in the declared region.

The return value will be either `DMA_MEMORY_MAP` or `DMA_MEMORY_IO` and must correspond to a passed in flag (i.e. no returning `DMA_MEMORY_IO` if only `DMA_MEMORY_MAP` were passed in) for success or zero for failure.

Note, for `DMA_MEMORY_IO` returns, all subsequent memory returned by `dma_alloc_coherent()` may no longer be accessed directly, but instead must be accessed using the correct bus functions. If your driver isn't prepared to handle this contingency, it should not specify `DMA_MEMORY_IO` in the input flags.

As a simplification for the platforms, only *\*one\** such region of memory may be declared per device.

For reasons of efficiency, most platforms choose to track the declared region only at the granularity of a page. For smaller allocations, you should use the `dma_pool()` API.

```
void
dma_release_declared_memory(struct device *dev)
```

Remove the memory region previously declared from the system. This API performs *\*no\** in-use checking for this region and will return unconditionally having removed all the required structures. It is the driver's job to ensure that no parts of this memory region are currently in use.

```
void *
dma_mark_declared_memory_occupied(struct device *dev,
                                   dma_addr_t device_addr, size_t size)
```

This is used to occupy specific regions of the declared space (`dma_alloc_coherent()` will hand out the first free region it finds).

`device_addr` is the *\*device\** address of the region requested.

`size` is the size (and should be a page-sized multiple).

## DMA-API.txt

The return value will be either a pointer to the processor virtual address of the memory, or an error (via PTR\_ERR()) if any part of the region is occupied.

### Part III - Debug drivers use of the DMA-API

---

The DMA-API as described above has some constraints. DMA addresses must be released with the corresponding function with the same size for example. With the advent of hardware IOMMUs it becomes more and more important that drivers do not violate those constraints. In the worst case such a violation can result in data corruption up to destroyed filesystems.

To debug drivers and find bugs in the usage of the DMA-API checking code can be compiled into the kernel which will tell the developer about those violations. If your architecture supports it you can select the "Enable debugging of DMA-API usage" option in your kernel configuration. Enabling this option has a performance impact. Do not enable it in production kernels.

If you boot the resulting kernel will contain code which does some bookkeeping about what DMA memory was allocated for which device. If this code detects an error it prints a warning message with some details into your kernel log. An example warning message may look like this:

```
-----[ cut here ]-----
WARNING: at /data2/repos/linux-2.6-iommu/lib/dma-debug.c:448
        check_unmap+0x203/0x490()
Hardware name:
forcedeth 0000:00:08.0: DMA-API: device driver frees DMA memory with wrong
        function [device address=0x00000000640444be] [size=66 bytes] [mapped as
single] [unmapped as page]
Modules linked in: nfsd exportfs bridge stp llc r8169
Pid: 0, comm: swapper Tainted: G           W 2.6.28-dmatest-09289-g8bb99c0 #1
Call Trace:
<IRQ>  [<ffffffff80240b22>] warn_slowpath+0xf2/0x130
[<ffffffff80647b70>] _spin_unlock+0x10/0x30
[<ffffffff80537e75>] usb_hcd_link_urb_to_ep+0x75/0xc0
[<ffffffff80647c22>] _spin_unlock_irqrestore+0x12/0x40
[<ffffffff8055347f>] ohci_urb_enqueue+0x19f/0x7c0
[<ffffffff80252f96>] queue_work+0x56/0x60
[<ffffffff80237e10>] enqueue_task_fair+0x20/0x50
[<ffffffff80539279>] usb_hcd_submit_urb+0x379/0xbc0
[<ffffffff803b78c3>] cpumask_next_and+0x23/0x40
[<ffffffff80235177>] find_busiest_group+0x207/0x8a0
[<ffffffff8064784f>] _spin_lock_irqsave+0x1f/0x50
[<ffffffff803c7ea3>] check_unmap+0x203/0x490
[<ffffffff803c8259>] debug_dma_unmap_page+0x49/0x50
[<ffffffff80485f26>] nv_tx_done_optimized+0xc6/0x2c0
[<ffffffff80486c13>] nv_nic_irq_optimized+0x73/0x2b0
[<ffffffff8026df84>] handle_IRQ_event+0x34/0x70
[<ffffffff8026ffe9>] handle_edge_irq+0xc9/0x150
[<ffffffff8020e3ab>] do_IRQ+0xcb/0x1c0
[<ffffffff8020c093>] ret_from_intr+0x0/0xa
<EOI> <4>---[ end trace f6435a98e2a38c0e ]---
```

The driver developer can find the driver and the device including a stacktrace

of the DMA-API call which caused this warning.

Per default only the first error will result in a warning message. All other errors will only silently counted. This limitation exist to prevent the code from flooding your kernel log. To support debugging a device driver this can be disabled via debugfs. See the debugfs interface documentation below for details.

The debugfs directory for the DMA-API debugging code is called `dma-api/`. In this directory the following files can currently be found:

<code>dma-api/all_errors</code>	This file contains a numeric value. If this value is not equal to zero the debugging code will print a warning for every error it finds into the kernel log. Be careful with this option, as it can easily flood your logs.
<code>dma-api/disabled</code>	This read-only file contains the character 'Y' if the debugging code is disabled. This can happen when it runs out of memory or if it was disabled at boot time
<code>dma-api/error_count</code>	This file is read-only and shows the total numbers of errors found.
<code>dma-api/num_errors</code>	The number in this file shows how many warnings will be printed to the kernel log before it stops. This number is initialized to one at system boot and be set by writing into this file
<code>dma-api/min_free_entries</code>	This read-only file can be read to get the minimum number of free <code>dma_debug_entries</code> the allocator has ever seen. If this value goes down to zero the code will disable itself because it is not longer reliable.
<code>dma-api/num_free_entries</code>	The current number of free <code>dma_debug_entries</code> in the allocator.
<code>dma-api/driver-filter</code>	You can write a name of a driver into this file to limit the debug output to requests from that particular driver. Write an empty string to that file to disable the filter and see all errors again.

If you have this code compiled into your kernel it will be enabled by default. If you want to boot without the bookkeeping anyway you can provide 'dma\_debug=off' as a boot parameter. This will disable DMA-API debugging. Notice that you can not enable it again at runtime. You have to reboot to do so.

If you want to see debug messages only for a special device driver you can

#### DMA-API.txt

specify the `dma_debug_driver=<drivername>` parameter. This will enable the driver filter at boot time. The debug code will only print errors for that driver afterwards. This filter can be disabled or changed later using `debugfs`.

When the code disables itself at runtime this is most likely because it ran out of `dma_debug_entries`. These entries are preallocated at boot. The number of preallocated entries is defined per architecture. If it is too low for you boot with '`dma_debug_entries=<your_desired_number>`' to overwrite the architectural default.