

```

                                debugobjects.templ.txt
<?xml version="1.0" encoding="UTF-8"?>
<!DOCTYPE book PUBLIC "-//OASIS//DTD DocBook XML V4.1.2//EN"
    "http://www.oasis-open.org/docbook/xml/4.1.2/docbookx.dtd" []>

<book id="debug-objects-guide">
  <bookinfo>
    <title>Debug objects life time</title>

    <authorgroup>
      <author>
        <firstname>Thomas</firstname>
        <surname>Gleixner</surname>
        <affiliation>
          <address>
            <email>tglx@linutronix.de</email>
          </address>
        </affiliation>
      </author>
    </authorgroup>

    <copyright>
      <year>2008</year>
      <holder>Thomas Gleixner</holder>
    </copyright>

    <legalnotice>
      <para>
        This documentation is free software; you can redistribute
        it and/or modify it under the terms of the GNU General Public
        License version 2 as published by the Free Software Foundation.
      </para>

      <para>
        This program is distributed in the hope that it will be
        useful, but WITHOUT ANY WARRANTY; without even the implied
        warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE.
        See the GNU General Public License for more details.
      </para>

      <para>
        You should have received a copy of the GNU General Public
        License along with this program; if not, write to the Free
        Software Foundation, Inc., 59 Temple Place, Suite 330, Boston,
        MA 02111-1307 USA
      </para>

      <para>
        For more details see the file COPYING in the source
        distribution of Linux.
      </para>
    </legalnotice>
  </bookinfo>

  <toc></toc>

  <chapter id="intro">

```

<title>Introduction</title>

<para>

debugobjects is a generic infrastructure to track the life time of kernel objects and validate the operations on those.

</para>

<para>

debugobjects is useful to check for the following error patterns:

<itemizedlist>

<listitem><para>Activation of uninitialized objects</para></listitem>

<listitem><para>Initialization of active objects</para></listitem>

<listitem><para>Usage of freed/destroyed objects</para></listitem>

</itemizedlist>

</para>

<para>

debugobjects is not changing the data structure of the real object so it can be compiled in with a minimal runtime impact and enabled on demand with a kernel command line option.

</para>

</chapter>

<chapter id="howto">

<title>Howto use debugobjects</title>

<para>

A kernel subsystem needs to provide a data structure which describes the object type and add calls into the debug code at appropriate places. The data structure to describe the object type needs at minimum the name of the object type. Optional functions can and should be provided to fixup detected problems so the kernel can continue to work and the debug information can be retrieved from a live system instead of hard core debugging with serial consoles and stack trace transcripts from the monitor.

</para>

<para>

The debug calls provided by debugobjects are:

<itemizedlist>

<listitem><para>debug_object_init</para></listitem>

<listitem><para>debug_object_init_on_stack</para></listitem>

<listitem><para>debug_object_activate</para></listitem>

<listitem><para>debug_object_deactivate</para></listitem>

<listitem><para>debug_object_destroy</para></listitem>

<listitem><para>debug_object_free</para></listitem>

</itemizedlist>

Each of these functions takes the address of the real object and a pointer to the object type specific debug description structure.

</para>

<para>

Each detected error is reported in the statistics and a limited number of errors are printk'ed including a full stack trace.

</para>

<para>

The statistics are available via /sys/kernel/debug/debug_objects/stats.

They provide information about the number of warnings and the number of successful fixups along with information about the usage of the internal tracking objects and the state of the

```

                                debugobjects.templ.txt
    internal tracking objects pool.
</para>
</chapter>
<chapter id="debugfunctions">
  <title>Debug functions</title>
  <sect1 id="prototypes">
    <title>Debug object function reference</title>
!Elib/debugobjects.c
  </sect1>
  <sect1 id="debug_object_init">
    <title>debug_object_init</title>
    <para>
      This function is called whenever the initialization function
      of a real object is called.
    </para>
    <para>
      When the real object is already tracked by debugobjects it is
      checked, whether the object can be initialized. Initializing
      is not allowed for active and destroyed objects. When
      debugobjects detects an error, then it calls the fixup_init
      function of the object type description structure if provided
      by the caller. The fixup function can correct the problem
      before the real initialization of the object happens. E.g. it
      can deactivate an active object in order to prevent damage to
      the subsystem.
    </para>
    <para>
      When the real object is not yet tracked by debugobjects,
      debugobjects allocates a tracker object for the real object
      and sets the tracker object state to ODEBUG_STATE_INIT. It
      verifies that the object is not on the callers stack. If it is
      on the callers stack then a limited number of warnings
      including a full stack trace is printk'ed. The calling code
      must use debug_object_init_on_stack() and remove the object
      before leaving the function which allocated it. See next
      section.
    </para>
  </sect1>

  <sect1 id="debug_object_init_on_stack">
    <title>debug_object_init_on_stack</title>
    <para>
      This function is called whenever the initialization function
      of a real object which resides on the stack is called.
    </para>
    <para>
      When the real object is already tracked by debugobjects it is
      checked, whether the object can be initialized. Initializing
      is not allowed for active and destroyed objects. When
      debugobjects detects an error, then it calls the fixup_init
      function of the object type description structure if provided
      by the caller. The fixup function can correct the problem
      before the real initialization of the object happens. E.g. it
      can deactivate an active object in order to prevent damage to
      the subsystem.
    </para>
  </sect1>

```

<para>
When the real object is not yet tracked by debugobjects debugobjects allocates a tracker object for the real object and sets the tracker object state to ODEBUG_STATE_INIT. It verifies that the object is on the callers stack.
</para>
<para>
An object which is on the stack must be removed from the tracker by calling debug_object_free() before the function which allocates the object returns. Otherwise we keep track of stale objects.
</para>
</sect1>

<sect1 id="debug_object_activate">
<title>debug_object_activate</title>
<para>
This function is called whenever the activation function of a real object is called.
</para>
<para>
When the real object is already tracked by debugobjects it is checked, whether the object can be activated. Activating is not allowed for active and destroyed objects. When debugobjects detects an error, then it calls the fixup_activate function of the object type description structure if provided by the caller. The fixup function can correct the problem before the real activation of the object happens. E.g. it can deactivate an active object in order to prevent damage to the subsystem.
</para>
<para>
When the real object is not yet tracked by debugobjects then the fixup_activate function is called if available. This is necessary to allow the legitimate activation of statically allocated and initialized objects. The fixup function checks whether the object is valid and calls the debug_objects_init() function to initialize the tracking of this object.
</para>
<para>
When the activation is legitimate, then the state of the associated tracker object is set to ODEBUG_STATE_ACTIVE.
</para>
</sect1>

<sect1 id="debug_object_deactivate">
<title>debug_object_deactivate</title>
<para>
This function is called whenever the deactivation function of a real object is called.
</para>
<para>
When the real object is tracked by debugobjects it is checked, whether the object can be deactivated. Deactivating is not allowed for untracked or destroyed objects.
</para>

```
<para>
  When the deactivation is legitimate, then the state of the
  associated tracker object is set to ODEBUG_STATE_INACTIVE.
</para>
</sect1>

<sect1 id="debug_object_destroy">
  <title>debug_object_destroy</title>
  <para>
    This function is called to mark an object destroyed. This is
    useful to prevent the usage of invalid objects, which are
    still available in memory: either statically allocated objects
    or objects which are freed later.
  </para>
  <para>
    When the real object is tracked by debugobjects it is checked,
    whether the object can be destroyed. Destruction is not
    allowed for active and destroyed objects. When debugobjects
    detects an error, then it calls the fixup_destroy function of
    the object type description structure if provided by the
    caller. The fixup function can correct the problem before the
    real destruction of the object happens. E.g. it can deactivate
    an active object in order to prevent damage to the subsystem.
  </para>
  <para>
    When the destruction is legitimate, then the state of the
    associated tracker object is set to ODEBUG_STATE_DESTROYED.
  </para>
</sect1>

<sect1 id="debug_object_free">
  <title>debug_object_free</title>
  <para>
    This function is called before an object is freed.
  </para>
  <para>
    When the real object is tracked by debugobjects it is checked,
    whether the object can be freed. Free is not allowed for
    active objects. When debugobjects detects an error, then it
    calls the fixup_free function of the object type description
    structure if provided by the caller. The fixup function can
    correct the problem before the real free of the object
    happens. E.g. it can deactivate an active object in order to
    prevent damage to the subsystem.
  </para>
  <para>
    Note that debug_object_free removes the object from the
    tracker. Later usage of the object is detected by the other
    debug checks.
  </para>
</sect1>
</chapter>
<chapter id="fixupfunctions">
  <title>Fixup functions</title>
  <sect1 id="debug_obj_descr">
    <title>Debug object type description structure</title>
```

```

!include/linux/debugobjects.h
</sect1>
<sect1 id="fixup_init">
  <title>fixup_init</title>
  <para>
    This function is called from the debug code whenever a problem
    in debug_object_init is detected. The function takes the
    address of the object and the state which is currently
    recorded in the tracker.
  </para>
  <para>
    Called from debug_object_init when the object state is:
    <itemizedlist>
      <listitem><para>ODEBUG_STATE_ACTIVE</para></listitem>
    </itemizedlist>
  </para>
  <para>
    The function returns 1 when the fixup was successful,
    otherwise 0. The return value is used to update the
    statistics.
  </para>
  <para>
    Note, that the function needs to call the debug_object_init()
    function again, after the damage has been repaired in order to
    keep the state consistent.
  </para>
</sect1>

<sect1 id="fixup_activate">
  <title>fixup_activate</title>
  <para>
    This function is called from the debug code whenever a problem
    in debug_object_activate is detected.
  </para>
  <para>
    Called from debug_object_activate when the object state is:
    <itemizedlist>
      <listitem><para>ODEBUG_STATE_NOTAVAILABLE</para></listitem>
      <listitem><para>ODEBUG_STATE_ACTIVE</para></listitem>
    </itemizedlist>
  </para>
  <para>
    The function returns 1 when the fixup was successful,
    otherwise 0. The return value is used to update the
    statistics.
  </para>
  <para>
    Note that the function needs to call the debug_object_activate()
    function again after the damage has been repaired in order to
    keep the state consistent.
  </para>
  <para>
    The activation of statically initialized objects is a special
    case. When debug_object_activate() has no tracked object for
    this object address then fixup_activate() is called with
    object state ODEBUG_STATE_NOTAVAILABLE. The fixup function

```

```

                                debugobjects.templ.txt
    needs to check whether this is a legitimate case of a
    statically initialized object or not. In case it is it calls
    debug_object_init() and debug_object_activate() to make the
    object known to the tracker and marked active. In this case
    the function should return 0 because this is not a real fixup.
</para>
</sect1>

<sect1 id="fixup_destroy">
  <title>fixup_destroy</title>
  <para>
    This function is called from the debug code whenever a problem
    in debug_object_destroy is detected.
  </para>
  <para>
    Called from debug_object_destroy when the object state is:
    <itemizedlist>
      <listitem><para>ODEBUG_STATE_ACTIVE</para></listitem>
    </itemizedlist>
  </para>
  <para>
    The function returns 1 when the fixup was successful,
    otherwise 0. The return value is used to update the
    statistics.
  </para>
</sect1>
<sect1 id="fixup_free">
  <title>fixup_free</title>
  <para>
    This function is called from the debug code whenever a problem
    in debug_object_free is detected. Further it can be called
    from the debug checks in kfree/vfree, when an active object is
    detected from the debug_check_no_obj_freed() sanity checks.
  </para>
  <para>
    Called from debug_object_free() or debug_check_no_obj_freed()
    when the object state is:
    <itemizedlist>
      <listitem><para>ODEBUG_STATE_ACTIVE</para></listitem>
    </itemizedlist>
  </para>
  <para>
    The function returns 1 when the fixup was successful,
    otherwise 0. The return value is used to update the
    statistics.
  </para>
</sect1>
</chapter>
<chapter id="bugs">
  <title>Known Bugs And Assumptions</title>
  <para>
    None (knock on wood).
  </para>
</chapter>
</book>

```