

mtdnand.tmpl.txt

```
<?xml version="1.0" encoding="UTF-8"?>
<!DOCTYPE book PUBLIC "-//OASIS//DTD DocBook XML V4.1.2//EN"
    "http://www.oasis-open.org/docbook/xml/4.1.2/docbookx.dtd" []>

<book id="MTD-NAND-Guide">
  <bookinfo>
    <title>MTD NAND Driver Programming Interface</title>

    <authorgroup>
      <author>
        <firstname>Thomas</firstname>
        <surname>Gleixner</surname>
        <affiliation>
          <address>
            <email>tglx@linutronix.de</email>
          </address>
        </affiliation>
      </author>
    </authorgroup>

    <copyright>
      <year>2004</year>
      <holder>Thomas Gleixner</holder>
    </copyright>

    <legalnotice>
      <para>
        This documentation is free software; you can redistribute
        it and/or modify it under the terms of the GNU General Public
        License version 2 as published by the Free Software Foundation.
      </para>

      <para>
        This program is distributed in the hope that it will be
        useful, but WITHOUT ANY WARRANTY; without even the implied
        warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE.
        See the GNU General Public License for more details.
      </para>

      <para>
        You should have received a copy of the GNU General Public
        License along with this program; if not, write to the Free
        Software Foundation, Inc., 59 Temple Place, Suite 330, Boston,
        MA 02111-1307 USA
      </para>

      <para>
        For more details see the file COPYING in the source
        distribution of Linux.
      </para>
    </legalnotice>
  </bookinfo>

  <toc></toc>

  <chapter id="intro">
```

<title>Introduction</title>

<para>

The generic NAND driver supports almost all NAND and AG-AND based chips and connects them to the Memory Technology Devices (MTD) subsystem of the Linux Kernel.

</para>

<para>

This documentation is provided for developers who want to implement board drivers or filesystem drivers suitable for NAND devices.

</para>

</chapter>

<chapter id="bugs">

<title>Known Bugs And Assumptions</title>

<para>

None.

</para>

</chapter>

<chapter id="dochints">

<title>Documentation hints</title>

<para>

The function and structure docs are autogenerated. Each function and struct member has a short description which is marked with an [XXX]

identifier.

The following chapters explain the meaning of those identifiers.

</para>

<sect1 id="Function_identifiers_XXX">

<title>Function identifiers [XXX]</title>

<para>

The functions are marked with [XXX] identifiers in the short comment. The identifiers explain the usage and scope of the functions. Following identifiers are used:

</para>

<itemizedlist>

<listitem><para>

[MTD Interface]</para><para>

These functions provide the interface to the MTD kernel API. They are not replacable and provide functionality which is complete hardware independent.

</para></listitem>

<listitem><para>

[NAND Interface]</para><para>

These functions are exported and provide the interface to the

NAND kernel API.

</para></listitem>

<listitem><para>

[GENERIC]</para><para>

Generic functions are not replacable and provide functionality which is complete hardware independent.

</para></listitem>

<listitem><para>

[DEFAULT]</para><para>

Default functions provide hardware related functionality which

is suitable

for most of the implementations. These functions can be replaced

by the board driver if necessary. Those functions are called via pointers in the NAND chip description structure. The board driver can set the functions which should be replaced by board dependent functions before calling nand_scan(). If the function pointer is NULL on entry to nand_scan() then the pointer is set to the default function which is suitable for the detected chip type.

</para></listitem>

</itemizedlist>

</sect1>

<sect1 id="Struct_member_identifiers_XXX">

<title>Struct member identifiers [XXX]</title>

<para>

The struct members are marked with [XXX] identifiers in the comment. The identifiers explain the usage and scope of the members. Following identifiers are used:

</para>

<itemizedlist>

<listitem><para>

[INTERN]</para><para>

These members are for NAND driver internal use only and must not be modified. Most of these values are calculated from the chip geometry information which is evaluated during nand_scan().

</para></listitem>

<listitem><para>

[REPLACEABLE]</para><para>

Replaceable members hold hardware related functions which can be provided by the board driver. The board driver can set the functions which should be replaced by board dependent functions before calling nand_scan(). If the function pointer is NULL on entry to nand_scan() then the pointer is set to the default function which is suitable for the detected chip type.

</para></listitem>

<listitem><para>

[BOARDSPECIFIC]</para><para>

Board specific members hold hardware related information which must be provided by the board driver. The board driver must set the function pointers and datafields before calling nand_scan().

</para></listitem>

<listitem><para>

[OPTIONAL]</para><para>

Optional members can hold information relevant for the board driver. The generic NAND driver code does not use this information.

```

                                mtdnand.tmpl.txt
                                </para></listitem>
                                </itemizedlist>
                                </sect1>
                                </chapter>

<chapter id="basicboarddriver">
  <title>Basic board driver</title>
  <para>
    For most boards it will be sufficient to provide just the
    basic functions and fill out some really board dependent
    members in the nand chip description structure.
  </para>
  <sect1 id="Basic_defines">
    <title>Basic defines</title>
    <para>
      At least you have to provide a mtd structure and
      a storage for the ioremap'ed chip address.
      You can allocate the mtd structure using kmalloc
      or you can allocate it statically.
      In case of static allocation you have to allocate
      a nand_chip structure too.
    </para>
    <para>
      Kmalloc based example
    </para>
    <programlisting>
static struct mtd_info *board_mtd;
static void __iomem *baseaddr;
    </programlisting>
    <para>
      Static example
    </para>
    <programlisting>
static struct mtd_info board_mtd;
static struct nand_chip board_chip;
static void __iomem *baseaddr;
    </programlisting>
  </sect1>
  <sect1 id="Partition_defines">
    <title>Partition defines</title>
    <para>
      If you want to divide your device into partitions, then
      enable the configuration switch CONFIG_MTD_PARTITIONS
and define
      a partitioning scheme suitable to your board.
    </para>
    <programlisting>
#define NUM_PARTITIONS 2
static struct mtd_partition partition_info[] = {
  { .name = "Flash partition 1",
    .offset = 0,
    .size = 8 * 1024 * 1024 },
  { .name = "Flash partition 2",
    .offset = MTDPART_OFS_NEXT,
    .size = MTDPART_SIZ_FULL },
};

```

```

                                mtdnand.tmpl.txt
        </programlisting>
</sect1>
<sect1 id="Hardware_control_functions">
    <title>Hardware control function</title>
    <para>
        The hardware control function provides access to the
        control pins of the NAND chip(s).
        The access can be done by GPIO pins or by address lines.
        If you use address lines, make sure that the timing
        requirements are met.
    </para>
    <para>
        <emphasis>GPIO based example</emphasis>
    </para>
    <programlisting>
static void board_hwcontrol(struct mtd_info *mtd, int cmd)
{
    switch(cmd) {
        case NAND_CTL_SETCLE: /* Set CLE pin high */ break;
        case NAND_CTL_CLRCLE: /* Set CLE pin low */ break;
        case NAND_CTL_SETALE: /* Set ALE pin high */ break;
        case NAND_CTL_CLRALE: /* Set ALE pin low */ break;
        case NAND_CTL_SETNCE: /* Set nCE pin low */ break;
        case NAND_CTL_CLRNCE: /* Set nCE pin high */ break;
    }
}

    </programlisting>
    <para>
        <emphasis>Address lines based example.</emphasis> It's
        assumed that the
        nCE pin is driven by a chip select decoder.
    </para>
    <programlisting>
static void board_hwcontrol(struct mtd_info *mtd, int cmd)
{
    struct nand_chip *this = (struct nand_chip *) mtd->priv;
    switch(cmd) {
        case NAND_CTL_SETCLE: this->IO_ADDR_W |= CLE_ADRR_BIT; break;
        case NAND_CTL_CLRCLE: this->IO_ADDR_W &= ~CLE_ADRR_BIT;
break;
        case NAND_CTL_SETALE: this->IO_ADDR_W |= ALE_ADRR_BIT; break;
        case NAND_CTL_CLRALE: this->IO_ADDR_W &= ~ALE_ADRR_BIT;
break;
    }
}

    </programlisting>
</sect1>
<sect1 id="Device_ready_function">
    <title>Device ready function</title>
    <para>
        If the hardware interface has the ready busy pin of the
        NAND chip connected to a
        GPIO or other accesible I/O pin, this function is used
        to read back the state of the
        pin. The function has no arguments and should return 0,
        if the device is busy (R/B pin

```

```

                                mtdnand.tmpl.txt
                                is low) and 1, if the device is ready (R/B pin is high).
                                If the hardware interface does not give access to the
ready busy pin, then
                                the function must not be defined and the function
pointer this->dev_ready is set to NULL.
                                </para>
                                </sect1>
                                <sect1 id="Init_function">
                                <title>Init function</title>
                                <para>
board                                The init function allocates memory and sets up all the
                                specific parameters and function pointers. When
everything                                is set up nand_scan() is called. This function tries to
                                detect and identify then chip. If a chip is found all
the                                internal data fields are initialized accordingly.
                                The structure(s) have to be zeroed out first and then
filled with the necceary                                information about the device.
                                </para>
                                <programlisting>
static int __init board_init (void)
{
    struct nand_chip *this;
    int err = 0;

    /* Allocate memory for MTD device structure and private data */
    board_mtd = kzalloc(sizeof(struct mtd_info) + sizeof(struct nand_chip),
GFP_KERNEL);
    if (!board_mtd) {
        printk ("Unable to allocate NAND MTD device structure.\n");
        err = -ENOMEM;
        goto out;
    }

    /* map physical address */
    baseaddr = ioremap(CHIP_PHYSICAL_ADDRESS, 1024);
    if (!baseaddr) {
        printk("Ioremap to access NAND chip failed\n");
        err = -EIO;
        goto out_mtd;
    }

    /* Get pointer to private data */
    this = (struct nand_chip *) ();
    /* Link the private data with the MTD structure */
    board_mtd->priv = this;

    /* Set address of NAND IO lines */
    this->IO_ADDR_R = baseaddr;
    this->IO_ADDR_W = baseaddr;
    /* Reference hardware control function */
    this->hwcontrol = board_hwcontrol;
    /* Set command delay time, see datasheet for correct value */

```

```

                                mtdnand.tmpl.txt
this->chip_delay = CHIP_DEPENDEND_COMMAND_DELAY;
/* Assign the device ready function, if available */
this->dev_ready = board_dev_ready;
this->eccmode = NAND_ECC_SOFT;

/* Scan to find existence of the device */
if (nand_scan (board_mtd, 1)) {
    err = -ENXIO;
    goto out_ior;
}

add_mtd_partitions(board_mtd, partition_info, NUM_PARTITIONS);
goto out;

out_ior:
    iounmap(baseaddr);
out_mtd:
    kfree (board_mtd);
out:
    return err;
}
module_init(board_init);
</programlisting>
</sect1>
<sect1 id="Exit_function">
    <title>Exit function</title>
    <para>
        The exit function is only necessary if the driver is
        compiled as a module. It releases all resources which
        are held by the chip driver and unregisters the
        partitions
        in the MTD layer.
    </para>
</sect1>
</programlisting>
#ifdef MODULE
static void __exit board_cleanup (void)
{
    /* Release resources, unregister device */
    nand_release (board_mtd);

    /* unmap physical address */
    iounmap(baseaddr);

    /* Free the MTD device structure */
    kfree (board_mtd);
}
module_exit(board_cleanup);
#endif
</programlisting>
</sect1>
</chapter>

<chapter id="boarddriversadvanced">
    <title>Advanced board driver functions</title>
    <para>
        This chapter describes the advanced functionality of the NAND

```

mtdnand.tmpl.txt

board driver. For a list of functions which can be overridden by the driver see the documentation of the nand_chip structure.

</para>

<sect1 id="Multiple_chip_control">

<title>Multiple chip control</title>

<para>

The nand driver can control chip arrays. Therefore the board driver must provide an own select_chip function.

This

function must (de)select the requested chip.

The function pointer in the nand_chip structure must be set before calling nand_scan(). The maxchip parameter of nand_scan() defines the maximum number of chips to scan for. Make sure that the select_chip function can handle the requested number of chips.

</para>

<para>

The nand driver concatenates the chips to one virtual chip and provides this virtual chip to the MTD layer.

</para>

<para>

<emphasis>Note: The driver can only handle linear chip

arrays

of equally sized chips. There is no support for parallel arrays which extend the buswidth.</emphasis>

</para>

<para>

<emphasis>GPIO based example</emphasis>

</para>

<programlisting>

```
static void board_select_chip (struct mtd_info *mtd, int chip)
{
    /* Deselect all chips, set all nCE pins high */
    GPIO (BOARD_NAND_NCE) |= 0xff;
    if (chip >= 0)
        GPIO (BOARD_NAND_NCE) &= ~ (1 << chip);
}
```

</programlisting>

<para>

<emphasis>Address lines based example.</emphasis>

Its assumed that the nCE pins are connected to an address decoder.

</para>

<programlisting>

```
static void board_select_chip (struct mtd_info *mtd, int chip)
{
    struct nand_chip *this = (struct nand_chip *) mtd->priv;

    /* Deselect all chips */
    this->IO_ADDR_R &= ~BOARD_NAND_ADDR_MASK;
    this->IO_ADDR_W &= ~BOARD_NAND_ADDR_MASK;
    switch (chip) {
    case 0:
        this->IO_ADDR_R |= BOARD_NAND_ADDR_CHIP0;
        this->IO_ADDR_W |= BOARD_NAND_ADDR_CHIP0;
```



```

                                mtdnand.tmpl.txt
                                break;
                                ....
                                case n:
                                    this->IO_ADDR_R |= BOARD_NAND_ADDR_CHIPn;
                                    this->IO_ADDR_W |= BOARD_NAND_ADDR_CHIPn;
                                    break;
                                }
                                }

                                </programlisting>
                                </sect1>
                                <sect1 id="Hardware_ECC_support">
                                    <title>Hardware ECC support</title>
                                    <sect2 id="Functions_and_constants">
                                        <title>Functions and constants</title>
                                        <para>
of                                The nand driver supports three different types
                                hardware ECC.
                                <itemizedlist>
                                <listitem><para>NAND_ECC_HW3_256</para><para>
                                Hardware ECC generator providing 3 bytes ECC per
                                256 byte.
                                </para> </listitem>
                                <listitem><para>NAND_ECC_HW3_512</para><para>
                                Hardware ECC generator providing 3 bytes ECC per
                                512 byte.
                                </para> </listitem>
                                <listitem><para>NAND_ECC_HW6_512</para><para>
                                Hardware ECC generator providing 6 bytes ECC per
                                512 byte.
                                </para> </listitem>
                                <listitem><para>NAND_ECC_HW8_512</para><para>
                                Hardware ECC generator providing 6 bytes ECC per
                                512 byte.
                                </para> </listitem>
                                </itemizedlist>
                                If your hardware generator has a different
functionality                    add it at the appropriate place in nand_base.c
                                </para>
                                <para>
                                The board driver must provide following
functions:                        <itemizedlist>
                                <listitem><para>enable_hwecc</para><para>
                                This function is called before reading / writing
to                                the chip. Reset or initialize the hardware
                                generator
                                in this function. The function is called with an
                                argument which let you distinguish between read
                                and write operations.
                                </para> </listitem>
                                <listitem><para>calculate_ecc</para><para>
                                This function is called after read / write from
/ to

```

set
below.

for
respectively 2
is
generator
software
provided
code.

Reed-Solomon
The syndrome
syndrome
generic

the data
work. This
ECC. The
longer
layout and
managed by
oob-layout
for

mtdnand.tmpl.txt
the chip. Transfer the ECC from the hardware to
the buffer. If the option NAND_HWECC_SYNDROME is
then the function is only called on write. See

</para> </listitem>
<listitem><para>correct_data</para><para>
In case of an ECC error this function is called
error detection and correction. Return 1
in case the error can be corrected. If the error
not correctable return -1. If your hardware
matches the default algorithm of the nand_ecc
generator then use the correction function
by nand_ecc instead of implementing duplicated

</para> </listitem>
</itemizedlist>

</para>

</sect2>

<sect2 id="Hardware_ECC_with_syndrome_calculation">

<title>Hardware ECC with syndrome calculation</title>

<para>

Many hardware ECC implementations provide
codes and calculate an error syndrome on read.
must be converted to a standard Reed-Solomon
before calling the error correction code in the
Reed-Solomon library.

</para>

<para>

The ECC bytes must be placed immediately after
bytes in order to make the syndrome generator
is contrary to the usual layout used by software
separation of data and out of band area is not
possible. The nand driver code handles this
the remaining free bytes in the oob area are
the autoplacement code. Provide a matching
in this case. See rts_from4.c and diskonchip.c
implementation reference. In those cases we must

also
layout is
positions.

mtdnand.tmpl.txt

use bad block tables on FLASH, because the ECC
interfering with the bad block marker

See bad block table support for details.

```
</para>
</sect2>
</sect1>
<sect1 id="Bad_Block_table_support">
  <title>Bad block table support</title>
  <para>
```

Most NAND chips mark the bad blocks at a defined position in the spare area. Those blocks must not be erased under any circumstances as the bad block information would be lost. It is possible to check the bad block mark each time when the blocks are accessed by reading the spare area of the first page in the block. This is time consuming so a bad block table is used.

```
</para>
<para>
```

The nand driver supports various types of bad block tables.

```
<itemizedlist>
```

```
<listitem><para>Per device</para><para>
```

The bad block table contains all bad block information of the device which can consist of multiple chips.

```
</para> </listitem>
```

```
<listitem><para>Per chip</para><para>
```

A bad block table is used per chip and contains the bad block information for this particular chip.

```
</para> </listitem>
```

```
<listitem><para>Fixed offset</para><para>
```

The bad block table is located at a fixed offset in the chip (device). This applies to various DiskOnChip devices.

```
</para> </listitem>
```

```
<listitem><para>Automatic placed</para><para>
```

The bad block table is automatically placed and detected either at the end or at the beginning of a chip (device)

```
</para> </listitem>
```

```
<listitem><para>Mirrored tables</para><para>
```

The bad block table is mirrored on the chip (device) to allow updates of the bad block table without data loss.

```
</para> </listitem>
```

```
</itemizedlist>
```

```
</para>
<para>
```

nand_scan() calls the function nand_default_bbt(). nand_default_bbt() selects appropriate default bad block table descriptors depending on the chip

information

which was retrieved by nand_scan().

```
</para>
```

<para>

The standard policy is scanning the device for bad blocks and build a ram based bad block table which allows faster access than always checking the bad block information on the flash chip itself.

</para>

<sect2 id="Flash_based_tables">

<title>Flash based tables</title>

<para>

block table in FLASH.

no factory marked

blocks. The marker pattern

So in case of

chip this block

Therefore we scan the

good blocks and

before erasing any

It may be desired or necceary to keep a bad

For AG-AND chips this is mandatory, as they have

bad blocks. They have factory marked good

is erased when the block is erased to be reused.

powerloss before writing the pattern back to the

would be lost and added to the bad blocks.

chip(s) when we detect them the first time for

store this information in a bad block table

of the blocks.

</para>

<para>

procteted against

memory bad block

are allowed

The blocks in which the tables are stored are

accidental access by marking them bad in the

table. The bad block table management functions

to circumvernt this protection.

</para>

<para>

block table support

option field of

nand_scan(). For AG-AND

table functionality

options are

chip</para></listitem>

block</para></listitem>

of the chip</para></listitem>

The simplest way to activate the FLASH based bad

is to set the option NAND_USE_FLASH_BBT in the

the nand chip structure before calling

chips is this done by default.

This activates the default FLASH based bad block

of the NAND driver. The default bad block table

<itemizedlist>

<listitem><para>Store bad block table per

<listitem><para>Use 2 bits per

<listitem><para>Automatic placement at the end

<listitem><para>Use mirrored tables with version

```

mtdnand.templ.txt
numbers</para></listitem>
the chip</para></listitem>
</itemizedlist>
</para>
</sect2>
<sect2 id="User_defined_tables">
  <title>User defined tables</title>
  <para>
    User defined tables are created by filling out a
    nand_bbt_descr structure and storing the pointer
    nand_chip structure member bbt_td before calling
    nand_scan().
    If a mirror table is necessary a second
    structure must be
    created and a pointer to this structure must be
    stored
    in bbt_md inside the nand_chip structure. If the
    member is set to NULL then only the main table
    is used
    and no scan for the mirrored table is performed.
  </para>
  <para>
    The most important field in the nand_bbt_descr
    structure
    is the options field. The options define most of
    the
    table properties. Use the predefined constants
    from
    nand.h to define the options.
    <itemizedlist>
    <listitem><para>Number of bits per block</para>
    <para>The supported number of bits is 1, 2, 4,
    8.</para></listitem>
    <listitem><para>Table per chip</para>
    <para>Setting the constant NAND_BBT_PERCHIP
    a bad block table is managed for each chip in a
    chip array.
    If this option is not set then a per device bad
    block table
    is used.</para></listitem>
    <listitem><para>Table location is
    absolute</para>
    <para>Use the option constant NAND_BBT_ABSPAGE
    and
    define the absolute page number where the bad
    block
    table starts in the field pages. If you have
    selected bad block
    tables per chip and you have a multi chip array
    then the start page
    must be given for each chip in the chip array.
  </para>
</sect2>

```

Note: there is no scan

fields

uninitialized</para></listitem>

detected</para>

first or the last good

NAND_BBT_LASTBLOCK to place

(device). The

pattern which

the block which

the pattern

pattern has to be

must be given

structure. For mirrored

mandatory.</para></listitem>

the table creation

Usually this is done only

the table write support.

in case a block has

function block_markbad

table. If the write

FLASH.</para>

mirrored tables with

the table version control.

mirrored tables with write

the bad block

mtdnand.templ.txt

for a table ident pattern performed, so the

pattern, veroffs, offs, len can be left

<listitem><para>Table location is automatically

<para>The table can either be located in the

blocks of the chip (device). Set

the bad block table at the end of the chip

bad block tables are marked and identified by a

is stored in the spare area of the first page in

holds the bad block table. Store a pointer to

in the pattern field. Further the length of the

stored in len and the offset in the spare area

in the offs member of the nand_bbt_descr

bad block tables different patterns are

<listitem><para>Table creation</para>

<para>Set the option NAND_BBT_CREATE to enable

if no table can be found during the scan.

once if a new chip is found. </para></listitem>

<listitem><para>Table write support</para>

<para>Set the option NAND_BBT_WRITE to enable

This allows the update of the bad block table(s)

to be marked bad due to wear. The MTD interface

is calling the update function of the bad block

support is enabled then the table is updated on

<para>

Note: Write support should only be enabled for

version control.

</para></listitem>

<listitem><para>Table version control</para>

<para>Set the option NAND_BBT_VERSION to enable

It's highly recommended to enable this for

support. It makes sure that the risk of losing

```

mtdnand.templ.txt
information about the
The version is stored in
device. The position of
veroffs in the bad block table
write</para>
table does contain
NAND_BBT_SAVECONTENT. When
block is read the bad
and everything is
bad block table
ignored and erased.
reserved for
blocks is defined
description structure.
a reasonable number.
scanned for the bad
mtdnand.templ.txt
table information is reduced to the loss of the
one worn out block which should be marked bad.
4 consecutive bytes in the spare area of the
the version number is defined by the member
descriptor.</para></listitem>
<listitem><para>Save block contents on
<para>
In case that the block which holds the bad block
other useful information, set the option
the bad block table is written then the whole
block table is updated and the block is erased
written back. If this option is not set only the
is written and everything else in the block is
</para></listitem>
<listitem><para>Number of reserved blocks</para>
<para>
For automatic placement some blocks must be
bad block table storage. The number of reserved
in the maxblocks member of the bad block table
Reserving 4 blocks for mirrored tables should be
This also limits the number of blocks which are
block table ident pattern.
</para></listitem>
</itemizedlist>
</para>
</sect2>
</sect1>
<sect1 id="Spare_area_placement">
<title>Spare area (auto)placement</title>
<para>
The nand driver implements different possibilities for
placement of filesystem data in the spare area,
<itemizedlist>
<listitem><para>Placement defined by fs
driver</para></listitem>
<listitem><para>Automatic placement</para></listitem>
</itemizedlist>
The default placement function is automatic placement.
The
nand driver has built in default placement schemes for

```

the

various chiptypes. If due to hardware ECC functionality the default placement does not fit then the board driver can provide a own placement scheme.

</para>
<para>

which

File system drivers can provide a own placement scheme is used instead of the default placement scheme.

</para>
<para>

structure

Placement schemes are defined by a nand_oobinfo

<programlisting>

```
struct nand_oobinfo {
    int    useecc;
    int    eccbytes;
    int    eccpos[24];
    int    oobfree[8][2];
};
```

</programlisting>

<itemizedlist>

<listitem><para>useecc</para><para>

function. The header

The useecc member controls the ecc and placement

select ecc and

file include/mtd/mtd-abi.h contains constants to

complete. This is

placement. MTD_NANDECC_OFF switches off the ecc

diagnosis only.

not recommended and available for testing and

placement, MTD_NANDECC_AUTOPLACE

MTD_NANDECC_PLACE selects caller defined

selects automatic placement.

</para></listitem>

<listitem><para>eccbytes</para><para>

bytes per page.

The eccbytes member defines the number of ecc

</para></listitem>

<listitem><para>eccpos</para><para>

spare area where

The eccpos array holds the byte offsets in the

the ecc codes are placed.

</para></listitem>

<listitem><para>oobfree</para><para>

area which can be

The oobfree array defines the areas in the spare

given in the format

used for automatic placement. The information is

usable area, size the

{offset, size}. offset defines the start of the

defined. The list is terminated

length in bytes. More than one area can be

by an {0, 0} entry.

</para></listitem>


```

                                mtdnand.templ.txt
                                </itemizedlist>
                                </para>
                                <sect2 id="Placement_defined_by_fs_driver">
                                <title>Placement defined by fs driver</title>
                                <para>
                                The calling function provides a pointer to a
                                structure which defines the ecc placement. For
                                caller must provide a spare area buffer along
                                data buffer. The spare area buffer size is
                                (number of pages) * (size of spare area). For reads the buffer size
                                is (number of pages) * ((size of spare area) +
                                (number of ecc steps per page) * sizeof (int)). The driver
                                stores the result of the ecc check for each tuple in the
                                spare buffer.
                                The storage sequence is
                                </para>
                                <para>
                                &lt;spare data page 0&gt;&lt;ecc result
                                0&gt;...&lt;ecc result n&gt;
                                </para>
                                <para>
                                ...
                                </para>
                                <para>
                                &lt;spare data page n&gt;&lt;ecc result
                                0&gt;...&lt;ecc result n&gt;
                                </para>
                                <para>
                                This is a legacy mode used by YAFFS1.
                                </para>
                                <para>
                                If the spare area buffer is NULL then only the
                                ECC placement is done according to the given scheme in the
                                nand_oobinfo structure.
                                </para>
                                </sect2>
                                <sect2 id="Automatic_placement">
                                <title>Automatic placement</title>
                                <para>
                                Automatic placement uses the built in defaults
                                to place the ecc bytes in the spare area. If filesystem data
                                have to be stored / read into the spare area then the calling
                                function must provide a buffer. The buffer size per page is determined
                                by the oobfree array in the nand_oobinfo structure.

```

```

                                mtdnand.tmpl.txt
                                </para>
                                <para>
ECC placement is                If the spare area buffer is NULL then only the
                                done according to the default builtin scheme.
                                </para>
                                </sect2>
                                <sect2 id="User_space_placement_selection">
                                <title>User space placement selection</title>
                                <para>
an internal                    All non ecc functions like mtd->read and mtd->write use
                                structure, which can be set by an ioctl. This structure
is preset                      to the autoplacement default.
                                <programlisting>
                                ioctl (fd, MEMSETOOBSEL, oobsel);
                                </programlisting>
                                oobsel is a pointer to a user supplied structure of type
match the                     nand_oobconfig. The contents of this structure must
                                criteria of the filesystem, which will be used. See an
example in utils/nandwrite.c.
                                </para>
                                </sect2>
                                </sect1>
                                <sect1 id="Spare_area_autoplacement_default">
                                <title>Spare area autoplacement default schemes</title>
                                <sect2 id="pagesize_256">
                                <title>256 byte pagesize</title>
<informaltable><tgroup cols="3"><tbody>
<row>
<entry>Offset</entry>
<entry>Content</entry>
<entry>Comment</entry>
</row>
<row>
<entry>0x00</entry>
<entry>ECC byte 0</entry>
<entry>Error correction code byte 0</entry>
</row>
<row>
<entry>0x01</entry>
<entry>ECC byte 1</entry>
<entry>Error correction code byte 1</entry>
</row>
<row>
<entry>0x02</entry>
<entry>ECC byte 2</entry>
<entry>Error correction code byte 2</entry>
</row>
<row>
<entry>0x03</entry>
<entry>Autoplace 0</entry>
<entry></entry>
</row>

```

```

<row>
<entry>0x04</entry>
<entry>Autoplace 1</entry>
<entry></entry>
</row>
<row>
<entry>0x05</entry>
<entry>Bad block marker</entry>
<entry>If any bit in this byte is zero, then this block is bad.
This applies only to the first page in a block. In the remaining
pages this byte is reserved</entry>
</row>
<row>
<entry>0x06</entry>
<entry>Autoplace 2</entry>
<entry></entry>
</row>
<row>
<entry>0x07</entry>
<entry>Autoplace 3</entry>
<entry></entry>
</row>
</tbody></tgroup></informaltable>
</sect2>
<sect2 id="pagesize_512">
<title>512 byte pagesize</title>
<informaltable><tgroup cols="3"><tbody>
<row>
<entry>Offset</entry>
<entry>Content</entry>
<entry>Comment</entry>
</row>
<row>
<entry>0x00</entry>
<entry>ECC byte 0</entry>
<entry>Error correction code byte 0 of the lower 256 Byte data in
this page</entry>
</row>
<row>
<entry>0x01</entry>
<entry>ECC byte 1</entry>
<entry>Error correction code byte 1 of the lower 256 Bytes of data
in this page</entry>
</row>
<row>
<entry>0x02</entry>
<entry>ECC byte 2</entry>
<entry>Error correction code byte 2 of the lower 256 Bytes of data
in this page</entry>
</row>
<row>
<entry>0x03</entry>
<entry>ECC byte 3</entry>
<entry>Error correction code byte 0 of the upper 256 Bytes of data
in this page</entry>
</row>

```

```

<row>
<entry>0x04</entry>
<entry>reserved</entry>
<entry>reserved</entry>
</row>
<row>
<entry>0x05</entry>
<entry>Bad block marker</entry>
<entry>If any bit in this byte is zero, then this block is bad.
This applies only to the first page in a block. In the remaining
pages this byte is reserved</entry>
</row>
<row>
<entry>0x06</entry>
<entry>ECC byte 4</entry>
<entry>Error correction code byte 1 of the upper 256 Bytes of data
in this page</entry>
</row>
<row>
<entry>0x07</entry>
<entry>ECC byte 5</entry>
<entry>Error correction code byte 2 of the upper 256 Bytes of data
in this page</entry>
</row>
<row>
<entry>0x08 - 0x0F</entry>
<entry>Autoplace 0 - 7</entry>
<entry></entry>
</row>
</tbody></tgroup></informaltable>
</sect2>
<sect2 id="pagesize_2048">
<title>2048 byte pagesize</title>
<informaltable><tgroup cols="3"><tbody>
<row>
<entry>Offset</entry>
<entry>Content</entry>
<entry>Comment</entry>
</row>
<row>
<entry>0x00</entry>
<entry>Bad block marker</entry>
<entry>If any bit in this byte is zero, then this block is bad.
This applies only to the first page in a block. In the remaining
pages this byte is reserved</entry>
</row>
<row>
<entry>0x01</entry>
<entry>Reserved</entry>
<entry>Reserved</entry>
</row>
<row>
<entry>0x02-0x27</entry>
<entry>Autoplace 0 - 37</entry>
<entry></entry>
</row>

```

```

<row>
<entry>0x28</entry>
<entry>ECC byte 0</entry>
<entry>Error correction code byte 0 of the first 256 Byte data in
this page</entry>
</row>
<row>
<entry>0x29</entry>
<entry>ECC byte 1</entry>
<entry>Error correction code byte 1 of the first 256 Bytes of data
in this page</entry>
</row>
<row>
<entry>0x2A</entry>
<entry>ECC byte 2</entry>
<entry>Error correction code byte 2 of the first 256 Bytes data in
this page</entry>
</row>
<row>
<entry>0x2B</entry>
<entry>ECC byte 3</entry>
<entry>Error correction code byte 0 of the second 256 Bytes of data
in this page</entry>
</row>
<row>
<entry>0x2C</entry>
<entry>ECC byte 4</entry>
<entry>Error correction code byte 1 of the second 256 Bytes of data
in this page</entry>
</row>
<row>
<entry>0x2D</entry>
<entry>ECC byte 5</entry>
<entry>Error correction code byte 2 of the second 256 Bytes of data
in this page</entry>
</row>
<row>
<entry>0x2E</entry>
<entry>ECC byte 6</entry>
<entry>Error correction code byte 0 of the third 256 Bytes of data
in this page</entry>
</row>
<row>
<entry>0x2F</entry>
<entry>ECC byte 7</entry>
<entry>Error correction code byte 1 of the third 256 Bytes of data
in this page</entry>
</row>
<row>
<entry>0x30</entry>
<entry>ECC byte 8</entry>
<entry>Error correction code byte 2 of the third 256 Bytes of data
in this page</entry>
</row>
<row>
<entry>0x31</entry>

```

```

<entry>ECC byte 9</entry>
<entry>Error correction code byte 0 of the fourth 256 Bytes of data
in this page</entry>
</row>
<row>
<entry>0x32</entry>
<entry>ECC byte 10</entry>
<entry>Error correction code byte 1 of the fourth 256 Bytes of data
in this page</entry>
</row>
<row>
<entry>0x33</entry>
<entry>ECC byte 11</entry>
<entry>Error correction code byte 2 of the fourth 256 Bytes of data
in this page</entry>
</row>
<row>
<entry>0x34</entry>
<entry>ECC byte 12</entry>
<entry>Error correction code byte 0 of the fifth 256 Bytes of data
in this page</entry>
</row>
<row>
<entry>0x35</entry>
<entry>ECC byte 13</entry>
<entry>Error correction code byte 1 of the fifth 256 Bytes of data
in this page</entry>
</row>
<row>
<entry>0x36</entry>
<entry>ECC byte 14</entry>
<entry>Error correction code byte 2 of the fifth 256 Bytes of data
in this page</entry>
</row>
<row>
<entry>0x37</entry>
<entry>ECC byte 15</entry>
<entry>Error correction code byte 0 of the sixth 256 Bytes of data
in this page</entry>
</row>
<row>
<entry>0x38</entry>
<entry>ECC byte 16</entry>
<entry>Error correction code byte 1 of the sixth 256 Bytes of data
in this page</entry>
</row>
<row>
<entry>0x39</entry>
<entry>ECC byte 17</entry>
<entry>Error correction code byte 2 of the sixth 256 Bytes of data
in this page</entry>
</row>
<row>
<entry>0x3A</entry>
<entry>ECC byte 18</entry>
<entry>Error correction code byte 0 of the seventh 256 Bytes of

```

data in this page</entry>
</row>
<row>
<entry>0x3B</entry>
<entry>ECC byte 19</entry>
<entry>Error correction code byte 1 of the seventh 256 Bytes of data in this page</entry>
</row>
<row>
<entry>0x3C</entry>
<entry>ECC byte 20</entry>
<entry>Error correction code byte 2 of the seventh 256 Bytes of data in this page</entry>
</row>
<row>
<entry>0x3D</entry>
<entry>ECC byte 21</entry>
<entry>Error correction code byte 0 of the eighth 256 Bytes of data in this page</entry>
</row>
<row>
<entry>0x3E</entry>
<entry>ECC byte 22</entry>
<entry>Error correction code byte 1 of the eighth 256 Bytes of data in this page</entry>
</row>
<row>
<entry>0x3F</entry>
<entry>ECC byte 23</entry>
<entry>Error correction code byte 2 of the eighth 256 Bytes of data in this page</entry>
</row>

</tbody></tgroup></informaltable>
</sect2>

</sect1>

</chapter>

<chapter id="filesystems">

<title>Filesystem support</title>

<para>

The NAND driver provides all necessary functions for a filesystem via the MTD interface.

</para>

<para>

Filesystems must be aware of the NAND peculiarities and restrictions. One major restriction of NAND Flash is, that you cannot write as often as you want to a page. The consecutive writes to a page, before erasing it again, are restricted to 1-3 writes, depending on the manufacturers specifications. This applies similar to the spare area.

</para>

<para>

Therefore NAND aware filesystems must either write in page size

chunks

up to

or hold a writebuffer to collect smaller writes until they sum
pagesize. Available NAND aware filesystems: JFFS2, YAFFS.

</para>

<para>

The spare area usage to store filesystem data is controlled by
the spare area placement functionality which is described in one
of the earlier chapters.

</para>

</chapter>

<chapter id="tools">

<title>Tools</title>

<para>

The MTD project provides a couple of helpful tools to handle
NAND Flash.

<itemizedlist>

<listitem><para>flasherase, flasheraseall: Erase and format
FLASH partitions</para></listitem>

<listitem><para>nandwrite: write filesystem images to NAND
FLASH</para></listitem>

<listitem><para>nanddump: dump the contents of a NAND FLASH
partitions</para></listitem>

</itemizedlist>

</para>

<para>

These tools are aware of the NAND restrictions. Please use those
tools instead of complaining about errors which are caused by non NAND
aware access methods.

</para>

</chapter>

<chapter id="defines">

<title>Constants</title>

<para>

This chapter describes the constants which might be relevant for a driver
developer.

</para>

<sect1 id="Chip_option_constants">

<title>Chip option constants</title>

<sect2 id="Constants_for_chip_id_table">

<title>Constants for chip id table</title>

<para>

These constants are defined in nand.h. They are ored together to
describe the chip functionality.

<programlisting>

/* Chip can not auto increment pages */

#define NAND_NO_AUTOINCR 0x00000001

/* Buswidth is 16 bit */

#define NAND_BUSWIDTH_16 0x00000002

/* Device supports partial programming without padding */

#define NAND_NO_PADDING 0x00000004


```

                                mtdnand.tmpl.txt
/* Chip has cache program function */
#define NAND_CACHEPRG          0x00000008
/* Chip has copy back function */
#define NAND_COPYBACK          0x00000010
/* AND Chip which has 4 banks and a confusing page / block
 * assignment. See Renesas datasheet for further information */
#define NAND_IS_AND             0x00000020
/* Chip has a array of 4 pages which can be read without
 * additional ready /busy waits */
#define NAND_4PAGE_ARRAY        0x00000040
    </programlisting>
    </para>
</sect2>
<sect2 id="Constants_for_runtime_options">
    <title>Constants for runtime options</title>
    <para>
        These constants are defined in nand.h. They are ored together to
describe
        the functionality.
    </programlisting>
/* Use a flash based bad block table. This option is parsed by the
 * default bad block table function (nand_default_bbt). */
#define NAND_USE_FLASH_BBT      0x00010000
/* The hw ecc generator provides a syndrome instead a ecc value on read
 * This can only work if we have the ecc bytes directly behind the
 * data bytes. Applies for DOC and AG-AND Renesas HW Reed Solomon generators */
#define NAND_HW ECC_SYNDROME    0x00020000
    </programlisting>
    </para>
</sect2>
</sect1>

<sect1 id="EEC_selection_constants">
    <title>ECC selection constants</title>
    <para>
        Use these constants to select the ECC algorithm.
    </programlisting>
/* No ECC. Usage is not recommended ! */
#define NAND_ECC_NONE           0
/* Software ECC 3 byte ECC per 256 Byte data */
#define NAND_ECC_SOFT           1
/* Hardware ECC 3 byte ECC per 256 Byte data */
#define NAND_ECC_HW3_256        2
/* Hardware ECC 3 byte ECC per 512 Byte data */
#define NAND_ECC_HW3_512        3
/* Hardware ECC 6 byte ECC per 512 Byte data */
#define NAND_ECC_HW6_512        4
/* Hardware ECC 6 byte ECC per 512 Byte data */
#define NAND_ECC_HW8_512        6
    </programlisting>
    </para>
</sect1>

<sect1 id="Hardware_control_related_constants">
    <title>Hardware control related constants</title>
    <para>

```

mtdnand.tmpl.txt

These constants describe the requested hardware access function when the boardspecific hardware control function is called

<programlisting>

```
/* Select the chip by setting nCE to low */
#define NAND_CTL_SETNCE      1
/* Deselect the chip by setting nCE to high */
#define NAND_CTL_CLRNCE     2
/* Select the command latch by setting CLE to high */
#define NAND_CTL_SETCLE     3
/* Deselect the command latch by setting CLE to low */
#define NAND_CTL_CLRCLE     4
/* Select the address latch by setting ALE to high */
#define NAND_CTL_SETALE     5
/* Deselect the address latch by setting ALE to low */
#define NAND_CTL_CLRALE     6
/* Set write protection by setting WP to high. Not used! */
#define NAND_CTL_SETWP      7
/* Clear write protection by setting WP to low. Not used! */
#define NAND_CTL_CLRWP      8
```

</programlisting>

</para>

</sect1>

<sect1 id="Bad_block_table_constants">

<title>Bad block table related constants</title>

<para>

These constants describe the options used for bad block table descriptors.

<programlisting>

```
/* Options for the bad block table descriptors */
```

```
/* The number of bits used per block in the bbt on the device */
#define NAND_BBT_NRBITS_MSK    0x0000000F
#define NAND_BBT_1BIT         0x00000001
#define NAND_BBT_2BIT         0x00000002
#define NAND_BBT_4BIT         0x00000004
#define NAND_BBT_8BIT         0x00000008
/* The bad block table is in the last good block of the device */
#define NAND_BBT_LASTBLOCK    0x00000010
/* The bbt is at the given page, else we must scan for the bbt */
#define NAND_BBT_ABSPAGE      0x00000020
/* The bbt is at the given page, else we must scan for the bbt */
#define NAND_BBT_SEARCH       0x00000040
/* bbt is stored per chip on multichip devices */
#define NAND_BBT_PERCHIP      0x00000080
/* bbt has a version counter at offset veroffs */
#define NAND_BBT_VERSION      0x00000100
/* Create a bbt if none exists */
#define NAND_BBT_CREATE       0x00000200
/* Search good / bad pattern through all pages of a block */
#define NAND_BBT_SCANALLPAGES 0x00000400
/* Scan block empty during good / bad block scan */
#define NAND_BBT_SCANEMPTY    0x00000800
/* Write bbt if necessary */
#define NAND_BBT_WRITE        0x00001000
/* Read and write back block contents when writing bbt */
```

```

                                mtdnand.tmpl.txt
#define NAND_BBT_SAVECONTENT 0x00002000
    </programlisting>
    </para>
</sect1>

</chapter>

<chapter id="structs">
    <title>Structures</title>
    <para>
        This chapter contains the autogenerated documentation of the structures
which are
        used in the NAND driver and might be relevant for a driver developer. Each
        struct member has a short description which is marked with an [XXX]
identifier.
        See the chapter "Documentation hints" for an explanation.
    </para>
!include/linux/mtd/nand.h
</chapter>

<chapter id="pubfunctions">
    <title>Public Functions Provided</title>
    <para>
        This chapter contains the autogenerated documentation of the NAND kernel
API functions
        which are exported. Each function has a short description which is marked
with an [XXX] identifier.
        See the chapter "Documentation hints" for an explanation.
    </para>
!Edrivers/mtd/nand/nand_base.c
!Edrivers/mtd/nand/nand_bbt.c
!Edrivers/mtd/nand/nand_ecc.c
</chapter>

<chapter id="intfunctions">
    <title>Internal Functions Provided</title>
    <para>
        This chapter contains the autogenerated documentation of the NAND driver
internal functions.
        Each function has a short description which is marked with an [XXX]
identifier.
        See the chapter "Documentation hints" for an explanation.
        The functions marked with [DEFAULT] might be relevant for a board driver
developer.
    </para>
!Idrivers/mtd/nand/nand_base.c
!Idrivers/mtd/nand/nand_bbt.c
<!-- No internal functions for kernel-doc:
X!Idrivers/mtd/nand/nand_ecc.c
-->
</chapter>

<chapter id="credits">
    <title>Credits</title>
    <para>

```

mtdnand.tmpl.txt

The following people have contributed to the NAND driver:

<orderedlist>

<listitem><para>Steven J.

Hill<email>sjhill@realitydiluted.com</email></para></listitem>

<listitem><para>David

Woodhouse<email>dwmw2@infradead.org</email></para></listitem>

<listitem><para>Thomas

Gleixner<email>tglx@linutronix.de</email></para></listitem>

</orderedlist>

A lot of users have provided bugfixes, improvements and helping hands for testing.

Thanks a lot.

</para>

<para>

The following people have contributed to this document:

<orderedlist>

<listitem><para>Thomas

Gleixner<email>tglx@linutronix.de</email></para></listitem>

</orderedlist>

</para>

</chapter>

</book>