

EDAC - Error Detection And Correction

Written by Doug Thompson <dougthompson@xmission.com>

7 Dec 2005

17 Jul 2007 Updated

(c) Mauro Carvalho Chehab <mchehab@redhat.com>

05 Aug 2009 Nehalem interface

EDAC is maintained and written by:

Doug Thompson, Dave Jiang, Dave Peterson et al,
original author: Thayne Harbaugh,

Contact:

website: bluesmoke.sourceforge.net

mailing list: bluesmoke-devel@lists.sourceforge.net

"bluesmoke" was the name for this device driver when it was "out-of-tree" and maintained at sourceforge.net. When it was pushed into 2.6.16 for the first time, it was renamed to 'EDAC'.

The bluesmoke project at sourceforge.net is now utilized as a 'staging area' for EDAC development, before it is sent upstream to kernel.org

At the bluesmoke/EDAC project site is a series of quilt patches against recent kernels, stored in a SVN repository. For easier downloading, there is also a tarball snapshot available.

EDAC PURPOSE

The 'edac' kernel module goal is to detect and report errors that occur within the computer system running under linux.

MEMORY

In the initial release, memory Correctable Errors (CE) and Uncorrectable Errors (UE) are the primary errors being harvested. These types of errors are harvested by the 'edac_mc' class of device.

Detecting CE events, then harvesting those events and reporting them, CAN be a predictor of future UE events. With CE events, the system can continue to operate, but with less safety. Preventive maintenance and proactive part replacement of memory DIMMs exhibiting CEs can reduce the likelihood of the dreaded UE events and system 'panics'.

NON-MEMORY

A new feature for EDAC, the edac_device class of device, was added in the 2.6.23 version of the kernel.

This new device type allows for non-memory type of ECC hardware detectors to have their states harvested and presented to userspace via the sysfs

interface.

Some architectures have ECC detectors for L1, L2 and L3 caches, along with DMA engines, fabric switches, main data path switches, interconnections, and various other hardware data paths. If the hardware reports it, then a edac_device device probably can be constructed to harvest and present that to userspace.

PCI BUS SCANNING

In addition, PCI Bus Parity and SERR Errors are scanned for on PCI devices in order to determine if errors are occurring on data transfers.

The presence of PCI Parity errors must be examined with a grain of salt. There are several add-in adapters that do NOT follow the PCI specification with regards to Parity generation and reporting. The specification says the vendor should tie the parity status bits to 0 if they do not intend to generate parity. Some vendors do not do this, and thus the parity bit can "float" giving false positives.

In the kernel there is a PCI device attribute located in sysfs that is checked by the EDAC PCI scanning code. If that attribute is set, PCI parity/error scanning is skipped for that device. The attribute is:

broken_parity_status

as is located in /sys/devices/pci<XXX>/0000:XX:YY.Z directories for PCI devices.

FUTURE HARDWARE SCANNING

EDAC will have future error detectors that will be integrated with EDAC or added to it, in the following list:

MCE	Machine Check Exception
MCA	Machine Check Architecture
NMI	NMI notification of ECC errors
MSRs	Machine Specific Register error cases
	and other mechanisms.

These errors are usually bus errors, ECC errors, thermal throttling and the like.

EDAC VERSIONING

EDAC is composed of a "core" module (edac_core.ko) and several Memory Controller (MC) driver modules. On a given system, the CORE is loaded and one MC driver will be loaded. Both the CORE and the MC driver (or edac_device driver) have individual versions that reflect current release level of their respective modules.

Thus, to "report" on what version a system is running, one must report both

the CORE's and the MC driver's versions.

LOADING

If 'edac' was statically linked with the kernel then no loading is necessary. If 'edac' was built as modules then simply modprobe the 'edac' pieces that you need. You should be able to modprobe hardware-specific modules and have the dependencies load the necessary core modules.

Example:

```
$> modprobe amd76x_edac
```

loads both the amd76x_edac.ko memory controller module and the edac_mc.ko core module.

EDAC sysfs INTERFACE

EDAC presents a 'sysfs' interface for control, reporting and attribute reporting purposes.

EDAC lives in the /sys/devices/system/edac directory.

Within this directory there currently reside 2 'edac' components:

```
mc      memory controller(s) system
pci     PCI control and status system
```

Memory Controller (mc) Model

First a background on the memory controller's model abstracted in EDAC. Each 'mc' device controls a set of DIMM memory modules. These modules are laid out in a Chip-Select Row (csrowX) and Channel table (chX). There can be multiple csrows and multiple channels.

Memory controllers allow for several csrows, with 8 csrows being a typical value.

Yet, the actual number of csrows depends on the electrical "loading" of a given motherboard, memory controller and DIMM characteristics.

Dual channels allows for 128 bit data transfers to the CPU from memory. Some newer chipsets allow for more than 2 channels, like Fully Buffered DIMMs (FB-DIMMs). The following example will assume 2 channels:

	Channel 0	Channel 1
csrow0	DIMM_A0	DIMM_B0
csrow1	DIMM_A0	DIMM_B0

csrow2	DIMM_A1	DIMM_B1
csrow3	DIMM_A1	DIMM_B1

In the above example table there are 4 physical slots on the motherboard for memory DIMMs:

```
DIMM_A0
DIMM_B0
DIMM_A1
DIMM_B1
```

Labels for these slots are usually silk screened on the motherboard. Slots labeled 'A' are channel 0 in this example. Slots labeled 'B' are channel 1. Notice that there are two csrows possible on a physical DIMM. These csrows are allocated their csrow assignment based on the slot into which the memory DIMM is placed. Thus, when 1 DIMM is placed in each Channel, the csrows cross both DIMMs.

Memory DIMMs come single or dual "ranked". A rank is a populated csrow. Thus, 2 single ranked DIMMs, placed in slots DIMM_A0 and DIMM_B0 above will have 1 csrow, csrow0. csrow1 will be empty. On the other hand, when 2 dual ranked DIMMs are similarly placed, then both csrow0 and csrow1 will be populated. The pattern repeats itself for csrow2 and csrow3.

The representation of the above is reflected in the directory tree in EDAC's sysfs interface. Starting in directory /sys/devices/system/edac/mc each memory controller will be represented by its own 'mcX' directory, where 'X' is the index of the MC.

```
.... /edac/mc/
      |
      |->mc0
      |->mc1
      |->mc2
      |
      ....
```

Under each 'mcX' directory each 'csrowX' is again represented by a 'csrowX', where 'X' is the csrow index:

```
... /mc/mc0/
      |
      |->csrow0
      |->csrow2
      |->csrow3
      |
      ....
```

Notice that there is no csrow1, which indicates that csrow0 is composed of a single ranked DIMMs. This should also apply in both Channels, in order to have dual-channel mode be operational. Since both csrow2 and csrow3 are populated, this indicates a dual ranked

edac.txt

set of DIMMs for channels 0 and 1.

Within each of the 'mcX' and 'csrowX' directories are several EDAC control and attribute files.

'mcX' DIRECTORIES

In 'mcX' directories are EDAC control and attribute files for this 'X' instance of the memory controllers:

Counter reset control file:

'reset_counters'

This write-only control file will zero all the statistical counters for UE and CE errors. Zeroing the counters will also reset the timer indicating how long since the last counter zero. This is useful for computing errors/time. Since the counters are always reset at driver initialization time, no module/kernel parameter is available.

RUN TIME: echo "anything" >/sys/devices/system/edac/mc/mc0/counter_reset

This resets the counters on memory controller 0

Seconds since last counter reset control file:

'seconds_since_reset'

This attribute file displays how many seconds have elapsed since the last counter reset. This can be used with the error counters to measure error rates.

Memory Controller name attribute file:

'mc_name'

This attribute file displays the type of memory controller that is being utilized.

Total memory managed by this memory controller attribute file:

'size_mb'

This attribute file displays, in count of megabytes, of memory that this instance of memory controller manages.

Total Uncorrectable Errors count attribute file:

edac.txt

'ue_count'

This attribute file displays the total count of uncorrectable errors that have occurred on this memory controller. If panic_on_ue is set this counter will not have a chance to increment, since EDAC will panic the system.

Total UE count that had no information attribute fileY:

'ue_noinfo_count'

This attribute file displays the number of UEs that have occurred with no information as to which DIMM slot is having errors.

Total Correctable Errors count attribute file:

'ce_count'

This attribute file displays the total count of correctable errors that have occurred on this memory controller. This count is very important to examine. CEs provide early indications that a DIMM is beginning to fail. This count field should be monitored for non-zero values and report such information to the system administrator.

Total Correctable Errors count attribute file:

'ce_noinfo_count'

This attribute file displays the number of CEs that have occurred wherewith no informations as to which DIMM slot is having errors. Memory is handicapped, but operational, yet no information is available to indicate which slot the failing memory is in. This count field should be also be monitored for non-zero values.

Device Symlink:

'device'

Symlink to the memory controller device.

Sdram memory scrubbing rate:

'sdram_scrub_rate'

Read/Write attribute file that controls memory scrubbing. The scrubbing rate is set by writing a minimum bandwidth in bytes/sec to the attribute file. The rate will be translated to an internal value that gives at least the specified rate.

Reading the file will return the actual scrubbing rate employed.

edac.txt

If configuration fails or memory scrubbing is not implemented, the value of the attribute file will be -1.

'csrowX' DIRECTORIES

In the 'csrowX' directories are EDAC control and attribute files for this 'X' instance of csrow:

Total Uncorrectable Errors count attribute file:

'ue_count'

This attribute file displays the total count of uncorrectable errors that have occurred on this csrow. If panic_on_ue is set this counter will not have a chance to increment, since EDAC will panic the system.

Total Correctable Errors count attribute file:

'ce_count'

This attribute file displays the total count of correctable errors that have occurred on this csrow. This count is very important to examine. CEs provide early indications that a DIMM is beginning to fail. This count field should be monitored for non-zero values and report such information to the system administrator.

Total memory managed by this csrow attribute file:

'size_mb'

This attribute file displays, in count of megabytes, of memory that this csrow contains.

Memory Type attribute file:

'mem_type'

This attribute file will display what type of memory is currently on this csrow. Normally, either buffered or unbuffered memory.
Examples:

Registered-DDR
Unbuffered-DDR

EDAC Mode of operation attribute file:

edac.txt

'edac_mode'

This attribute file will display what type of Error detection and correction is being utilized.

Device type attribute file:

'dev_type'

This attribute file will display what type of DRAM device is being utilized on this DIMM.

Examples:

x1
x2
x4
x8

Channel 0 CE Count attribute file:

'ch0_ce_count'

This attribute file will display the count of CEs on this DIMM located in channel 0.

Channel 0 UE Count attribute file:

'ch0_ue_count'

This attribute file will display the count of UEs on this DIMM located in channel 0.

Channel 0 DIMM Label control file:

'ch0_dimm_label'

This control file allows this DIMM to have a label assigned to it. With this label in the module, when errors occur the output can provide the DIMM label in the system log. This becomes vital for panic events to isolate the cause of the UE event.

DIMM Labels must be assigned after booting, with information that correctly identifies the physical slot with its silk screen label. This information is currently very motherboard specific and determination of this information must occur in userland at this time.

Channel 1 CE Count attribute file:

'ch1_ce_count'

edac.txt

This attribute file will display the count of CEs on this DIMM located in channel 1.

Channel 1 UE Count attribute file:

'chl_ue_count'

This attribute file will display the count of UEs on this DIMM located in channel 0.

Channel 1 DIMM Label control file:

'chl_dimm_label'

This control file allows this DIMM to have a label assigned to it. With this label in the module, when errors occur the output can provide the DIMM label in the system log. This becomes vital for panic events to isolate the cause of the UE event.

DIMM Labels must be assigned after booting, with information that correctly identifies the physical slot with its silk screen label. This information is currently very motherboard specific and determination of this information must occur in userland at this time.

SYSTEM LOGGING

If logging for UEs and CEs are enabled then system logs will have error notices indicating errors that have been detected:

EDAC MC0: CE page 0x283, offset 0xce0, grain 8, syndrome 0x6ec3, row 0, channel 1 "DIMM_B1": amd76x_edac

EDAC MC0: CE page 0x1e5, offset 0xfb0, grain 8, syndrome 0xb741, row 0, channel 1 "DIMM_B1": amd76x_edac

The structure of the message is:

the memory controller	(MC0)
Error type	(CE)
memory page	(0x283)
offset in the page	(0xce0)
the byte granularity	(grain 8)
or resolution of the error	
the error syndrome	(0xb741)
memory row	(row 0)
memory channel	(channel 1)
DIMM label, if set prior	(DIMM B1
and then an optional, driver-specific message that may	
have additional information.	

Both UEs and CEs with no info will lack all but memory controller,

edac.txt

error type, a notice of "no info" and then an optional, driver-specific error message.

PCI Bus Parity Detection

On Header Type 00 devices the primary status is looked at for any parity error regardless of whether Parity is enabled on the device. (The spec indicates parity is generated in some cases). On Header Type 01 bridges, the secondary status register is also looked at to see if parity occurred on the bus on the other side of the bridge.

SYSFS CONFIGURATION

Under /sys/devices/system/edac/pci are control and attribute files as follows:

Enable/Disable PCI Parity checking control file:

'check_pci_parity'

This control file enables or disables the PCI Bus Parity scanning operation. Writing a 1 to this file enables the scanning. Writing a 0 to this file disables the scanning.

Enable:

echo "1" >/sys/devices/system/edac/pci/check_pci_parity

Disable:

echo "0" >/sys/devices/system/edac/pci/check_pci_parity

Parity Count:

'pci_parity_count'

This attribute file will display the number of parity errors that have been detected.

MODULE PARAMETERS

Panic on UE control file:

'edac_mc_panic_on_ue'

An uncorrectable error will cause a machine panic. This is usually desirable. It is a bad idea to continue when an uncorrectable error occurs - it is indeterminate what was uncorrected and the operating system context might be so mangled that continuing will lead to further

edac.txt

corruption. If the kernel has MCE configured, then EDAC will never notice the UE.

LOAD TIME: module/kernel parameter: edac_mc_panic_on_ue=[0|1]

RUN TIME: echo "1" >
/sys/module/edac_core/parameters/edac_mc_panic_on_ue

Log UE control file:

'edac_mc_log_ue'

Generate kernel messages describing uncorrectable errors. These errors are reported through the system message log system. UE statistics will be accumulated even when UE logging is disabled.

LOAD TIME: module/kernel parameter: edac_mc_log_ue=[0|1]

RUN TIME: echo "1" > /sys/module/edac_core/parameters/edac_mc_log_ue

Log CE control file:

'edac_mc_log_ce'

Generate kernel messages describing correctable errors. These errors are reported through the system message log system. CE statistics will be accumulated even when CE logging is disabled.

LOAD TIME: module/kernel parameter: edac_mc_log_ce=[0|1]

RUN TIME: echo "1" > /sys/module/edac_core/parameters/edac_mc_log_ce

Polling period control file:

'edac_mc_poll_msec'

The time period, in milliseconds, for polling for error information. Too small a value wastes resources. Too large a value might delay necessary handling of errors and might lose valuable information for locating the error. 1000 milliseconds (once each second) is the current default. Systems which require all the bandwidth they can get, may increase this.

LOAD TIME: module/kernel parameter: edac_mc_poll_msec=[0|1]

RUN TIME: echo "1000" >
/sys/module/edac_core/parameters/edac_mc_poll_msec

Panic on PCI PARITY Error:

'panic_on_pci_parity'

edac.txt

This control files enables or disables panicking when a parity error has been detected.

module/kernel parameter: edac_panic_on_pci_pe=[0|1]

Enable:

echo "1" > /sys/module/edac_core/parameters/edac_panic_on_pci_pe

Disable:

echo "0" > /sys/module/edac_core/parameters/edac_panic_on_pci_pe

=====

EDAC_DEVICE type of device

In the header file, edac_core.h, there is a series of edac_device structures and APIs for the EDAC_DEVICE.

User space access to an edac_device is through the sysfs interface.

At the location /sys/devices/system/edac (sysfs) new edac_device devices will appear.

There is a three level tree beneath the above 'edac' directory. For example, the 'test_device_edac' device (found at the bluesmoke.sourceforge.net website) installs itself as:

/sys/devices/system/edac/test-instance

in this directory are various controls, a symlink and one or more 'instance' directories.

The standard default controls are:

log_ce	boolean to log CE events
log_ue	boolean to log UE events
panic_on_ue	boolean to 'panic' the system if an UE is encountered (default off, can be set true via startup script)
poll_msec	time period between POLL cycles for events

The test_device_edac device adds at least one of its own custom control:

test_bits	which in the current test driver does nothing but show how it is installed. A ported driver can add one or more such controls and/or attributes for specific uses. One out-of-tree driver uses controls here to allow for ERROR INJECTION operations to hardware injection registers
-----------	---

The symlink points to the 'struct dev' that is registered for this edac_device.

INSTANCES

One or more instance directories are present. For the 'test_device_edac' case:

test-instance0

In this directory there are two default counter attributes, which are totals of counter in deeper subdirectories.

ce_count	total of CE events of subdirectories
ue_count	total of UE events of subdirectories

BLOCKS

At the lowest directory level is the 'block' directory. There can be 0, 1 or more blocks specified in each instance.

test-block0

In this directory the default attributes are:

ce_count	which is counter of CE events for this 'block' of hardware being monitored
ue_count	which is counter of UE events for this 'block' of hardware being monitored

The 'test_device_edac' device adds 4 attributes and 1 control:

test-block-bits-0	for every POLL cycle this counter is incremented
test-block-bits-1	every 10 cycles, this counter is bumped once, and test-block-bits-0 is set to 0
test-block-bits-2	every 100 cycles, this counter is bumped once, and test-block-bits-1 is set to 0
test-block-bits-3	every 1000 cycles, this counter is bumped once, and test-block-bits-2 is set to 0
reset-counters	writing ANY thing to this control will reset all the above counters.

Use of the 'test_device_edac' driver should any others to create their own unique drivers for their hardware systems.

The 'test_device_edac' sample driver is located at the bluesmoke.sourceforge.net project site for EDAC.

NEHALEM USAGE OF EDAC APIs

This chapter documents some EXPERIMENTAL mappings for EDAC API to handle

edac.txt

Nehalem EDAC driver. They will likely be changed on future versions of the driver.

Due to the way Nehalem exports Memory Controller data, some adjustments were done at i7core_edac driver. This chapter will cover those differences

- 1) On Nehalem, there are one Memory Controller per Quick Patch Interconnect (QPI). At the driver, the term "socket" means one QPI. This is associated with a physical CPU socket.

Each MC have 3 physical read channels, 3 physical write channels and 3 logic channels. The driver currently sees it as just 3 channels. Each channel can have up to 3 DIMMs.

The minimum known unit is DIMMs. There are no information about csrows. As EDAC API maps the minimum unit is csrows, the driver sequentially maps channel/dimm into different csrows.

For example, suposing the following layout:

```
Ch0 phy rd0, wr0 (0x063f4031): 2 ranks, UDIMMs
  dimm 0 1024 Mb offset: 0, bank: 8, rank: 1, row: 0x4000, col: 0x400
  dimm 1 1024 Mb offset: 4, bank: 8, rank: 1, row: 0x4000, col: 0x400
Ch1 phy rd1, wr1 (0x063f4031): 2 ranks, UDIMMs
  dimm 0 1024 Mb offset: 0, bank: 8, rank: 1, row: 0x4000, col: 0x400
Ch2 phy rd3, wr3 (0x063f4031): 2 ranks, UDIMMs
  dimm 0 1024 Mb offset: 0, bank: 8, rank: 1, row: 0x4000, col: 0x400
```

The driver will map it as:

```
csrow0: channel 0, dimm0
csrow1: channel 0, dimm1
csrow2: channel 1, dimm0
csrow3: channel 2, dimm0
```

exports one

DIMM per csrow.

Each QPI is exported as a different memory controller.

- 2) Nehalem MC has the hability to generate errors. The driver implements this functionality via some error injection nodes:

For injecting a memory error, there are some sysfs nodes, under /sys/devices/system/edac/mc/mc?/:

inject_addrmatch/*:

Controls the error injection mask register. It is possible to specify several characteristics of the address to match an error code:

```
dimm = the affected dimm. Numbers are relative to a channel;
rank = the memory rank;
channel = the channel that will generate an error;
bank = the affected bank;
page = the page address;
column (or col) = the address column.
```

each of the above values can be set to "any" to match any valid value.

At driver init, all values are set to any.

edac.txt

For example, to generate an error at rank 1 of dimm 2, for any channel, any bank, any page, any column:

```
echo 2 >/sys/devices/system/edac/mc/mc0/inject_addrmatch/dimm
echo 1 >/sys/devices/system/edac/mc/mc0/inject_addrmatch/rank
```

To return to the default behaviour of matching any, you can do:

```
echo any >/sys/devices/system/edac/mc/mc0/inject_addrmatch/dimm
echo any >/sys/devices/system/edac/mc/mc0/inject_addrmatch/rank
```

inject_eccmask:

specifies what bits will have troubles,

inject_section:

specifies what ECC cache section will get the error:

```
3 for both
2 for the highest
1 for the lowest
```

inject_type:

specifies the type of error, being a combination of the following bits:

```
bit 0 - repeat
bit 1 - ecc
bit 2 - parity
```

inject_enable starts the error generation when something different than 0 is written.

All inject vars can be read. root permission is needed for write.

Datasheet states that the error will only be generated after a write on an address that matches inject_addrmatch. It seems, however, that reading will also produce an error.

For example, the following code will generate an error for any write access at socket 0, on any DIMM/address on channel 2:

```
echo 2 >/sys/devices/system/edac/mc/mc0/inject_addrmatch/channel
echo 2 >/sys/devices/system/edac/mc/mc0/inject_type
echo 64 >/sys/devices/system/edac/mc/mc0/inject_eccmask
echo 3 >/sys/devices/system/edac/mc/mc0/inject_section
echo 1 >/sys/devices/system/edac/mc/mc0/inject_enable
dd if=/dev/mem of=/dev/null seek=16k bs=4k count=1 >& /dev/null
```

For socket 1, it is needed to replace "mc0" by "mc1" at the above commands.

The generated error message will look like:

```
EDAC MC0: UE row 0, channel-a= 0 channel-b= 0 labels "-": NON_FATAL (addr =
0x0075b980, socket=0, Dimm=0, Channel=2, syndrome=0x00000040, count=1,
Err=8c0000400001009f:4000080482 (read error: read ECC error))
```

3) Nehalem specific Corrected Error memory counters

Nehalem have some registers to count memory errors. The driver uses those registers to report Corrected Errors on devices with Registered Dimms.

edac.txt

However, those counters don't work with Unregistered Dimms. As the chipset offers some counters that also work with UDIMMS (but with a worse level of granularity than the default ones), the driver exposes those registers for UDIMM memories.

They can be read by looking at the contents of all_channel_counts/

```
$ for i in /sys/devices/system/edac/mc/mc0/all_channel_counts/*; do echo $i;
cat $i; done
/sys/devices/system/edac/mc/mc0/all_channel_counts/udimm0
0
/sys/devices/system/edac/mc/mc0/all_channel_counts/udimm1
0
/sys/devices/system/edac/mc/mc0/all_channel_counts/udimm2
0
```

What happens here is that errors on different csrows, but at the same dimm number will increment the same counter.

So, in this memory mapping:

```
csrow0: channel 0, dimm0
csrow1: channel 0, dimm1
csrow2: channel 1, dimm0
csrow3: channel 2, dimm0
```

The hardware will increment udimm0 for an error at the first dimm at either csrow0, csrow2 or csrow3;

The hardware will increment udimm1 for an error at the second dimm at either csrow0, csrow2 or csrow3;

The hardware will increment udimm2 for an error at the third dimm at either csrow0, csrow2 or csrow3;

4) Standard error counters

The standard error counters are generated when an mcelog error is received by the driver. Since, with udimm, this is counted by software, it is possible that some errors could be lost. With rdimm's, they displays the contents of the registers