

Fault injection capabilities infrastructure

See also `drivers/md/faulty.c` and "every_nth" module option for `scsi_debug`.

Available fault injection capabilities

o `failslab`

injects slab allocation failures. (`kmalloc()`, `kmem_cache_alloc()`, ...)

o `fail_page_alloc`

injects page allocation failures. (`alloc_pages()`, `get_free_pages()`, ...)

o `fail_make_request`

injects disk IO errors on devices permitted by setting
`/sys/block/<device>/make-it-fail` or
`/sys/block/<device>/<partition>/make-it-fail`. (`generic_make_request()`)

Configure fault-injection capabilities behavior

o debugfs entries

`fault-inject-debugfs` kernel module provides some debugfs entries for runtime configuration of fault-injection capabilities.

- `/sys/kernel/debug/fail*/probability:`

likelihood of failure injection, in percent.
Format: `<percent>`

Note that one-failure-per-hundred is a very high error rate for some testcases. Consider setting `probability=100` and configure `/sys/kernel/debug/fail*/interval` for such testcases.

- `/sys/kernel/debug/fail*/interval:`

specifies the interval between failures, for calls to `should_fail()` that pass all the other tests.

Note that if you enable this, by setting `interval>1`, you will probably want to set `probability=100`.

- `/sys/kernel/debug/fail*/times:`

specifies how many times failures may happen at most.
A value of `-1` means "no limit".

- `/sys/kernel/debug/fail*/space:`

specifies an initial resource "budget", decremented by "size"

fault-injection.txt

on each call to `should_fail(size)`. Failure injection is suppressed until "space" reaches zero.

- `/sys/kernel/debug/fail*/verbose`

Format: { 0 | 1 | 2 }

specifies the verbosity of the messages when failure is injected. '0' means no messages; '1' will print only a single log line per failure; '2' will print a call trace too -- useful to debug the problems revealed by fault injection.

- `/sys/kernel/debug/fail*/task-filter:`

Format: { 'Y' | 'N' }

A value of 'N' disables filtering by process (default).

Any positive value limits failures to only processes indicated by `/proc/<pid>/make-it-fail==1`.

- `/sys/kernel/debug/fail*/require-start:`

- `/sys/kernel/debug/fail*/require-end:`

- `/sys/kernel/debug/fail*/reject-start:`

- `/sys/kernel/debug/fail*/reject-end:`

specifies the range of virtual addresses tested during stacktrace walking. Failure is injected only if some caller in the walked stacktrace lies within the required range, and none lies within the rejected range.

Default required range is `[0, ULONG_MAX)` (whole of virtual address space).

Default rejected range is `[0, 0)`.

- `/sys/kernel/debug/fail*/stacktrace-depth:`

specifies the maximum stacktrace depth walked during search for a caller within `[require-start, require-end)` OR `[reject-start, reject-end)`.

- `/sys/kernel/debug/fail_page_alloc/ignore-gfp-highmem:`

Format: { 'Y' | 'N' }

default is 'N', setting it to 'Y' won't inject failures into highmem/user allocations.

- `/sys/kernel/debug/failslab/ignore-gfp-wait:`

- `/sys/kernel/debug/fail_page_alloc/ignore-gfp-wait:`

Format: { 'Y' | 'N' }

default is 'N', setting it to 'Y' will inject failures only into non-sleep allocations (GFP_ATOMIC allocations).

- `/sys/kernel/debug/fail_page_alloc/min-order:`

specifies the minimum page allocation order to be injected failures.

o Boot option

fault-injection.txt

In order to inject faults while debugfs is not available (early boot time), use the boot option:

```
failslab=  
fail_page_alloc=  
fail_make_request=<interval>,<probability>,<space>,<times>
```

How to add new fault injection capability

- o #include <linux/fault-inject.h>

- o define the fault attributes

```
DECLARE_FAULT_INJECTION(name);
```

Please see the definition of struct fault_attr in fault-inject.h for details.

- o provide a way to configure fault attributes

- boot option

If you need to enable the fault injection capability from boot time, you can provide boot option to configure it. There is a helper function for it:

```
setup_fault_attr(attr, str);
```

- debugfs entries

failslab, fail_page_alloc, and fail_make_request use this way.
Helper functions:

```
init_fault_attr_dentries(entries, attr, name);  
void cleanup_fault_attr_dentries(entries);
```

- module parameters

If the scope of the fault injection capability is limited to a single kernel module, it is better to provide module parameters to configure the fault attributes.

- o add a hook to insert failures

Upon should_fail() returning true, client code should inject a failure.

```
should_fail(attr, size);
```

Application Examples

- o Inject slab allocation failures into module init/exit code

```
#!/bin/bash
```

fault-injection.txt

```
FAULTYPE=failslab
echo Y > /sys/kernel/debug/$FAULTYPE/task-filter
echo 10 > /sys/kernel/debug/$FAULTYPE/probability
echo 100 > /sys/kernel/debug/$FAULTYPE/interval
echo -1 > /sys/kernel/debug/$FAULTYPE/times
echo 0 > /sys/kernel/debug/$FAULTYPE/space
echo 2 > /sys/kernel/debug/$FAULTYPE/verbose
echo 1 > /sys/kernel/debug/$FAULTYPE/ignore-gfp-wait

faulty_system()
{
    bash -c "echo 1 > /proc/self/make-it-fail && exec $*"
}

if [ $# -eq 0 ]
then
    echo "Usage: $0 modulename [ modulename ... ]"
    exit 1
fi

for m in $*
do
    echo inserting $m...
    faulty_system modprobe $m

    echo removing $m...
    faulty_system modprobe -r $m
done

-----

o Inject page allocation failures only for a specific module

#!/bin/bash

FAULTYPE=fail_page_alloc
module=$1

if [ -z $module ]
then
    echo "Usage: $0 <modulename>"
    exit 1
fi

modprobe $module

if [ ! -d /sys/module/$module/sections ]
then
    echo Module $module is not loaded
    exit 1
fi

cat /sys/module/$module/sections/.text >
/sys/kernel/debug/$FAULTYPE/require-start
cat /sys/module/$module/sections/.data > /sys/kernel/debug/$FAULTYPE/require-end
```

```
                                fault-injection.txt
echo N > /sys/kernel/debug/$FAILTYPE/task-filter
echo 10 > /sys/kernel/debug/$FAILTYPE/probability
echo 100 > /sys/kernel/debug/$FAILTYPE/interval
echo -1 > /sys/kernel/debug/$FAILTYPE/times
echo 0 > /sys/kernel/debug/$FAILTYPE/space
echo 2 > /sys/kernel/debug/$FAILTYPE/verbose
echo 1 > /sys/kernel/debug/$FAILTYPE/ignore-gfp-wait
echo 1 > /sys/kernel/debug/$FAILTYPE/ignore-gfp-highmem
echo 10 > /sys/kernel/debug/$FAILTYPE/stacktrace-depth

trap "echo 0 > /sys/kernel/debug/$FAILTYPE/probability" SIGINT SIGTERM EXIT

echo "Injecting errors into the module $module... (interrupt to stop)"
sleep 1000000
```