

=====

## ASYNCHRONOUS OPERATIONS HANDLING

=====

By: David Howells <dhowells@redhat.com>

Contents:

- (\*) Overview.
- (\*) Operation record initialisation.
- (\*) Parameters.
- (\*) Procedure.
- (\*) Asynchronous callback.

=====

### OVERVIEW

=====

FS-Cache has an asynchronous operations handling facility that it uses for its data storage and retrieval routines. Its operations are represented by `fscache_operation` structs, though these are usually embedded into some other structure.

This facility is available to and expected to be used by the cache backends, and FS-Cache will create operations and pass them off to the appropriate cache backend for completion.

To make use of this facility, `<linux/fscache-cache.h>` should be `#included`.

=====

### OPERATION RECORD INITIALISATION

=====

An operation is recorded in an `fscache_operation` struct:

```
struct fscache_operation {
    union {
        struct work_struct fast_work;
        struct slow_work slow_work;
    };
    unsigned long          flags;
    fscache_operation_processor_t processor;
    ...
};
```

Someone wanting to issue an operation should allocate something with this struct embedded in it. They should initialise it by calling:

```
void fscache_operation_init(struct fscache_operation *op,
                           fscache_operation_release_t release);
```

第 1 页

with the operation to be initialised and the release function to use.

The `op->flags` parameter should be set to indicate the CPU time provision and the exclusivity (see the Parameters section).

The `op->fast_work`, `op->slow_work` and `op->processor` flags should be set as appropriate for the CPU time provision (see the Parameters section).

`FSCACHE_OP_WAITING` may be set in `op->flags` prior to each submission of the operation and waited for afterwards.

## PARAMETERS

There are a number of parameters that can be set in the operation record's flag parameter. There are three options for the provision of CPU time in these operations:

- (1) The operation may be done synchronously (`FSCACHE_OP_MYTHREAD`). A thread may decide it wants to handle an operation itself without deferring it to another thread.

This is, for example, used in read operations for calling `readpages()` on the backing filesystem in `CacheFiles`. Although `readpages()` does an asynchronous data fetch, the determination of whether pages exist is done synchronously – and the netfs does not proceed until this has been determined.

If this option is to be used, `FSCACHE_OP_WAITING` must be set in `op->flags` before submitting the operation, and the operating thread must wait for it to be cleared before proceeding:

```
wait_on_bit(&op->flags, FSCACHE_OP_WAITING,
            fscache_wait_bit, TASK_UNINTERRUPTIBLE);
```

- (2) The operation may be fast asynchronous (`FSCACHE_OP_FAST`), in which case it will be given to `keventd` to process. Such an operation is not permitted to sleep on I/O.

This is, for example, used by `CacheFiles` to copy data from a backing fs page to a netfs page after the backing fs has read the page in.

If this option is used, `op->fast_work` and `op->processor` must be initialised before submitting the operation:

```
INIT_WORK(&op->fast_work, do_some_work);
```

- (3) The operation may be slow asynchronous (`FSCACHE_OP_SLOW`), in which case it will be given to the slow work facility to process. Such an operation is permitted to sleep on I/O.

operations.txt

This is, for example, used by FS-Cache to handle background writes of pages that have just been fetched from a remote server.

If this option is used, `op->slow_work` and `op->processor` must be initialised before submitting the operation:

```
fscache_operation_init_slow(op, processor)
```

Furthermore, operations may be one of two types:

- (1) Exclusive (FSCACHE\_OP\_EXCLUSIVE). Operations of this type may not run in conjunction with any other operation on the object being operated upon.

An example of this is the attribute change operation, in which the file being written to may need truncation.

- (2) Shareable. Operations of this type may be running simultaneously. It's up to the operation implementation to prevent interference between other operations running at the same time.

=====  
PROCEDURE  
=====

Operations are used through the following procedure:

- (1) The submitting thread must allocate the operation and initialise it itself. Normally this would be part of a more specific structure with the generic `op` embedded within.
- (2) The submitting thread must then submit the operation for processing using one of the following two functions:

```
int fscache_submit_op(struct fscache_object *object,  
                     struct fscache_operation *op);
```

```
int fscache_submit_exclusive_op(struct fscache_object *object,  
                               struct fscache_operation *op);
```

The first function should be used to submit non-exclusive ops and the second to submit exclusive ones. The caller must still set the `FSCACHE_OP_EXCLUSIVE` flag.

If successful, both functions will assign the operation to the specified object and return 0. `-ENOBUFFS` will be returned if the object specified is permanently unavailable.

The operation manager will defer operations on an object that is still undergoing lookup or creation. The operation will also be deferred if an operation of conflicting exclusivity is in progress on the object.

If the operation is asynchronous, the manager will retain a reference to it, so the caller should put their reference to it by passing it to:

operations.txt

```
void fscache_put_operation(struct fscache_operation *op);
```

- (3) If the submitting thread wants to do the work itself, and has marked the operation with `FSCACHE_OP_MYTHREAD`, then it should monitor `FSCACHE_OP_WAITING` as described above and check the state of the object if necessary (the object might have died whilst the thread was waiting).

When it has finished doing its processing, it should call `fscache_put_operation()` on it.

- (4) The operation holds an effective lock upon the object, preventing other exclusive ops conflicting until it is released. The operation can be enqueued for further immediate asynchronous processing by adjusting the CPU time provisioning option if necessary, eg:

```
op->flags &= ~FSCACHE_OP_TYPE;
op->flags |= ~FSCACHE_OP_FAST;
```

and calling:

```
void fscache_enqueue_operation(struct fscache_operation *op)
```

This can be used to allow other things to have use of the worker thread pools.

#### =====

#### ASYNCHRONOUS CALLBACK

#### =====

When used in asynchronous mode, the worker thread pool will invoke the processor method with a pointer to the operation. This should then get at the container struct by using `container_of()`:

```
static void fscache_write_op(struct fscache_operation *_op)
{
    struct fscache_storage *op =
        container_of(_op, struct fscache_storage, op);
    ...
}
```

The caller holds a reference on the operation, and will invoke `fscache_put_operation()` when the processor function returns. The processor function is at liberty to call `fscache_enqueue_operation()` or to take extra references.