

configs.txt

configs - Userspace-driven kernel object configuration.

Joel Becker <joel.becker@oracle.com>

Updated: 31 March 2005

Copyright (c) 2005 Oracle Corporation,
Joel Becker <joel.becker@oracle.com>

[What is configs?]

configs is a ram-based filesystem that provides the converse of sysfs's functionality. Where sysfs is a filesystem-based view of kernel objects, configs is a filesystem-based manager of kernel objects, or config_items.

With sysfs, an object is created in kernel (for example, when a device is discovered) and it is registered with sysfs. Its attributes then appear in sysfs, allowing userspace to read the attributes via readdir(3)/read(2). It may allow some attributes to be modified via write(2). The important point is that the object is created and destroyed in kernel, the kernel controls the lifecycle of the sysfs representation, and sysfs is merely a window on all this.

A configs config_item is created via an explicit userspace operation: mkdir(2). It is destroyed via rmdir(2). The attributes appear at mkdir(2) time, and can be read or modified via read(2) and write(2). As with sysfs, readdir(3) queries the list of items and/or attributes. symlink(2) can be used to group items together. Unlike sysfs, the lifetime of the representation is completely driven by userspace. The kernel modules backing the items must respond to this.

Both sysfs and configs can and should exist together on the same system. One is not a replacement for the other.

[Using configs]

configs can be compiled as a module or into the kernel. You can access it by doing

```
mount -t configs none /config
```

The configs tree will be empty unless client modules are also loaded. These are modules that register their item types with configs as subsystems. Once a client subsystem is loaded, it will appear as a subdirectory (or more than one) under /config. Like sysfs, the configs tree is always there, whether mounted on /config or not.

An item is created via mkdir(2). The item's attributes will also appear at this time. readdir(3) can determine what the attributes are, read(2) can query their default values, and write(2) can store new values. Like sysfs, attributes should be ASCII text files, preferably with only one value per file. The same efficiency caveats from sysfs apply. Don't mix more than one attribute in one attribute file.

configs.txt

Like sysfs, configs expects write(2) to store the entire buffer at once. When writing to configs attributes, userspace processes should first read the entire file, modify the portions they wish to change, and then write the entire buffer back. Attribute files have a maximum size of one page (PAGE_SIZE, 4096 on i386).

When an item needs to be destroyed, remove it with rmdir(2). An item cannot be destroyed if any other item has a link to it (via symlink(2)). Links can be removed via unlink(2).

[Configuring FakeNBD: an Example]

Imagine there's a Network Block Device (NBD) driver that allows you to access remote block devices. Call it FakeNBD. FakeNBD uses configs for its configuration. Obviously, there will be a nice program that sysadmins use to configure FakeNBD, but somehow that program has to tell the driver about it. Here's where configs comes in.

When the FakeNBD driver is loaded, it registers itself with configs. readdir(3) sees this just fine:

```
# ls /config
fakenbd
```

A fakenbd connection can be created with mkdir(2). The name is arbitrary, but likely the tool will make some use of the name. Perhaps it is a uuid or a disk name:

```
# mkdir /config/fakenbd/disk1
# ls /config/fakenbd/disk1
target device rw
```

The target attribute contains the IP address of the server FakeNBD will connect to. The device attribute is the device on the server. Predictably, the rw attribute determines whether the connection is read-only or read-write.

```
# echo 10.0.0.1 > /config/fakenbd/disk1/target
# echo /dev/sda1 > /config/fakenbd/disk1/device
# echo 1 > /config/fakenbd/disk1/rw
```

That's it. That's all there is. Now the device is configured, via the shell no less.

[Coding With configs]

Every object in configs is a config_item. A config_item reflects an object in the subsystem. It has attributes that match values on that object. configs handles the filesystem representation of that object and its attributes, allowing the subsystem to ignore all but the basic show/store interaction.

Items are created and destroyed inside a config_group. A group is a collection of items that share the same attributes and operations. Items are created by mkdir(2) and removed by rmdir(2), but configs

handles that. The group has a set of operations to perform these tasks

A subsystem is the top level of a client module. During initialization, the client module registers the subsystem with configs, the subsystem appears as a directory at the top of the configs filesystem. A subsystem is also a config_group, and can do everything a config_group can.

[struct config_item]

```
struct config_item {
    char                *ci_name;
    char                ci_namebuf[UOBJ_NAME_LEN];
    struct kref         ci_kref;
    struct list_head    ci_entry;
    struct config_item  *ci_parent;
    struct config_group *ci_group;
    struct config_item_type *ci_type;
    struct dentry       *ci_dentry;
};

void config_item_init(struct config_item *);
void config_item_init_type_name(struct config_item *,
                                const char *name,
                                struct config_item_type *type);
struct config_item *config_item_get(struct config_item *);
void config_item_put(struct config_item *);
```

Generally, struct config_item is embedded in a container structure, a structure that actually represents what the subsystem is doing. The config_item portion of that structure is how the object interacts with configs.

Whether statically defined in a source file or created by a parent config_group, a config_item must have one of the _init() functions called on it. This initializes the reference count and sets up the appropriate fields.

All users of a config_item should have a reference on it via config_item_get(), and drop the reference when they are done via config_item_put().

By itself, a config_item cannot do much more than appear in configs. Usually a subsystem wants the item to display and/or store attributes, among other things. For that, it needs a type.

[struct config_item_type]

```
struct configs_item_operations {
    void (*release)(struct config_item *);
    ssize_t (*show_attribute)(struct config_item *,
                              struct configs_attribute *,
                              char *);
    ssize_t (*store_attribute)(struct config_item *,
                              struct configs_attribute *,
                              const char *, size_t);
```

```

                                configfs.txt
int (*allow_link)(struct config_item *src,
                  struct config_item *target);
int (*drop_link)(struct config_item *src,
                 struct config_item *target);
};

struct config_item_type {
    struct module                                *ct_owner;
    struct configfs_item_operations              *ct_item_ops;
    struct configfs_group_operations            *ct_group_ops;
    struct configfs_attribute                   **ct_attrs;
};

```

The most basic function of a `config_item_type` is to define what operations can be performed on a `config_item`. All items that have been allocated dynamically will need to provide the `ct_item_ops->release()` method. This method is called when the `config_item`'s reference count reaches zero. Items that wish to display an attribute need to provide the `ct_item_ops->show_attribute()` method. Similarly, storing a new attribute value uses the `store_attribute()` method.

[struct configfs_attribute]

```

struct configfs_attribute {
    char                *ca_name;
    struct module        *ca_owner;
    mode_t               ca_mode;
};

```

When a `config_item` wants an attribute to appear as a file in the item's `configfs` directory, it must define a `configfs_attribute` describing it. It then adds the attribute to the NULL-terminated array `config_item_type->ct_attrs`. When the item appears in `configfs`, the attribute file will appear with the `configfs_attribute->ca_name` filename. `configfs_attribute->ca_mode` specifies the file permissions.

If an attribute is readable and the `config_item` provides a `ct_item_ops->show_attribute()` method, that method will be called whenever userspace asks for a `read(2)` on the attribute. The converse will happen for `write(2)`.

[struct config_group]

A `config_item` cannot live in a vacuum. The only way one can be created is via `mkdir(2)` on a `config_group`. This will trigger creation of a child item.

```

struct config_group {
    struct config_item      cg_item;
    struct list_head        cg_children;
    struct configfs_subsystem *cg_subsys;
    struct config_group     **default_groups;
};

void config_group_init(struct config_group *group);
void config_group_init_type_name(struct config_group *group,

```

configs.txt

```
const char *name,  
struct config_item_type *type);
```

The `config_group` structure contains a `config_item`. Properly configuring that item means that a group can behave as an item in its own right. However, it can do more: it can create child items or groups. This is accomplished via the group operations specified on the group's `config_item_type`.

```
struct configs_group_operations {  
    struct config_item *(*make_item)(struct config_group *group,  
                                     const char *name);  
    struct config_group *(*make_group)(struct config_group *group,  
                                       const char *name);  
    int (*commit_item)(struct config_item *item);  
    void (*disconnect_notify)(struct config_group *group,  
                              struct config_item *item);  
    void (*drop_item)(struct config_group *group,  
                      struct config_item *item);  
};
```

A group creates child items by providing the `ct_group_ops->make_item()` method. If provided, this method is called from `mkdir(2)` in the group's directory. The subsystem allocates a new `config_item` (or more likely, its container structure), initializes it, and returns it to configs. Configs will then populate the filesystem tree to reflect the new item.

If the subsystem wants the child to be a group itself, the subsystem provides `ct_group_ops->make_group()`. Everything else behaves the same, using the group `_init()` functions on the group.

Finally, when userspace calls `rmdir(2)` on the item or group, `ct_group_ops->drop_item()` is called. As a `config_group` is also a `config_item`, it is not necessary for a separate `drop_group()` method. The subsystem must `config_item_put()` the reference that was initialized upon item allocation. If a subsystem has no work to do, it may omit the `ct_group_ops->drop_item()` method, and configs will call `config_item_put()` on the item on behalf of the subsystem.

IMPORTANT: `drop_item()` is void, and as such cannot fail. When `rmdir(2)` is called, configs WILL remove the item from the filesystem tree (assuming that it has no children to keep it busy). The subsystem is responsible for responding to this. If the subsystem has references to the item in other threads, the memory is safe. It may take some time for the item to actually disappear from the subsystem's usage. But it is gone from configs.

When `drop_item()` is called, the item's linkage has already been torn down. It no longer has a reference on its parent and has no place in the item hierarchy. If a client needs to do some cleanup before this teardown happens, the subsystem can implement the `ct_group_ops->disconnect_notify()` method. The method is called after configs has removed the item from the filesystem view but before the item is removed from its parent group. Like `drop_item()`,

configs.txt

disconnect_notify() is void and cannot fail. Client subsystems should not drop any references here, as they still must do it in drop_item().

A config_group cannot be removed while it still has child items. This is implemented in the configs rmdir(2) code. ->drop_item() will not be called, as the item has not been dropped. rmdir(2) will fail, as the directory is not empty.

[struct configs_subsystem]

A subsystem must register itself, usually at module_init time. This tells configs to make the subsystem appear in the file tree.

```
struct configs_subsystem {
    struct config_group    su_group;
    struct mutex           su_mutex;
};

int configs_register_subsystem(struct configs_subsystem *subsys);
void configs_unregister_subsystem(struct configs_subsystem *subsys);
```

A subsystem consists of a toplevel config_group and a mutex. The group is where child config_items are created. For a subsystem, this group is usually defined statically. Before calling configs_register_subsystem(), the subsystem must have initialized the group via the usual group_init() functions, and it must also have initialized the mutex.

When the register call returns, the subsystem is live, and it will be visible via configs. At that point, mkdir(2) can be called and the subsystem must be ready for it.

[An Example]

The best example of these basic concepts is the simple_children subsystem/group and the simple_child item in configs_example_explicit.c and configs_example_macros.c. It shows a trivial object displaying and storing an attribute, and a simple group creating and destroying these children.

The only difference between configs_example_explicit.c and configs_example_macros.c is how the attributes of the childless item are defined. The childless item has extended attributes, each with their own show()/store() operation. This follows a convention commonly used in sysfs. configs_example_explicit.c creates these attributes by explicitly defining the structures involved. Conversely configs_example_macros.c uses some convenience macros from configs.h to define the attributes. These macros are similar to their sysfs counterparts.

[Hierarchy Navigation and the Subsystem Mutex]

There is an extra bonus that configs provides. The config_groups and config_items are arranged in a hierarchy due to the fact that they appear in a filesystem. A subsystem is NEVER to touch the filesystem parts, but the subsystem might be interested in this hierarchy. For this reason, the hierarchy is mirrored via the config_group->cg_children

configs.txt

and config_item->ci_parent structure members.

A subsystem can navigate the cg_children list and the ci_parent pointer to see the tree created by the subsystem. This can race with configs' management of the hierarchy, so configs uses the subsystem mutex to protect modifications. Whenever a subsystem wants to navigate the hierarchy, it must do so under the protection of the subsystem mutex.

A subsystem will be prevented from acquiring the mutex while a newly allocated item has not been linked into this hierarchy. Similarly, it will not be able to acquire the mutex while a dropping item has not yet been unlinked. This means that an item's ci_parent pointer will never be NULL while the item is in configs, and that an item will only be in its parent's cg_children list for the same duration. This allows a subsystem to trust ci_parent and cg_children while they hold the mutex.

[Item Aggregation Via symlink(2)]

configs provides a simple group via the group->item parent/child relationship. Often, however, a larger environment requires aggregation outside of the parent/child connection. This is implemented via symlink(2).

A config_item may provide the ct_item_ops->allow_link() and ct_item_ops->drop_link() methods. If the ->allow_link() method exists, symlink(2) may be called with the config_item as the source of the link. These links are only allowed between configs config_items. Any symlink(2) attempt outside the configs filesystem will be denied.

When symlink(2) is called, the source config_item's ->allow_link() method is called with itself and a target item. If the source item allows linking to target item, it returns 0. A source item may wish to reject a link if it only wants links to a certain type of object (say, in its own subsystem).

When unlink(2) is called on the symbolic link, the source item is notified via the ->drop_link() method. Like the ->drop_item() method, this is a void function and cannot return failure. The subsystem is responsible for responding to the change.

A config_item cannot be removed while it links to any other item, nor can it be removed while an item links to it. Dangling symlinks are not allowed in configs.

[Automatically Created Subgroups]

A new config_group may want to have two types of child config_items. While this could be codified by magic names in ->make_item(), it is much more explicit to have a method whereby userspace sees this divergence.

Rather than have a group where some items behave differently than others, configs provides a method whereby one or many subgroups are automatically created inside the parent at its creation. Thus, mkdir("parent") results in "parent", "parent/subgroup1", up through

configs.txt

```
"parent/subgroupN".  Items of type 1 can now be created in
"parent/subgroup1", and items of type N can be created in
"parent/subgroupN".
```

These automatic subgroups, or default groups, do not preclude other children of the parent group. If `ct_group_ops->make_group()` exists, other child groups can be created on the parent group directly.

A configs subsystem specifies default groups by filling in the NULL-terminated array `default_groups` on the `config_group` structure. Each group in that array is populated in the configs tree at the same time as the parent group. Similarly, they are removed at the same time as the parent. No extra notification is provided. When a `->drop_item()` method call notifies the subsystem the parent group is going away, it also means every default group child associated with that parent group.

As a consequence of this, `default_groups` cannot be removed directly via `rmdir(2)`. They also are not considered when `rmdir(2)` on the parent group is checking for children.

[Dependant Subsystems]

Sometimes other drivers depend on particular configs items. For example, `ocfs2` mounts depend on a heartbeat region item. If that region item is removed with `rmdir(2)`, the `ocfs2` mount must BUG or go readonly. Not happy.

configs provides two additional API calls: `configs_depend_item()` and `configs_undepend_item()`. A client driver can call `configs_depend_item()` on an existing item to tell configs that it is depended on. configs will then return `-EBUSY` from `rmdir(2)` for that item. When the item is no longer depended on, the client driver calls `configs_undepend_item()` on it.

These API cannot be called underneath any configs callbacks, as they will conflict. They can block and allocate. A client driver probably shouldn't calling them of its own gumption. Rather it should be providing an API that external subsystems call.

How does this work? Imagine the `ocfs2` mount process. When it mounts, it asks for a heartbeat region item. This is done via a call into the heartbeat code. Inside the heartbeat code, the region item is looked up. Here, the heartbeat code calls `configs_depend_item()`. If it succeeds, then heartbeat knows the region is safe to give to `ocfs2`. If it fails, it was being torn down anyway, and heartbeat can gracefully pass up an error.

[Committable Items]

NOTE: Committable items are currently unimplemented.

Some `config_items` cannot have a valid initial state. That is, no default values can be specified for the item's attributes such that the item can do its work. Userspace must configure one or more attributes, after which the subsystem can start whatever entity this item represents.

configfs.txt

Consider the FakeNBD device from above. Without a target address *and* a target device, the subsystem has no idea what block device to import. The simple example assumes that the subsystem merely waits until all the appropriate attributes are configured, and then connects. This will, indeed, work, but now every attribute store must check if the attributes are initialized. Every attribute store must fire off the connection if that condition is met.

Far better would be an explicit action notifying the subsystem that the `config_item` is ready to go. More importantly, an explicit action allows the subsystem to provide feedback as to whether the attributes are initialized in a way that makes sense. `configfs` provides this as `committable` items.

`configfs` still uses only normal filesystem operations. An item is committed via `rename(2)`. The item is moved from a directory where it can be modified to a directory where it cannot.

Any group that provides the `ct_group_ops->commit_item()` method has `committable` items. When this group appears in `configfs`, `mkdir(2)` will not work directly in the group. Instead, the group will have two subdirectories: "live" and "pending". The "live" directory does not support `mkdir(2)` or `rmdir(2)` either. It only allows `rename(2)`. The "pending" directory does allow `mkdir(2)` and `rmdir(2)`. An item is created in the "pending" directory. Its attributes can be modified at will. Userspace commits the item by renaming it into the "live" directory. At this point, the subsystem receives the `->commit_item()` callback. If all required attributes are filled to satisfaction, the method returns zero and the item is moved to the "live" directory.

As `rmdir(2)` does not work in the "live" directory, an item must be shutdown, or "uncommitted". Again, this is done via `rename(2)`, this time from the "live" directory back to the "pending" one. The subsystem is notified by the `ct_group_ops->uncommit_object()` method.