

Optimized MPEG Filesystem (OMFS)

Overview

OMFS is a filesystem created by SonicBlue for use in the ReplayTV DVR and Rio Karma MP3 player. The filesystem is extent-based, utilizing block sizes from 2k to 8k, with hash-based directories. This filesystem driver may be used to read and write disks from these devices.

Note, it is not recommended that this FS be used in place of a general filesystem for your own streaming media device. Native Linux filesystems will likely perform better.

More information is available at:

<http://linux-karma.sf.net/>

Various utilities, including mkomfs and omfsck, are included with omfsprogs, available at:

<http://bobcopeland.com/karma/>

Instructions are included in its README.

Options

OMFS supports the following mount-time options:

```
uid=n      - make all files owned by specified user
gid=n      - make all files owned by specified group
umask=xxx  - set permission umask to xxx
fmask=xxx  - set umask to xxx for files
dmask=xxx  - set umask to xxx for directories
```

Disk format

OMFS discriminates between "sysblocks" and normal data blocks. The sysblock group consists of super block information, file metadata, directory structures, and extents. Each sysblock has a header containing CRCs of the entire sysblock, and may be mirrored in successive blocks on the disk. A sysblock may have a smaller size than a data block, but since they are both addressed by the same 64-bit block number, any remaining space in the smaller sysblock is unused.

Sysblock header information:

```
struct omfs_header {
    __be64 h_self;                /* FS block where this is located */
    __be32 h_body_size;           /* size of useful data after header */
    __be16 h_crc;                 /* crc-ccitt of body_size bytes */
    char h_fill1[2];
    u8 h_version;                 /* version, always 1 */
}
```

```

                                omfs.txt
char h_type;                    /* OMFS_INODE_X */
u8 h_magic;                     /* OMFS_IMAGIC */
u8 h_check_xor;                 /* XOR of header bytes before this */
__be32 h_fill2;
};

```

Files and directories are both represented by omfs_inode:

```

struct omfs_inode {
    struct omfs_header i_head;    /* header */
    __be64 i_parent;              /* parent containing this inode */
    __be64 i_sibling;             /* next inode in hash bucket */
    __be64 i_ctime;               /* ctime, in milliseconds */
    char i_fill1[35];
    char i_type;                  /* OMFS_[DIR, FILE] */
    __be32 i_fill2;
    char i_fill3[64];
    char i_name[OMFS_NAMELEN];    /* filename */
    __be64 i_size;                /* size of file, in bytes */
};

```

Directories in OMFS are implemented as a large hashtable. Filenames are hashed then prepended into the bucket list beginning at OMFS_DIR_START. Lookup requires hashing the filename, then seeking across i_sibling pointers until a match is found on i_name. Empty buckets are represented by block pointers with all-1s (~0).

A file is an omfs_inode structure followed by an extent table beginning at OMFS_EXTENT_START:

```

struct omfs_extent_entry {
    __be64 e_cluster;             /* start location of a set of blocks */
    __be64 e_blocks;              /* number of blocks after e_cluster */
};

struct omfs_extent {
    __be64 e_next;                /* next extent table location */
    __be32 e_extent_count;        /* total # extents in this table */
    __be32 e_fill;
    struct omfs_extent_entry e_entry; /* start of extent entries */
};

```

Each extent holds the block offset followed by number of blocks allocated to the extent. The final extent in each table is a terminator with e_cluster being ~0 and e_blocks being ones' complement of the total number of blocks in the table.

If this table overflows, a continuation inode is written and pointed to by e_next. These have a header but lack the rest of the inode structure.