

POHMELFS network protocol.

Basic structure used in network communication is following command:

```
struct netfs_cmd
{
    __u16      cmd;      /* Command number */
    __u16      csize;    /* Attached crypto information size */
    __u16      cpad;     /* Attached padding size */
    __u16      ext;      /* External flags */
    __u32      size;     /* Size of the attached data */
    __u32      trans;    /* Transaction id */
    __u64      id;       /* Object ID to operate on. Used for
feedback. */
    __u64      start;    /* Start of the object. */
    __u64      iv;       /* IV sequence */
    __u8       data[0];
};
```

Commands can be embedded into transaction command (which in turn has own command), so one can extend protocol as needed without breaking backward compatibility as long as old commands are supported. All string lengths include tail 0 byte.

All commands are transferred over the network in big-endian. CPU endianness is used at the end peers.

@cmd - command number, which specifies command to be processed. Following commands are used currently:

```
NETFS_READDIR    = 1,      /* Read directory for given inode number */
NETFS_READ_PAGE,  /* Read data page from the server */
NETFS_WRITE_PAGE, /* Write data page to the server */
NETFS_CREATE,     /* Create directory entry */
NETFS_REMOVE,     /* Remove directory entry */
NETFS_LOOKUP,     /* Lookup single object */
NETFS_LINK,       /* Create a link */
NETFS_TRANS,      /* Transaction */
NETFS_OPEN,       /* Open intent */
NETFS_INODE_INFO, /* Metadata cache coherency synchronization
message */
NETFS_PAGE_CACHE, /* Page cache invalidation message */
NETFS_READ_PAGES, /* Read multiple contiguous pages in one go */
NETFS_RENAME,     /* Rename object */
NETFS_CAPABILITIES, /* Capabilities of the client, for example
supported crypto */
NETFS_LOCK,       /* Distributed lock message */
NETFS_XATTR_SET,  /* Set extended attribute */
NETFS_XATTR_GET,  /* Get extended attribute */
```

@ext - external flags. Used by different commands to specify some extra arguments

like partial size of the embedded objects or creation flags.

@size - size of the attached data. For NETFS_READ_PAGE and NETFS_READ_PAGES no

data is attached,
but size of the requested data is incorporated here. It does not include
size of the command
header (struct netfs_cmd) itself.

@id - id of the object this command operates on. Each command can use it for own purpose.

@start - start of the object this command operates on. Each command can use it for own purpose.

@csize, @cpad - size and padding size of the (attached if needed) crypto information.

Command specifications.

@NETFS_READDIR

This command is used to sync content of the remote dir to the client.

@ext - length of the path to object.

@size - the same.

@id - local inode number of the directory to read.

@start - zero.

@NETFS_READ_PAGE

This command is used to read data from remote server.

Data size does not exceed local page cache size.

@id - inode number.

@start - first byte offset.

@size - number of bytes to read plus length of the path to object.

@ext - object path length.

@NETFS_CREATE

Used to create object.

It does not require that all directories on top of the object were already created, it will create them automatically. Each object has associated @netfs_path_entry data structure, which contains creation mode (permissions and type) and length of the name as long as name itself.

@start - 0

@size - size of the all data structures needed to create a path

@id - local inode number

@ext - 0

@NETFS_REMOVE

Used to remove object.

@ext - length of the path to object.

@size - the same.

@id - local inode number.

@start - zero.

network_protocol.txt

@NETFS_LOOKUP

Lookup information about object on server.

@ext - length of the path to object.

@size - the same.

@id - local inode number of the directory to look object in.

@start - local inode number of the object to look at.

@NETFS_LINK

Create hard or symlink.

Command is sent as "object_path|target_path".

@size - size of the above string.

@id - parent local inode number.

@start - 1 for symlink, 0 for hardlink.

@ext - size of the "object_path" above.

@NETFS_TRANS

Transaction header.

@size - incorporates all embedded command sizes including their header sizes.

@start - transaction generation number - unique id used to find transaction.

@ext - transaction flags. Unused at the moment.

@id - 0.

@NETFS_OPEN

Open intent for given transaction.

@id - local inode number.

@start - 0.

@size - path length to the object.

@ext - open flags (O_RDONLY and so on).

@NETFS_INODE_INFO

Metadata update command.

It is sent to servers when attributes of the object are changed and received when data or metadata were updated. It operates with the following structure:

```
struct netfs_inode_info
{
    unsigned int    mode;
    unsigned int    nlink;
    unsigned int    uid;
    unsigned int    gid;
    unsigned int    blocksize;
    unsigned int    padding;
    __u64           ino;
    __u64           blocks;
    __u64           rdev;
    __u64           size;
    __u64           version;
}
```

network_protocol.txt

};

It effectively mirrors stat(2) returned data.

@ext - path length to the object.
@size - the same plus size of the netfs_inode_info structure.
@id - local inode number.
@start - 0.

@NETFS_PAGE_CACHE

Command is only received by clients. It contains information about page to be marked as not up-to-date.

@id - client's inode number.
@start - last byte of the page to be invalidated. If it is not equal to current inode size, it will be vmtruncated().
@size - 0
@ext - 0

@NETFS_READ_PAGES

Used to read multiple contiguous pages in one go.

@start - first byte of the contiguous region to read.
@size - contains of two fields: lower 8 bits are used to represent page cache shift used by client, another 3 bytes are used to get number of pages.
@id - local inode number.
@ext - path length to the object.

@NETFS_RENAME

Used to rename object.

Attached data is formed into following string: "old_path|new_path".

@id - local inode number.
@start - parent inode number.
@size - length of the above string.
@ext - length of the old path part.

@NETFS_CAPABILITIES

Used to exchange crypto capabilities with server.

If crypto capabilities are not supported by server, then client will disable it or fail (if 'crypto_fail_unsupported' mount options was specified).

@id - superblock index. Used to specify crypto information for group of servers.
@size - size of the attached capabilities structure.
@start - 0.
@size - 0.
@scsize - 0.

@NETFS_LOCK

Used to send lock request/release messages. Although it sends byte range request

network_protocol.txt

and is capable of flushing pages based on that, it is not used, since all Linux filesystems lock the whole inode.

@id - lock generation number.

@start - start of the locked range.

@size - size of the locked range.

@ext - lock type: read/write. Not used actually. 15'th bit is used to determine, if it is lock request (1) or release (0).

@NETFS_XATTR_SET

@NETFS_XATTR_GET

Used to set/get extended attributes for given inode.

@id - attribute generation number or xattr setting type

@start - size of the attribute (request or attached)

@size - name length, path len and data size for given attribute

@ext - path length for given object