

USING VFAT

To use the vfat filesystem, use the filesystem type 'vfat'. i.e.
`mount -t vfat /dev/fd0 /mnt`

No special partition formatter is required. mkdosfs will work fine if you want to format from within Linux.

VFAT MOUNT OPTIONS

`uid=###` -- Set the owner of all files on this filesystem.
 The default is the uid of current process.

`gid=###` -- Set the group of all files on this filesystem.
 The default is the gid of current process.

`umask=###` -- The permission mask (for files and directories, see `umask(1)`).
 The default is the umask of current process.

`dmask=###` -- The permission mask for the directory.
 The default is the umask of current process.

`fmask=###` -- The permission mask for files.
 The default is the umask of current process.

`allow_utime=###` -- This option controls the permission check of mtime/atime.

20 - If current process is in group of file's group ID,
 you can change timestamp.

2 - Other users can change timestamp.

The default is set from ``dmask'` option. (If the directory is writable, `utime(2)` is also allowed. I.e. `~dmask & 022`)

Normally `utime(2)` checks current process is owner of the file, or it has `CAP_FOWNER` capability. But FAT filesystem doesn't have uid/gid on disk, so normal check is too unflexible. With this option you can relax it.

`codepage=###` -- Sets the codepage number for converting to shortname characters on FAT filesystem.
 By default, `FAT_DEFAULT_CODEPAGE` setting is used.

`iocharset=<name>` -- Character set to use for converting between the encoding is used for user visible filename and 16 bit Unicode characters. Long filenames are stored on disk in Unicode format, but Unix for the most part doesn't know how to deal with Unicode.
 By default, `FAT_DEFAULT_IOCHARSET` setting is used.

There is also an option of doing UTF-8 translations with the `utf8` option.

NOTE: "iocharset=utf8" is not recommended. If unsure, you should consider the following option instead.

utf8=<bool> -- UTF-8 is the filesystem safe version of Unicode that is used by the console. It can be enabled for the filesystem with this option. If 'uni_xlate' gets set, UTF-8 gets disabled.

uni_xlate=<bool> -- Translate unhandled Unicode characters to special escaped sequences. This would let you backup and restore filenames that are created with any Unicode characters. Until Linux supports Unicode for real, this gives you an alternative. Without this option, a '?' is used when no translation is possible. The escape character is ':' because it is otherwise illegal on the vfat filesystem. The escape sequence that gets used is ':' and the four digits of hexadecimal unicode.

nonumtail=<bool> -- When creating 8.3 aliases, normally the alias will end in '~1' or tilde followed by some number. If this option is set, then if the filename is "longfilename.txt" and "longfile.txt" does not currently exist in the directory, 'longfile.txt' will be the short alias instead of 'longfi~1.txt'.

usefree -- Use the "free clusters" value stored on FSINFO. It'll be used to determine number of free clusters without scanning disk. But it's not used by default, because recent Windows don't update it correctly in some case. If you are sure the "free clusters" on FSINFO is correct, by this option you can avoid scanning disk.

quiet -- Stops printing certain warning messages.

check=s|r|n -- Case sensitivity checking setting.
s: strict, case sensitive
r: relaxed, case insensitive
n: normal, default setting, currently case insensitive

nocase -- This was deprecated for vfat. Use shortname=win95 instead.

shortname=lower|win95|winnt|mixed
-- Shortname display/create setting.
lower: convert to lowercase for display,
emulate the Windows 95 rule for create.
win95: emulate the Windows 95 rule for display/create.
winnt: emulate the Windows NT rule for display/create.
mixed: emulate the Windows NT rule for display,
emulate the Windows 95 rule for create.
Default setting is 'mixed'.

tz=UTC -- Interpret timestamps as UTC rather than local time. This option disables the conversion of timestamps between local time (as used by Windows on FAT) and UTC (which Linux uses internally). This is particularly useful when mounting devices (like digital cameras) that are set to UTC in order to avoid the pitfalls of

vfat.txt

local time.

- showexec -- If set, the execute permission bits of the file will be allowed only if the extension part of the name is .EXE, .COM, or .BAT. Not set by default.
- debug -- Can be set, but unused by the current implementation.
- sys_immutable -- If set, ATTR_SYS attribute on FAT is handled as IMMUTABLE flag on Linux. Not set by default.
- flush -- If set, the filesystem will try to flush to disk more early than normal. Not set by default.
- rodir -- FAT has the ATTR_RO (read-only) attribute. On Windows, the ATTR_RO of the directory will just be ignored, and is used only by applications as a flag (e.g. it's set for the customized folder).

If you want to use ATTR_RO as read-only flag even for the directory, set this option.

- errors=panic|continue|remount-ro
 -- specify FAT behavior on critical errors: panic, continue without doing anything or remount the partition in read-only mode (default behavior).

<bool>: 0, 1, yes, no, true, false

TODO

- * Need to get rid of the raw scanning stuff. Instead, always use a get next directory entry approach. The only thing left that uses raw scanning is the directory renaming code.

POSSIBLE PROBLEMS

- * vfat_valid_longname does not properly checked reserved names.
- * When a volume name is the same as a directory name in the root directory of the filesystem, the directory name sometimes shows up as an empty file.
- * autoconv option does not work correctly.

BUG REPORTS

If you have trouble with the VFAT filesystem, mail bug reports to chaffee@bmrc.cs.berkeley.edu. Please specify the filename and the operation that gave you trouble.

TEST SUITE

If you plan to make any modifications to the vfat filesystem, please get the test suite that comes with the vfat distribution at

<http://bmrc.berkeley.edu/people/chaffee/vfat.html>

This tests quite a few parts of the vfat filesystem and additional tests for new features or untested features would be appreciated.

NOTES ON THE STRUCTURE OF THE VFAT FILESYSTEM

(This documentation was provided by Galen C. Hunt <gchunt@cs.rochester.edu> and lightly annotated by Gordon Chaffee).

This document presents a very rough, technical overview of my knowledge of the extended FAT file system used in Windows NT 3.5 and Windows 95. I don't guarantee that any of the following is correct, but it appears to be so.

The extended FAT file system is almost identical to the FAT file system used in DOS versions up to and including 6.223410239847 :-). The significant change has been the addition of long file names. These names support up to 255 characters including spaces and lower case characters as opposed to the traditional 8.3 short names.

Here is the description of the traditional FAT entry in the current Windows 95 filesystem:

```
struct directory { // Short 8.3 names
    unsigned char name[8];           // file name
    unsigned char ext[3];            // file extension
    unsigned char attr;              // attribute byte
    unsigned char lcase;             // Case for base and extension
    unsigned char ctime_ms;          // Creation time, milliseconds
    unsigned char ctime[2];          // Creation time
    unsigned char cdate[2];          // Creation date
    unsigned char adate[2];          // Last access date
    unsigned char reserved[2];       // reserved values (ignored)
    unsigned char time[2];           // time stamp
    unsigned char date[2];           // date stamp
    unsigned char start[2];          // starting cluster number
    unsigned char size[4];           // size of the file
};
```

The lcase field specifies if the base and/or the extension of an 8.3 name should be capitalized. This field does not seem to be used by Windows 95 but it is used by Windows NT. The case of filenames is not completely compatible from Windows NT to Windows 95. It is not completely compatible in the reverse direction, however. Filenames that fit in the 8.3 namespace and are written on Windows NT to be lowercase will show up as uppercase on Windows 95.

Note that the "start" and "size" values are actually little endian integer values. The descriptions of the fields in this structure are public knowledge and can be found elsewhere.

With the extended FAT system, Microsoft has inserted extra directory entries for any files with extended names. (Any name which legally fits within the old 8.3 encoding scheme does not have extra entries.) I call these extra entries slots. Basically, a slot is a specially formatted directory entry which holds up to 13 characters of

a file's extended name. Think of slots as additional labeling for the directory entry of the file to which they correspond. Microsoft prefers to refer to the 8.3 entry for a file as its alias and the extended slot directory entries as the file name.

The C structure for a slot directory entry follows:

```
struct slot { // Up to 13 characters of a long name
    unsigned char id;           // sequence number for slot
    unsigned char name0_4[10];  // first 5 characters in name
    unsigned char attr;         // attribute byte
    unsigned char reserved;     // always 0
    unsigned char alias_checksum; // checksum for 8.3 alias
    unsigned char name5_10[12]; // 6 more characters in name
    unsigned char start[2];     // starting cluster number
    unsigned char name11_12[4]; // last 2 characters in name
};
```

If the layout of the slots looks a little odd, it's only because of Microsoft's efforts to maintain compatibility with old software. The slots must be disguised to prevent old software from panicking. To this end, a number of measures are taken:

- 1) The attribute byte for a slot directory entry is always set to 0x0f. This corresponds to an old directory entry with attributes of "hidden", "system", "read-only", and "volume label". Most old software will ignore any directory entries with the "volume label" bit set. Real volume label entries don't have the other three bits set.
- 2) The starting cluster is always set to 0, an impossible value for a DOS file.

Because the extended FAT system is backward compatible, it is possible for old software to modify directory entries. Measures must be taken to ensure the validity of slots. An extended FAT system can verify that a slot does in fact belong to an 8.3 directory entry by the following:

- 1) Positioning. Slots for a file always immediately proceed their corresponding 8.3 directory entry. In addition, each slot has an id which marks its order in the extended file name. Here is a very abbreviated view of an 8.3 directory entry and its corresponding long name slots for the file "My Big File.Extension which is long":

```
<proceeding files...>
<slot #3, id = 0x43, characters = "h is long">
<slot #2, id = 0x02, characters = "xtension whic">
<slot #1, id = 0x01, characters = "My Big File.E">
<directory entry, name = "MYBIGFIL.EXT">
```

Note that the slots are stored from last to first. Slots are numbered from 1 to N. The Nth slot is or'ed with 0x40 to mark it as the last one.

vfat.txt

- 2) Checksum. Each slot has an "alias_checksum" value. The checksum is calculated from the 8.3 name using the following algorithm:

```
for (sum = i = 0; i < 11; i++) {  
    sum = (((sum&1)<<7)|((sum&0xfe)>>1)) + name[i]  
}
```

- 3) If there is free space in the final slot, a Unicode NULL (0x0000) is stored after the final character. After that, all unused characters in the final slot are set to Unicode 0xFFFF.

Finally, note that the extended name is stored in Unicode. Each Unicode character takes two bytes.