

Naming and data format standards for sysfs files

The libensors library offers an interface to the raw sensors data through the sysfs interface. Since lm-sensors 3.0.0, libensors is completely chip-independent. It assumes that all the kernel drivers implement the standard sysfs interface described in this document. This makes adding or updating support for any given chip very easy, as libensors, and applications using it, do not need to be modified. This is a major improvement compared to lm-sensors 2.

Note that motherboards vary widely in the connections to sensor chips. There is no standard that ensures, for example, that the second temperature sensor is connected to the CPU, or that the second fan is on the CPU. Also, some values reported by the chips need some computation before they make full sense. For example, most chips can only measure voltages between 0 and +4V. Other voltages are scaled back into that range using external resistors. Since the values of these resistors can change from motherboard to motherboard, the conversions cannot be hard coded into the driver and have to be done in user space.

For this reason, even if we aim at a chip-independent libensors, it will still require a configuration file (e.g. /etc/sensors.conf) for proper values conversion, labeling of inputs and hiding of unused inputs.

An alternative method that some programs use is to access the sysfs files directly. This document briefly describes the standards that the drivers follow, so that an application program can scan for entries and access this data in a simple and consistent way. That said, such programs will have to implement conversion, labeling and hiding of inputs. For this reason, it is still not recommended to bypass the library.

Each chip gets its own directory in the sysfs /sys/devices tree. To find all sensor chips, it is easier to follow the device symlinks from /sys/class/hwmon/hwmon*.

Up to lm-sensors 3.0.0, libensors looks for hardware monitoring attributes in the "physical" device directory. Since lm-sensors 3.0.1, attributes found in the hwmon "class" device directory are also supported. Complex drivers (e.g. drivers for multifunction chips) may want to use this possibility to avoid namespace pollution. The only drawback will be that older versions of libensors won't support the driver in question.

All sysfs values are fixed point numbers.

There is only one value per file, unlike the older /proc specification. The common scheme for files naming is: <type><number>_<item>. Usual types for sensor chips are "in" (voltage), "temp" (temperature) and "fan" (fan). Usual items are "input" (measured value), "max" (high threshold, "min" (low threshold). Numbering usually starts from 1, except for voltages which start from 0 (because most data sheets use this). A number is always used for elements that can be present more than once, even if there is a single element of the given type on the specific chip. Other files do not refer to a specific element, so they have a simple name, and no number.

sysfs-interface..txt

Alarms are direct indications read from the chips. The drivers do NOT make comparisons of readings to thresholds. This allows violations between readings to be caught and alarmed. The exact definition of an alarm (for example, whether a threshold must be met or must be exceeded to cause an alarm) is chip-dependent.

When setting values of hwmon sysfs attributes, the string representation of the desired value must be written, note that strings which are not a number are interpreted as 0! For more on how written strings are interpreted see the "sysfs attribute writes interpretation" section at the end of this file.

[0-*] denotes any positive number starting from 0
[1-*] denotes any positive number starting from 1
RO read only value
WO write only value
RW read/write value

Read/write values may be read-only for some chips, depending on the hardware implementation.

All entries (except name) are optional, and should only be created in a given driver if the chip has the feature.

* Global attributes *

name The chip name.
This should be a short, lowercase string, not containing spaces nor dashes, representing the chip name. This is the only mandatory attribute.
I2C devices get this attribute created automatically.
RO

update_rate The rate at which the chip will update readings.
Unit: millisecond
RW
Some devices have a variable update rate. This attribute can be used to change the update rate to the desired frequency.

* Voltages *

in[0-*]_min Voltage min value.
Unit: millivolt
RW

in[0-*]_max Voltage max value.
Unit: millivolt
RW

sysfs-interface..txt

in[0-*]_input Voltage input value.
Unit: millivolt
RO
Voltage measured on the chip pin.
Actual voltage depends on the scaling resistors on the motherboard, as recommended in the chip datasheet.
This varies by chip and by motherboard.
Because of this variation, values are generally NOT scaled by the chip driver, and must be done by the application.
However, some drivers (notably lm87 and via686a) do scale, because of internal resistors built into a chip.
These drivers will output the actual voltage. Rule of thumb: drivers should report the voltage values at the "pins" of the chip.

in[0-*]_label Suggested voltage channel label.
Text string
Should only be created if the driver has hints about what this voltage channel is being used for, and user-space doesn't. In all other cases, the label is provided by user-space.
RO

cpu[0-*]_vid CPU core reference voltage.
Unit: millivolt
RO
Not always correct.

vrn Voltage Regulator Module version number.
RW (but changing it should no more be necessary)
Originally the VRM standard version multiplied by 10, but now an arbitrary number, as not all standards have a version number.
Affects the way the driver calculates the CPU core reference voltage from the vid pins.

Also see the Alarms section for status flags associated with voltages.

* Fans *

fan[1-*]_min Fan minimum value
Unit: revolution/min (RPM)
RW

fan[1-*]_max Fan maximum value
Unit: revolution/min (RPM)
Only rarely supported by the hardware.
RW

fan[1-*]_input Fan input value.
Unit: revolution/min (RPM)
RO

sysfs-interface..txt

fan[1-*]_div Fan divisor.
Integer value in powers of two (1, 2, 4, 8, 16, 32, 64, 128).
RW
Some chips only support values 1, 2, 4 and 8.
Note that this is actually an internal clock divisor, which affects the measurable speed range, not the read value.

fan[1-*]_target
Desired fan speed
Unit: revolution/min (RPM)
RW
Only makes sense if the chip supports closed-loop fan speed control based on the measured fan speed.

fan[1-*]_label Suggested fan channel label.
Text string
Should only be created if the driver has hints about what this fan channel is being used for, and user-space doesn't. In all other cases, the label is provided by user-space.
RO

Also see the Alarms section for status flags associated with fans.

* PWM *

pwm[1-*] Pulse width modulation fan control.
Integer value in the range 0 to 255
RW
255 is max or 100%.

pwm[1-*]_enable
Fan speed control method:
0: no fan speed control (i.e. fan at full speed)
1: manual fan speed control enabled (using pwm[1-*])
2+: automatic fan speed control enabled
Check individual chip documentation files for automatic mode details.
RW

pwm[1-*]_mode 0: DC mode (direct current)
1: PWM mode (pulse-width modulation)
RW

pwm[1-*]_freq Base PWM frequency in Hz.
Only possibly available when pwmN_mode is PWM, but not always present even then.
RW

pwm[1-*]_auto_channels_temp
Select which temperature channels affect this PWM output in auto mode. Bitfield, 1 is temp1, 2 is temp2, 4 is temp3 etc... Which values are possible depend on the chip used.

RW

pwm[1-*]_auto_point[1-*]_pwm
 pwm[1-*]_auto_point[1-*]_temp
 pwm[1-*]_auto_point[1-*]_temp_hyst

Define the PWM vs temperature curve. Number of trip points is chip-dependent. Use this for chips which associate trip points to PWM output channels.

RW

temp[1-*]_auto_point[1-*]_pwm
 temp[1-*]_auto_point[1-*]_temp
 temp[1-*]_auto_point[1-*]_temp_hyst

Define the PWM vs temperature curve. Number of trip points is chip-dependent. Use this for chips which associate trip points to temperature channels.

RW

There is a third case where trip points are associated to both PWM output channels and temperature channels: the PWM values are associated to PWM output channels while the temperature values are associated to temperature channels. In that case, the result is determined by the mapping between temperature inputs and PWM outputs. When several temperature inputs are mapped to a given PWM output, this leads to several candidate PWM values. The actual result is up to the chip, but in general the highest candidate value (fastest fan speed) wins.

* Temperatures *

temp[1-*]_type Sensor type selection.

Integers 1 to 6

RW

1: PII/Celeron Diode

2: 3904 transistor

3: thermal diode

4: thermistor

5: AMD AMDSI

6: Intel PECI

Not all types are supported by all chips

temp[1-*]_max Temperature max value.

Unit: millidegree Celsius (or millivolt, see below)

RW

temp[1-*]_min Temperature min value.

Unit: millidegree Celsius

RW

temp[1-*]_max_hyst

Temperature hysteresis value for max limit.

Unit: millidegree Celsius

Must be reported as an absolute temperature, NOT a delta from the max value.

RW

temp[1-*]_input Temperature input value.
Unit: millidegree Celsius
RO

temp[1-*]_crit Temperature critical value, typically greater than
corresponding temp_max values.
Unit: millidegree Celsius
RW

temp[1-*]_crit_hyst Temperature hysteresis value for critical limit.
Unit: millidegree Celsius
Must be reported as an absolute temperature, NOT a delta
from the critical value.
RW

temp[1-*]_offset Temperature offset which is added to the temperature reading
by the chip.
Unit: millidegree Celsius
Read/Write value.

temp[1-*]_label Suggested temperature channel label.
Text string
Should only be created if the driver has hints about what
this temperature channel is being used for, and user-space
doesn't. In all other cases, the label is provided by
user-space.
RO

temp[1-*]_lowest Historical minimum temperature
Unit: millidegree Celsius
RO

temp[1-*]_highest Historical maximum temperature
Unit: millidegree Celsius
RO

temp[1-*]_reset_history Reset temp_lowest and temp_highest
WO

temp_reset_history Reset temp_lowest and temp_highest for all sensors
WO

Some chips measure temperature using external thermistors and an ADC, and report the temperature measurement as a voltage. Converting this voltage back to a temperature (or the other way around for limits) requires mathematical functions not available in the kernel, so the conversion must occur in user space. For these chips, all temp* files described above should contain values expressed in millivolt instead of millidegree

Celsius. In other words, such temperature channels are handled as voltage channels by the driver.

Also see the Alarms section for status flags associated with temperatures.

 * Currents *

Note that no known chip provides current measurements as of writing, so this part is theoretical, so to say.

curr[1-*]_max Current max value
 Unit: milliampere
 RW

curr[1-*]_min Current min value.
 Unit: milliampere
 RW

curr[1-*]_input Current input value
 Unit: milliampere
 RO

 * Power *

power[1-*]_average Average power use
 Unit: microWatt
 RO

power[1-*]_average_interval Power use averaging interval. A poll
 notification is sent to this file if the
 hardware changes the averaging interval.
 Unit: milliseconds
 RW

power[1-*]_average_interval_max Maximum power use averaging interval
 Unit: milliseconds
 RO

power[1-*]_average_interval_min Minimum power use averaging interval
 Unit: milliseconds
 RO

power[1-*]_average_highest Historical average maximum power use
 Unit: microWatt
 RO

power[1-*]_average_lowest Historical average minimum power use
 Unit: microWatt
 RO

power[1-*]_average_max A poll notification is sent to

	sysfs-interface..txt
	power[1-*]_average when power use rises above this value. Unit: microWatt RW
power[1-*]_average_min	A poll notification is sent to power[1-*]_average when power use sinks below this value. Unit: microWatt RW
power[1-*]_input	Instantaneous power use Unit: microWatt RO
power[1-*]_input_highest	Historical maximum power use Unit: microWatt RO
power[1-*]_input_lowest	Historical minimum power use Unit: microWatt RO
power[1-*]_reset_history	Reset input_highest, input_lowest, average_highest and average_lowest. WO
power[1-*]_accuracy	Accuracy of the power meter. Unit: Percent RO
power[1-*]_alarm	1 if the system is drawing more power than the cap allows; 0 otherwise. A poll notification is sent to this file when the power use exceeds the cap. This file only appears if the cap is known to be enforced by hardware. RO
power[1-*]_cap	If power use rises above this limit, the system should take action to reduce power use. A poll notification is sent to this file if the cap is changed by the hardware. The *_cap files only appear if the cap is known to be enforced by hardware. Unit: microWatt RW
power[1-*]_cap_hyst	Margin of hysteresis built around capping and notification. Unit: microWatt RW
power[1-*]_cap_max	Maximum cap that can be set. Unit: microWatt RO

sysfs-interface..txt

power[1-*]_cap_min Minimum cap that can be set.
 Unit: microWatt
 RO

* Energy *

energy[1-*]_input Cumulative energy use
 Unit: microJoule
 RO

* Alarms *

Each channel or limit may have an associated alarm file, containing a boolean value. 1 means than an alarm condition exists, 0 means no alarm.

Usually a given chip will either use channel-related alarms, or limit-related alarms, not both. The driver should just reflect the hardware implementation.

in[0-*]_alarm
fan[1-*]_alarm
temp[1-*]_alarm
 Channel alarm
 0: no alarm
 1: alarm
 RO

OR

in[0-*]_min_alarm
in[0-*]_max_alarm
fan[1-*]_min_alarm
fan[1-*]_max_alarm
temp[1-*]_min_alarm
temp[1-*]_max_alarm
temp[1-*]_crit_alarm
 Limit alarm
 0: no alarm
 1: alarm
 RO

Each input channel may have an associated fault file. This can be used to notify open diodes, unconnected fans etc. where the hardware supports it. When this boolean has value 1, the measurement for that channel should not be trusted.

in[0-*]_fault
fan[1-*]_fault
temp[1-*]_fault
 Input fault condition
 0: no fault occured

sysfs-interface..txt

1: fault condition
RO

Some chips also offer the possibility to get beeped when an alarm occurs:

beep_enable Master beep enable
 0: no beeps
 1: beeps
 RW

in[0-*]_beep
fan[1-*]_beep
temp[1-*]_beep
 Channel beep
 0: disable
 1: enable
 RW

In theory, a chip could provide per-limit beep masking, but no such chip was seen so far.

Old drivers provided a different, non-standard interface to alarms and beeps. These interface files are deprecated, but will be kept around for compatibility reasons:

alarms Alarm bitmask.
 RO
 Integer representation of one to four bytes.
 A '1' bit means an alarm.
 Chips should be programmed for 'comparator' mode so that the alarm will 'come back' after you read the register if it is still valid.
 Generally a direct representation of a chip's internal alarm registers; there is no standard for the position of individual bits. For this reason, the use of this interface file for new drivers is discouraged. Use individual *_alarm and *_fault files instead.
 Bits are defined in kernel/include/sensors.h.

beep_mask Bitmask for beep.
 Same format as 'alarms' with the same bit locations, use discouraged for the same reason. Use individual *_beep files instead.
 RW

* Intrusion detection *

intrusion[0-*]_alarm
 Chassis intrusion detection
 0: OK
 1: intrusion detected
 RW
 Contrary to regular alarm flags which clear themselves

sysfs-interface..txt

automatically when read, this one sticks until cleared by the user. This is done by writing 0 to the file. Writing other values is unsupported.

```
intrusion[0-*]_beep
    Chassis intrusion beep
    0: disable
    1: enable
    RW
```

sysfs attribute writes interpretation

hwmon sysfs attributes always contain numbers, so the first thing to do is to convert the input to a number, there are 2 ways to do this depending whether the number can be negative or not:

```
unsigned long u = simple_strtoul(buf, NULL, 10);
long s = simple_strtol(buf, NULL, 10);
```

With `buf` being the buffer with the user input being passed by the kernel. Notice that we do not use the second argument of `strto[u]l`, and thus cannot tell when 0 is returned, if this was really 0 or is caused by invalid input. This is done deliberately as checking this everywhere would add a lot of code to the kernel.

Notice that it is important to always store the converted value in an unsigned long or long, so that no wrap around can happen before any further checking.

After the input string is converted to an (unsigned) long, the value should be checked if its acceptable. Be careful with further conversions on the value before checking it for validity, as these conversions could still cause a wrap around before the check. For example do not multiply the result, and only add/subtract if it has been divided before the add/subtract.

What to do if a value is found to be invalid, depends on the type of the sysfs attribute that is being set. If it is a continuous setting like a `tempX_max` or `inX_max` attribute, then the value should be clamped to its limits using `SENSORS_LIMIT(value, min_limit, max_limit)`. If it is not continuous like for example a `tempX_type`, then when an invalid value is written, `-EINVAL` should be returned.

Example1, `templ_max`, register is a signed 8 bit value (-128 - 127 degrees):

```
long v = simple_strtol(buf, NULL, 10) / 1000;
v = SENSORS_LIMIT(v, -128, 127);
/* write v to register */
```

Example2, fan divider setting, valid values 2, 4 and 8:

```
unsigned long v = simple_strtoul(buf, NULL, 10);

switch (v) {
case 2: v = 1; break;
case 4: v = 2; break;
```

```
sysfs-interface..txt
case 8: v = 3; break;
default:
    return -EINVAL;
}
/* write v to register */
```