

hdio.txt
Summary of HDIO_ ioctl calls.
=====

Edward A. Falk <efalk@google.com>

November, 2004

This document attempts to describe the ioctl(2) calls supported by the HD/IDE layer. These are by-and-large implemented (as of Linux 2.6) in drivers/ide/ide.c and drivers/block/scsi_ioctl.c

ioctl values are listed in <linux/hdreg.h>. As of this writing, they are as follows:

ioctls that pass argument pointers to user space:

HDIO_GETGEO	get device geometry
HDIO_GET_UNMASKINTR	get current unmask setting
HDIO_GET_MULTCOUNT	get current IDE blockmode setting
HDIO_GET_QDMA	get use-qdma flag
HDIO_SET_XFER	set transfer rate via proc
HDIO_OBSOLETE_IDENTITY	OBSOLETE, DO NOT USE
HDIO_GET_KEEPPSETTINGS	get keep-settings-on-reset flag
HDIO_GET_32BIT	get current io_32bit setting
HDIO_GET_NOWERR	get ignore-write-error flag
HDIO_GET_DMA	get use-dma flag
HDIO_GET_NICE	get nice flags
HDIO_GET_IDENTITY	get IDE identification info
HDIO_GET_WCACHE	get write cache mode on/off
HDIO_GET_ACOUSTIC	get acoustic value
HDIO_GET_ADDRESS	get sector addressing mode
HDIO_GET_BUSSTATE	get the bus state of the hwif
HDIO_TRISTATE_HWIF	execute a channel tristate
HDIO_DRIVE_RESET	execute a device reset
HDIO_DRIVE_TASKFILE	execute raw taskfile
HDIO_DRIVE_TASK	execute task and special drive command
HDIO_DRIVE_CMD	execute a special drive command
HDIO_DRIVE_CMD_AEB	HDIO_DRIVE_TASK

ioctls that pass non-pointer values:

HDIO_SET_MULTCOUNT	change IDE blockmode
HDIO_SET_UNMASKINTR	permit other irqs during I/O
HDIO_SET_KEEPPSETTINGS	keep ioctl settings on reset
HDIO_SET_32BIT	change io_32bit flags
HDIO_SET_NOWERR	change ignore-write-error flag
HDIO_SET_DMA	change use-dma flag
HDIO_SET_PIO_MODE	reconfig interface to new speed
HDIO_SCAN_HWIF	register and (re)scan interface
HDIO_SET_NICE	set nice flags
HDIO_UNREGISTER_HWIF	unregister interface
HDIO_SET_WCACHE	change write cache enable-disable
HDIO_SET_ACOUSTIC	change acoustic behavior
HDIO_SET_BUSSTATE	set the bus state of the hwif
HDIO_SET_QDMA	change use-qdma flag
HDIO_SET_ADDRESS	change lba addressing modes

hdio.txt

HDIO_SET_IDE_SCSI	Set scsi emulation mode on/off
HDIO_SET_SCSI_IDE	not implemented yet

The information that follows was determined from reading kernel source code. It is likely that some corrections will be made over time.

General:

Unless otherwise specified, all ioctl calls return 0 on success and -1 with errno set to an appropriate value on error.

Unless otherwise specified, all ioctl calls return -1 and set errno to EFAULT on a failed attempt to copy data to or from user address space.

Unless otherwise specified, all data structures and constants are defined in <linux/hdreg.h>

HDIO_GETGEO get device geometry

usage:

```
struct hd_geometry geom;
ioctl(fd, HDIO_GETGEO, &geom);
```

inputs: none

outputs:

hd_geometry structure containing:

heads	number of heads
sectors	number of sectors/track
cylinders	number of cylinders, mod 65536
start	starting sector of this partition.

error returns:

EINVAL	if the device is not a disk drive or floppy drive, or if the user passes a null pointer
--------	--

notes:

Not particularly useful with modern disk drives, whose geometry

hdio.txt

is a polite fiction anyway. Modern drives are addressed purely by sector number nowadays (lba addressing), and the drive geometry is an abstraction which is actually subject to change. Currently (as of Nov 2004), the geometry values are the "bios" values -- presumably the values the drive had when Linux first booted.

In addition, the cylinders field of the `hd_geometry` is an unsigned short, meaning that on most architectures, this `ioctl` will not return a meaningful value on drives with more than 65535 tracks.

The start field is unsigned long, meaning that it will not contain a meaningful value for disks over 219 Gb in size.

`HDIO_GET_UNMASKINTR` get current unmask setting

usage:

```
long val;
ioctl(fd, HDIO_GET_UNMASKINTR, &val);
```

inputs: none

outputs:

The value of the drive's current unmask setting

`HDIO_SET_UNMASKINTR` permit other irqs during I/O

usage:

```
unsigned long val;
ioctl(fd, HDIO_SET_UNMASKINTR, val);
```

inputs:

New value for unmask flag

outputs: none

error return:

<code>EINVAL</code>	(<code>bdev != bdev->bd_contains</code>) (not sure what this means)
<code>EACCES</code>	Access denied: requires <code>CAP_SYS_ADMIN</code>
<code>EINVAL</code>	value out of range [0 1]
<code>EBUSY</code>	Controller busy

`HDIO_GET_MULTCOUNT` get current IDE blockmode setting

usage:

hdio.txt

```
long val;  
ioctl(fd, HDIO_GET_MULTCOUNT, &val);
```

inputs: none

outputs:

The value of the current IDE block mode setting. This controls how many sectors the drive will transfer per interrupt.

HDIO_SET_MULTCOUNT change IDE blockmode

usage:

```
int val;  
ioctl(fd, HDIO_SET_MULTCOUNT, val);
```

inputs:

New value for IDE block mode setting. This controls how many sectors the drive will transfer per interrupt.

outputs: none

error return:

EINVAL	(bdev != bdev->bd_contains) (not sure what this means)
EACCES	Access denied: requires CAP_SYS_ADMIN
EINVAL	value out of range supported by disk.
EBUSY	Controller busy or blockmode already set.
EIO	Drive did not accept new block mode.

notes:

Source code comments read:

This is tightly woven into the driver->do_special cannot touch. DON'T do it again until a total personality rewrite is committed.

If blockmode has already been set, this ioctl will fail with EBUSY

HDIO_GET_QDMA get use-qdma flag

Not implemented, as of 2.6.8.1

HDIO_SET_XFER set transfer rate via proc

Not implemented, as of 2.6.8.1

HDIO_OBSOLETE_IDENTITY OBSOLETE, DO NOT USE

Same as HDIO_GET_IDENTITY (see below), except that it only returns the first 142 bytes of drive identity information.

HDIO_GET_IDENTITY get IDE identification info

usage:

```
unsigned char identity[512];
ioctl(fd, HDIO_GET_IDENTITY, identity);
```

inputs: none

outputs:

ATA drive identity information. For full description, see the IDENTIFY DEVICE and IDENTIFY PACKET DEVICE commands in the ATA specification.

error returns:

EINVAL	(bdev != bdev->bd_contains) (not sure what this means)
ENOMSG	IDENTIFY DEVICE information not available

notes:

Returns information that was obtained when the drive was probed. Some of this information is subject to change, and this ioctl does not re-probe the drive to update the information.

This information is also available from /proc/ide/hdX/identify

HDIO_GET_KEEPPSETTINGS get keep-settings-on-reset flag

usage:

```
long val;
ioctl(fd, HDIO_GET_KEEPPSETTINGS, &val);
```

inputs: none

outputs:

The value of the current "keep settings" flag

notes:

When set, indicates that kernel should restore settings after a drive reset.

HDIO_SET_KEESETTINGS keep ioctl settings on reset

usage:

```
long val;  
ioctl(fd, HDIO_SET_KEESETTINGS, val);
```

inputs:

New value for keep_settings flag

outputs: none

error return:

EINVAL	(bdev != bdev->bd_contains) (not sure what this means)
EACCES	Access denied: requires CAP_SYS_ADMIN
EINVAL	value out of range [0 1]
EBUSY	Controller busy

HDIO_GET_32BIT get current io_32bit setting

usage:

```
long val;  
ioctl(fd, HDIO_GET_32BIT, &val);
```

inputs: none

outputs:

The value of the current io_32bit setting

notes:

0=16-bit, 1=32-bit, 2,3 = 32bit+sync

HDIO_GET_NOWERR get ignore-write-error flag

usage:

```
long val;  
ioctl(fd, HDIO_GET_NOWERR, &val);
```

inputs: none

outputs:

The value of the current ignore-write-error flag

HDIO_GET_DMA get use-dma flag

hdio.txt

usage:

```
long val;
ioctl(fd, HDIO_GET_DMA, &val);
```

inputs: none

outputs:

The value of the current use-dma flag

HDIO_GET_NICE get nice flags

usage:

```
long nice;
ioctl(fd, HDIO_GET_NICE, &nice);
```

inputs: none

outputs:

The drive's "nice" values.

notes:

Per-drive flags which determine when the system will give more bandwidth to other devices sharing the same IDE bus.
See <linux/hdreg.h>, near symbol IDE_NICE_DSC_OVERLAP.

HDIO_SET_NICE set nice flags

usage:

```
unsigned long nice;
...
ioctl(fd, HDIO_SET_NICE, nice);
```

inputs:

bitmask of nice flags.

outputs: none

error returns:

EACCES	Access denied: requires CAP_SYS_ADMIN
EPERM	Flags other than DSC_OVERLAP and NICE_1 set.
EPERM	DSC_OVERLAP specified but not supported by drive

notes:

This ioctl sets the DSC_OVERLAP and NICE_1 flags from values provided by the user.

hdio.txt

Nice flags are listed in <linux/hdreg.h>, starting with IDE_NICE_DSC_OVERLAP. These values represent shifts.

HDIO_GET_WCACHE get write cache mode on|off

usage:

```
long val;
ioctl(fd, HDIO_GET_WCACHE, &val);
```

inputs: none

outputs:

The value of the current write cache mode

HDIO_GET_ACOUSTIC get acoustic value

usage:

```
long val;
ioctl(fd, HDIO_GET_ACOUSTIC, &val);
```

inputs: none

outputs:

The value of the current acoustic settings

notes:

See HDIO_SET_ACOUSTIC

HDIO_GET_ADDRESS

usage:

```
long val;
ioctl(fd, HDIO_GET_ADDRESS, &val);
```

inputs: none

outputs:

The value of the current addressing mode:

0 = 28-bit

1 = 48-bit

2 = 48-bit doing 28-bit

3 = 64-bit

HDIO_GET_BUSSTATE get the bus state of the hwif

usage:

```
long state;
ioctl(fd, HDIO_SCAN_HWIF, &state);
```

inputs: none

outputs:

Current power state of the IDE bus. One of BUSSTATE_OFF, BUSSTATE_ON, or BUSSTATE_TRISTATE

error returns:

EACCES Access denied: requires CAP_SYS_ADMIN

HDIO_SET_BUSSTATE set the bus state of the hwif

usage:

```
int state;
...
ioctl(fd, HDIO_SCAN_HWIF, state);
```

inputs:

Desired IDE power state. One of BUSSTATE_OFF, BUSSTATE_ON, or BUSSTATE_TRISTATE

outputs: none

error returns:

EACCES Access denied: requires CAP_SYS_RAWIO

EOPNOTSUPP Hardware interface does not support bus power control

HDIO_TRISTATE_HWIF execute a channel tristate

Not implemented, as of 2.6.8.1. See HDIO_SET_BUSSTATE

HDIO_DRIVE_RESET execute a device reset

usage:

```
int args[3]
...
ioctl(fd, HDIO_DRIVE_RESET, args);
```

hdio.txt

inputs: none

outputs: none

error returns:

EACCES	Access denied: requires CAP_SYS_ADMIN
ENXIO	No such device: phy dead or ctl_addr == 0
EIO	I/O error: reset timed out or hardware error

notes:

Execute a reset on the device as soon as the current IO operation has completed.

Executes an ATAPI soft reset if applicable, otherwise executes an ATA soft reset on the controller.

HDIO_DRIVE_TASKFILE

execute raw taskfile

Note: If you don't have a copy of the ANSI ATA specification handy, you should probably ignore this ioctl.

Execute an ATA disk command directly by writing the "taskfile" registers of the drive. Requires ADMIN and RAWIO access privileges.

usage:

```
struct {
    ide_task_request_t req_task;
    u8 outbuf[OUTPUT_SIZE];
    u8 inbuf[INPUT_SIZE];
} task;
memset(&task.req_task, 0, sizeof(task.req_task));
task.req_task.out_size = sizeof(task.outbuf);
task.req_task.in_size = sizeof(task.inbuf);
...
ioctl(fd, HDIO_DRIVE_TASKFILE, &task);
...
```

inputs:

(See below for details on memory area passed to ioctl.)

io_ports[8]	values to be written to taskfile registers
hob_ports[8]	high-order bytes, for extended commands.
out_flags	flags indicating which registers are valid
in_flags	flags indicating which registers should be returned
data_phase	see below
req_cmd	command type to be executed
out_size	size of output buffer
outbuf	buffer of data to be transmitted to disk
inbuf	buffer of data to be received from disk (see [1])

outputs:

io_ports[]	values returned in the taskfile registers
hob_ports[]	high-order bytes, for extended commands.
out_flags	flags indicating which registers are valid (see [2])
in_flags	flags indicating which registers should be returned
outbuf	buffer of data to be transmitted to disk (see [1])
inbuf	buffer of data to be received from disk

error returns:

EACCES	CAP_SYS_ADMIN or CAP_SYS_RAWIO privilege not set.
ENOMSG	Device is not a disk drive.
ENOMEM	Unable to allocate memory for task
EFAULT	req_cmd == TASKFILE_IN_OUT (not implemented as of 2.6.8)
EPERM	req_cmd == TASKFILE_MULTI_OUT and drive multi-count not yet set.
EIO	Drive failed the command.

notes:

[1] READ THE FOLLOWING NOTES *CAREFULLY*. THIS IOCTL IS FULL OF GOTCHAS. Extreme caution should be used with using this ioctl. A mistake can easily corrupt data or hang the system.

[2] Both the input and output buffers are copied from the user and written back to the user, even when not used.

[3] If one or more bits are set in out_flags and in_flags is zero, the following values are used for in_flags.all and written back into in_flags on completion.

- * IDE_TASKFILE_STD_IN_FLAGS | (IDE_HOB_STD_IN_FLAGS << 8)
if LBA48 addressing is enabled for the drive
- * IDE_TASKFILE_STD_IN_FLAGS
if CHS/LBA28

The association between in_flags.all and each enable bitfield flips depending on endianness; fortunately, TASKFILE only uses in_flags.b.data bit and ignores all other bits. The end result is that, on any endian machines, it has no effect other than modifying in_flags on completion.

[4] The default value of SELECT is (0xa0|DEV_bit|LBA_bit) except for four drives per port chipsets. For four drives per port chipsets, it's (0xa0|DEV_bit|LBA_bit) for the first pair and (0x80|DEV_bit|LBA_bit) for the second pair.

[5] The argument to the ioctl is a pointer to a region of memory containing a ide_task_request_t structure, followed by an optional buffer of data to be transmitted to the drive, followed by an optional buffer to receive data from the drive.

Command is passed to the disk drive via the ide_task_request_t structure, which contains these fields:

hdio.txt

io_ports[8]	values for the taskfile registers
hob_ports[8]	high-order bytes, for extended commands
out_flags	flags indicating which entries in the io_ports[] and hob_ports[] arrays contain valid values. Type ide_reg_valid_t.
in_flags	flags indicating which entries in the io_ports[] and hob_ports[] arrays are expected to contain valid values on return.
data_phase	See below
req_cmd	Command type, see below
out_size	output (user->drive) buffer size, bytes
in_size	input (drive->user) buffer size, bytes

When out_flags is zero, the following registers are loaded.

HOB_FEATURE	If the drive supports LBA48
HOB_NSECTOR	If the drive supports LBA48
HOB_SECTOR	If the drive supports LBA48
HOB_LCYL	If the drive supports LBA48
HOB_HCYL	If the drive supports LBA48
FEATURE	
NSECTOR	
SECTOR	
LCYL	
HCYL	
SELECT	First, masked with 0xE0 if LBA48, 0xEF otherwise; then, or'ed with the default value of SELECT.

If any bit in out_flags is set, the following registers are loaded.

HOB_DATA	If out_flags.b.data is set. HOB_DATA will travel on DD8-DD15 on little endian machines and on DD0-DD7 on big endian machines.
DATA	If out_flags.b.data is set. DATA will travel on DD0-DD7 on little endian machines and on DD8-DD15 on big endian machines.
HOB_NSECTOR	If out_flags.b.nsector_hob is set
HOB_SECTOR	If out_flags.b.sector_hob is set
HOB_LCYL	If out_flags.b.lcyl_hob is set
HOB_HCYL	If out_flags.b.hcyl_hob is set
FEATURE	If out_flags.b.feature is set
NSECTOR	If out_flags.b.nsector is set
SECTOR	If out_flags.b.sector is set
LCYL	If out_flags.b.lcyl is set
HCYL	If out_flags.b.hcyl is set
SELECT	Or'ed with the default value of SELECT and loaded regardless of out_flags.b.select.

Taskfile registers are read back from the drive into {io|hob}_ports[] after the command completes iff one of the following conditions is met; otherwise, the original values will be written back, unchanged.

hdio.txt

1. The drive fails the command (EIO).
2. One or more than one bits are set in out_flags.
3. The requested data_phase is TASKFILE_NO_DATA.

HOB_DATA	If in_flags.b.data is set. It will contain DD8-DD15 on little endian machines and DD0-DD7 on big endian machines.
DATA	If in_flags.b.data is set. It will contain DD0-DD7 on little endian machines and DD8-DD15 on big endian machines.
HOB_FEATURE	If the drive supports LBA48
HOB_NSECTOR	If the drive supports LBA48
HOB_SECTOR	If the drive supports LBA48
HOB_LCYL	If the drive supports LBA48
HOB_HCYL	If the drive supports LBA48
NSECTOR	
SECTOR	
LCYL	
HCYL	

The data_phase field describes the data transfer to be performed. Value is one of:

TASKFILE_IN	
TASKFILE_MULTI_IN	
TASKFILE_OUT	
TASKFILE_MULTI_OUT	
TASKFILE_IN_OUT	
TASKFILE_IN_DMA	
TASKFILE_IN_DMAQ	== IN_DMA (queueing not supported)
TASKFILE_OUT_DMA	
TASKFILE_OUT_DMAQ	== OUT_DMA (queueing not supported)
TASKFILE_P_IN	unimplemented
TASKFILE_P_IN_DMA	unimplemented
TASKFILE_P_IN_DMAQ	unimplemented
TASKFILE_P_OUT	unimplemented
TASKFILE_P_OUT_DMA	unimplemented
TASKFILE_P_OUT_DMAQ	unimplemented

The req_cmd field classifies the command type. It may be one of:

IDE_DRIVE_TASK_NO_DATA	
IDE_DRIVE_TASK_SET_XFER	unimplemented
IDE_DRIVE_TASK_IN	
IDE_DRIVE_TASK_OUT	unimplemented
IDE_DRIVE_TASK_RAW_WRITE	

[6] Do not access {in|out}_flags->all except for resetting all the bits. Always access individual bit fields. ->all value will flip depending on endianness. For the same reason, do not use IDE_{TASKFILE|HOB}_STD_{OUT|IN}_FLAGS constants defined in hdreg.h.

hdio.txt

HDIO_DRIVE_CMD

execute a special drive command

Note: If you don't have a copy of the ANSI ATA specification handy, you should probably ignore this ioctl.

usage:

```
u8 args[4+XFER_SIZE];
...
ioctl(fd, HDIO_DRIVE_CMD, args);
```

inputs:

Commands other than WIN_SMART

args[0]	COMMAND
args[1]	NSECTOR
args[2]	FEATURE
args[3]	NSECTOR

WIN_SMART

args[0]	COMMAND
args[1]	SECTOR
args[2]	FEATURE
args[3]	NSECTOR

outputs:

args[] buffer is filled with register values followed by any data returned by the disk.

args[0]	status
args[1]	error
args[2]	NSECTOR
args[3]	undefined
args[4+]	NSECTOR * 512 bytes of data returned by the command.

error returns:

EACCES	Access denied: requires CAP_SYS_RAWIO
ENOMEM	Unable to allocate memory for task
EIO	Drive reports error

notes:

[1] For commands other than WIN_SMART, args[1] should equal args[3]. SECTOR, LCYL and HCYL are undefined. For WIN_SMART, 0x4f and 0xc2 are loaded into LCYL and HCYL respectively. In both cases SELECT will contain the default value for the drive. Please refer to HDIO_DRIVE_TASKFILE notes for the default value of SELECT.

[2] If NSECTOR value is greater than zero and the drive sets DRQ when interrupting for the command, NSECTOR * 512 bytes are read from the device into the area following NSECTOR. In the above example, the area would be args[4..4+XFER_SIZE]. 16bit PIO is used regardless of HDIO_SET_32BIT setting.

hdio.txt

[3] If COMMAND == WIN_SETFEATURES && FEATURE == SETFEATURES_XFER && NSECTOR >= XFER_SW_DMA_0 && the drive supports any DMA mode, IDE driver will try to tune the transfer mode of the drive accordingly.

HDIO_DRIVE_TASK execute task and special drive command

Note: If you don't have a copy of the ANSI ATA specification handy, you should probably ignore this ioctl.

usage:

```
u8 args[7];
...
ioctl(fd, HDIO_DRIVE_TASK, args);
```

inputs:

Taskfile register values:

args[0]	COMMAND
args[1]	FEATURE
args[2]	NSECTOR
args[3]	SECTOR
args[4]	LCYL
args[5]	HCYL
args[6]	SELECT

outputs:

Taskfile register values:

args[0]	status
args[1]	error
args[2]	NSECTOR
args[3]	SECTOR
args[4]	LCYL
args[5]	HCYL
args[6]	SELECT

error returns:

EACCES	Access denied: requires CAP_SYS_RAWIO
ENOMEM	Unable to allocate memory for task
ENOMSG	Device is not a disk drive.
EIO	Drive failed the command.

notes:

[1] DEV bit (0x10) of SELECT register is ignored and the appropriate value for the drive is used. All other bits are used unaltered.

HDIO_DRIVE_CMD_AEB

HDIO_DRIVE_TASK

hdio.txt
Not implemented, as of 2.6.8.1

HDIO_SET_32BIT change io_32bit flags

usage:

```
int val;  
ioctl(fd, HDIO_SET_32BIT, val);
```

inputs:

New value for io_32bit flag

outputs: none

error return:

EINVAL	(bdev != bdev->bd_contains) (not sure what this means)
EACCES	Access denied: requires CAP_SYS_ADMIN
EINVAL	value out of range [0 3]
EBUSY	Controller busy

HDIO_SET_NOWERR change ignore-write-error flag

usage:

```
int val;  
ioctl(fd, HDIO_SET_NOWERR, val);
```

inputs:

New value for ignore-write-error flag. Used for ignoring
WRERR_STAT

outputs: none

error return:

EINVAL	(bdev != bdev->bd_contains) (not sure what this means)
EACCES	Access denied: requires CAP_SYS_ADMIN
EINVAL	value out of range [0 1]
EBUSY	Controller busy

HDIO_SET_DMA change use-dma flag

usage:

```
long val;  
ioctl(fd, HDIO_SET_DMA, val);
```

inputs:

New value for use-dma flag

hdio.txt

outputs: none

error return:

EINVAL	(bdev != bdev->bd_contains) (not sure what this means)
EACCES	Access denied: requires CAP_SYS_ADMIN
EINVAL	value out of range [0 1]
EBUSY	Controller busy

HDIO_SET_PIO_MODE reconfig interface to new speed

usage:

```
long val;  
ioctl(fd, HDIO_SET_PIO_MODE, val);
```

inputs:

New interface speed.

outputs: none

error return:

EINVAL	(bdev != bdev->bd_contains) (not sure what this means)
EACCES	Access denied: requires CAP_SYS_ADMIN
EINVAL	value out of range [0 255]
EBUSY	Controller busy

HDIO_SCAN_HWIF register and (re)scan interface

usage:

```
int args[3]  
...  
ioctl(fd, HDIO_SCAN_HWIF, args);
```

inputs:

args[0]	io address to probe
args[1]	control address to probe
args[2]	irq number

outputs: none

error returns:

EACCES	Access denied: requires CAP_SYS_RAWIO
EIO	Probe failed.

notes:

This ioctl initializes the addresses and irq for a disk controller, probes for drives, and creates /proc/ide interfaces as appropriate.

HDIO_UNREGISTER_HWIF unregister interface

usage:

```
int index;
ioctl(fd, HDIO_UNREGISTER_HWIF, index);
```

inputs:

index index of hardware interface to unregister

outputs: none

error returns:

EACCES Access denied: requires CAP_SYS_RAWIO

notes:

This ioctl removes a hardware interface from the kernel.

Currently (2.6.8) this ioctl silently fails if any drive on the interface is busy.

HDIO_SET_WCACHE change write cache enable-disable

usage:

```
int val;
ioctl(fd, HDIO_SET_WCACHE, val);
```

inputs:

New value for write cache enable

outputs: none

error return:

EINVAL (bdev != bdev->bd_contains) (not sure what this means)

EACCES Access denied: requires CAP_SYS_ADMIN

EINVAL value out of range [0 1]

EBUSY Controller busy

HDIO_SET_ACOUSTIC change acoustic behavior

usage:

```
int val;
ioctl(fd, HDIO_SET_ACOUSTIC, val);
```

inputs:

New value for drive acoustic settings

outputs: none

hdio.txt

error return:
EINVAL (bdev != bdev->bd_contains) (not sure what this means)
EACCES Access denied: requires CAP_SYS_ADMIN
EINVAL value out of range [0 254]
EBUSY Controller busy

HDIO_SET_QDMA change use-qdma flag

Not implemented, as of 2.6.8.1

HDIO_SET_ADDRESS change lba addressing modes

usage:

```
int val;  
ioctl(fd, HDIO_SET_ADDRESS, val);
```

inputs:

New value for addressing mode
0 = 28-bit
1 = 48-bit
2 = 48-bit doing 28-bit

outputs: none

error return:
EINVAL (bdev != bdev->bd_contains) (not sure what this means)
EACCES Access denied: requires CAP_SYS_ADMIN
EINVAL value out of range [0 2]
EBUSY Controller busy
EIO Drive does not support lba48 mode.

HDIO_SET_IDE_SCSI

usage:

```
long val;  
ioctl(fd, HDIO_SET_IDE_SCSI, val);
```

inputs:

New value for scsi emulation mode (?)

outputs: none

error return:
EINVAL (bdev != bdev->bd_contains) (not sure what this means)
EACCES Access denied: requires CAP_SYS_ADMIN
EINVAL value out of range [0 1]
EBUSY Controller busy

hdio.txt

HDIO_SET SCSI_IDE

Not implemented, as of 2.6.8.1