

GETTING STARTED WITH KMEMCHECK

Vegard Nossum <vegardno@ifi.uio.no>

Contents

- 0. Introduction
- 1. Downloading
- 2. Configuring and compiling
- 3. How to use
 - 3.1. Booting
 - 3.2. Run-time enable/disable
 - 3.3. Debugging
 - 3.4. Annotating false positives
- 4. Reporting errors
- 5. Technical description

0. Introduction

kmemcheck is a debugging feature for the Linux Kernel. More specifically, it is a dynamic checker that detects and warns about some uses of uninitialized memory.

Userspace programmers might be familiar with Valgrind's memcheck. The main difference between memcheck and kmemcheck is that memcheck works for userspace programs only, and kmemcheck works for the kernel only. The implementations are of course vastly different. Because of this, kmemcheck is not as accurate as memcheck, but it turns out to be good enough in practice to discover real programmer errors that the compiler is not able to find through static analysis.

Enabling kmemcheck on a kernel will probably slow it down to the extent that the machine will not be usable for normal workloads such as e.g. an interactive desktop. kmemcheck will also cause the kernel to use about twice as much memory as normal. For this reason, kmemcheck is strictly a debugging feature.

1. Downloading

As of version 2.6.31-rc1, kmemcheck is included in the mainline kernel.

2. Configuring and compiling

kmemcheck only works for the x86 (both 32- and 64-bit) platform. A number of configuration variables must have specific settings in order for the kmemcheck menu to even appear in "menuconfig". These are:

- o CONFIG_CC_OPTIMIZE_FOR_SIZE=n

This option is located under "General setup" / "Optimize for size".

Without this, gcc will use certain optimizations that usually lead to false positive warnings from kmemcheck. An example of this is a 16-bit field in a struct, where gcc may load 32 bits, then discard the upper 16 bits. kmemcheck sees only the 32-bit load, and may trigger a warning for the upper 16 bits (if they're uninitialized).

- o CONFIG_SLAB=y or CONFIG_SLUB=y

This option is located under "General setup" / "Choose SLAB allocator".

- o CONFIG_FUNCTION_TRACER=n

This option is located under "Kernel hacking" / "Tracers" / "Kernel Function Tracer"

When function tracing is compiled in, gcc emits a call to another function at the beginning of every function. This means that when the page fault handler is called, the ftrace framework will be called before kmemcheck has had a chance to handle the fault. If ftrace then modifies memory that was tracked by kmemcheck, the result is an endless recursive page fault.

- o CONFIG_DEBUG_PAGEALLOC=n

This option is located under "Kernel hacking" / "Debug page memory allocations".

In addition, I highly recommend turning on CONFIG_DEBUG_INFO=y. This is also located under "Kernel hacking". With this, you will be able to get line number information from the kmemcheck warnings, which is extremely valuable in debugging a problem. This option is not mandatory, however, because it slows down the compilation process and produces a much bigger kernel image.

Now the kmemcheck menu should be visible (under "Kernel hacking" / "kmemcheck: trap use of uninitialized memory"). Here follows a description of the kmemcheck configuration variables:

- o CONFIG_KMEMCHECK

This must be enabled in order to use kmemcheck at all...

- o CONFIG_KMEMCHECK_[DISABLED | ENABLED | ONESHOT]_BY_DEFAULT

This option controls the status of kmemcheck at boot-time. "Enabled" will enable kmemcheck right from the start, "disabled" will boot the kernel as normal (but with the kmemcheck code compiled in, so it can be enabled at run-time after the kernel has booted), and "one-shot" is a special mode which will turn kmemcheck off automatically after detecting the first use of uninitialized memory.

If you are using kmemcheck to actively debug a problem, then you probably want to choose "enabled" here.

The one-shot mode is mostly useful in automated test setups because it can prevent floods of warnings and increase the chances of the machine surviving in case something is really wrong. In other cases, the one-shot mode could actually be counter-productive because it would turn itself off at the very first error — in the case of a false positive too — and this would come in the way of debugging the specific problem you were interested in.

If you would like to use your kernel as normal, but with a chance to enable kmemcheck in case of some problem, it might be a good idea to choose "disabled" here. When kmemcheck is disabled, most of the run-time overhead is not incurred, and the kernel will be almost as fast as normal.

o CONFIG_KMEMCHECK_QUEUE_SIZE

Select the maximum number of error reports to store in an internal (fixed-size) buffer. Since errors can occur virtually anywhere and in any context, we need a temporary storage area which is guaranteed not to generate any other page faults when accessed. The queue will be emptied as soon as a tasklet may be scheduled. If the queue is full, new error reports will be lost.

The default value of 64 is probably fine. If some code produces more than 64 errors within an irqs-off section, then the code is likely to produce many, many more, too, and these additional reports seldom give any more information (the first report is usually the most valuable anyway).

This number might have to be adjusted if you are not using serial console or similar to capture the kernel log. If you are using the "dmesg" command to save the log, then getting a lot of kmemcheck warnings might overflow the kernel log itself, and the earlier reports will get lost in that way instead. Try setting this to 10 or so on such a setup.

o CONFIG_KMEMCHECK_SHADOW_COPY_SHIFT

Select the number of shadow bytes to save along with each entry of the error-report queue. These bytes indicate what parts of an allocation are initialized, uninitialized, etc. and will be displayed when an error is detected to help the debugging of a particular problem.

The number entered here is actually the logarithm of the number of bytes that will be saved. So if you pick for example 5 here, kmemcheck will save $2^5 = 32$ bytes.

The default value should be fine for debugging most problems. It also fits nicely within 80 columns.

o CONFIG_KMEMCHECK_PARTIAL_OK

This option (when enabled) works around certain GCC optimizations that produce 32-bit reads from 16-bit variables where the upper 16 bits are thrown away afterwards.

The default value (enabled) is recommended. This may of course hide some real errors, but disabling it would probably produce a lot of false positives.

o CONFIG_KMEMCHECK_BITOPS_OK

This option silences warnings that would be generated for bit-field accesses where not all the bits are initialized at the same time. This may also hide some real bugs.

This option is probably obsolete, or it should be replaced with the kmemcheck-/bitfield-annotations for the code in question. The default value is therefore fine.

Now compile the kernel as usual.

3. How to use

=====

3.1. Booting

=====

First some information about the command-line options. There is only one option specific to kmemcheck, and this is called "kmemcheck". It can be used to override the default mode as chosen by the CONFIG_KMEMCHECK_*_BY_DEFAULT option. Its possible settings are:

- o kmemcheck=0 (disabled)
- o kmemcheck=1 (enabled)
- o kmemcheck=2 (one-shot mode)

If SLUB debugging has been enabled in the kernel, it may take precedence over kmemcheck in such a way that the slab caches which are under SLUB debugging will not be tracked by kmemcheck. In order to ensure that this doesn't happen (even though it shouldn't by default), use SLUB's boot option "slub_debug", like this: slub_debug=--

In fact, this option may also be used for fine-grained control over SLUB vs. kmemcheck. For example, if the command line includes "kmemcheck=1 slub_debug=,dentry", then SLUB debugging will be used only for the "dentry" slab cache, and with kmemcheck tracking all the other caches. This is advanced usage, however, and is not generally recommended.

3.2. Run-time enable/disable

=====

When the kernel has booted, it is possible to enable or disable kmemcheck at run-time. WARNING: This feature is still experimental and may cause false positive warnings to appear. Therefore, try not to use this. If you find that it doesn't work properly (e.g. you see an unreasonable amount of warnings), I will be happy to take bug reports.

Use the file /proc/sys/kernel/kmemcheck for this purpose, e.g.:

kmemcheck.txt

called from a thousand places in the kernel, and is therefore not very useful. These inlined functions would not show up in the stack trace above, simply because the kernel doesn't load the extra debugging information. This technique can of course be used with ordinary kernel oopses as well.

In this case, it's the caller of `memcpy()` that is interesting, and it can be found in `include/asm-generic/siginfo.h`, line 287:

```
281 static inline void copy_siginfo(struct siginfo *to, struct siginfo *from)
282 {
283     if (from->si_code < 0)
284         memcpy(to, from, sizeof(*to));
285     else
286         /* _sigchld is currently the largest know union member */
287         memcpy(to, from, __ARCH_SI_PREAMBLE_SIZE +
288             sizeof(from->_sifields._sigchld));
288 }
```

Since this was a read (kmemcheck usually warns about reads only, though it can warn about writes to unallocated or freed memory as well), it was probably the "from" argument which contained some uninitialized bytes. Following the chain of calls, we move upwards to see where "from" was allocated or initialized, `kernel/signal.c`, line 380:

```
359 static void collect_signal(int sig, struct sigpending *list, siginfo_t
360 *info)
361 {
362     ...
363     list_for_each_entry(q, &list->list, list) {
364         if (q->info.si_signo == sig) {
365             if (first)
366                 goto still_pending;
367             first = q;
368         }
369     }
370     if (first) {
371         still_pending:
372         list_del_init(&first->list);
373         copy_siginfo(info, &first->info);
374         __sigqueue_free(first);
375     }
376 }
377 }
```

Here, it is `&first->info` that is being passed on to `copy_siginfo()`. The variable "first" was found on a list -- passed in as the second argument to `collect_signal()`. We continue our journey through the stack, to figure out where the item on "list" was allocated or initialized. We move to line 410:

```
395 static int __dequeue_signal(struct sigpending *pending, sigset_t *mask,
396 siginfo_t *info)
397 {
398     ...
399     collect_signal(sig, pending, info);
400     ...
401 }
```

kmemcheck.txt

Now we need to follow the "pending" pointer, since that is being passed on to collect_signal() as "list". At this point, we've run out of lines from the "addr2line" output. Not to worry, we just paste the next addresses from the kmemcheck stack dump, i.e.:

```
[<ffffffff8104f04e>] dequeue_signal+0x8e/0x170
[<ffffffff81050bd8>] get_signal_to_deliver+0x98/0x390
[<ffffffff8100b87d>] do_notify_resume+0xad/0x7d0
[<ffffffff8100c7b5>] int_signal+0x12/0x17
```

```
$ addr2line -e vmlinux -i fffffffff8104f04e fffffffff81050bd8 \
    fffffffff8100b87d fffffffff8100c7b5
kernel/signal.c:446
kernel/signal.c:1806
arch/x86/kernel/signal.c:805
arch/x86/kernel/signal.c:871
arch/x86/kernel/entry_64.S:694
```

Remember that since these addresses were found on the stack and not as the RIP value, they actually point to the `_next_instruction` (they are return addresses). This becomes obvious when we look at the code for line 446:

```
422 int dequeue_signal(struct task_struct *tsk, sigset_t *mask, siginfo_t *info)
423 {
...
431         signr = __dequeue_signal(&tsk->signal->shared_pending,
432                                 mask, info);
433         /*
434          * itimer signal ?
435          *
436          * itimers are process shared and we restart periodic
437          * itimers in the signal delivery path to prevent DoS
438          * attacks in the high resolution timer case. This is
439          * compliant with the old way of self restarting
440          * itimers, as the SIGALRM is a legacy signal and only
441          * queued once. Changing the restart behaviour to
442          * restart the timer in the signal dequeue path is
443          * reducing the timer noise on heavy loaded !highres
444          * systems too.
445          */
446         if (unlikely(signr == SIGALRM)) {
...
489 }
```

So instead of looking at 446, we should be looking at 431, which is the line that executes just before 446. Here we see that what we are looking for is `&tsk->signal->shared_pending`.

Our next task is now to figure out which function that puts items on this "shared_pending" list. A crude, but efficient tool, is git grep:

```
$ git grep -n 'shared_pending' kernel/
...
kernel/signal.c:828:     pending = group ? &t->signal->shared_pending :
&t->pending;
kernel/signal.c:1339:     pending = group ? &t->signal->shared_pending :
第 7 页
```

```
&t->pending;
...
```

There were more results, but none of them were related to list operations, and these were the only assignments. We inspect the line numbers more closely and find that this is indeed where items are being added to the list:

```
816 static int send_signal(int sig, struct siginfo *info, struct task_struct *t,
817                        int group)
818 {
...
828     pending = group ? &t->signal->shared_pending : &t->pending;
...
851     q = __sigqueue_alloc(t, GFP_ATOMIC, (sig < SIGRTMIN &&
852                                         (is_si_special(info) ||
853                                         info->si_code >= 0)));
854     if (q) {
855         list_add_tail(&q->list, &pending->list);
...
890 }
```

and:

```
1309 int send_sigqueue(struct sigqueue *q, struct task_struct *t, int group)
1310 {
...
1339     pending = group ? &t->signal->shared_pending : &t->pending;
1340     list_add_tail(&q->list, &pending->list);
...
1347 }
```

In the first case, the list element we are looking for, "q", is being returned from the function `__sigqueue_alloc()`, which looks like an allocation function. Let's take a look at it:

```
187 static struct sigqueue *__sigqueue_alloc(struct task_struct *t, gfp_t flags,
188                                           int override_rlimit)
189 {
190     struct sigqueue *q = NULL;
191     struct user_struct *user;
192
193     /*
194      * We won't get problems with the target's UID changing under us
195      * because changing it requires RCU be used, and if t != current,
196      * caller must be holding the RCU readlock (by way of a spinlock)
197      * we use RCU protection here
198      */
199     user = get_uid(__task_cred(t)->user);
200     atomic_inc(&user->sigpending);
201     if (override_rlimit ||
202         atomic_read(&user->sigpending) <=
203             t->signal->rlim[RLIMIT_SIGPENDING].rlim_cur)
204         q = kmem_cache_alloc(sigqueue_cachep, flags);
205     if (unlikely(q == NULL)) {
```



```

kmemcheck.txt
206         atomic_dec(&user->sigpending);
207         free_uid(user);
208     } else {
209         INIT_LIST_HEAD(&q->list);
210         q->flags = 0;
211         q->user = user;
212     }
213
214     return q;
215 }
```

We see that this function initializes `q->list`, `q->flags`, and `q->user`. It seems that now is the time to look at the definition of "struct sigqueue", e.g.:

```

14 struct sigqueue {
15     struct list_head list;
16     int flags;
17     siginfo_t info;
18     struct user_struct *user;
19 };

```

And, you might remember, it was a `memcpy()` on `&first->info` that caused the warning, so this makes perfect sense. It also seems reasonable to assume that it is the caller of `__sigqueue_alloc()` that has the responsibility of filling out (initializing) this member.

But just which fields of the struct were uninitialized? Let's look at kmemcheck's report again:

[illegible]

These first two lines are the memory dump of the memory object itself, and the shadow bytemap, respectively. The memory object itself is in this case `&first->info`. Just beware that the start of this dump is NOT the start of the object itself! The position of the caret (^) corresponds with the address of the read (`ffff88003e4a2024`).

The shadow bytemap dump legend is as follows:

```
i - initialized
u - uninitialized
a - unallocated (memory has been allocated by the slab layer, but has not
    yet been handed off to anybody)
f - freed (memory has been allocated by the slab layer, but has been freed
    by the previous owner)
```

In order to figure out where (relative to the start of the object) the uninitialized memory was located, we have to look at the disassembly. For that, we'll need the RIP address again:

```
RIP: 0010:[<ffffffff8104ede8>] [<ffffffff8104ede8>] dequeue signal+0xc8/0x190
```

```

                                kmemcheck.txt
$ objdump -d --no-show-raw-insn vmlinux | grep -C 8 ffffffff8104ede8:
ffffffff8104edc8:      mov     %r8, 0x8(%r8)
ffffffff8104edcc:      test    %r10d, %r10d
ffffffff8104edcf:      js      ffffffff8104ee88 <__dequeue_signal+0x168>
ffffffff8104edd5:      mov     %rax, %rdx
ffffffff8104edd8:      mov     $0xc, %ecx
ffffffff8104eddd:      mov     %r13, %rdi
ffffffff8104ede0:      mov     $0x30, %eax
ffffffff8104ede5:      mov     %rdx, %rsi
ffffffff8104ede8:      rep movsl %ds: (%rsi), %es: (%rdi)
ffffffff8104edea:      test     $0x2, %al
ffffffff8104edec:      je       ffffffff8104edf0 <__dequeue_signal+0xd0>
ffffffff8104edee:      movsw    %ds: (%rsi), %es: (%rdi)
ffffffff8104edf0:      test     $0x1, %al
ffffffff8104edf2:      je       ffffffff8104edf5 <__dequeue_signal+0xd5>
ffffffff8104edf4:      movsb     %ds: (%rsi), %es: (%rdi)
ffffffff8104edf5:      mov      %r8, %rdi
ffffffff8104edf8:      callq    ffffffff8104de60 <__sigqueue_free>

```

As expected, it's the "rep movsl" instruction from the memcpy() that causes the warning. We know about REP MOVSL that it uses the register RCX to count the number of remaining iterations. By taking a look at the register dump again (from the kmemcheck report), we can figure out how many bytes were left to copy:

```
RAX: 00000000000000030 RBX: ffff88003d4ea968 RCX: 0000000000000009
```

By looking at the disassembly, we also see that %ecx is being loaded with the value \$0xc just before (ffffffff8104edd8), so we are very lucky. Keep in mind that this is the number of iterations, not bytes. And since this is a "long" operation, we need to multiply by 4 to get the number of bytes. So this means that the uninitialized value was encountered at $4 * (0xc - 0x9) = 12$ bytes from the start of the object.

We can now try to figure out which field of the "struct siginfo" that was not initialized. This is the beginning of the struct:

```

40 typedef struct siginfo {
41     int si_signo;
42     int si_errno;
43     int si_code;
44
45     union {
46
47
48
49
50
51
52
53
54
55
56
57
58
59
60
61
62
63
64
65
66
67
68
69
70
71
72
73
74
75
76
77
78
79
80
81
82
83
84
85
86
87
88
89
90
91
92     } _sifields;
93 } siginfo_t;

```

On 64-bit, the int is 4 bytes long, so it must be the union member that has not been initialized. We can verify this using gdb:

```


$ gdb vmlinux
(gdb) p &((struct siginfo *) 0)->_sifields
$1 = (union {...} *) 0x10

```

Actually, it seems that the union member is located at offset 0x10 -- which

kmemcheck.txt

means that gcc has inserted 4 bytes of padding between the members `si_code` and `_sifields`. We can now get a fuller picture of the memory dump:



This allows us to realize another important fact: `si_code` contains the value `0x80`. Remember that x86 is little endian, so the first 4 bytes `"80000000"` are really the number `0x00000080`. With a bit of research, we find that this is actually the constant `SI_KERNEL` defined in `include/asm-generic/siginfo.h`:

```
144 #define SI_KERNEL      0x80          /* sent by the kernel from somewhere
    */
```

This macro is used in exactly one place in the x86 kernel: In `send_signal()` in `kernel/signal.c`:

```

816 static int send_signal(int sig, struct siginfo *info, struct task_struct *t,
817                        int group)
818 {
...
828     pending = group ? &t->signal->shared_pending : &t->pending;
...
851     q = __sigqueue_alloc(t, GFP_ATOMIC, (sig < SIGRTMIN &&
852                                         (is_si_special(info) ||
853                                         info->si_code >= 0)));
854     if (q) {
855         list_add_tail(&q->list, &pending->list);
856         switch ((unsigned long) info) {
...
865             case (unsigned long) SEND_SIG_PRIV:
866                 q->info.si_signo = sig;
867                 q->info.si_errno = 0;
868                 q->info.si_code = SI_KERNEL;
869                 q->info.si_pid = 0;
870                 q->info.si_uid = 0;
871                 break;
...
890 }

```

Not only does this match with the `.si_code` member, it also matches the place we found earlier when looking for where `siginfo_t` objects are enqueued on the "shared pending" list.

So to sum up: It seems that it is the padding introduced by the compiler between two struct fields that is uninitialized, and this gets reported when we do a `memcpy()` on the struct. This means that we have identified a false positive warning.

Normally, kmemcheck will not report uninitialized accesses in memcpy() calls

kmemcheck.txt

when both the source and destination addresses are tracked. (Instead, we copy the shadow bytemap as well). In this case, the destination address clearly was not tracked. We can dig a little deeper into the stack trace from above:

```
arch/x86/kernel/signal.c:805
arch/x86/kernel/signal.c:871
arch/x86/kernel/entry_64.S:694
```

And we clearly see that the destination siginfo object is located on the stack:

```
782 static void do_signal(struct pt_regs *regs)
783 {
784     struct k_sigaction ka;
785     siginfo_t info;
...
804     signr = get_signal_to_deliver(&info, &ka, regs, NULL);
...
854 }
```

And this &info is what eventually gets passed to copy_siginfo() as the destination argument.

Now, even though we didn't find an actual error here, the example is still a good one, because it shows how one would go about to find out what the report was all about.

3.4. Annotating false positives

There are a few different ways to make annotations in the source code that will keep kmemcheck from checking and reporting certain allocations. Here they are:

o __GFP_NOTRACK_FALSE_POSITIVE

This flag can be passed to kmalloc() or kmem_cache_alloc() (therefore also to other functions that end up calling one of these) to indicate that the allocation should not be tracked because it would lead to a false positive report. This is a "big hammer" way of silencing kmemcheck; after all, even if the false positive pertains to particular field in a struct, for example, we will now lose the ability to find (real) errors in other parts of the same struct.

Example:

```
/* No warnings will ever trigger on accessing any part of x */
x = kmalloc(sizeof *x, GFP_KERNEL | __GFP_NOTRACK_FALSE_POSITIVE);
```

o kmemcheck_bitfield_begin(name)/kmemcheck_bitfield_end(name) and kmemcheck_annotate_bitfield(ptr, name)

The first two of these three macros can be used inside struct definitions to signal, respectively, the beginning and end of a bitfield. Additionally, this will assign the bitfield a name, which

kmemcheck.txt
is given as an argument to the macros.

Having used these markers, one can later use `kmemcheck_annotate_bitfield()` at the point of allocation, to indicate which parts of the allocation is part of a bitfield.

Example:

```
struct foo {
    int x;

    kmemcheck_bitfield_begin(flags);
    int flag_a:1;
    int flag_b:1;
    kmemcheck_bitfield_end(flags);

    int y;
};

struct foo *x = kmalloc(sizeof *x);

/* No warnings will trigger on accessing the bitfield of x */
kmemcheck_annotate_bitfield(x, flags);
```

Note that `kmemcheck_annotate_bitfield()` can be used even before the return value of `kmalloc()` is checked — in other words, passing NULL as the first argument is legal (and will do nothing).

4. Reporting errors

As we have seen, `kmemcheck` will produce false positive reports. Therefore, it is not very wise to blindly post `kmemcheck` warnings to mailing lists and maintainers. Instead, I encourage maintainers and developers to find errors in their own code. If you get a warning, you can try to work around it, try to figure out if it's a real error or not, or simply ignore it. Most developers know their own code and will quickly and efficiently determine the root cause of a `kmemcheck` report. This is therefore also the most efficient way to work with `kmemcheck`.

That said, we (the `kmemcheck` maintainers) will always be on the lookout for false positives that we can annotate and silence. So whatever you find, please drop us a note privately! Kernel configs and steps to reproduce (if available) are of course a great help too.

Happy hacking!

5. Technical description

`kmemcheck` works by marking memory pages non-present. This means that whenever somebody attempts to access the page, a page fault is generated. The page fault handler notices that the page was in fact only hidden, and so it calls on the `kmemcheck` code to make further investigations.

When the investigations are completed, kmemcheck "shows" the page by marking it present (as it would be under normal circumstances). This way, the interrupted code can continue as usual.

But after the instruction has been executed, we should hide the page again, so that we can catch the next access too! Now kmemcheck makes use of a debugging feature of the processor, namely single-stepping. When the processor has finished the one instruction that generated the memory access, a debug exception is raised. From here, we simply hide the page again and continue execution, this time with the single-stepping feature turned off.

kmemcheck requires some assistance from the memory allocator in order to work. The memory allocator needs to

1. Tell kmemcheck about newly allocated pages and pages that are about to be freed. This allows kmemcheck to set up and tear down the shadow memory for the pages in question. The shadow memory stores the status of each byte in the allocation proper, e.g. whether it is initialized or uninitialized.
2. Tell kmemcheck which parts of memory should be marked uninitialized. There are actually a few more states, such as "not yet allocated" and "recently freed".

If a slab cache is set up using the SLAB_NOTRACK flag, it will never return memory that can take page faults because of kmemcheck.

If a slab cache is NOT set up using the SLAB_NOTRACK flag, callers can still request memory with the `__GFP_NOTRACK` or `__GFP_NOTRACK_FALSE_POSITIVE` flags. This does not prevent the page faults from occurring, however, but marks the object in question as being initialized so that no warnings will ever be produced for this object.

Currently, the SLAB and SLUB allocators are supported by kmemcheck.