

can.txt

Readme file for the Controller Area Network Protocol Family (aka Socket CAN)

This file contains

- 1 Overview / What is Socket CAN
- 2 Motivation / Why using the socket API
- 3 Socket CAN concept
  - 3.1 receive lists
  - 3.2 local loopback of sent frames
  - 3.3 network security issues (capabilities)
  - 3.4 network problem notifications
- 4 How to use Socket CAN
  - 4.1 RAW protocol sockets with can\_filters (SOCK\_RAW)
    - 4.1.1 RAW socket option CAN\_RAW\_FILTER
    - 4.1.2 RAW socket option CAN\_RAW\_ERR\_FILTER
    - 4.1.3 RAW socket option CAN\_RAW\_LOOPBACK
    - 4.1.4 RAW socket option CAN\_RAW\_RECV\_OWN\_MSGS
  - 4.2 Broadcast Manager protocol sockets (SOCK\_DGRAM)
  - 4.3 connected transport protocols (SOCK\_SEQPACKET)
  - 4.4 unconnected transport protocols (SOCK\_DGRAM)
- 5 Socket CAN core module
  - 5.1 can.ko module params
  - 5.2 procfs content
  - 5.3 writing own CAN protocol modules
- 6 CAN network drivers
  - 6.1 general settings
  - 6.2 local loopback of sent frames
  - 6.3 CAN controller hardware filters
  - 6.4 The virtual CAN driver (vcan)
  - 6.5 The CAN network device driver interface
    - 6.5.1 Netlink interface to set/get devices properties
    - 6.5.2 Setting the CAN bit-timing
    - 6.5.3 Starting and stopping the CAN network device
  - 6.6 supported CAN hardware
- 7 Socket CAN resources
- 8 Credits

---

## 1. Overview / What is Socket CAN

---

The socketcan package is an implementation of CAN protocols (Controller Area Network) for Linux. CAN is a networking technology which has widespread use in automation, embedded devices, and

automotive fields. While there have been other CAN implementations for Linux based on character devices, Socket CAN uses the Berkeley socket API, the Linux network stack and implements the CAN device drivers as network interfaces. The CAN socket API has been designed as similar as possible to the TCP/IP protocols to allow programmers, familiar with network programming, to easily learn how to use CAN sockets.

## 2. Motivation / Why using the socket API

---

There have been CAN implementations for Linux before Socket CAN so the question arises, why we have started another project. Most existing implementations come as a device driver for some CAN hardware, they are based on character devices and provide comparatively little functionality. Usually, there is only a hardware-specific device driver which provides a character device interface to send and receive raw CAN frames, directly to/from the controller hardware. Queueing of frames and higher-level transport protocols like ISO-TP have to be implemented in user space applications. Also, most character-device implementations support only one single process to open the device at a time, similar to a serial interface. Exchanging the CAN controller requires employment of another device driver and often the need for adaption of large parts of the application to the new driver's API.

Socket CAN was designed to overcome all of these limitations. A new protocol family has been implemented which provides a socket interface to user space applications and which builds upon the Linux network layer, so to use all of the provided queueing functionality. A device driver for CAN controller hardware registers itself with the Linux network layer as a network device, so that CAN frames from the controller can be passed up to the network layer and on to the CAN protocol family module and also vice-versa. Also, the protocol family module provides an API for transport protocol modules to register, so that any number of transport protocols can be loaded or unloaded dynamically. In fact, the can core module alone does not provide any protocol and cannot be used without loading at least one additional protocol module. Multiple sockets can be opened at the same time, on different or the same protocol module and they can listen/send frames on different or the same CAN IDs. Several sockets listening on the same interface for frames with the same CAN ID are all passed the same received matching CAN frames. An application wishing to communicate using a specific transport protocol, e.g. ISO-TP, just selects that protocol when opening the socket, and then can read and write application data byte streams, without having to deal with CAN-IDs, frames, etc.

Similar functionality visible from user-space could be provided by a character device, too, but this would lead to a technically inelegant solution for a couple of reasons:

- \* Intricate usage. Instead of passing a protocol argument to `socket(2)` and using `bind(2)` to select a CAN interface and CAN ID, an application would have to do all these operations using `ioctl(2)`s.

- \* Code duplication. A character device cannot make use of the Linux network queueing code, so all that code would have to be duplicated for CAN networking.
- \* Abstraction. In most existing character-device implementations, the hardware-specific device driver for a CAN controller directly provides the character device for the application to work with. This is at least very unusual in Unix systems for both, char and block devices. For example you don't have a character device for a certain UART of a serial interface, a certain sound chip in your computer, a SCSI or IDE controller providing access to your hard disk or tape streamer device. Instead, you have abstraction layers which provide a unified character or block device interface to the application on the one hand, and a interface for hardware-specific device drivers on the other hand. These abstractions are provided by subsystems like the tty layer, the audio subsystem or the SCSI and IDE subsystems for the devices mentioned above.

The easiest way to implement a CAN device driver is as a character device without such a (complete) abstraction layer, as is done by most existing drivers. The right way, however, would be to add such a layer with all the functionality like registering for certain CAN IDs, supporting several open file descriptors and (de)multiplexing CAN frames between them, (sophisticated) queueing of CAN frames, and providing an API for device drivers to register with. However, then it would be no more difficult, or may be even easier, to use the networking framework provided by the Linux kernel, and this is what Socket CAN does.

The use of the networking framework of the Linux kernel is just the natural and most appropriate way to implement CAN for Linux.

### 3. Socket CAN concept

---

As described in chapter 2 it is the main goal of Socket CAN to provide a socket interface to user space applications which builds upon the Linux network layer. In contrast to the commonly known TCP/IP and ethernet networking, the CAN bus is a broadcast-only(!) medium that has no MAC-layer addressing like ethernet. The CAN-identifier (`can_id`) is used for arbitration on the CAN-bus. Therefore the CAN-IDs have to be chosen uniquely on the bus. When designing a CAN-ECU network the CAN-IDs are mapped to be sent by a specific ECU. For this reason a CAN-ID can be treated best as a kind of source address.

#### 3.1 receive lists


The network transparent access of multiple applications leads to the problem that different applications may be interested in the same CAN-IDs from the same CAN network interface. The Socket CAN core module - which implements the protocol family CAN - provides several high efficient receive lists for this reason. If e.g. a user space application opens a CAN RAW socket, the raw protocol module itself requests the (range of) CAN-IDs from the Socket CAN core that are requested by the user. The subscription and unsubscription of CAN-IDs can be done for specific CAN interfaces or for all(!) known

can.txt

CAN interfaces with the `can_rx_(un)register()` functions provided to CAN protocol modules by the SocketCAN core (see chapter 5). To optimize the CPU usage at runtime the receive lists are split up into several specific lists per device that match the requested filter complexity for a given use-case.

### 3.2 local loopback of sent frames

As known from other networking concepts the data exchanging applications may run on the same or different nodes without any change (except for the according addressing information):



To ensure that application A receives the same information in the example (2) as it would receive in example (1) there is need for some kind of local loopback of the sent CAN frames on the appropriate node.

The Linux network devices (by default) just can handle the transmission and reception of media dependent frames. Due to the arbitration on the CAN bus the transmission of a low prio CAN-ID may be delayed by the reception of a high prio CAN frame. To reflect the correct\* traffic on the node the loopback of the sent data has to be performed right after a successful transmission. If the CAN network interface is not capable of performing the loopback for some reason the SocketCAN core can do this task as a fallback solution. See chapter 6.2 for details (recommended).

The loopback functionality is enabled by default to reflect standard networking behaviour for CAN applications. Due to some requests from the RT-SocketCAN group the loopback optionally may be disabled for each separate socket. See `sockopts` from the CAN RAW sockets in chapter 4.1.

\* = you really like to have this when you're running analyser tools like 'candump' or 'cansniffer' on the (same) node.

### 3.3 network security issues (capabilities)

The Controller Area Network is a local field bus transmitting only broadcast messages without any routing and security concepts. In the majority of cases the user application has to deal with raw CAN frames. Therefore it might be reasonable NOT to restrict the CAN access only to the user root, as known from other networks. Since the currently implemented CAN\_RAW and CAN\_BCM sockets can only send and receive frames to/from CAN interfaces it does not affect security of others networks to allow all users to access the CAN. To enable non-root users to access CAN\_RAW and CAN\_BCM protocol sockets the Kconfig options CAN\_RAW\_USER and/or CAN\_BCM\_USER may be selected at kernel compile time.

### 3.4 network problem notifications

The use of the CAN bus may lead to several problems on the physical and media access control layer. Detecting and logging of these lower layer problems is a vital requirement for CAN users to identify hardware issues on the physical transceiver layer as well as arbitration problems and error frames caused by the different ECUs. The occurrence of detected errors are important for diagnosis and have to be logged together with the exact timestamp. For this reason the CAN interface driver can generate so called Error Frames that can optionally be passed to the user application in the same way as other CAN frames. Whenever an error on the physical layer or the MAC layer is detected (e.g. by the CAN controller) the driver creates an appropriate error frame. Error frames can be requested by the user application using the common CAN filter mechanisms. Inside this filter definition the (interested) type of errors may be selected. The reception of error frames is disabled by default. The format of the CAN error frame is briefly described in the Linux header file "include/linux/can/error.h".

### 4. How to use Socket CAN

---

Like TCP/IP, you first need to open a socket for communicating over a CAN network. Since Socket CAN implements a new protocol family, you need to pass PF\_CAN as the first argument to the socket(2) system call. Currently, there are two CAN protocols to choose from, the raw socket protocol and the broadcast manager (BCM). So to open a socket, you would write

```
s = socket(PF_CAN, SOCK_RAW, CAN_RAW);
```

and

```
s = socket(PF_CAN, SOCK_DGRAM, CAN_BCM);
```

respectively. After the successful creation of the socket, you would normally use the bind(2) system call to bind the socket to a CAN interface (which is different from TCP/IP due to different addressing – see chapter 3). After binding (CAN\_RAW) or connecting (CAN\_BCM) the socket, you can read(2) and write(2) from/to the socket or use send(2), sendto(2), sendmsg(2) and the recv\* counterpart operations on the socket as usual. There are also CAN specific socket options described below.

The basic CAN frame structure and the sockaddr structure are defined in include/linux/can.h:

```
struct can_frame {
    canid_t can_id; /* 32 bit CAN_ID + EFF/RTR/ERR flags */
    __u8 can_dlc; /* data length code: 0 .. 8 */
    __u8 data[8] __attribute__((aligned(8)));
};
```

The alignment of the (linear) payload data[] to a 64bit boundary allows the user to define own structs and unions to easily access the

can.txt

CAN payload. There is no given byteorder on the CAN bus by default. A read(2) system call on a CAN\_RAW socket transfers a struct can\_frame to the user space.

The sockaddr\_can structure has an interface index like the PF\_PACKET socket, that also binds to a specific interface:

```
struct sockaddr_can {
    sa_family_t can_family;
    int         can_ifindex;
    union {
        /* transport protocol class address info (e.g. ISOTP) */
        struct { canid_t rx_id, tx_id; } tp;

        /* reserved for future CAN protocols address information */
    } can_addr;
};
```

To determine the interface index an appropriate ioctl() has to be used (example for CAN\_RAW sockets without error checking):

```
int s;
struct sockaddr_can addr;
struct ifreq ifr;

s = socket(PF_CAN, SOCK_RAW, CAN_RAW);

strcpy(ifr.ifr_name, "can0");
ioctl(s, SIOCGIFINDEX, &ifr);

addr.can_family = AF_CAN;
addr.can_ifindex = ifr.ifr_ifindex;

bind(s, (struct sockaddr *)&addr, sizeof(addr));

(..)
```

To bind a socket to all(!) CAN interfaces the interface index must be 0 (zero). In this case the socket receives CAN frames from every enabled CAN interface. To determine the originating CAN interface the system call recvfrom(2) may be used instead of read(2). To send on a socket that is bound to 'any' interface sendto(2) is needed to specify the outgoing interface.

Reading CAN frames from a bound CAN\_RAW socket (see above) consists of reading a struct can\_frame:

```
struct can_frame frame;

nbytes = read(s, &frame, sizeof(struct can_frame));

if (nbytes < 0) {
    perror("can raw socket read");
    return 1;
}
```

can.txt

```
/* paranoid check ... */
if (nbytes < sizeof(struct can_frame)) {
    fprintf(stderr, "read: incomplete CAN frame\n");
    return 1;
}
```

```
/* do something with the received CAN frame */
```

Writing CAN frames can be done similarly, with the write(2) system call:

```
nbytes = write(s, &frame, sizeof(struct can_frame));
```

When the CAN interface is bound to 'any' existing CAN interface (addr.can\_ifindex = 0) it is recommended to use recvfrom(2) if the information about the originating CAN interface is needed:

```
struct sockaddr_can addr;
struct ifreq ifr;
socklen_t len = sizeof(addr);
struct can_frame frame;

nbytes = recvfrom(s, &frame, sizeof(struct can_frame),
                 0, (struct sockaddr*)&addr, &len);

/* get interface name of the received CAN frame */
ifr.ifr_ifindex = addr.can_ifindex;
ioctl(s, SIOCGIFNAME, &ifr);
printf("Received a CAN frame from interface %s", ifr.ifr_name);
```

To write CAN frames on sockets bound to 'any' CAN interface the outgoing interface has to be defined certainly.

```
strcpy(ifr.ifr_name, "can0");
ioctl(s, SIOCGIFINDEX, &ifr);
addr.can_ifindex = ifr.ifr_ifindex;
addr.can_family = AF_CAN;

nbytes = sendto(s, &frame, sizeof(struct can_frame),
               0, (struct sockaddr*)&addr, sizeof(addr));
```

#### 4.1 RAW protocol sockets with can\_filters (SOCK\_RAW)

Using CAN\_RAW sockets is extensively comparable to the commonly known access to CAN character devices. To meet the new possibilities provided by the multi user SocketCAN approach, some reasonable defaults are set at RAW socket binding time:

- The filters are set to exactly one filter receiving everything
- The socket only receives valid data frames (=> no error frames)
- The loopback of sent CAN frames is enabled (see chapter 3.2)
- The socket does not receive its own sent frames (in loopback mode)

These default settings may be changed before or after binding the socket. To use the referenced definitions of the socket options for CAN\_RAW sockets, include <linux/can/raw.h>.

#### 4.1.1 RAW socket option CAN\_RAW\_FILTER

The reception of CAN frames using CAN\_RAW sockets can be controlled by defining 0 .. n filters with the CAN\_RAW\_FILTER socket option.

The CAN filter structure is defined in include/linux/can.h:

```
struct can_filter {
    canid_t can_id;
    canid_t can_mask;
};
```

A filter matches, when

```
<received_can_id> & mask == can_id & mask
```

which is analogous to known CAN controllers hardware filter semantics. The filter can be inverted in this semantic, when the CAN\_INV\_FILTER bit is set in can\_id element of the can\_filter structure. In contrast to CAN controller hardware filters the user may set 0 .. n receive filters for each open socket separately:

```
struct can_filter rfilter[2];

rfilter[0].can_id    = 0x123;
rfilter[0].can_mask  = CAN_SFF_MASK;
rfilter[1].can_id    = 0x200;
rfilter[1].can_mask  = 0x700;

setsockopt(s, SOL_CAN_RAW, CAN_RAW_FILTER, &rfilter, sizeof(rfilter));
```

To disable the reception of CAN frames on the selected CAN\_RAW socket:

```
setsockopt(s, SOL_CAN_RAW, CAN_RAW_FILTER, NULL, 0);
```

To set the filters to zero filters is quite obsolete as not read data causes the raw socket to discard the received CAN frames. But having this 'send only' use-case we may remove the receive list in the Kernel to save a little (really a very little!) CPU usage.

#### 4.1.2 RAW socket option CAN\_RAW\_ERR\_FILTER

As described in chapter 3.4 the CAN interface driver can generate so called Error Frames that can optionally be passed to the user application in the same way as other CAN frames. The possible errors are divided into different error classes that may be filtered using the appropriate error mask. To register for every possible error condition CAN\_ERR\_MASK can be used as value for the error mask. The values for the error mask are defined in linux/can/error.h .

```
can_err_mask_t err_mask = ( CAN_ERR_TX_TIMEOUT | CAN_ERR_BUSOFF );

setsockopt(s, SOL_CAN_RAW, CAN_RAW_ERR_FILTER,
    &err_mask, sizeof(err_mask));
```

#### 4.1.3 RAW socket option CAN\_RAW\_LOOPBACK



To meet multi user needs the local loopback is enabled by default (see chapter 3.2 for details). But in some embedded use-cases (e.g. when only one application uses the CAN bus) this loopback functionality can be disabled (separately for each socket):

```
int loopback = 0; /* 0 = disabled, 1 = enabled (default) */

setsockopt(s, SOL_CAN_RAW, CAN_RAW_LOOPBACK, &loopback, sizeof(loopback));
```

#### 4.1.4 RAW socket option CAN\_RAW\_RECV\_OWN\_MSGS

When the local loopback is enabled, all the sent CAN frames are looped back to the open CAN sockets that registered for the CAN frames' CAN-ID on this given interface to meet the multi user needs. The reception of the CAN frames on the same socket that was sending the CAN frame is assumed to be unwanted and therefore disabled by default. This default behaviour may be changed on demand:

```
int recv_own_msgs = 1; /* 0 = disabled (default), 1 = enabled */

setsockopt(s, SOL_CAN_RAW, CAN_RAW_RECV_OWN_MSGS,
           &recv_own_msgs, sizeof(recv_own_msgs));
```

#### 4.2 Broadcast Manager protocol sockets (SOCK\_DGRAM)

#### 4.3 connected transport protocols (SOCK\_SEQPACKET)

#### 4.4 unconnected transport protocols (SOCK\_DGRAM)

### 5. Socket CAN core module

---

The Socket CAN core module implements the protocol family PF\_CAN. CAN protocol modules are loaded by the core module at runtime. The core module provides an interface for CAN protocol modules to subscribe needed CAN IDs (see chapter 3.1).

#### 5.1 can.ko module params

- stats\_timer: To calculate the Socket CAN core statistics (e.g. current/maximum frames per second) this 1 second timer is invoked at can.ko module start time by default. This timer can be disabled by usingstattimer=0 on the module commandline.
- debug: (removed since SocketCAN SVN r546)

#### 5.2 procfs content

As described in chapter 3.1 the Socket CAN core uses several filter lists to deliver received CAN frames to CAN protocol modules. These receive lists, their filters and the count of filter matches can be checked in the appropriate receive list. All entries contain the device and a protocol module identifier:

```
foo@bar:~$ cat /proc/net/can/rcvlist_all
```

can.txt

```
receive list 'rx_all':
(vcan3: no entry)
(vcan2: no entry)
(vcan1: no entry)
device    can_id    can_mask    function    userdata    matches    ident
vcan0      000      000000000    f88e6370    f6c6f400          0    raw
(any: no entry)
```

In this example an application requests any CAN traffic from vcan0.

```
rcvlist_all - list for unfiltered entries (no filter operations)
rcvlist_eff - list for single extended frame (EFF) entries
rcvlist_err - list for error frames masks
rcvlist_fil - list for mask/value filters
rcvlist_inv - list for mask/value filters (inverse semantic)
rcvlist_sff - list for single standard frame (SFF) entries
```

Additional procfs files in /proc/net/can

```
stats      - Socket CAN core statistics (rx/tx frames, match ratios, ...)
reset_stats - manual statistic reset
version     - prints the Socket CAN core version and the ABI version
```

### 5.3 writing own CAN protocol modules

To implement a new protocol in the protocol family PF\_CAN a new protocol has to be defined in include/linux/can.h .  
The prototypes and definitions to use the Socket CAN core can be accessed by including include/linux/can/core.h .  
In addition to functions that register the CAN protocol and the CAN device notifier chain there are functions to subscribe CAN frames received by CAN interfaces and to send CAN frames:

```
can_rx_register - subscribe CAN frames from a specific interface
can_rx_unregister - unsubscribe CAN frames from a specific interface
can_send        - transmit a CAN frame (optional with local loopback)
```

For details see the kernel doc documentation in net/can/af\_can.c or the source code of net/can/raw.c or net/can/bcm.c .

## 6. CAN network drivers

---

Writing a CAN network device driver is much easier than writing a CAN character device driver. Similar to other known network device drivers you mainly have to deal with:

- TX: Put the CAN frame from the socket buffer to the CAN controller.
- RX: Put the CAN frame from the CAN controller to the socket buffer.

See e.g. at Documentation/networking/netdevices.txt . The differences for writing CAN network device driver are described below:

### 6.1 general settings

```

                                can.txt
dev->type  = ARPHRD_CAN; /* the netdevice hardware type */
dev->flags = IFF_NOARP;  /* CAN has no arp */

dev->mtu   = sizeof(struct can_frame);

```

The struct can\_frame is the payload of each socket buffer in the protocol family PF\_CAN.

## 6.2 local loopback of sent frames

As described in chapter 3.2 the CAN network device driver should support a local loopback functionality similar to the local echo e.g. of tty devices. In this case the driver flag IFF\_ECHO has to be set to prevent the PF\_CAN core from locally echoing sent frames (aka loopback) as fallback solution:

```
dev->flags = (IFF_NOARP | IFF_ECHO);
```

## 6.3 CAN controller hardware filters

To reduce the interrupt load on deep embedded systems some CAN controllers support the filtering of CAN IDs or ranges of CAN IDs. These hardware filter capabilities vary from controller to controller and have to be identified as not feasible in a multi-user networking approach. The use of the very controller specific hardware filters could make sense in a very dedicated use-case, as a filter on driver level would affect all users in the multi-user system. The high efficient filter sets inside the PF\_CAN core allow to set different multiple filters for each socket separately. Therefore the use of hardware filters goes to the category 'handmade tuning on deep embedded systems'. The author is running a MPC603e @133MHz with four SJA1000 CAN controllers from 2002 under heavy bus load without any problems ...

## 6.4 The virtual CAN driver (vcan)

Similar to the network loopback devices, vcan offers a virtual local CAN interface. A full qualified address on CAN consists of

- a unique CAN Identifier (CAN ID)
- the CAN bus this CAN ID is transmitted on (e.g. can0)

so in common use cases more than one virtual CAN interface is needed.

The virtual CAN interfaces allow the transmission and reception of CAN frames without real CAN controller hardware. Virtual CAN network devices are usually named 'vcanX', like vcan0 vcan1 vcan2 ... When compiled as a module the virtual CAN driver module is called vcan.ko

Since Linux Kernel version 2.6.24 the vcan driver supports the Kernel netlink interface to create vcan network devices. The creation and removal of vcan network devices can be managed with the ip(8) tool:

- Create a virtual CAN network interface:  
\$ ip link add type vcan

can.txt

- Create a virtual CAN network interface with a specific name 'vcan42':  
\$ ip link add dev vcan42 type vcan
- Remove a (virtual CAN) network interface 'vcan42':  
\$ ip link del vcan42

## 6.5 The CAN network device driver interface

The CAN network device driver interface provides a generic interface to setup, configure and monitor CAN network devices. The user can then configure the CAN device, like setting the bit-timing parameters, via the netlink interface using the program "ip" from the "IPROUTE2" utility suite. The following chapter describes briefly how to use it. Furthermore, the interface uses a common data structure and exports a set of common functions, which all real CAN network device drivers should use. Please have a look to the SJA1000 or MSCAN driver to understand how to use them. The name of the module is can-dev.ko.

### 6.5.1 Netlink interface to set/get devices properties

The CAN device must be configured via netlink interface. The supported netlink message types are defined and briefly described in "include/linux/can/netlink.h". CAN link support for the program "ip" of the IPROUTE2 utility suite is available and it can be used as shown below:

- Setting CAN device properties:

```
$ ip link set can0 type can help
Usage: ip link set DEVICE type can
      [ bitrate BITRATE [ sample-point SAMPLE-POINT] ] |
      [ tq TQ prop-seg PROP_SEG phase-seg1 PHASE-SEG1
        phase-seg2 PHASE-SEG2 [ sjw SJW ] ]

      [ loopback { on | off } ]
      [ listen-only { on | off } ]
      [ triple-sampling { on | off } ]

      [ restart-ms TIME-MS ]
      [ restart ]
```

```
Where: BITRATE      := { 1..1000000 }
       SAMPLE-POINT := { 0.000..0.999 }
       TQ           := { NUMBER }
       PROP-SEG     := { 1..8 }
       PHASE-SEG1   := { 1..8 }
       PHASE-SEG2   := { 1..8 }
       SJW          := { 1..4 }
       RESTART-MS   := { 0 | NUMBER }
```

- Display CAN device details and statistics:

```
$ ip -details -statistics link show can0
2: can0: <NOARP,UP,LOWER_UP,ECHO> mtu 16 qdisc pfifo_fast state UP qlen 10
link/can
can <TRIPLE-SAMPLING> state ERROR-ACTIVE restart-ms 100
```

```

                                can.txt
bitrate 125000 sample_point 0.875
tq 125 prop-seg 6 phase-seg1 7 phase-seg2 2 sjw 1
sjal000: tseg1 1..16 tseg2 1..8 sjw 1..4 brp 1..64 brp-inc 1
clock 8000000
re-started bus-errors arbit-lost error-warn error-pass bus-off
41          17457      0          41          42          41
RX: bytes  packets  errors  dropped  overrun  mcast
140859     17608     17457    0        0        0
TX: bytes  packets  errors  dropped  carrier  collsns
861        112      0        41      0        0

```

More info to the above output:

"<TRIPLE-SAMPLING>"

Shows the list of selected CAN controller modes: LOOPBACK, LISTEN-ONLY, or TRIPLE-SAMPLING.

"state ERROR-ACTIVE"

The current state of the CAN controller: "ERROR-ACTIVE", "ERROR-WARNING", "ERROR-PASSIVE", "BUS-OFF" or "STOPPED"

"restart-ms 100"

Automatic restart delay time. If set to a non-zero value, a restart of the CAN controller will be triggered automatically in case of a bus-off condition after the specified delay time in milliseconds. By default it's off.

"bitrate 125000 sample\_point 0.875"

Shows the real bit-rate in bits/sec and the sample-point in the range 0.000..0.999. If the calculation of bit-timing parameters is enabled in the kernel (CONFIG\_CAN\_CALC\_BITTIMING=y), the bit-timing can be defined by setting the "bitrate" argument. Optionally the "sample-point" can be specified. By default it's 0.000 assuming CIA-recommended sample-points.

"tq 125 prop-seg 6 phase-seg1 7 phase-seg2 2 sjw 1"

Shows the time quanta in ns, propagation segment, phase buffer segment 1 and 2 and the synchronisation jump width in units of tq. They allow to define the CAN bit-timing in a hardware independent format as proposed by the Bosch CAN 2.0 spec (see chapter 8 of <http://www.semiconductors.bosch.de/pdf/can2spec.pdf>).

"sjal000: tseg1 1..16 tseg2 1..8 sjw 1..4 brp 1..64 brp-inc 1  
clock 8000000"

Shows the bit-timing constants of the CAN controller, here the "sjal000". The minimum and maximum values of the time segment 1 and 2, the synchronisation jump width in units of tq, the bitrate pre-scaler and the CAN system clock frequency in Hz. These constants could be used for user-defined (non-standard) bit-timing calculation algorithms in user-space.

"re-started bus-errors arbit-lost error-warn error-pass bus-off"

Shows the number of restarts, bus and arbitration lost errors, and the state changes to the error-warning, error-passive and bus-off state. RX overrun errors are listed in the "overrun" field of the standard network statistics.

### 6.5.2 Setting the CAN bit-timing

The CAN bit-timing parameters can always be defined in a hardware independent format as proposed in the Bosch CAN 2.0 specification specifying the arguments "tq", "prop\_seg", "phase\_seg1", "phase\_seg2" and "sjw":

```
$ ip link set canX type can tq 125 prop-seg 6 \
                             phase-seg1 7 phase-seg2 2 sjw 1
```

If the kernel option `CONFIG_CAN_CALC_BITTIMING` is enabled, CIA recommended CAN bit-timing parameters will be calculated if the bit-rate is specified with the argument "bitrate":

```
$ ip link set canX type can bitrate 125000
```

Note that this works fine for the most common CAN controllers with standard bit-rates but may *\*fail\** for exotic bit-rates or CAN system clock frequencies. Disabling `CONFIG_CAN_CALC_BITTIMING` saves some space and allows user-space tools to solely determine and set the bit-timing parameters. The CAN controller specific bit-timing constants can be used for that purpose. They are listed by the following command:

```
$ ip -details link show can0
...
    sja1000: clock 8000000 tseg1 1..16 tseg2 1..8 sjw 1..4 brp 1..64 brp-inc 1
```

### 6.5.3 Starting and stopping the CAN network device

A CAN network device is started or stopped as usual with the command "ifconfig canX up/down" or "ip link set canX up/down". Be aware that you *\*must\** define proper bit-timing parameters for real CAN devices before you can start it to avoid error-prone default settings:

```
$ ip link set canX up type can bitrate 125000
```

A device may enter the "bus-off" state if too much errors occurred on the CAN bus. Then no more messages are received or sent. An automatic bus-off recovery can be enabled by setting the "restart-ms" to a non-zero value, e.g.:

```
$ ip link set canX type can restart-ms 100
```

Alternatively, the application may realize the "bus-off" condition by monitoring CAN error frames and do a restart when appropriate with the command:

```
$ ip link set canX type can restart
```

Note that a restart will also create a CAN error frame (see also chapter 3.4).

## 6.6 Supported CAN hardware

can.txt

Please check the "Kconfig" file in "drivers/net/can" to get an actual list of the support CAN hardware. On the Socket CAN project website (see chapter 7) there might be further drivers available, also for older kernel versions.

## 7. Socket CAN resources

---

You can find further resources for Socket CAN like user space tools, support for old kernel versions, more drivers, mailing lists, etc. at the BerliOS OSS project website for Socket CAN:

<http://developer.berlios.de/projects/socketcan>

If you have questions, bug fixes, etc., don't hesitate to post them to the Socketcan-Users mailing list. But please search the archives first.

## 8. Credits

---

Oliver Hartkopp (PF\_CAN core, filters, drivers, bcm, SJA1000 driver)  
Urs Thuermann (PF\_CAN core, kernel integration, socket interfaces, raw, vcan)  
Jan Kizka (RT-SocketCAN core, Socket-API reconciliation)  
Wolfgang Grandegger (RT-SocketCAN core & drivers, Raw Socket-API reviews,  
CAN device driver interface, MSCAN driver)  
Robert Schwebel (design reviews, PTXdist integration)  
Marc Kleine-Budde (design reviews, Kernel 2.6 cleanups, drivers)  
Benedikt Spranger (reviews)  
Thomas Gleixner (LKML reviews, coding style, posting hints)  
Andrey Volkov (kernel subtree structure, ioctls, MSCAN driver)  
Matthias Brukner (first SJA1000 CAN netdevice implementation Q2/2003)  
Klaus Hitschler (PEAK driver integration)  
Uwe Koppe (CAN netdevices with PF\_PACKET approach)  
Michael Schulze (driver layer loopback requirement, RT CAN drivers review)  
Pavel Pisa (Bit-timing calculation)  
Sascha Hauer (SJA1000 platform driver)  
Sebastian Haas (SJA1000 EMS PCI driver)  
Markus Plessing (SJA1000 EMS PCI driver)  
Per Dalen (SJA1000 Kvaser PCI driver)  
Sam Ravnborg (reviews, coding style, kbuild help)