

## HOWTO for the linux packet generator

---

Date: 041221

Enable CONFIG\_NET\_PKTGEN to compile and build pktgen.o either in kernel or as module. Module is preferred. insmod pktgen if needed. Once running pktgen creates a thread on each CPU where each thread has affinity to its CPU. Monitoring and controlling is done via /proc. Easiest to select a suitable a sample script and configure.

On a dual CPU:

```
ps aux | grep pkt
root    129  0.3  0.0    0   0 ?        SW    2003 523:20 [pktgen/0]
root    130  0.3  0.0    0   0 ?        SW    2003 509:50 [pktgen/1]
```

For monitoring and control pktgen creates:

```
/proc/net/pktgen/pgctrl
/proc/net/pktgen/kpktgend_X
/proc/net/pktgen/ethX
```

### Viewing threads

---

```
/proc/net/pktgen/kpktgend_0
Name: kpktgend_0  max_before_softirq: 10000
Running:
Stopped: eth1
Result: OK: max_before_softirq=10000
```

Most important the devices assigned to thread. Note! A device can only belong to one thread.

### Viewing devices

---

Parm section holds configured info. Current hold running stats.  
Result is printed after run or after interruption. Example:

```
/proc/net/pktgen/eth1
```

```
Params: count 10000000  min_pkt_size: 60  max_pkt_size: 60
       frags: 0  delay: 0  clone_skb: 1000000  ifname: eth1
       flows: 0  flowlen: 0
       dst_min: 10.10.11.2  dst_max:
       src_min:   src_max:
       src_mac: 00:00:00:00:00:00  dst_mac: 00:04:23:AC:FD:82
       udp_src_min: 9  udp_src_max: 9  udp_dst_min: 9  udp_dst_max: 9
       src_mac_count: 0  dst_mac_count: 0
       Flags:
```

Current:

```

                                pktgen.txt
pkts-sofar: 10000000  errors: 39664
started: 1103053986245187us  stopped: 1103053999346329us  idle: 880401us
seq_num: 10000011  cur_dst_mac_offset: 0  cur_src_mac_offset: 0
cur_saddr: 0x10a0a0a  cur_daddr: 0x20b0a0a
cur_udp_dst: 9  cur_udp_src: 9
flows: 0
Result: OK: 13101142(c12220741+d880401) usec, 10000000 (60byte,0frags)
       763292pps 390Mb/sec (390805504bps) errors: 39664

```

## Configuring threads and devices

=====

This is done via the /proc interface easiest done via pgset in the scripts

### Examples:

```

pgset "clone_skb 1"      sets the number of copies of the same packet
pgset "clone_skb 0"      use single SKB for all transmits
pgset "pkt_size 9014"    sets packet size to 9014
pgset "frags 5"          packet will consist of 5 fragments
pgset "count 200000"     sets number of packets to send, set to zero
                        for continuous sends until explicitly stopped.

pgset "delay 5000"       adds delay to hard_start_xmit(). nanoseconds

pgset "dst 10.0.0.1"     sets IP destination address
                        (BEWARE! This generator is very aggressive!)

pgset "dst_min 10.0.0.1" Same as dst
pgset "dst_max 10.0.0.254" Set the maximum destination IP.
pgset "src_min 10.0.0.1"  Set the minimum (or only) source IP.
pgset "src_max 10.0.0.254" Set the maximum source IP.
pgset "dst6 fec0::1"      IPV6 destination address
pgset "src6 fec0::2"      IPV6 source address
pgset "dstmac 00:00:00:00:00:00" sets MAC destination address
pgset "srcmac 00:00:00:00:00:00" sets MAC source address

pgset "queue_map_min 0"  Sets the min value of tx queue interval
pgset "queue_map_max 7"  Sets the max value of tx queue interval, for multiqueue
devices

                        To select queue 1 of a given device,
                        use queue_map_min=1 and queue_map_max=1

pgset "src_mac_count 1"  Sets the number of MACs we'll range through.
                        The 'minimum' MAC is what you set with srcmac.

pgset "dst_mac_count 1"  Sets the number of MACs we'll range through.
                        The 'minimum' MAC is what you set with dstmac.

pgset "flag [name]"      Set a flag to determine behaviour.  Current flags
are: IPSRC_RND #IP Source is random (between min/max),
      IPDST_RND, UDPSRC_RND,
      UDPDST_RND, MACSRC_RND, MACDST_RND
      MPLS_RND, VID_RND, SVID_RND
      QUEUE_MAP_RND # queue map random
      QUEUE_MAP_CPU # queue map mirrors

```

smp\_processor\_id()

## pktgen.txt

```
pgset "udp_src_min 9"    set UDP source port min, If < udp_src_max, then
                           cycle through the port range.

pgset "udp_src_max 9"    set UDP source port max.
pgset "udp_dst_min 9"    set UDP destination port min, If < udp_dst_max, then
                           cycle through the port range.
pgset "udp_dst_max 9"    set UDP destination port max.

pgset "mpls 0001000a,0002000a,0000000a" set MPLS labels (in this example
                                                outer label=16,middle label=32,
                                                inner label=0 (IPv4 NULL)) Note that
                                                there must be no spaces between the
                                                arguments. Leading zeros are required.
                                                Do not set the bottom of stack bit,
                                                that's done automatically. If you do
                                                set the bottom of stack bit, that
                                                indicates that you want to randomly
                                                generate that address and the flag
                                                MPLS_RND will be turned on. You
                                                can have any mix of random and fixed
                                                labels in the label stack.

pgset "mpls 0"           turn off mpls (or any invalid argument works too!)

pgset "vlan_id 77"       set VLAN ID 0-4095
pgset "vlan_p 3"         set priority bit 0-7 (default 0)
pgset "vlan_cfi 0"       set canonical format identifier 0-1 (default 0)

pgset "svlan_id 22"      set SVLAN ID 0-4095
pgset "svlan_p 3"        set priority bit 0-7 (default 0)
pgset "svlan_cfi 0"      set canonical format identifier 0-1 (default 0)

pgset "vlan_id 9999"     > 4095 remove vlan and svlan tags
pgset "svlan 9999"       > 4095 remove svlan tag

pgset "tos XX"           set former IPv4 TOS field (e.g. "tos 28" for AF11 no
ECN, default 00)
pgset "traffic_class XX" set former IPv6 TRAFFIC CLASS (e.g. "traffic_class B8"
for EF no ECN, default 00)

pgset stop               aborts injection. Also, ^C aborts generator.
```

## Example scripts

=====

A collection of small tutorial scripts for pktgen is in examples dir.

```
pktgen.conf-1-1          # 1 CPU 1 dev
pktgen.conf-1-2          # 1 CPU 2 dev
pktgen.conf-2-1          # 2 CPU's 1 dev
pktgen.conf-2-2          # 2 CPU's 2 dev
pktgen.conf-1-1-rdos     # 1 CPU 1 dev w. route DoS
                        第 3 页
```

	pktgen.txt
pktgen.conf-1-1-ipv6	# 1 CPU 1 dev ipv6
pktgen.conf-1-1-ipv6-rdos	# 1 CPU 1 dev ipv6 w. route DoS
pktgen.conf-1-1-flows	# 1 CPU 1 dev multiple flows.

Run in shell: ./pktgen.conf-X-Y It does all the setup including sending.

## Interrupt affinity

=====

Note when adding devices to a specific CPU there good idea to also assign /proc/irq/XX/smp\_affinity so the TX-interrupts gets bound to the same CPU. as this reduces cache bouncing when freeing skb's.

## Current commands and configuration options

### \*\* Pgcontrol commands:

start  
stop

### \*\* Thread commands:

add\_device  
rem\_device\_all  
max\_before\_softirq

### \*\* Device commands:

count  
clone\_skb  
debug

frags  
delay

src\_mac\_count  
dst\_mac\_count

pkt\_size  
min\_pkt\_size  
max\_pkt\_size

mpls

udp\_src\_min  
udp\_src\_max

udp\_dst\_min  
udp\_dst\_max

flag  
IPSRC\_RND  
TXSIZE\_RND

pktgen.txt

IPDST\_RND  
UDPSRC\_RND  
UDPDST\_RND  
MACSRC\_RND  
MACDST\_RND

dst\_min  
dst\_max

src\_min  
src\_max

dst\_mac  
src\_mac

clear\_counters

dst6  
src6

flows  
flowlen

References:

<ftp://robur.slu.se/pub/Linux/net-development/pktgen-testing/>  
<ftp://robur.slu.se/pub/Linux/net-development/pktgen-testing/examples/>

Paper from Linux-Kongress in Erlangen 2004.

[ftp://robur.slu.se/pub/Linux/net-development/pktgen-testing/pktgen\\_paper.pdf](ftp://robur.slu.se/pub/Linux/net-development/pktgen-testing/pktgen_paper.pdf)

Thanks to:

Grant Grundler for testing on IA-64 and parisc, Harald Welte, Lennert Buytenhek  
Stephen Hemminger, Andi Kleen, Dave Miller and many others.

Good luck with the linux net-development.