Overview ____

This readme tries to provide some background on the hows and whys of RDS, and will hopefully help you find your way around the code.

In addition, please see this email about RDS origins: http://oss.oracle.com/pipermail/rds-devel/2007-November/000228.html

RDS Architecture

RDS provides reliable, ordered datagram delivery by using a single reliable connection between any two nodes in the cluster. This allows applications to use a single socket to talk to any other process in the cluster — so in a cluster with N processes you need N sockets, in contrast to N*N if you use a connection—oriented socket transport like TCP.

RDS is not Infiniband-specific; it was designed to support different transports. The current implementation used to support RDS over TCP as well as IB. Work is in progress to support RDS over iWARP, and using DCE to guarantee no dropped packets on Ethernet, it may be possible to use RDS over UDP in the future.

The high-level semantics of RDS from the application's point of view are

* Addressing

RDS uses IPv4 addresses and 16bit port numbers to identify the end point of a connection. All socket operations that involve passing addresses between kernel and user space generally use a struct sockaddr in.

The fact that IPv4 addresses are used does not mean the underlying transport has to be IP-based. In fact, RDS over IB uses a reliable IB connection; the IP address is used exclusively to locate the remote node's GID (by ARPing for the given IP).

The port space is entirely independent of UDP, TCP or any other protocol.

* Socket interface

RDS sockets work *mostly* as you would expect from a BSD socket. The next section will cover the details. At any rate, all $\rm I/O$ is performed through the standard BSD socket API. Some additions like zerocopy support are implemented through control messages, while other extensions use the getsockopt/setsockopt calls.

Sockets must be bound before you can send or receive data. This is needed because binding also selects a transport and attaches it to the socket. Once bound, the transport assignment does not change. RDS will tolerate IPs moving around (eg in a active-active HA scenario), but only as long as the address doesn't move to a different transport.

* sysctls

RDS supports a number of sysctls in /proc/sys/net/rds

Socket Interface

AF RDS, PF RDS, SOL RDS

These constants haven't been assigned yet, because RDS isn't in mainline yet. Currently, the kernel module assigns some constant and publishes it to user space through two sysctl files /proc/sys/net/rds/pf rds

/proc/sys/net/rds/pt_rds /proc/sys/net/rds/sol rds

fd = socket(PF_RDS, SOCK_SEQPACKET, 0);
 This creates a new, unbound RDS socket.

setsockopt(SOL_SOCKET): send and receive buffer size
RDS honors the send and receive buffer size socket options.
You are not allowed to queue more than SO_SNDSIZE bytes to
a socket. A message is queued when sendmsg is called, and
it leaves the queue when the remote system acknowledges
its arrival.

The SO_RCVSIZE option controls the maximum receive queue length. This is a soft limit rather than a hard limit - RDS will continue to accept and queue incoming messages, even if that takes the queue length over the limit. However, it will also mark the port as "congested" and send a congestion update to the source node. The source node is supposed to throttle any processes sending to this congested port.

bind(fd, &sockaddr in, ...)

This binds the socket to a local IP address and port, and a transport.

sendmsg(fd, ...)

Sends a message to the indicated recipient. The kernel will transparently establish the underlying reliable connection if it isn't up yet.

An attempt to send a message that exceeds SO_SNDSIZE will return with -EMSGSIZE

An attempt to send a message that would take the total number of queued bytes over the $SO_SNDSIZE$ threshold will return EAGAIN.

An attempt to send a message to a destination that is marked as "congested" will return ENOBUFS.

recvmsg(fd, ...)

Receives a message that was queued to this socket. The sockets recv queue accounting is adjusted, and if the queue length drops below SO_SNDSIZE, the port is marked uncongested, and a congestion update is sent to all peers.

第 2 页

Applications can ask the RDS kernel module to receive notifications via control messages (for instance, there is a notification when a congestion update arrived, or when a RDMA operation completes). These notifications are received through the msg.msg_control buffer of struct msghdr. The format of the messages is described in manpages.

pol1(fd)

RDS supports the poll interface to allow the application to implement async I/O.

POLLIN handling is pretty straightforward. When there's an incoming message queued to the socket, or a pending notification, we signal POLLIN.

POLLOUT is a little harder. Since you can essentially send to any destination, RDS will always signal POLLOUT as long as there's room on the send queue (ie the number of bytes queued is less than the sendbuf size).

However, the kernel will refuse to accept messages to a destination marked congested — in this case you will loop forever if you rely on poll to tell you what to do. This isn't a trivial problem, but applications can deal with this — by using congestion notifications, and by checking for ENOBUFS errors returned by sendmsg.

setsockopt(SOL_RDS, RDS_CANCEL_SENT_TO, &sockaddr_in)
 This allows the application to discard all messages queued to a
 specific destination on this particular socket.

This allows the application to cancel outstanding messages if it detects a timeout. For instance, if it tried to send a message, and the remote host is unreachable, RDS will keep trying forever. The application may decide it's not worth it, and cancel the operation. In this case, it would use RDS_CANCEL_SENT_TO to nuke any pending messages.

RDMA for RDS

see rds-rdma(7) manpage (available in rds-tools)

Congestion Notifications

see rds(7) manpage

RDS Protocol

Message header

第 3 页

The message header is a 'struct rds header' (see rds.h): Fields: h sequence: per-packet sequence number piggybacked acknowledgment of last packet received length of data, not including header h sport: source port h dport: destination port CONG_BITMAP - this is a congestion update bitmap ACK REQUIRED - receiver must ack this packet RETRANSMITTED - packet has previously been sent h credit: indicate to other end of connection that it has more credits available (i.e. there is more send room) h padding[4]: unused, for future use h csum: header checksum h exthdr: optional data can be passed here. This is currently used for passing RDMA-related information.

ACK and retransmit handling

One might think that with reliable IB connections you wouldn't need to ack messages that have been received. The problem is that IB hardware generates an ack message before it has DMAed the message into memory. This creates a potential message loss if the HCA is disabled for any reason between when it sends the ack and before the message is DMAed and processed. This is only a potential issue if another HCA is available for fail-over.

Sending an ack immediately would allow the sender to free the sent message from their send queue quickly, but could cause excessive traffic to be used for acks. RDS piggybacks acks on sent data packets. Ack-only packets are reduced by only allowing one to be in flight at a time, and by the sender only asking for acks when its send buffers start to fill up. All retransmissions are also acked.

Flow Control

RDS's IB transport uses a credit-based mechanism to verify that there is space in the peer's receive buffers for more data. This eliminates the need for hardware retries on the connection.

Congestion

Messages waiting in the receive queue on the receiving socket 第 4 页

are accounted against the sockets SO_RCVBUF option value. Only the payload bytes in the message are accounted for. If the number of bytes queued equals or exceeds revbuf then the socket is congested. All sends attempted to this socket's address should return block or return -EWOULDBLOCK.

Applications are expected to be reasonably tuned such that this situation very rarely occurs. An application encountering this "back-pressure" is considered a bug.

This is implemented by having each node maintain bitmaps which indicate which ports on bound addresses are congested. As the bitmap changes it is sent through all the connections which terminate in the local address of the bitmap which changed.

The bitmaps are allocated as connections are brought up. This avoids allocation in the interrupt handling path which queues sages on sockets. The dense bitmaps let transports send the entire bitmap on any bitmap change reasonably efficiently. This is much easier to implement than some finer-grained communication of per-port congestion. The sender does a very inexpensive bit test to test if the port it's about to send to is congested or not.

RDS Transport Layer

As mentioned above, RDS is not IB-specific. Its code is divided into a general RDS layer and a transport layer.

The general layer handles the socket API, congestion handling, loopback, stats, usermem pinning, and the connection state machine.

The transport layer handles the details of the transport. The IB transport, for example, handles all the queue pairs, work requests, CM event handlers, and other Infiniband details.

RDS Kernel Structures

struct rds message

aka possibly "rds_outgoing", the generic RDS layer copies data to be sent and sets header fields as needed, based on the socket API. This is then queued for the individual connection and sent by the connection's transport.

struct rds incoming

a generic struct referring to incoming data that can be handed from the transport to the general code and queued by the general code while the socket is awoken. It is then passed back to the transport code to handle the actual copy-to-user.

struct rds socket

per-socket information struct rds_connection per-connection information

```
struct rds_transport
pointers to transport-specific functions
struct rds_statistics
non-transport-specific statistics
struct rds_cong_map
wraps the raw congestion bitmap, contains rbnode, waitq, etc.
```

Connection management

Connections may be in UP, DOWN, CONNECTING, DISCONNECTING, and ERROR states.

The first time an attempt is made by an RDS socket to send data to a node, a connection is allocated and connected. That connection is then maintained forever — if there are transport errors, the connection will be dropped and re-established.

Dropping a connection while packets are queued will cause queued or partially-sent datagrams to be retransmitted when the connection is re-established.

The send path

```
rds sendmsg()
  struct rds_message built from incoming data
  CMSGs parsed (e.g. RDMA ops)
  transport connection alloced and connected if not already
  rds message placed on send queue
  send worker awoken
rds_send_worker()
  calls rds send xmit() until queue is empty
rds send xmit()
  transmits congestion map if one is pending
  may set ACK REQUIRED
  calls transport to send either non-RDMA or RDMA message
  (RDMA ops never retransmitted)
rds ib xmit()
  allocs work requests from send ring
  adds any new send credits available to peer (h_credits)
  maps the rds_message's sg list
  piggybacks ack
  populates work requests
  post send to connection's queue pair
```

The recv path

```
rds_ib_recv_cq_comp_handler()
  looks at write completions
  unmaps recv buffer from device
  no errors, call rds_ib_process_recv()
  refill recv ring
rds ib process recv()
```

validate header checksum copy header to rds_ib_incoming struct if start of a new datagram add to ibinc's fraglist if competed datagram: update cong map if datagram was cong update call rds_recv_incoming() otherwise note if ack is required rds_recv_incoming() drop duplicate packets respond to pings find the sock associated with this datagram add to sock queue wake up sock do some congestion calculations rds recvmsg copy data into user iovec handle CMSGs return to application