

The existing interfaces for getting network packages time stamped are:

*** SO_TIMESTAMP**

Generate time stamp for each incoming packet using the (not necessarily monotonous!) system time. Result is returned via `recv_msg()` in a control message as `timeval` (usec resolution).

*** SO_TIMESTAMPNS**

Same time stamping mechanism as `SO_TIMESTAMP`, but returns result as `timespec` (nsec resolution).

*** IP_MULTICAST_LOOP + SO_TIMESTAMP[NS]**

Only for multicasts: approximate send time stamp by receiving the looped packet and using its receive time stamp.

The following interface complements the existing ones: receive time stamps can be generated and returned for arbitrary packets and much closer to the point where the packet is really sent. Time stamps can be generated in software (as before) or in hardware (if the hardware has such a feature).

SO_TIMESTAMPING:

Instructs the socket layer which kind of information is wanted. The parameter is an integer with some of the following bits set. Setting other bits is an error and doesn't change the current state.

<code>SOF_TIMESTAMPING_TX_HARDWARE:</code>	try to obtain send time stamp in hardware
<code>SOF_TIMESTAMPING_TX_SOFTWARE:</code>	if <code>SOF_TIMESTAMPING_TX_HARDWARE</code> is off or fails, then do it in software
<code>SOF_TIMESTAMPING_RX_HARDWARE:</code>	return the original, unmodified time stamp as generated by the hardware
<code>SOF_TIMESTAMPING_RX_SOFTWARE:</code>	if <code>SOF_TIMESTAMPING_RX_HARDWARE</code> is off or fails, then do it in software
<code>SOF_TIMESTAMPING_RAW_HARDWARE:</code>	return original raw hardware time stamp
<code>SOF_TIMESTAMPING_SYS_HARDWARE:</code>	return hardware time stamp transformed to the system time base
<code>SOF_TIMESTAMPING_SOFTWARE:</code>	return system time stamp generated in software

`SOF_TIMESTAMPING_TX/RX` determine how time stamps are generated.
`SOF_TIMESTAMPING_RAW/SYS` determine how they are reported in the following control message:

```
struct scm_timestamping {
    struct timespec systime;
    struct timespec hwtimetrans;
    struct timespec hwtimeraw;
};
```

`recvmsg()` can be used to get this control message for regular incoming packets. For send time stamps the outgoing packet is looped back to the socket's error queue with the send time stamp(s) attached. It can be received with `recvmsg(flags=MSG_ERRQUEUE)`. The call returns the original outgoing packet data including all headers prepended down to and including the link layer, the `scm_timestamping` control message and

timestamping.txt

a sock_extended_err control message with ee_errno==ENOMSG and ee_origin==SO_EE_ORIGIN_TIMESTAMPING. A socket with such a pending bounced packet is ready for reading as far as select() is concerned. If the outgoing packet has to be fragmented, then only the first fragment is time stamped and returned to the sending socket.

All three values correspond to the same event in time, but were generated in different ways. Each of these values may be empty (= all zero), in which case no such value was available. If the application is not interested in some of these values, they can be left blank to avoid the potential overhead of calculating them.

sysptime is the value of the system time at that moment. This corresponds to the value also returned via SO_TIMESTAMP[NS]. If the time stamp was generated by hardware, then this field is empty. Otherwise it is filled in if SOF_TIMESTAMPING_SOFTWARE is set.

hwtime_raw is the original hardware time stamp. Filled in if SOF_TIMESTAMPING_RAW_HARDWARE is set. No assumptions about its relation to system time should be made.

hwtime_trans is the hardware time stamp transformed so that it corresponds as good as possible to system time. This correlation is not perfect; as a consequence, sorting packets received via different NICs by their hwtime_trans may differ from the order in which they were received. hwtime_trans may be non-monotonic even for the same NIC. Filled in if SOF_TIMESTAMPING_SYS_HARDWARE is set. Requires support by the network device and will be empty without that support.

SIOCSHWTSTAMP:

Hardware time stamping must also be initialized for each device driver that is expected to do hardware time stamping. The parameter is defined in /include/linux/net_timestamp.h as:

```
struct hwtimestamp_config {
    int flags;          /* no flags defined right now, must be zero */
    int tx_type;        /* HWTSTAMP_TX_* */
    int rx_filter;      /* HWTSTAMP_FILTER_* */
};
```

Desired behavior is passed into the kernel and to a specific device by calling ioctl(SIOCSHWTSTAMP) with a pointer to a struct ifreq whose ifr_data points to a struct hwtimestamp_config. The tx_type and rx_filter are hints to the driver what it is expected to do. If the requested fine-grained filtering for incoming packets is not supported, the driver may time stamp more than just the requested types of packets.

A driver which supports hardware time stamping shall update the struct with the actual, possibly more permissive configuration. If the requested packets cannot be time stamped, then nothing should be changed and ERANGE shall be returned (in contrast to EINVAL, which indicates that SIOCSHWTSTAMP is not supported at all).

timestamping.txt

Only a processes with admin rights may change the configuration. User space is responsible to ensure that multiple processes don't interfere with each other and that the settings are reset.

```
/* possible values for hwtstamp_config->tx_type */
enum {
    /*
     * no outgoing packet will need hardware time stamping;
     * should a packet arrive which asks for it, no hardware
     * time stamping will be done
     */
    HWTSTAMP_TX_OFF,

    /*
     * enables hardware time stamping for outgoing packets;
     * the sender of the packet decides which are to be
     * time stamped by setting SOF_TIMESTAMPING_TX_SOFTWARE
     * before sending the packet
     */
    HWTSTAMP_TX_ON,
};

/* possible values for hwtstamp_config->rx_filter */
enum {
    /* time stamp no incoming packet at all */
    HWTSTAMP_FILTER_NONE,

    /* time stamp any incoming packet */
    HWTSTAMP_FILTER_ALL,

    /* return value: time stamp all packets requested plus some others */
    HWTSTAMP_FILTER_SOME,

    /* PTP v1, UDP, any kind of event packet */
    HWTSTAMP_FILTER_PTP_V1_L4_EVENT,

    /* for the complete list of values, please check
     * the include file /include/linux/net_tstamp.h
     */
};
```

DEVICE IMPLEMENTATION

A driver which supports hardware time stamping must support the SIOCSHWTSTAMP ioctl and update the supplied struct hwtstamp_config with the actual values as described in the section on SIOCSHWTSTAMP.

Time stamps for received packets must be stored in the skb. To get a pointer to the shared time stamp structure of the skb call skb_hwtstamps(). Then set the time stamps in the structure:

```
struct skb_shared_hwtstamps {
    /* hardware time stamp transformed into duration
     * since arbitrary point in time
```

timestamping.txt

```
    */  
    ktime_t hwtstamp;  
    ktime_t syststamp; /* hwtstamp transformed to system time base */  
};
```

Time stamps for outgoing packets are to be generated as follows:

- In `hard_start_xmit()`, check if `skb_tx(skb)->hardware` is set non-zero. If yes, then the driver is expected to do hardware time stamping.
- If this is possible for the `skb` and requested, then declare that the driver is doing the time stamping by setting the field `skb_tx(skb)->in_progress` non-zero. You might want to keep a pointer to the associated `skb` for the next step and not free the `skb`. A driver not supporting hardware time stamping doesn't do that. A driver must never touch `sk_buff::tstamp`! It is used to store software generated time stamps by the network subsystem.
- As soon as the driver has sent the packet and/or obtained a hardware time stamp for it, it passes the time stamp back by calling `skb_hwtstamp_tx()` with the original `skb`, the raw hardware time stamp. `skb_hwtstamp_tx()` clones the original `skb` and adds the timestamps, therefore the original `skb` has to be freed now. If obtaining the hardware time stamp somehow fails, then the driver should not fall back to software time stamping. The rationale is that this would occur at a later time in the processing pipeline than other software time stamping and therefore could lead to unexpected deltas between time stamps.
- If the driver did not call `set_skb_tx(skb)->in_progress`, then `dev_hard_start_xmit()` checks whether software time stamping is wanted as fallback and potentially generates the time stamp.