

The PCI Express Advanced Error Reporting Driver Guide HOWTO

T. Long Nguyen <tom.l.nguyen@intel.com>

Yanmin Zhang <yanmin.zhang@intel.com>

07/29/2006

## 1. Overview

### 1.1 About this guide

This guide describes the basics of the PCI Express Advanced Error Reporting (AER) driver and provides information on how to use it, as well as how to enable the drivers of endpoint devices to conform with PCI Express AER driver.

### 1.2 Copyright (C) Intel Corporation 2006.

### 1.3 What is the PCI Express AER Driver?

PCI Express error signaling can occur on the PCI Express link itself or on behalf of transactions initiated on the link. PCI Express defines two error reporting paradigms: the baseline capability and the Advanced Error Reporting capability. The baseline capability is required of all PCI Express components providing a minimum defined set of error reporting requirements. Advanced Error Reporting capability is implemented with a PCI Express advanced error reporting extended capability structure providing more robust error reporting.

The PCI Express AER driver provides the infrastructure to support PCI Express Advanced Error Reporting capability. The PCI Express AER driver provides three basic functions:

- Gathers the comprehensive error information if errors occurred.
- Reports error to the users.
- Performs error recovery actions.

AER driver only attaches root ports which support PCI-Express AER capability.

## 2. User Guide

### 2.1 Include the PCI Express AER Root Driver into the Linux Kernel

The PCI Express AER Root driver is a Root Port service driver attached to the PCI Express Port Bus driver. If a user wants to use it, the driver has to be compiled. Option CONFIG\_PCIEAER supports this capability. It depends on CONFIG\_PCIEPORTBUS, so pls. set CONFIG\_PCIEPORTBUS=y and CONFIG\_PCIEAER = y.

### 2.2 Load PCI Express AER Root Driver

There is a case where a system has AER support in BIOS. Enabling the AER Root driver and having AER support in BIOS may result unpredictable behavior. To avoid this conflict, a successful load of the AER Root driver requires ACPI \_OSC support in the BIOS to allow the AER Root driver to request for native control of AER. See the PCI FW 3.0 Specification for

pcieaer-howto.txt

details regarding OSC usage. Currently, lots of firmwares don't provide \_OSC support while they use PCI Express. To support such firmwares, `forceload`, a parameter of type bool, could enable AER to continue to be initiated although firmwares have no \_OSC support. To enable the workaround, pls. add `aerdriver.forceload=y` to kernel boot parameter line when booting kernel. Note that `forceload=n` by default.

`nosourceid`, another parameter of type bool, can be used when broken hardware (mostly chipsets) has root ports that cannot obtain the reporting source ID. `nosourceid=n` by default.

### 2.3 AER error output

When a PCI-E AER error is captured, an error message will be outputted to console. If it's a correctable error, it is outputted as a warning. Otherwise, it is printed as an error. So users could choose different log level to filter out correctable error messages.

Below shows an example:

```
0000:50:00.0: PCIe Bus Error: severity=Uncorrected (Fatal), type=Transaction
Layer, id=0500(Requester ID)
0000:50:00.0:   device [8086:0329] error status/mask=00100000/00000000
0000:50:00.0:   [20] Unsupported Request      (First)
0000:50:00.0:   TLP Header: 04000001 00200a03 05010000 00050100
```

In the example, 'Requester ID' means the ID of the device who sends the error message to root port. Pls. refer to pci express specs for other fields.

## 3. Developer Guide

To enable AER aware support requires a software driver to configure the AER capability structure within its device and to provide callbacks.

To support AER better, developers need understand how AER does work firstly.

PCI Express errors are classified into two types: correctable errors and uncorrectable errors. This classification is based on the impacts of those errors, which may result in degraded performance or function failure.

Correctable errors pose no impacts on the functionality of the interface. The PCI Express protocol can recover without any software intervention or any loss of data. These errors are detected and corrected by hardware. Unlike correctable errors, uncorrectable errors impact functionality of the interface. Uncorrectable errors can cause a particular transaction or a particular PCI Express link to be unreliable. Depending on those error conditions, uncorrectable errors are further classified into non-fatal errors and fatal errors. Non-fatal errors cause the particular transaction to be unreliable, but the PCI Express link itself is fully functional. Fatal errors, on the other hand, cause the link to be unreliable.

When AER is enabled, a PCI Express device will automatically send an error message to the PCIe root port above it when the device captures

an error. The Root Port, upon receiving an error reporting message, internally processes and logs the error message in its PCI Express capability structure. Error information being logged includes storing the error reporting agent's requestor ID into the Error Source Identification Registers and setting the error bits of the Root Error Status Register accordingly. If AER error reporting is enabled in Root Error Command Register, the Root Port generates an interrupt if an error is detected.

Note that the errors as described above are related to the PCI Express hierarchy and links. These errors do not include any device specific errors because device specific errors will still get sent directly to the device driver.

### 3.1 Configure the AER capability structure

AER aware drivers of PCI Express component need change the device control registers to enable AER. They also could change AER registers, including mask and severity registers. Helper function `pci_enable_pcie_error_reporting` could be used to enable AER. See section 3.3.

### 3.2. Provide callbacks

#### 3.2.1 callback `reset_link` to reset pci express link

This callback is used to reset the pci express physical link when a fatal error happens. The root port aer service driver provides a default `reset_link` function, but different upstream ports might have different specifications to reset pci express link, so all upstream ports should provide their own `reset_link` functions.

In struct `pcie_port_service_driver`, a new pointer, `reset_link`, is added.

```
pci_ers_result_t (*reset_link) (struct pci_dev *dev);
```

Section 3.2.2.2 provides more detailed info on when to call `reset_link`.

#### 3.2.2 PCI error-recovery callbacks

The PCI Express AER Root driver uses error callbacks to coordinate with downstream device drivers associated with a hierarchy in question when performing error recovery actions.

Data struct `pci_driver` has a pointer, `err_handler`, to point to `pci_error_handlers` who consists of a couple of callback function pointers. AER driver follows the rules defined in `pci-error-recovery.txt` except pci express specific parts (e.g. `reset_link`). Pls. refer to `pci-error-recovery.txt` for detailed definitions of the callbacks.

Below sections specify when to call the error callback functions.

##### 3.2.2.1 Correctable errors

Correctable errors pose no impacts on the functionality of the interface. The PCI Express protocol can recover without any software intervention or any loss of data. These errors do not require any recovery actions. The AER driver clears the device's correctable error status register accordingly and logs these errors.

### 3.2.2.2 Non-correctable (non-fatal and fatal) errors

If an error message indicates a non-fatal error, performing link reset at upstream is not required. The AER driver calls `error_detected(dev, pci_channel_io_normal)` to all drivers associated within a hierarchy in question. for example,

EndPoint<==>DownstreamPort B<==>UpstreamPort A<==>RootPort.

If Upstream port A captures an AER error, the hierarchy consists of Downstream port B and EndPoint.

A driver may return `PCI_ERS_RESULT_CAN_RECOVER`, `PCI_ERS_RESULT_DISCONNECT`, or `PCI_ERS_RESULT_NEED_RESET`, depending on whether it can recover or the AER driver calls `mmio_enabled` as next.

If an error message indicates a fatal error, kernel will broadcast `error_detected(dev, pci_channel_io_frozen)` to all drivers within a hierarchy in question. Then, performing link reset at upstream is necessary. As different kinds of devices might use different approaches to reset link, AER port service driver is required to provide the function to reset link. Firstly, kernel looks for if the upstream component has an aer driver. If it has, kernel uses the `reset_link` callback of the aer driver. If the upstream component has no aer driver and the port is downstream port, we will perform a hot reset as the default by setting the Secondary Bus Reset bit of the Bridge Control register associated with the downstream port. As for upstream ports, they should provide their own aer service drivers with `reset_link` function. If `error_detected` returns `PCI_ERS_RESULT_CAN_RECOVER` and `reset_link` returns `PCI_ERS_RESULT_RECOVERED`, the error handling goes to `mmio_enabled`.

## 3.3 helper functions

3.3.1 `int pci_enable_pcie_error_reporting(struct pci_dev *dev);`  
`pci_enable_pcie_error_reporting` enables the device to send error messages to root port when an error is detected. Note that devices don't enable the error reporting by default, so device drivers need call this function to enable it.

3.3.2 `int pci_disable_pcie_error_reporting(struct pci_dev *dev);`  
`pci_disable_pcie_error_reporting` disables the device to send error messages to root port when an error is detected.

3.3.3 `int pci_cleanup_aer_uncorrect_error_status(struct pci_dev *dev);`  
`pci_cleanup_aer_uncorrect_error_status` cleansups the uncorrectable error status register.

## 3.4 Frequent Asked Questions

Q: What happens if a PCI Express device driver does not provide an

pcieaer-howto.txt

error recovery handler (pci\_driver->err\_handler is equal to NULL)?

A: The devices attached with the driver won't be recovered. If the error is fatal, kernel will print out warning messages. Please refer to section 3 for more information.

Q: What happens if an upstream port service driver does not provide callback reset\_link?

A: Fatal error recovery will fail if the errors are reported by the upstream ports who are attached by the service driver.

Q: How does this infrastructure deal with driver that is not PCI Express aware?

A: This infrastructure calls the error callback functions of the driver when an error happens. But if the driver is not aware of PCI Express, the device might not report its own errors to root port.

Q: What modifications will that driver need to make it compatible with the PCI Express AER Root driver?

A: It could call the helper functions to enable AER in devices and cleanup uncorrectable status register. Pls. refer to section 3.3.

#### 4. Software error injection

Debugging PCIe AER error recovery code is quite difficult because it is hard to trigger real hardware errors. Software based error injection can be used to fake various kinds of PCIe errors.

First you should enable PCIe AER software error injection in kernel configuration, that is, following item should be in your .config.

CONFIG\_PCIEAER\_INJECT=y or CONFIG\_PCIEAER\_INJECT=m

After reboot with new kernel or insert the module, a device file named /dev/aer\_inject should be created.

Then, you need a user space tool named aer-inject, which can be gotten from:

<http://www.kernel.org/pub/linux/utils/pci/aer-inject/>

More information about aer-inject can be found in the document comes with its source code.