

## Run-time Power Management Framework for I/O Devices

(C) 2009 Rafael J. Wysocki <rjw@sisk.pl>, Novell Inc.

## 1. Introduction

Support for run-time power management (run-time PM) of I/O devices is provided at the power management core (PM core) level by means of:

- \* The power management workqueue `pm_wq` in which bus types and device drivers can put their PM-related work items. It is strongly recommended that `pm_wq` be used for queuing all work items related to run-time PM, because this allows them to be synchronized with system-wide power transitions (suspend to RAM, hibernation and resume from system sleep states). `pm_wq` is declared in `include/linux/pm_runtime.h` and defined in `kernel/power/main.c`.
- \* A number of run-time PM fields in the 'power' member of 'struct device' (which is of the type 'struct dev\_pm\_info', defined in `include/linux/pm.h`) that can be used for synchronizing run-time PM operations with one another.
- \* Three device run-time PM callbacks in 'struct dev\_pm\_ops' (defined in `include/linux/pm.h`).
- \* A set of helper functions defined in `drivers/base/power/runtime.c` that can be used for carrying out run-time PM operations in such a way that the synchronization between them is taken care of by the PM core. Bus types and device drivers are encouraged to use these functions.

The run-time PM callbacks present in 'struct dev\_pm\_ops', the device run-time PM fields of 'struct dev\_pm\_info' and the core helper functions provided for run-time PM are described below.

## 2. Device Run-time PM Callbacks

There are three device run-time PM callbacks defined in 'struct dev\_pm\_ops':

```
struct dev_pm_ops {
    ...
    int (*runtime_suspend)(struct device *dev);
    int (*runtime_resume)(struct device *dev);
    int (*runtime_idle)(struct device *dev);
    ...
};
```

The `->runtime_suspend()`, `->runtime_resume()` and `->runtime_idle()` callbacks are executed by the PM core for either the bus type, or device type (if the bus type's callback is not defined), or device class (if the bus type's and device type's callbacks are not defined) of given device. The bus type, device type and device class callbacks are referred to as subsystem-level callbacks in what follows.

The subsystem-level suspend callback is entirely responsible for handling the suspend of the device as appropriate, which may, but need not include executing the device driver's own `->runtime_suspend()` callback (from the PM core's point of view it is not necessary to implement a `->runtime_suspend()` callback in a device driver as long as the subsystem-level suspend callback

knows what to do to handle the device).

- \* Once the subsystem-level suspend callback has completed successfully for given device, the PM core regards the device as suspended, which need not mean that the device has been put into a low power state. It is supposed to mean, however, that the device will not process data and will not communicate with the CPU(s) and RAM until the subsystem-level resume callback is executed for it. The run-time PM status of a device after successful execution of the subsystem-level suspend callback is 'suspended'.
- \* If the subsystem-level suspend callback returns -EBUSY or -EAGAIN, the device's run-time PM status is 'active', which means that the device must be fully operational afterwards.
- \* If the subsystem-level suspend callback returns an error code different from -EBUSY or -EAGAIN, the PM core regards this as a fatal error and will refuse to run the helper functions described in Section 4 for the device, until the status of it is directly set either to 'active', or to 'suspended' (the PM core provides special helper functions for this purpose).

In particular, if the driver requires remote wake-up capability (i.e. hardware mechanism allowing the device to request a change of its power state, such as PCI PME) for proper functioning and device\_run\_wake() returns 'false' for the device, then ->runtime\_suspend() should return -EBUSY. On the other hand, if device\_run\_wake() returns 'true' for the device and the device is put into a low power state during the execution of the subsystem-level suspend callback, it is expected that remote wake-up will be enabled for the device. Generally, remote wake-up should be enabled for all input devices put into a low power state at run time.

The subsystem-level resume callback is entirely responsible for handling the resume of the device as appropriate, which may, but need not include executing the device driver's own ->runtime\_resume() callback (from the PM core's point of view it is not necessary to implement a ->runtime\_resume() callback in a device driver as long as the subsystem-level resume callback knows what to do to handle the device).

- \* Once the subsystem-level resume callback has completed successfully, the PM core regards the device as fully operational, which means that the device must be able to complete I/O operations as needed. The run-time PM status of the device is then 'active'.
- \* If the subsystem-level resume callback returns an error code, the PM core regards this as a fatal error and will refuse to run the helper functions described in Section 4 for the device, until its status is directly set either to 'active' or to 'suspended' (the PM core provides special helper functions for this purpose).

The subsystem-level idle callback is executed by the PM core whenever the device appears to be idle, which is indicated to the PM core by two counters, the device's usage counter and the counter of 'active' children of the device.

- \* If any of these counters is decreased using a helper function provided by the PM core and it turns out to be equal to zero, the other counter is checked. If that counter also is equal to zero, the PM core executes the subsystem-level idle callback with the device as an argument.

The action performed by a subsystem-level idle callback is totally dependent on the subsystem in question, but the expected and recommended action is to check if the device can be suspended (i.e. if all of the conditions necessary for suspending the device are satisfied) and to queue up a suspend request for the device in that case. The value returned by this callback is ignored by the PM core.

The helper functions provided by the PM core, described in Section 4, guarantee that the following constraints are met with respect to the bus type's run-time PM callbacks:

- (1) The callbacks are mutually exclusive (e.g. it is forbidden to execute `->runtime_suspend()` in parallel with `->runtime_resume()` or with another instance of `->runtime_suspend()` for the same device) with the exception that `->runtime_suspend()` or `->runtime_resume()` can be executed in parallel with `->runtime_idle()` (although `->runtime_idle()` will not be started while any of the other callbacks is being executed for the same device).
- (2) `->runtime_idle()` and `->runtime_suspend()` can only be executed for 'active' devices (i.e. the PM core will only execute `->runtime_idle()` or `->runtime_suspend()` for the devices the run-time PM status of which is 'active').
- (3) `->runtime_idle()` and `->runtime_suspend()` can only be executed for a device the usage counter of which is equal to zero and either the counter of 'active' children of which is equal to zero, or the 'power.ignore\_children' flag of which is set.
- (4) `->runtime_resume()` can only be executed for 'suspended' devices (i.e. the PM core will only execute `->runtime_resume()` for the devices the run-time PM status of which is 'suspended').

Additionally, the helper functions provided by the PM core obey the following rules:

- \* If `->runtime_suspend()` is about to be executed or there's a pending request to execute it, `->runtime_idle()` will not be executed for the same device.
- \* A request to execute or to schedule the execution of `->runtime_suspend()` will cancel any pending requests to execute `->runtime_idle()` for the same device.
- \* If `->runtime_resume()` is about to be executed or there's a pending request to execute it, the other callbacks will not be executed for the same device.
- \* A request to execute `->runtime_resume()` will cancel any pending or scheduled requests to execute the other callbacks for the same device.

### 3. Run-time PM Device Fields

The following device run-time PM fields are present in 'struct dev\_pm\_info', as defined in include/linux/pm.h:

```
struct timer_list suspend_timer;  
- timer used for scheduling (delayed) suspend request
```

## runtime\_pm.txt

unsigned long timer\_expires;  
- timer expiration time, in jiffies (if this is different from zero, the timer is running and will expire at that time, otherwise the timer is not running)

struct work\_struct work;  
- work structure used for queuing up requests (i.e. work items in pm\_wq)

wait\_queue\_head\_t wait\_queue;  
- wait queue used if any of the helper functions needs to wait for another one to complete

spinlock\_t lock;  
- lock used for synchronisation

atomic\_t usage\_count;  
- the usage counter of the device

atomic\_t child\_count;  
- the count of 'active' children of the device

unsigned int ignore\_children;  
- if set, the value of child\_count is ignored (but still updated)

unsigned int disable\_depth;  
- used for disabling the helper functions (they work normally if this is equal to zero); the initial value of it is 1 (i.e. run-time PM is initially disabled for all devices)

unsigned int runtime\_error;  
- if set, there was a fatal error (one of the callbacks returned error code as described in Section 2), so the helper functions will not work until this flag is cleared; this is the error code returned by the failing callback

unsigned int idle\_notification;  
- if set, ->runtime\_idle() is being executed

unsigned int request\_pending;  
- if set, there's a pending request (i.e. a work item queued up into pm\_wq)

enum rpm\_request request;  
- type of request that's pending (valid if request\_pending is set)

unsigned int deferred\_resume;  
- set if ->runtime\_resume() is about to be run while ->runtime\_suspend() is being executed for that device and it is not practical to wait for the suspend to complete; means "start a resume as soon as you've suspended"

unsigned int run\_wake;  
- set if the device is capable of generating run-time wake-up events

enum rpm\_status runtime\_status;  
- the run-time PM status of the device; this field's initial value is RPM\_SUSPENDED, which means that each device is initially regarded by the

runtime\_pm.txt

PM core as 'suspended', regardless of its real hardware status

unsigned int runtime\_auto;

- if set, indicates that the user space has allowed the device driver to power manage the device at run time via the /sys/devices/.../power/control interface; it may only be modified with the help of the pm\_runtime\_allow() and pm\_runtime\_forbid() helper functions

All of the above fields are members of the 'power' member of 'struct device'.

#### 4. Run-time PM Device Helper Functions

The following run-time PM helper functions are defined in drivers/base/power/runtime.c and include/linux/pm\_runtime.h:

void pm\_runtime\_init(struct device \*dev);

- initialize the device run-time PM fields in 'struct dev\_pm\_info'

void pm\_runtime\_remove(struct device \*dev);

- make sure that the run-time PM of the device will be disabled after removing the device from device hierarchy

int pm\_runtime\_idle(struct device \*dev);

- execute the subsystem-level idle callback for the device; returns 0 on success or error code on failure, where -EINPROGRESS means that ->runtime\_idle() is already being executed

int pm\_runtime\_suspend(struct device \*dev);

- execute the subsystem-level suspend callback for the device; returns 0 on success, 1 if the device's run-time PM status was already 'suspended', or error code on failure, where -EAGAIN or -EBUSY means it is safe to attempt to suspend the device again in future

int pm\_runtime\_resume(struct device \*dev);

- execute the subsystem-level resume callback for the device; returns 0 on success, 1 if the device's run-time PM status was already 'active' or error code on failure, where -EAGAIN means it may be safe to attempt to resume the device again in future, but 'power.runtime\_error' should be checked additionally

int pm\_request\_idle(struct device \*dev);

- submit a request to execute the subsystem-level idle callback for the device (the request is represented by a work item in pm\_wq); returns 0 on success or error code if the request has not been queued up

int pm\_schedule\_suspend(struct device \*dev, unsigned int delay);

- schedule the execution of the subsystem-level suspend callback for the device in future, where 'delay' is the time to wait before queuing up a suspend work item in pm\_wq, in milliseconds (if 'delay' is zero, the work item is queued up immediately); returns 0 on success, 1 if the device's PM run-time status was already 'suspended', or error code if the request hasn't been scheduled (or queued up if 'delay' is 0); if the execution of ->runtime\_suspend() is already scheduled and not yet expired, the new value of 'delay' will be used as the time to wait

int pm\_request\_resume(struct device \*dev);

runtime\_pm.txt

- submit a request to execute the subsystem-level resume callback for the device (the request is represented by a work item in pm\_wq); returns 0 on success, 1 if the device's run-time PM status was already 'active', or error code if the request hasn't been queued up

void pm\_runtime\_get\_noresume(struct device \*dev);

- increment the device's usage counter

int pm\_runtime\_get(struct device \*dev);

- increment the device's usage counter, run pm\_request\_resume(dev) and return its result

int pm\_runtime\_get\_sync(struct device \*dev);

- increment the device's usage counter, run pm\_runtime\_resume(dev) and return its result

void pm\_runtime\_put\_noidle(struct device \*dev);

- decrement the device's usage counter

int pm\_runtime\_put(struct device \*dev);

- decrement the device's usage counter, run pm\_request\_idle(dev) and return its result

int pm\_runtime\_put\_sync(struct device \*dev);

- decrement the device's usage counter, run pm\_runtime\_idle(dev) and return its result

void pm\_runtime\_enable(struct device \*dev);

- enable the run-time PM helper functions to run the device bus type's run-time PM callbacks described in Section 2

int pm\_runtime\_disable(struct device \*dev);

- prevent the run-time PM helper functions from running subsystem-level run-time PM callbacks for the device, make sure that all of the pending run-time PM operations on the device are either completed or canceled; returns 1 if there was a resume request pending and it was necessary to execute the subsystem-level resume callback for the device to satisfy that request, otherwise 0 is returned

void pm\_suspend\_ignore\_children(struct device \*dev, bool enable);

- set/unset the power.ignore\_children flag of the device

int pm\_runtime\_set\_active(struct device \*dev);

- clear the device's 'power.runtime\_error' flag, set the device's run-time PM status to 'active' and update its parent's counter of 'active' children as appropriate (it is only valid to use this function if 'power.runtime\_error' is set or 'power.disable\_depth' is greater than zero); it will fail and return error code if the device has a parent which is not active and the 'power.ignore\_children' flag of which is unset

void pm\_runtime\_set\_suspended(struct device \*dev);

- clear the device's 'power.runtime\_error' flag, set the device's run-time PM status to 'suspended' and update its parent's counter of 'active' children as appropriate (it is only valid to use this function if 'power.runtime\_error' is set or 'power.disable\_depth' is greater than zero)

## runtime\_pm.txt

```
bool pm_runtime_suspended(struct device *dev);
- return true if the device's runtime PM status is 'suspended', or false
  otherwise

void pm_runtime_allow(struct device *dev);
- set the power.runtime_auto flag for the device and decrease its usage
  counter (used by the /sys/devices/.../power/control interface to
  effectively allow the device to be power managed at run time)

void pm_runtime_forbid(struct device *dev);
- unset the power.runtime_auto flag for the device and increase its usage
  counter (used by the /sys/devices/.../power/control interface to
  effectively prevent the device from being power managed at run time)
```

It is safe to execute the following helper functions from interrupt context:

```
pm_request_idle()
pm_schedule_suspend()
pm_request_resume()
pm_runtime_get_noresume()
pm_runtime_get()
pm_runtime_put_noidle()
pm_runtime_put()
pm_suspend_ignore_children()
pm_runtime_set_active()
pm_runtime_set_suspended()
pm_runtime_enable()
```

## 5. Run-time PM Initialization, Device Probing and Removal

Initially, the run-time PM is disabled for all devices, which means that the majority of the run-time PM helper functions described in Section 4 will return -EAGAIN until `pm_runtime_enable()` is called for the device.

In addition to that, the initial run-time PM status of all devices is 'suspended', but it need not reflect the actual physical state of the device. Thus, if the device is initially active (i.e. it is able to process I/O), its run-time PM status must be changed to 'active', with the help of `pm_runtime_set_active()`, before `pm_runtime_enable()` is called for the device.

However, if the device has a parent and the parent's run-time PM is enabled, calling `pm_runtime_set_active()` for the device will affect the parent, unless the parent's 'power.ignore\_children' flag is set. Namely, in that case the parent won't be able to suspend at run time, using the PM core's helper functions, as long as the child's status is 'active', even if the child's run-time PM is still disabled (i.e. `pm_runtime_enable()` hasn't been called for the child yet or `pm_runtime_disable()` has been called for it). For this reason, once `pm_runtime_set_active()` has been called for the device, `pm_runtime_enable()` should be called for it too as soon as reasonably possible or its run-time PM status should be changed back to 'suspended' with the help of `pm_runtime_set_suspended()`.

If the default initial run-time PM status of the device (i.e. 'suspended') reflects the actual state of the device, its bus type's or its driver's `->probe()` callback will likely need to wake it up using one of the PM core's

helper functions described in Section 4. In that case, `pm_runtime_resume()` should be used. Of course, for this purpose the device's run-time PM has to be enabled earlier by calling `pm_runtime_enable()`.

If the device bus type's or driver's `->probe()` or `->remove()` callback runs `pm_runtime_suspend()` or `pm_runtime_idle()` or their asynchronous counterparts, they will fail returning `-EAGAIN`, because the device's usage counter is incremented by the core before executing `->probe()` and `->remove()`. Still, it may be desirable to suspend the device as soon as `->probe()` or `->remove()` has finished, so the PM core uses `pm_runtime_idle_sync()` to invoke the subsystem-level idle callback for the device at that time.

The user space can effectively disallow the driver of the device to power manage it at run time by changing the value of its `/sys/devices/.../power/control` attribute to "on", which causes `pm_runtime_forbid()` to be called. In principle, this mechanism may also be used by the driver to effectively turn off the run-time power management of the device until the user space turns it on. Namely, during the initialization the driver can make sure that the run-time PM status of the device is 'active' and call `pm_runtime_forbid()`. It should be noted, however, that if the user space has already intentionally changed the value of `/sys/devices/.../power/control` to "auto" to allow the driver to power manage the device at run time, the driver may confuse it by using `pm_runtime_forbid()` this way.

## 6. Run-time PM and System Sleep

Run-time PM and system sleep (i.e., system suspend and hibernation, also known as suspend-to-RAM and suspend-to-disk) interact with each other in a couple of ways. If a device is active when a system sleep starts, everything is straightforward. But what should happen if the device is already suspended?

The device may have different wake-up settings for run-time PM and system sleep. For example, remote wake-up may be enabled for run-time suspend but disallowed for system sleep (`device_may_wakeup(dev)` returns 'false'). When this happens, the subsystem-level system suspend callback is responsible for changing the device's wake-up setting (it may leave that to the device driver's system suspend routine). It may be necessary to resume the device and suspend it again in order to do so. The same is true if the driver uses different power levels or other settings for run-time suspend and system sleep.

During system resume, devices generally should be brought back to full power, even if they were suspended before the system sleep began. There are several reasons for this, including:

- \* The device might need to switch power levels, wake-up settings, etc.
- \* Remote wake-up events might have been lost by the firmware.
- \* The device's children may need the device to be at full power in order to resume themselves.
- \* The driver's idea of the device state may not agree with the device's physical state. This can happen during resume from hibernation.
- \* The device might need to be reset.



runtime\_pm.txt

\* Even though the device was suspended, if its usage counter was  $> 0$  then most likely it would need a run-time resume in the near future anyway.

\* Always going back to full power is simplest.

If the device was suspended before the sleep began, then its run-time PM status will have to be updated to reflect the actual post-system sleep status. The way to do this is:

```
pm_runtime_disable(dev);
pm_runtime_set_active(dev);
pm_runtime_enable(dev);
```

The PM core always increments the run-time usage counter before calling the `->prepare()` callback and decrements it after calling the `->complete()` callback. Hence disabling run-time PM temporarily like this will not cause any run-time suspend callbacks to be lost.

## 7. Generic subsystem callbacks

Subsystems may wish to conserve code space by using the set of generic power management callbacks provided by the PM core, defined in `driver/base/power/generic_ops.c`:

```
int pm_generic_runtime_idle(struct device *dev);
- invoke the ->runtime_idle() callback provided by the driver of this
  device, if defined, and call pm_runtime_suspend() for this device if the
  return value is 0 or the callback is not defined

int pm_generic_runtime_suspend(struct device *dev);
- invoke the ->runtime_suspend() callback provided by the driver of this
  device and return its result, or return -EINVAL if not defined

int pm_generic_runtime_resume(struct device *dev);
- invoke the ->runtime_resume() callback provided by the driver of this
  device and return its result, or return -EINVAL if not defined

int pm_generic_suspend(struct device *dev);
- if the device has not been suspended at run time, invoke the ->suspend()
  callback provided by its driver and return its result, or return 0 if not
  defined

int pm_generic_resume(struct device *dev);
- invoke the ->resume() callback provided by the driver of this device and,
  if successful, change the device's runtime PM status to 'active'

int pm_generic_freeze(struct device *dev);
- if the device has not been suspended at run time, invoke the ->freeze()
  callback provided by its driver and return its result, or return 0 if not
  defined

int pm_generic_thaw(struct device *dev);
- if the device has not been suspended at run time, invoke the ->thaw()
  callback provided by its driver and return its result, or return 0 if not
  defined
```

runtime\_pm.txt

```
int pm_generic_poweroff(struct device *dev);
```

- if the device has not been suspended at run time, invoke the `->poweroff()` callback provided by its driver and return its result, or return 0 if not defined

```
int pm_generic_restore(struct device *dev);
```

- invoke the `->restore()` callback provided by the driver of this device and, if successful, change the device's runtime PM status to 'active'

These functions can be assigned to the `->runtime_idle()`, `->runtime_suspend()`, `->runtime_resume()`, `->suspend()`, `->resume()`, `->freeze()`, `->thaw()`, `->poweroff()`, or `->restore()` callback pointers in the subsystem-level `dev_pm_ops` structures.

If a subsystem wishes to use all of them at the same time, it can simply assign the `GENERIC_SUBSYS_PM_OPS` macro, defined in `include/linux/pm.h`, to its `dev_pm_ops` structure pointer.

Device drivers that wish to use the same function as a system suspend, freeze, poweroff and run-time suspend callback, and similarly for system resume, thaw, restore, and run-time resume, can achieve this with the help of the `UNIVERSAL_DEV_PM_OPS` macro defined in `include/linux/pm.h` (possibly setting its last argument to `NULL`).