

## Using RCU to Protect Read-Mostly Arrays

Although RCU is more commonly used to protect linked lists, it can also be used to protect arrays. Three situations are as follows:

1. Hash Tables
2. Static Arrays
3. Resizeable Arrays

Each of these situations are discussed below.

## Situation 1: Hash Tables

Hash tables are often implemented as an array, where each array entry has a linked-list hash chain. Each hash chain can be protected by RCU as described in the listRCU.txt document. This approach also applies to other array-of-list situations, such as radix trees.

## Situation 2: Static Arrays

Static arrays, where the data (rather than a pointer to the data) is located in each array element, and where the array is never resized, have not been used with RCU. Rik van Riel recommends using seqlock in this situation, which would also have minimal read-side overhead as long as updates are rare.

Quick Quiz: Why is it so important that updates be rare when using seqlock?

## Situation 3: Resizeable Arrays

Use of RCU for resizeable arrays is demonstrated by the `grow_ary()` function used by the System V IPC code. The array is used to map from semaphore, message-queue, and shared-memory IDs to the data structure that represents the corresponding IPC construct. The `grow_ary()` function does not acquire any locks; instead its caller must hold the `ids->sem` semaphore.

The `grow_ary()` function, shown below, does some limit checks, allocates a new `ipc_id_ary`, copies the old to the new portion of the new, initializes the remainder of the new, updates the `ids->entries` pointer to point to the new array, and invokes `ipc_rcu_putref()` to free up the old array. Note that `rcu_assign_pointer()` is used to update the `ids->entries` pointer, which includes any memory barriers required on whatever architecture you are running on.

```
static int grow_ary(struct ipc_ids* ids, int newsize)
{
    struct ipc_id_ary* new;
    struct ipc_id_ary* old;
```

```

                                arrayRCU.txt

int i;
int size = ids->entries->size;

if(newsize > IPCMNI)
    newsize = IPCMNI;
if(newsize <= size)
    return newsize;

new = ipc_rcu_alloc(sizeof(struct kern_ipc_perm *)*newsize +
                    sizeof(struct ipc_id_ary));
if(new == NULL)
    return size;
new->size = newsize;
memcpy(new->p, ids->entries->p,
        sizeof(struct kern_ipc_perm *)*size +
        sizeof(struct ipc_id_ary));
for(i=size;i<newsize;i++) {
    new->p[i] = NULL;
}
old = ids->entries;

/*
 * Use rcu_assign_pointer() to make sure the memcpied
 * contents of the new array are visible before the new
 * array becomes visible.
 */
rcu_assign_pointer(ids->entries, new);

ipc_rcu_putref(old);
return newsize;
}

```

The `ipc_rcu_putref()` function decrements the array's reference count and then, if the reference count has dropped to zero, uses `call_rcu()` to free the array after a grace period has elapsed.

The array is traversed by the `ipc_lock()` function. This function indexes into the array under the protection of `rcu_read_lock()`, using `rcu_dereference()` to pick up the pointer to the array so that it may later safely be dereferenced — memory barriers are required on the Alpha CPU. Since the size of the array is stored with the array itself, there can be no array-size mismatches, so a simple check suffices. The pointer to the structure corresponding to the desired IPC object is placed in "out", with NULL indicating a non-existent entry. After acquiring "out->lock", the "out->deleted" flag indicates whether the IPC object is in the process of being deleted, and, if not, the pointer is returned.

```

struct kern_ipc_perm* ipc_lock(struct ipc_ids* ids, int id)
{
    struct kern_ipc_perm* out;
    int lid = id % SEQ_MULTIPLIER;
    struct ipc_id_ary* entries;

    rcu_read_lock();
    entries = rcu_dereference(ids->entries);

```

```

                                arrayRCU.txt
if(lid >= entries->size) {
    rcu_read_unlock();
    return NULL;
}
out = entries->p[lid];
if(out == NULL) {
    rcu_read_unlock();
    return NULL;
}
spin_lock(&out->lock);

/* ipc_rmid() may have already freed the ID while ipc_lock
 * was spinning: here verify that the structure is still valid
 */
if (out->deleted) {
    spin_unlock(&out->lock);
    rcu_read_unlock();
    return NULL;
}
return out;
}

```

Answer to Quick Quiz:

The reason that it is important that updates be rare when using seqlock is that frequent updates can livelock readers. One way to avoid this problem is to assign a seqlock for each array entry rather than to the entire array.