

Reference-count design for elements of lists/arrays protected by RCU.

Reference counting on elements of lists which are protected by traditional reader/writer spinlocks or semaphores are straightforward:

```

1.          2.
add()      search_and_reference()
{          {
    alloc_object      read_lock(&list_lock);
    ...              search_for_element
    atomic_set(&el->rc, 1);      atomic_inc(&el->rc);
    write_lock(&list_lock);      ...
    add_element      read_unlock(&list_lock);
    ...              ...
    write_unlock(&list_lock);    }
}

3.          4.
release_referenced()      delete()
{                          {
    ...                  write_lock(&list_lock);
    atomic_dec(&el->rc, refunc);      ...
    ...                  delete_element
    ...                  write_unlock(&list_lock);
    ...                  ...
    ...                  if (atomic_dec_and_test(&el->rc))
    ...                  kfree(el);
    ...                  ...
    ...                  }
}

```

If this list/array is made lock free using RCU as in changing the `write_lock()` in `add()` and `delete()` to `spin_lock()` and changing `read_lock()` in `search_and_reference()` to `rcu_read_lock()`, the `atomic_inc()` in `search_and_reference()` could potentially hold reference to an element which has already been deleted from the list/array. Use `atomic_inc_not_zero()` in this scenario as follows:

```

1.          2.
add()      search_and_reference()
{          {
    alloc_object      rcu_read_lock();
    ...              search_for_element
    atomic_set(&el->rc, 1);      if (!atomic_inc_not_zero(&el->rc)) {
    spin_lock(&list_lock);      rcu_read_unlock();
    ...                  return FAIL;
    add_element      }
    ...              ...
    spin_unlock(&list_lock);      rcu_read_unlock();
}

3.          4.
release_referenced()      delete()
{                          {
    ...                  spin_lock(&list_lock);
    if (atomic_dec_and_test(&el->rc))      ...
        call_rcu(&el->head, el_free);      delete_element
    ...                  spin_unlock(&list_lock);
    ...

```

rcuref.txt

```
}
    ...
    if (atomic_dec_and_test(&el->rc))
        call_rcu(&el->head, el_free);
    ...
}
```

Sometimes, a reference to the element needs to be obtained in the update (write) stream. In such cases, `atomic_inc_not_zero()` might be overkill, since we hold the update-side spinlock. One might instead use `atomic_inc()` in such cases.