

Debugging390.txt

Debugging on Linux for s/390 & z/Architecture by

Denis Joseph Barrow (djbarrow@de.ibm.com, barrow_dj@yahoo.com)
Copyright (C) 2000-2001 IBM Deutschland Entwicklung GmbH, IBM

Corporation

Best viewed with fixed width fonts

Overview of Document:

=====

This document is intended to give a good overview of how to debug Linux for s/390 & z/Architecture. It isn't intended as a complete reference & not a tutorial on the fundamentals of C & assembly. It doesn't go into 390 IO in any detail. It is intended to complement the documents in the reference section below & any other worthwhile references you get.

It is intended like the Enterprise Systems Architecture/390 Reference Summary to be printed out & used as a quick cheat sheet self help style reference when problems occur.

Contents

=====

Register Set
Address Spaces on Intel Linux
Address Spaces on Linux for s/390 & z/Architecture
The Linux for s/390 & z/Architecture Kernel Task Structure
Register Usage & Stackframes on Linux for s/390 & z/Architecture
A sample program with comments
Compiling programs for debugging on Linux for s/390 & z/Architecture
Figuring out gcc compile errors
Debugging Tools
objdump
strace
Performance Debugging
Debugging under VM
s/390 & z/Architecture IO Overview
Debugging IO on s/390 & z/Architecture under VM
GDB on s/390 & z/Architecture
Stack chaining in gdb by hand
Examining core dumps
ldd
Debugging modules
The proc file system
Starting points for debugging scripting languages etc.
Dumptool & Lcrash
SysRq
References
Special Thanks

Register Set

=====

The current architectures have the following registers.

16 General purpose registers, 32 bit on s/390 64 bit on z/Architecture, r0-r15 or gpr0-gpr15 used for arithmetic & addressing.

Debugging390.txt

16 Control registers, 32 bit on s/390 64 bit on z/Architecture, (cr0-cr15 kernel usage only) used for memory management, interrupt control, debugging control etc.

16 Access registers (ar0-ar15) 32 bit on s/390 & z/Architecture not used by normal programs but potentially could be used as temporary storage. Their main purpose is their 1 to 1 association with general purpose registers and are used in the kernel for copying data between kernel & user address spaces. Access register 0 (& access register 1 on z/Architecture (needs 64 bit pointer)) is currently used by the pthread library as a pointer to the current running threads private area.

16 64 bit floating point registers (fp0-fp15) IEEE & HFP floating point format compliant on G5 upwards & a Floating point control reg (FPC)
4 64 bit registers (fp0,fp2,fp4 & fp6) HFP only on older machines.

Note:

Linux (currently) always uses IEEE & emulates G5 IEEE format on older machines, (provided the kernel is configured for this).

The PSW is the most important register on the machine it is 64 bit on s/390 & 128 bit on z/Architecture & serves the roles of a program counter (pc), condition code register, memory space designator. In IBM standard notation I am counting bit 0 as the MSB. It has several advantages over a normal program counter in that you can change address translation & program counter in a single instruction. To change address translation, e.g. switching address translation off requires that you have a logical=physical mapping for the address you are currently running at.

Bit		Value
s/390	z/Architecture	
0	0	Reserved (must be 0) otherwise specification exception occurs.
1	1	Program Event Recording 1 PER enabled, PER is used to facilitate debugging e.g. single stepping.
2-4	2-4	Reserved (must be 0).
5	5	Dynamic address translation 1=DAT on.
6	6	Input/Output interrupt Mask
7	7	External interrupt Mask used primarily for interprocessor signalling & clock interrupts.
8-11	8-11	PSW Key used for complex memory protection mechanism not used under linux
12	12	1 on s/390 0 on z/Architecture
13	13	Machine Check Mask 1=enable machine check interrupts

Debugging390.txt

14	14	Wait State set this to 1 to stop the processor except for interrupts & give time to other LPARS used in CPU idle in the kernel to increase overall usage of processor resources.
15	15	Problem state (if set to 1 certain instructions are disabled) all linux user programs run with this bit 1 (useful info for debugging under VM).
16-17	16-17	Address Space Control
		00 Primary Space Mode when DAT on The linux kernel currently runs in this mode, CR1 is affiliated with this mode & points to the primary segment table origin etc.
		01 Access register mode this mode is used in functions to copy data between kernel & user space.
		10 Secondary space mode not used in linux however CR7 the register affiliated with this mode is & this & normally CR13=CR7 to allow us to copy data between kernel & user space. We do this as follows: We set ar2 to 0 to designate its affiliated gpr (gpr2)to point to primary=kernel space. We set ar4 to 1 to designate its affiliated gpr (gpr4) to point to secondary=home=user space & then essentially do a memcpy(gpr2,gpr4,size) to copy data between the address spaces, the reason we use home space for the kernel & don't keep secondary space free is that code will not run in secondary space.
		11 Home Space Mode all user programs run in this mode. it is affiliated with CR13.
18-19	18-19	Condition codes (CC)
20	20	Fixed point overflow mask if 1=FPU exceptions for this event occur (normally 0)
21	21	Decimal overflow mask if 1=FPU exceptions for this event occur (normally 0)
22	22	Exponent underflow mask if 1=FPU exceptions for this event occur (normally 0)
23	23	Significance Mask if 1=FPU exceptions for this event occur (normally 0)
24-31	24-30	Reserved Must be 0.
	31	Extended Addressing Mode

32 Basic Addressing Mode
Used to set addressing mode
PSW 31 PSW 32
0 0 24 bit
0 1 31 bit
1 1 64 bit

32 1=31 bit addressing mode 0=24 bit addressing mode (for backward compatibility), linux always runs with this bit set to 1

33-64 Instruction address.
33-63 Reserved must be 0
64-127 Address
In 24 bits mode bits 64-103=0 bits 104-127 Address
In 31 bits mode bits 64-96=0 bits 97-127 Address
Note: unlike 31 bit mode on s/390 bit 96 must be zero when loading the address with LPSWE otherwise a specification exception occurs, LPSW is fully backward compatible.

Prefix Page(s)

This per cpu memory area is too intimately tied to the processor not to mention. It exists between the real addresses 0-4096 on s/390 & 0-8192 z/Architecture & is exchanged with a 1 page on s/390 or 2 pages on z/Architecture in absolute storage by the set prefix instruction in linux'es startup.
This page is mapped to a different prefix for each processor in an SMP configuration
(assuming the os designer is sane of course :-)).
Bytes 0-512 (200 hex) on s/390 & 0-512, 4096-4544, 4604-5119 currently on z/Architecture are used by the processor itself for holding such information as exception indications & entry points for exceptions.
Bytes after 0xc00 hex are used by linux for per processor globals on s/390 & z/Architecture
(there is a gap on z/Architecture too currently between 0xc00 & 1000 which linux uses).
The closest thing to this on traditional architectures is the interrupt vector table. This is a good thing & does simplify some of the kernel coding however it means that we now cannot catch stray NULL pointers in the kernel without hard coded checks.

Address Spaces on Intel Linux

=====

The traditional Intel Linux is approximately mapped as follows forgive the ascii art.

0xFFFFFFFF 4GB Himem

* *
* Kernel Space *

Debugging390.txt

```

*
*****
*****
User Space Himem (typically 0xC0000000 3GB ) * User Stack *
*
*
* Shared Libs * * Next
Process *
*
* <== * to
* <== * Run
* User Program *
* Data BSS *
* Text *
* Sections *
*
0x00000000 *****
*****

```

Now it is easy to see that on Intel it is quite easy to recognise a kernel address as being one greater than user space himem (in this case 0xC0000000). & addresses of less than this are the ones in the current running program on this processor (if an smp box).

If using the virtual machine (VM) as a debugger it is quite difficult to know which user process is running as the address space you are looking at could be from any process in the run queue.

The limitation of Intels addressing technique is that the linux kernel uses a very simple real address to virtual addressing technique of Real Address=Virtual Address-User Space Himem.

This means that on Intel the kernel linux can typically only address Himem=0xFFFFFFFF-0xC0000000=1GB & this is all the RAM these machines can typically use.

They can lower User Himem to 2GB or lower & thus be able to use 2GB of RAM however this shrinks the maximum size of User Space from 3GB to 2GB they have a no win limit of 4GB unless they go to 64 Bit.

On 390 our limitations & strengths make us slightly different. For backward compatibility we are only allowed use 31 bits (2GB) of our 32 bit addresses, however, we use entirely separate address spaces for the user & kernel.

This means we can support 2GB of non Extended RAM on s/390, & more with the Extended memory management swap device & currently 4TB of physical memory currently on z/Architecture.

Address Spaces on Linux for s/390 & z/Architecture

Our addressing scheme is as follows

```

Himem 0x7fffffff 2GB on s/390      *****
currently 0x3fffffff (2^42)-1 * User Stack *      *
on z/Architecture.               *****      *
                                * Shared Libs *      *
                                *****
                                *               *      *
                                * User Program *      *
                                *   Data BSS   *      *
                                *    Text     *      *
                                *   Sections  *      *
0x00000000                        *****      *

```

This also means that we need to look at the PSW problem state bit or the addressing mode to decide whether we are looking at user or kernel space.

Virtual Addresses on s/390 & z/Architecture

A virtual address on s/390 is made up of 3 parts

The SX (segment index, roughly corresponding to the PGD & PMD in linux terminology)
being bits 1-11.

The PX (page index, corresponding to the page table entry (pte) in linux terminology)
being bits 12-19.

The remaining bits BX (the byte index are the offset in the page)
i.e. bits 20 to 31.

On z/Architecture in linux we currently make up an address from 4 parts.

The region index bits (RX) 0-32 we currently use bits 22-32

The segment index (SX) being bits 33-43

The page index (PX) being bits 44-51

The byte index (BX) being bits 52-63

Notes:

1) s/390 has no PMD so the PMD is really the PGD also.
A lot of this stuff is defined in pgtable.h.

2) Also seeing as s/390's page indexes are only 1k in size
(bits 12-19 x 4 bytes per pte) we use 1 (page 4k)
to make the best use of memory by updating 4 segment indices
entries each time we mess with a PMD & use offsets
0, 1024, 2048 & 3072 in this page as for our segment indexes.
On z/Architecture our page indexes are now 2k in size
(bits 12-19 x 8 bytes per pte) we do a similar trick
but only mess with 2 segment indices each time we mess with
a PMD.

3) As z/Architecture supports up to a massive 5-level page table lookup we can only use 3 currently on Linux (as this is all the generic kernel currently supports) however this may change in future this allows us to access (according to my sums) 4TB of virtual storage per process i.e.
 $4096 * 512 \text{ (PTES)} * 1024 \text{ (PMDS)} * 2048 \text{ (PGD)} = 4398046511104 \text{ bytes,}$
 enough for another 2 or 3 of years I think :-).
 to do this we use a region-third-table designation type in our address space control registers.

The Linux for s/390 & z/Architecture Kernel Task Structure

Each process/thread under Linux for S390 has its own kernel task_struct defined in linux/include/linux/sched.h

The S390 on initialisation & resuming of a process on a cpu sets the `__LC_KERNEL_STACK` variable in the spare prefix area for this cpu (which we use for per-processor globals).

The kernel stack pointer is intimately tied with the task structure for each processor as follows.

```

                s/390
*****
* 1 page kernel stack *
*      ( 4K )         *
*****
* 1 page task_struct *
*      ( 4K )         *
8K aligned *****
    
```

```

                z/Architecture
*****
* 2 page kernel stack *
*      ( 8K )         *
*****
* 2 page task_struct *
*      ( 8K )         *
16K aligned *****
    
```

What this means is that we don't need to dedicate any register or global variable to point to the current running process & can retrieve it with the following very simple construct for s/390 & one very similar for z/Architecture.

```

static inline struct task_struct * get_current(void)
{
    struct task_struct *current;
    __asm__ ("lhi    %0, -8192\n\t"
            "nr      %0, 15"
            : "=r" (current) );
    return current;
}
    
```

i.e. just anding the current kernel stack pointer with the mask -8192. Thankfully because Linux doesn't have support for nested IO interrupts

& our devices have large buffers can survive interrupts being shut for short amounts of time we don't need a separate stack for interrupts.

Register Usage & Stackframes on Linux for s/390 & z/Architecture

Overview:

This is the code that gcc produces at the top & the bottom of each function. It usually is fairly consistent & similar from function to function & if you know its layout you can probably make some headway in finding the ultimate cause of a problem after a crash without a source level debugger.

Note: To follow stackframes requires a knowledge of C or Pascal & limited knowledge of one assembly language.

It should be noted that there are some differences between the s/390 & z/Architecture stack layouts as the z/Architecture stack layout didn't have to maintain compatibility with older linkage formats.

Glossary:

alloca:

This is a built in compiler function for runtime allocation of extra space on the callers stack which is obviously freed up on function exit (e.g. the caller may choose to allocate nothing of a buffer of 4k if required for temporary purposes), it generates very efficient code (a few cycles) when compared to alternatives like malloc.

automatics: These are local variables on the stack, i.e they aren't in registers & they aren't static.

back-chain:

This is a pointer to the stack pointer before entering a framed functions (see frameless function) prologue got by dereferencing the address of the current stack pointer, i.e. got by accessing the 32 bit value at the stack pointers current location.

base-pointer:

This is a pointer to the back of the literal pool which is an area just behind each procedure used to store constants in each function.

call-clobbered: The caller probably needs to save these registers if there is something of value in them, on the stack or elsewhere before making a call to another procedure so that it can restore it later.

epilogue:

The code generated by the compiler to return to the caller.

frameless-function

A frameless function in Linux for s390 & z/Architecture is one which doesn't need more than the register save area (96 bytes on s/390, 160 on z/Architecture)

given to it by the caller.

A frameless function never:

- 1) Sets up a back chain.
- 2) Calls alloca.
- 3) Calls other normal functions
- 4) Has automatics.

GOT-pointer:

This is a pointer to the global-offset-table in ELF (Executable Linkable Format, Linux's most common executable format), all globals & shared library objects are found using this pointer.

lazy-binding

ELF shared libraries are typically only loaded when routines in the shared library are actually first called at runtime. This is lazy binding.

procedure-linkage-table

This is a table found from the GOT which contains pointers to routines in other shared libraries which can't be called to by easier means.

prologue:

The code generated by the compiler to set up the stack frame.

outgoing-args:

This is extra area allocated on the stack of the calling function if the parameters for the callee's cannot all be put in registers, the same area can be reused by each function the caller calls.

routine-descriptor:

A COFF executable format based concept of a procedure reference actually being 8 bytes or more as opposed to a simple pointer to the routine. This is typically defined as follows
Routine Descriptor offset 0=Pointer to Function
Routine Descriptor offset 4=Pointer to Table of Contents
The table of contents/TOC is roughly equivalent to a GOT pointer.
& it means that shared libraries etc. can be shared between several environments each with their own TOC.

static-chain: This is used in nested functions a concept adopted from pascal by gcc not used in ansi C or C++ (although quite useful), basically it is a pointer used to reference local variables of enclosing functions. You might come across this stuff once or twice in your lifetime.

e. g.

The function below should return 11 though gcc may get upset & toss warnings about unused variables.

```
int FunctionA(int a)
{
    int b;
    FunctionC(int c)
    {
```

```

        b=c+1;
    }
    FunctionC(10);
    return(b);
}

```

s/390 & z/Architecture Register usage

r0	used by syscalls/assembly	call-clobbered
r1	used by syscalls/assembly	call-clobbered
r2	argument 0 / return value 0	call-clobbered
r3	argument 1 / return value 1 (if long long)	call-clobbered
r4	argument 2	call-clobbered
r5	argument 3	call-clobbered
r6	argument 4	saved
r7	pointer-to arguments 5 to ...	saved
r8	this & that	saved
r9	this & that	saved
r10	static-chain (if nested function)	saved
r11	frame-pointer (if function used alloca)	saved
r12	got-pointer	saved
r13	base-pointer	saved
r14	return-address	saved
r15	stack-pointer	saved
f0	argument 0 / return value (float/double)	call-clobbered
f2	argument 1	call-clobbered
f4	z/Architecture argument 2	saved
f6	z/Architecture argument 3	saved

The remaining floating points
f1, f3, f5 f7-f15 are call-clobbered.

Notes:

- 1) The only requirement is that registers which are used by the callee are saved, e.g. the compiler is perfectly capable of using r11 for purposes other than a frame a frame pointer if a frame pointer is not needed.
- 2) In functions with variable arguments e.g. printf the calling procedure is identical to one without variable arguments & the same number of parameters. However, the prologue of this function is somewhat more hairy owing to it having to move these parameters to the stack to get va_start, va_arg & va_end to work.
- 3) Access registers are currently unused by gcc but are used in the kernel. Possibilities exist to use them at the moment for temporary storage but it isn't recommended.
- 4) Only 4 of the floating point registers are used for parameter passing as older machines such as G3 only have only 4 & it keeps the stack frame compatible with other compilers. However with IEEE floating point emulation under linux on the older machines you are free to use the other 12.
- 5) A long long or double parameter cannot be have the first 4 bytes in a register & the second four bytes in the outgoing args area. It must be purely in the outgoing args area if crossing this boundary.

- 6) Floating point parameters are mixed with outgoing args on the outgoing args area in the order they are passed in as parameters.
- 7) Floating point arguments 2 & 3 are saved in the outgoing args area for z/Architecture

Stack Frame Layout

s/390	z/Architecture	
0	0	back chain (a 0 here signifies end of back chain)
4	8	eos (end of stack, not used on Linux for S390 used in other linkage formats)
8	16	glue used in other s/390 linkage formats for saved routine descriptors etc.
12	24	glue used in other s/390 linkage formats for saved routine descriptors etc.
16	32	scratch area
20	40	scratch area
24	48	saved r6 of caller function
28	56	saved r7 of caller function
32	64	saved r8 of caller function
36	72	saved r9 of caller function
40	80	saved r10 of caller function
44	88	saved r11 of caller function
48	96	saved r12 of caller function
52	104	saved r13 of caller function
56	112	saved r14 of caller function
60	120	saved r15 of caller function
64	128	saved f4 of caller function
72	132	saved f6 of caller function
80		undefined
96	160	outgoing args passed from caller to callee
96+x	160+x	possible stack alignment (8 bytes desirable)
96+x+y	160+x+y	alloca space of caller (if used)
96+x+y+z	160+x+y+z	automatics of caller (if used)
0		back-chain

A sample program with comments.

Comments on the function test

- 1) It didn't need to set up a pointer to the constant pool gpr13 as it isn't used (:-).
 - 2) This is a frameless function & no stack is bought.
 - 3) The compiler was clever enough to recognise that it could return the value in r2 as well as use it for the passed in parameter (:-).
 - 4) The basr (branch relative & save) trick works as follows the instruction has a special case with r0,r0 with some instruction operands is understood as the literal value 0, some risc architectures also do this). So now we are branching to the next address & the address new program counter is in r13, so now we subtract the size of the function prologue we have executed + the size of the literal pool to get to the top of the literal pool
- ```
0040037c int test(int b)
{
```

# Function prologue

# Debugging390. txt

```

below
40037c: 90 de f0 34 stm %r13,%r14,52(%r15) # Save registers r13
& r14
400380: 0d d0 basr %r13,%r0 # Set up pointer to
constant pool using
400382: a7 da ff fa ahi %r13,-6 # basr trick
return(5+b);
400386: a7 2a 00 05 ahi %r2,5 # Huge main program
add 5 to r2
Function epilogue
below
40038a: 98 de f0 34 lm %r13,%r14,52(%r15) # restore registers
r13 & 14
40038e: 07 fe br %r14 # return
}

```

## Comments on the function main

1) The compiler did this function optimally ( 8- )

Literal pool for main.

400390: ff ff ff ec .long 0xffffffff

main(int argc, char \*argv[])

```

{
Function prologue
below
400394: 90 bf f0 2c stm %r11,%r15,44(%r15) # Save necessary
registers
400398: 18 0f lr %r0,%r15 # copy stack pointer
to r0
40039a: a7 fa ff a0 ahi %r15,-96 # Make area for
callee saving
40039e: 0d d0 basr %r13,%r0 # Set up r13 to point
to
4003a0: a7 da ff f0 ahi %r13,-16 # literal pool
4003a4: 50 00 f0 00 st %r0,0(%r15) # Save backchain

return(test(5));
4003a8: 58 e0 d0 00 l %r14,0(%r13) # Main Program Below
address of test from
load relative
literal pool
4003ac: a7 28 00 05 lhi %r2,5 # Set first parameter
to 5
4003b0: 4d ee d0 00 bas %r14,0(%r14,%r13) # jump to test
setting r14 as return
address using
branch & save instruction.
Function Epilogue
below
4003b4: 98 bf f0 8c lm %r11,%r15,140(%r15) # Restore necessary
registers.
4003b8: 07 fe br %r14 # return to do
program exit
}

```

## Compiler updates

```

main(int argc, char *argv[])
{
 4004fc: 90 7f f0 1c stm %r7, %r15, 28(%r15)
 400500: a7 d5 00 04 bras %r13, 400508 <main+0xc>
 400504: 00 40 04 f4 .long 0x004004f4
 # compiler now puts constant pool in code to so it saves an instruction
 400508: 18 0f lr %r0, %r15
 40050a: a7 fa ff a0 ahi %r15, -96
 40050e: 50 00 f0 00 st %r0, 0(%r15)
 return(test(5));
 400512: 58 10 d0 00 l %r1, 0(%r13)
 400516: a7 28 00 05 lhi %r2, 5
 40051a: 0d e1 basr %r14, %r1
 # compiler adds 1 extra instruction to epilogue this is done to
 # avoid processor pipeline stalls owing to data dependencies on g5 &
 # above as register 14 in the old code was needed directly after being loaded
 # by the lm %r11, %r15, 140(%r15) for the br %14.
 40051c: 58 40 f0 98 l %r4, 152(%r15)
 400520: 98 7f f0 7c lm %r7, %r15, 124(%r15)
 400524: 07 f4 br %r4
}

```

Hartmut ( our compiler developer ) also has been threatening to take out the stack backchain in optimised code as this also causes pipeline stalls, you have been warned.

## 64 bit z/Architecture code disassembly

If you understand the stuff above you'll understand the stuff below too so I'll avoid repeating myself & just say that some of the instructions have g's on the end of them to indicate they are 64 bit & the stack offsets are a bigger, the only other difference you'll find between 32 & 64 bit is that we now use f4 & f6 for floating point arguments on 64 bit.

00000000800005b0 <test>:

```

int test(int b)
{
 return(5+b);
 800005b0: a7 2a 00 05 ahi %r2, 5
 800005b4: b9 14 00 22 lgfr %r2, %r2 # downcast to integer
 800005b8: 07 fe br %r14
 800005ba: 07 07 bcr 0, %r7

}

```

00000000800005bc <main>:

```

main(int argc, char *argv[])
{
 800005bc: eb bf f0 58 00 24 stmg %r11, %r15, 88(%r15)
}

```

#### Debugging390.txt

```
800005c2: b9 04 00 1f lgr %r1,%r15
800005c6: a7 fb ff 60 aghi %r15,-160
800005ca: e3 10 f0 00 00 24 stg %r1,0(%r15)
 return(test(5));
800005d0: a7 29 00 05 lghi %r2,5
brasl allows jumps > 64k & is overkill here bras would do fine
800005d4: c0 e5 ff ff ff ee brasl %r14,800005b0 <test>
800005da: e3 40 f1 10 00 04 lg %r4,272(%r15)
800005e0: eb bf f0 f8 00 04 lmg %r11,%r15,248(%r15)
800005e6: 07 f4 br %r4
}
```

#### Compiling programs for debugging on Linux for s/390 & z/Architecture

---

-gdwarf-2 now works it should be considered the default debugging format for s/390 & z/Architecture as it is more reliable for debugging shared libraries, normal -g debugging works much better now Thanks to the IBM java compiler developers bug reports.

This is typically done adding/appending the flags -g or -gdwarf-2 to the CFLAGS & LDFLAGS variables Makefile of the program concerned.

If using gdb & you would like accurate displays of registers & stack traces compile without optimisation i.e make sure that there is no -O2 or similar on the CFLAGS line of the Makefile & the emitted gcc commands, obviously this will produce worse code ( not advisable for shipment ) but it is an aid to the debugging process.

This aids debugging because the compiler will copy parameters passed in in registers onto the stack so backtracing & looking at passed in parameters will work, however some larger programs which use inline functions will not compile without optimisation.

Debugging with optimisation has since much improved after fixing some bugs, please make sure you are using gdb-5.0 or later developed after Nov' 2000.

#### Figuring out gcc compile errors

---

If you are getting a lot of syntax errors compiling a program & the problem isn't blatantly obvious from the source.

It often helps to just preprocess the file, this is done with the -E option in gcc.

What this does is that it runs through the very first phase of compilation ( compilation in gcc is done in several stages & gcc calls many programs to achieve its end result ) with the -E option gcc just calls the gcc preprocessor (cpp).

The c preprocessor does the following, it joins all the files #included together recursively ( #include files can #include other files ) & also the c file you wish to compile.

It puts a fully qualified path of the #included files in a comment & it does macro expansion.

This is useful for debugging because

1) You can double check whether the files you expect to be included are the ones

### Debugging390.txt

that are being included ( e.g. double check that you aren't going to the i386 asm directory ).

- 2) Check that macro definitions aren't clashing with typedefs,
- 3) Check that definitions aren't being used before they are being included.
- 4) Helps put the line emitting the error under the microscope if it contains macros.

For convenience the Linux kernel's makefile will do preprocessing automatically for you  
by suffixing the file you want built with .i ( instead of .o )

e.g.

from the linux directory type  
make arch/s390/kernel/signal.i  
this will build

```
s390-gcc -D__KERNEL__ -I/homel/barrow/linux/include -Wall -Wstrict-prototypes
-O2 -fomit-frame-pointer
-fno-strict-aliasing -D__SMP__ -pipe -fno-strength-reduce -E
arch/s390/kernel/signal.c
> arch/s390/kernel/signal.i
```

Now look at signal.i you should see something like.

```
1 "/homel/barrow/linux/include/asm/types.h" 1
typedef unsigned short umode_t;
typedef __signed__ char __s8;
typedef unsigned char __u8;
typedef __signed__ short __s16;
typedef unsigned short __u16;
```

If instead you are getting errors further down e.g.

unknown instruction:2515 "move.l" or better still unknown instruction:2515  
"Fixme not implemented yet, call Martin" you are probably are attempting to  
compile some code

meant for another architecture or code that is simply not implemented, with a  
fixme statement

stuck into the inline assembly code so that the author of the file now knows he  
has work to do.

To look at the assembly emitted by gcc just before it is about to call gas ( the  
gnu assembler )

use the -S option.

Again for your convenience the Linux kernel's Makefile will hold your hand &  
do all this donkey work for you also by building the file with the .s suffix.

e.g.

from the Linux directory type  
make arch/s390/kernel/signal.s

```
s390-gcc -D__KERNEL__ -I/homel/barrow/linux/include -Wall -Wstrict-prototypes
-O2 -fomit-frame-pointer
-fno-strict-aliasing -D__SMP__ -pipe -fno-strength-reduce -S
arch/s390/kernel/signal.c
-o arch/s390/kernel/signal.s
```

### Debugging390.txt

This will output something like, ( please note the constant pool & the useful comments in the prologue to give you a hand at interpreting it ).

```
.LC54: .string "misaligned (__u16 *) in __xchg\n"
.LC57: .string "misaligned (__u32 *) in __xchg\n"
.L$PG1: # Pool sys_sigsuspend
.LC192: .long -262401
.LC193: .long -1
.LC194: .long schedule-.L$PG1
.LC195: .long do_signal-.L$PG1
 .align 4
.globl sys_sigsuspend
 .type sys_sigsuspend,@function
sys_sigsuspend:
leaf function 0
automatics 16
outgoing args 0
need frame pointer 0
call alloca 0
has varargs 0
incoming args (stack) 0
function length 168
 STM 8,15,32(15)
 LR 0,15
 AHI 15,-112
 BASR 13,0
.L$C01: AHI 13,.L$PG1-.L$C01
 ST 0,0(15)
 LR 8,2
 N 5,.LC192-.L$PG1(13)
```

Adding -g to the above output makes the output even more useful  
e.g. typing  
make CC:="s390-gcc -g" kernel/sched.s

which compiles.

```
s390-gcc -g -D__KERNEL__ -I/home/barrow/linux-2.3/include -Wall
-Wstrict-prototypes -O2 -fomit-frame-pointer -fno-strict-aliasing -pipe
-fno-strength-reduce -S kernel/sched.c -o kernel/sched.s
```

also outputs stabs ( debugger ) info, from this info you can find out the offsets & sizes of various elements in structures.  
e.g. the stab for the structure

```
struct rlimit {
 unsigned long rlim_cur;
 unsigned long rlim_max;
};
```

is

```
.stabs "rlimit:T(151,2)=s8rlim_cur:(0,5),0,32;rlim_max:(0,5),32,32;;",128,0,0,0
```



## Debugging390.txt

from this stab you can see that  
rlimit\_cur starts at bit offset 0 & is 32 bits in size  
rlimit\_max starts at bit offset 32 & is 32 bits in size.

### Debugging Tools:

#### objdump

This is a tool with many options the most useful being ( if compiled with -g).  
objdump --source <victim program or object file> > <victims debug listing >

The whole kernel can be compiled like this ( Doing this will make a 17MB kernel  
& a 200 MB listing ) however you have to strip it before building the image  
using the strip command to make it a more reasonable size to boot it.

A source/assembly mixed dump of the kernel can be done with the line  
objdump --source vmlinux > vmlinux.lst  
Also, if the file isn't compiled -g, this will output as much debugging  
information  
as it can (e.g. function names). This is very slow as it spends lots  
of time searching for debugging info. The following self explanatory line should  
be used  
instead if the code isn't compiled -g, as it is much faster:  
objdump --disassemble-all --syms vmlinux > vmlinux.lst

As hard drive space is valuable most of us use the following approach.

- 1) Look at the emitted psw on the console to find the crash address in the kernel.
- 2) Look at the file System.map ( in the linux directory ) produced when building

the kernel to find the closest address less than the current PSW to find the offending function.

- 3) use grep or similar to search the source tree looking for the source file with this function if you don't know where it is.
  - 4) rebuild this object file with -g on, as an example suppose the file was ( /arch/s390/kernel/signal.o )
  - 5) Assuming the file with the erroneous function is signal.c Move to the base of the Linux source tree.
  - 6) rm /arch/s390/kernel/signal.o
  - 7) make /arch/s390/kernel/signal.o
  - 8) watch the gcc command line emitted
  - 9) type it in again or alternatively cut & paste it on the console adding the -g option.
  - 10) objdump --source arch/s390/kernel/signal.o > signal.lst
- This will output the source & the assembly intermixed, as the snippet below shows  
This will unfortunately output addresses which aren't the same  
as the kernel ones you should be able to get around the mental arithmetic  
by playing with the --adjust-vma parameter to objdump.

```
static inline void spin_lock(spinlock_t *lp)
{
 a0: 18 34 lr %r3,%r4
 a2: a7 3a 03 bc ahi %r3,956
 __asm__ __volatile__(" lhi 1,-1\n"
 a6: a7 18 ff ff lhi %r1,-1
 aa: 1f 00 slr %r0,%r0
 ac: ba 01 30 00 cs %r0,%r1,0(%r3)
 b0: a7 44 ff fd jm aa <sys_sigsuspend+0x2e>
 saveset = current->blocked;
 b4: d2 07 f0 68 mvc 104(8,%r15),972(%r4)
 b8: 43 cc
 return (set->sig[0] & mask) != 0;
}
```

6) If debugging under VM go down to that section in the document for more info.

I now have a tool which takes the pain out of --adjust-vma  
& you are able to do something like  
make /arch/s390/kernel/traps.lst  
& it automatically generates the correctly relocated entries for  
the text segment in traps.lst.  
This tool is now standard in linux distro's in scripts/makelst

strace:

Q. What is it ?

A. It is a tool for intercepting calls to the kernel & logging them  
to a file & on the screen.

Q. What use is it ?

A. You can use it to find out what files a particular program opens.

#### Example 1

If you wanted to know does ping work but didn't have the source  
strace ping -c 1 127.0.0.1  
& then look at the man pages for each of the syscalls below,  
( In fact this is sometimes easier than looking at some spaghetti  
source which conditionally compiles for several architectures ).  
Not everything that it throws out needs to make sense immediately.

Just looking quickly you can see that it is making up a RAW socket  
for the ICMP protocol.  
Doing an alarm(10) for a 10 second timeout  
& doing a gettimeofday call before & after each read to see  
how long the replies took, & writing some text to stdout so the user  
has an idea what is going on.

```
socket(PF_INET, SOCK_RAW, IPPROTO_ICMP) = 3
getuid() = 0
setuid(0) = 0
```

### Debugging390.txt

```
stat("/usr/share/locale/C/libc.cat", 0xbffff134) = -1 ENOENT (No such file or
directory)
stat("/usr/share/locale/libc/C", 0xbffff134) = -1 ENOENT (No such file or
directory)
stat("/usr/local/share/locale/C/libc.cat", 0xbffff134) = -1 ENOENT (No such file
or directory)
getpid() = 353
setsockopt(3, SOL_SOCKET, SO_BROADCAST, [1], 4) = 0
setsockopt(3, SOL_SOCKET, SO_RCVBUF, [49152], 4) = 0
fstat(1, {st_mode=S_IFCHR|0620, st_rdev=makedev(3, 1), ...}) = 0
mmap(0, 4096, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_ANONYMOUS, -1, 0) =
0x40008000
ioctl(1, TCGETS, {B9600 opost isig icanon echo ...}) = 0
write(1, "PING 127.0.0.1 (127.0.0.1): 56 d"... , 42PING 127.0.0.1 (127.0.0.1): 56
data bytes
) = 42
sigaction(SIGINT, {0x8049ba0, [], SA_RESTART}, {SIG_DFL}) = 0
sigaction(SIGALRM, {0x8049600, [], SA_RESTART}, {SIG_DFL}) = 0
gettimeofday({948904719, 138951}, NULL) = 0
sendto(3, "\10\0D\201a\1\0\0\17#\2178\307\36"... , 64, 0, {sin_family=AF_INET,
sin_port=htons(0), sin_addr=inet_addr("127.0.0.1")}, 16) = 64
sigaction(SIGALRM, {0x8049600, [], SA_RESTART}, {0x8049600, [], SA_RESTART}) = 0
sigaction(SIGALRM, {0x8049ba0, [], SA_RESTART}, {0x8049600, [], SA_RESTART}) = 0
alarm(10) = 0
recvfrom(3, "E\0\0T\0005\0\0@1|r\177\0\0\1\177"... , 192, 0,
{sin_family=AF_INET, sin_port=htons(50882), sin_addr=inet_addr("127.0.0.1")},
[16]) = 84
gettimeofday({948904719, 160224}, NULL) = 0
recvfrom(3, "E\0\0T\0006\0\0\377\1\275p\177\0"... , 192, 0,
{sin_family=AF_INET, sin_port=htons(50882), sin_addr=inet_addr("127.0.0.1")},
[16]) = 84
gettimeofday({948904719, 166952}, NULL) = 0
write(1, "64 bytes from 127.0.0.1: icmp_se"... ,
5764 bytes from 127.0.0.1: icmp_seq=0 ttl=255 time=28.0 ms
```

### Example 2

```
strace passwd 2>&1 | grep open
produces the following output
open("/etc/ld.so.cache", O_RDONLY) = 3
open("/opt/kde/lib/libc.so.5", O_RDONLY) = -1 ENOENT (No such file or directory)
open("/lib/libc.so.5", O_RDONLY) = 3
open("/dev", O_RDONLY) = 3
open("/var/run/utmp", O_RDONLY) = 3
open("/etc/passwd", O_RDONLY) = 3
open("/etc/shadow", O_RDONLY) = 3
open("/etc/login.defs", O_RDONLY) = 4
open("/dev/tty", O_RDONLY) = 4
```

The 2>&1 is done to redirect stderr to stdout & grep is then filtering this input through the pipe for each line containing the string open.

### Example 3

## Debugging390.txt

Getting sophisticated  
telnetd crashes & I don't know why

### Steps

---

- 1) Replace the following line in /etc/inetd.conf  
telnet stream tcp nowait root /usr/sbin/in.telnetd -h  
with  
telnet stream tcp nowait root /blah
- 2) Create the file /blah with the following contents to start tracing telnetd  
#!/bin/bash  
/usr/bin/strace -o/t1 -f /usr/sbin/in.telnetd -h
- 3) chmod 700 /blah to make it executable only to root
- 4)  
killall -HUP inetd  
or ps aux | grep inetd  
get inetd's process id  
& kill -HUP inetd to restart it.

### Important options

---

-o is used to tell strace to output to a file in our case t1 in the root directory  
-f is to follow children i.e.  
e.g in our case above telnetd will start the login process & subsequently a shell like bash.  
You will be able to tell which is which from the process ID's listed on the left hand side of the strace output.  
-p<pid> will tell strace to attach to a running process, yup this can be done provided  
it isn't being traced or debugged already & you have enough privileges,  
the reason 2 processes cannot trace or debug the same program is that strace becomes the parent process of the one being debugged & processes ( unlike people )  
can have only one parent.

However the file /t1 will get big quite quickly  
to test it telnet 127.0.0.1

now look at what files in.telnetd execve'd  
413 execve("/usr/sbin/in.telnetd", ["/usr/sbin/in.telnetd", "-h"], [/\* 17 vars \*/]) = 0  
414 execve("/bin/login", ["/bin/login", "-h", "localhost", "-p"], [/\* 2 vars \*/]) = 0

Whey it worked!.

### Other hints:

---

If the program is not very interactive ( i.e. not much keyboard input )  
& is crashing in one architecture but not in another you can do  
an strace of both programs under as identical a scenario as you can

## Debugging390.txt

on both architectures outputting to a file then.  
do a diff of the two traces using the diff program  
i.e.  
diff output1 output2  
& maybe you'll be able to see where the call paths differed, this  
is possibly near the cause of the crash.

### More info

---

Look at man pages for strace & the various syscalls  
e.g. man strace, man alarm, man socket.

### Performance Debugging

---

gcc is capable of compiling in profiling code just add the -p option  
to the CFLAGS, this obviously affects program size & performance.  
This can be used by the gprof gnu profiling tool or the  
gcov the gnu code coverage tool ( code coverage is a means of testing  
code quality by checking if all the code in an executable is exercised by  
a tester ).

### Using top to find out where processes are sleeping in the kernel

---

To do this copy the System.map from the root directory where  
the linux kernel was built to the /boot directory on your  
linux machine.  
Start top  
Now type fU<return>  
You should see a new field called WCHAN which  
tells you where each process is sleeping here is a typical output.

```
6:59pm up 41 min, 1 user, load average: 0.00, 0.00, 0.00
28 processes: 27 sleeping, 1 running, 0 zombie, 0 stopped
CPU states: 0.0% user, 0.1% system, 0.0% nice, 99.8% idle
Mem: 254900K av, 45976K used, 208924K free, 0K shrd, 28636K buff
Swap: 0K av, 0K used, 0K free, 8620K cached
```

| PID        | USER | PRI | NI | SIZE | RSS  | SHARE | WCHAN     | STAT | LIB | %CPU | %MEM | TIME     |
|------------|------|-----|----|------|------|-------|-----------|------|-----|------|------|----------|
| COMMAND    |      |     |    |      |      |       |           |      |     |      |      |          |
| 750        | root | 12  | 0  | 848  | 848  | 700   | do_select | S    | 0   | 0.1  | 0.3  | 0:00     |
| in.telnetd |      |     |    |      |      |       |           |      |     |      |      |          |
| 767        | root | 16  | 0  | 1140 | 1140 | 964   |           | R    | 0   | 0.1  | 0.4  | 0:00 top |
| 1          | root | 8   | 0  | 212  | 212  | 180   | do_select | S    | 0   | 0.0  | 0.0  | 0:00     |
| init       |      |     |    |      |      |       |           |      |     |      |      |          |
| 2          | root | 9   | 0  | 0    | 0    | 0     | down_inte | SW   | 0   | 0.0  | 0.0  | 0:00     |
| kmcheck    |      |     |    |      |      |       |           |      |     |      |      |          |

### The time command

---

Another related command is the time command which gives you an indication  
of where a process is spending the majority of its time.  
e.g.  
time ping -c 5 nc  
outputs

## Debugging390. txt

real 0m4.054s  
user 0m0.010s  
sys 0m0.010s

### Debugging under VM

=====

#### Notes

-----

Addresses & values in the VM debugger are always hex never decimal  
Address ranges are of the format <HexValue1>-<HexValue2> or  
<HexValue1>.<HexValue2>  
e.g. The address range 0x2000 to 0x3000 can be described as 2000-3000 or  
2000.1000

The VM Debugger is case insensitive.

VM's strengths are usually other debuggers weaknesses you can get at any  
resource  
no matter how sensitive e.g. memory management resources, change address  
translation  
in the PSW. For kernel hacking you will reap dividends if you get good at it.

The VM Debugger displays operators but not operands, probably because some  
of it was written when memory was expensive & the programmer was probably proud  
that  
it fitted into 2k of memory & the programmers & didn't want to shock hardcore  
VM'ers by  
changing the interface :-), also the debugger displays useful information on the  
same line &  
the author of the code probably felt that it was a good idea not to go over  
the 80 columns on the screen.

As some of you are probably in a panic now this isn't as unintuitive as it may  
seem  
as the 390 instructions are easy to decode mentally & you can make a good guess  
at a lot  
of them as all the operands are nibble ( half byte aligned ) & if you have an  
objdump listing  
also it is quite easy to follow, if you don't have an objdump listing keep a  
copy of  
the s/390 Reference Summary & look at between pages 2 & 7 or alternatively the  
s/390 principles of operation.  
e.g. even I can guess that  
0001AFF8' LR 180F CC 0  
is a ( load register ) lr r0,r15

Also it is very easy to tell the length of a 390 instruction from the 2 most  
significant  
bits in the instruction ( not that this info is really useful except if you are  
trying to  
make sense of a hexdump of code ).  
Here is a table

| Bits | Instruction Length |
|------|--------------------|
| 00   | 2 Bytes            |

## Debugging390.txt

|    |         |
|----|---------|
| 01 | 4 Bytes |
| 10 | 4 Bytes |
| 11 | 6 Bytes |

The debugger also displays other useful info on the same line such as the addresses being operated on destination addresses of branches & condition codes.  
e. g.

```
00019736' AHI A7DAFF0E CC 1
000198BA' BRC A7840004 -> 000198C2' CC 0
000198CE' STM 900EF068 >> 0FA95E78 CC 2
```

## Useful VM debugger commands

---

I suppose I'd better mention this before I start to list the current active traces do

Q TR

there can be a maximum of 255 of these per set  
( more about trace sets later ).

To stop traces issue a

TR END.

To delete a particular breakpoint issue

TR DEL <breakpoint number>

The PA1 key drops to CP mode so you can issue debugger commands,

Doing alt c (on my 3270 console at least ) clears the screen.

hitting b <enter> comes back to the running operating system  
from cp mode ( in our case linux ).

It is typically useful to add shortcuts to your profile.exec file  
if you have one ( this is roughly equivalent to autoexec.bat in DOS ).  
file here are a few from mine.

```
/* this gives me command history on issuing f12 */
```

```
set pf12 retrieve
```

```
/* this continues */
```

```
set pf8 imm b
```

```
/* goes to trace set a */
```

```
set pf1 imm tr goto a
```

```
/* goes to trace set b */
```

```
set pf2 imm tr goto b
```

```
/* goes to trace set c */
```

```
set pf3 imm tr goto c
```

## Instruction Tracing

---

Setting a simple breakpoint

TR I PSWA <address>

To debug a particular function try

TR I R <function address range>

TR I on its own will single step.

TR I DATA <MNEMONIC> <OPTIONAL RANGE> will trace for particular mnemonics

e.g.

TR I DATA 4D R 0197BC.4000

will trace for BAS'es ( opcode 4D ) in the range 0197BC.4000

if you were inclined you could add traces for all branch instructions & suffix them with the run prefix so you would have a backtrace on screen when a program crashes.

TR BR <INTO OR FROM> will trace branches into or out of an address.

e.g.

TR BR INTO 0 is often quite useful if a program is getting awkward & deciding to branch to 0 & crashing as this will stop at the address before in jumps to 0.

TR I R <address range> RUN cmd d g

single steps a range of addresses but stays running & displays the gprs on each step.

## Displaying & modifying Registers

---

D G will display all the gprs

Adding a extra G to all the commands is necessary to access the full 64 bit content in VM on z/Architecture obviously this isn't required for access registers

as these are still 32 bit.

e.g. DGG instead of DG

D X will display all the control registers

D AR will display all the access registers

D AR4-7 will display access registers 4 to 7

CPU ALL D G will display the GRPS of all CPUS in the configuration

D PSW will display the current PSW

st PSW 2000 will put the value 2000 into the PSW & cause crash your machine.

D PREFIX displays the prefix offset

## Displaying Memory

---

To display memory mapped using the current PSW's mapping try

D <range>

To make VM display a message each time it hits a particular address & continue try

D I<range> will disassemble/display a range of instructions.

ST addr 32 bit word will store a 32 bit aligned address

D T<range> will display the EBCDIC in an address ( if you are that way inclined )

D R<range> will display real addresses ( without DAT ) but with prefixing.

There are other complex options to display if you need to get at say home space but are in primary space the easiest thing to do is to temporarily modify the PSW to the other addressing mode, display the stuff & then restore it.

## Hints

---

If you want to issue a debugger command without halting your virtual machine



with the

PA1 key try prefixing the command with #CP e.g.

#cp tr i pswa 2000

also suffixing most debugger commands with RUN will cause them not to stop just display the mnemonic at the current instruction on the console.

If you have several breakpoints you want to put into your program &

you get fed up of cross referencing with System.map

you can do the following trick for several symbols.

grep do\_signal System.map

which emits the following among other things

0001f4e0 T do\_signal

now you can do

TR I PSWA 0001f4e0 cmd msg \* do\_signal

This sends a message to your own console each time do\_signal is entered.

( As an aside I wrote a perl script once which automatically generated a REXX script with breakpoints on every kernel procedure, this isn't a good idea because there are thousands of these routines & VM can only set 255 breakpoints at a time so you nearly had to spend as long pruning the file down as you would entering the msg's by hand ),however, the trick might be useful for a single object file.

On linux'es 3270 emulator x3270 there is a very useful option under the file ment

Save Screens In File this is very good of keeping a copy of traces.

From CMS help <command name> will give you online help on a particular command.

e.g.

HELP DISPLAY

Also CP has a file called profile.exec which automatically gets called on startup of CMS ( like autoexec.bat ), keeping on a DOS analogy session

CP has a feature similar to doskey, it may be useful for you to

use profile.exec to define some keystrokes.

e.g.

SET PF9 IMM B

This does a single step in VM on pressing F8.

SET PF10

This sets up the ^ key.

which can be used for ^c (ctrl-c), ^z (ctrl-z) which can't be typed directly into some 3270 consoles.

SET PF11 ^-

This types the starting keystrokes for a sysrq see SysRq below.

SET PF12 RETRIEVE

This retrieves command history on pressing F12.

Sometimes in VM the display is set up to scroll automatically this can be very annoying if there are messages you wish to look at

to stop this do

TERM MORE 255 255

This will nearly stop automatic screen updates, however it will

cause a denial of service if lots of messages go to the 3270 console,

so it would be foolish to use this as the default on a production machine.

Tracing particular processes

-----  
The kernel's text segment is intentionally at an address in memory that it will very seldom collide with text segments of user programs ( thanks Martin ), this simplifies debugging the kernel.

However it is quite common for user processes to have addresses which collide this can make debugging a particular process under VM painful under normal circumstances as the process may change when doing a  
TR I R <address range>.

Thankfully after reading VM's online help I figured out how to debug I particular process.

Your first problem is to find the STD ( segment table designation ) of the program you wish to debug.

There are several ways you can do this here are a few  
1) objdump --syms <program to be debugged> | grep main

To get the address of main in the program.

tr i pswa <address of main>

Start the program, if VM drops to CP on what looks like the entry point of the main function this is most likely the process you wish to debug.

Now do a D X13 or D XG13 on z/Architecture.

On 31 bit the STD is bits 1-19 ( the ST0 segment table origin ) & 25-31 ( the STL segment table length ) of CR13.

now type

TR I R STD <CR13's value> 0.7fffffff

e.g.

TR I R STD 8F32E1FF 0.7fffffff

Another very useful variation is

TR STORE INTO STD <CR13's value> <address range>  
for finding out when a particular variable changes.

An alternative way of finding the STD of a currently running process is to do the following, ( this method is more complex but could be quite convenient if you aren't updating the kernel much & so your kernel structures will stay constant for a reasonable period of time ).

grep task /proc/<pid>/status

from this you should see something like

task: 0f160000 ksp: 0f161de8 pt\_regs: 0f161f68

This now gives you a pointer to the task structure.

Now make CC:="s390-gcc -g" kernel/sched.s

To get the task\_struct stabinfo.

( task\_struct is defined in include/linux/sched.h ).

Now we want to look at

task->active\_mm->pgd

on my machine the active\_mm in the task structure stab is

active\_mm: (4, 12), 672, 32

its offset is 672/8=84=0x54

the pgd member in the mm\_struct stab is

pgd: (4, 6)=\*(29, 5), 96, 32

so its offset is 96/8=12=0xc

so we'll

hexdump -s 0xf160054 /dev/mem | more

i.e. task\_struct+active\_mm offset

to look at the active\_mm member

## Debugging390.txt

```
f160054 0fee cc60 0019 e334 0000 0000 0000 0011
hexdump -s 0x0feecc6c /dev/mem | more
i.e. active_mm+pgd offset
feecc6c 0f2c 0000 0000 0001 0000 0001 0000 0010
we get something like
now do
TR I R STD <pgd|0x7f> 0.7fffffff
i.e. the 0x7f is added because the pgd only
gives the page table origin & we need to set the low bits
to the maximum possible segment table length.
TR I R STD 0f2c007f 0.7fffffff
on z/Architecture you'll probably need to do
TR I R STD <pgd|0x7> 0.ffffffffffffffff
to set the TableType to 0x1 & the Table length to 3.
```

## Tracing Program Exceptions

---

If you get a crash which says something like  
illegal operation or specification exception followed by a register dump  
You can restart linux & trace these using the tr prog <range or value> trace  
option.

The most common ones you will normally be tracing for is  
1=operation exception  
2=privileged operation exception  
4=protection exception  
5=addressing exception  
6=specification exception  
10=segment translation exception  
11=page translation exception

The full list of these is on page 22 of the current s/390 Reference Summary.  
e.g.  
tr prog 10 will trace segment translation exceptions.  
tr prog on its own will trace all program interruption codes.

## Trace Sets

---

On starting VM you are initially in the INITIAL trace set.  
You can do a Q TR to verify this.  
If you have a complex tracing situation where you wish to wait for instance  
till a driver is open before you start tracing IO, but know in your  
heart that you are going to have to make several runs through the code till you  
have a clue whats going on.

What you can do is  
TR I PSWA <Driver open address>  
hit b to continue till breakpoint  
reach the breakpoint  
now do your  
TR GOTO B  
TR IO 7c08-7c09 inst int run

or whatever the IO channels you wish to trace are & hit b

To get back to the initial trace set do

TR GOTO INITIAL

& the TR I PSWA <Driver open address> will be the only active breakpoint again.

#### Tracing linux syscalls under VM

---

Syscalls are implemented on Linux for S390 by the Supervisor call instruction (SVC) there 256

possibilities of these as the instruction is made up of a 0xA opcode & the second byte being

the syscall number. They are traced using the simple command.

TR SVC <Optional value or range>

the syscalls are defined in linux/arch/s390/include/asm/unistd.h

e.g. to trace all file opens just do

TR SVC 5 ( as this is the syscall number of open )

#### SMP Specific commands

---

To find out how many cpus you have

Q CPUS displays all the CPU's available to your virtual machine

To find the cpu that the current cpu VM debugger commands are being directed at do

Q CPU to change the current cpu VM debugger commands are being directed at do

CPU <desired cpu no>

On a SMP guest issue a command to all CPUs try prefixing the command with cpu all.

To issue a command to a particular cpu try cpu <cpu number> e.g.

CPU 01 TR I R 2000.3000

If you are running on a guest with several cpus & you have a IO related problem & cannot follow the flow of code but you know it isn't smp related.

from the bash prompt issue

shutdown -h now or halt.

do a Q CPUS to find out how many cpus you have

detach each one of them from cp except cpu 0

by issuing a

DETACH CPU 01-(number of cpus in configuration)

& boot linux again.

TR SIGP will trace inter processor signal processor instructions.

DEFINE CPU 01-(number in configuration)

will get your guests cpus back.

#### Help for displaying ascii textstrings

---

On the very latest VM Nucleus'es VM can now display ascii

( thanks Neale for the hint ) by doing

D TX<lowaddr>.<len>

e.g.

D TX0.100

Alternatively

## Debugging390.txt

=====

Under older VM debuggers ( I love EBDIC too ) you can use this little program I wrote which will convert a command line of hex digits to ascii text which can be compiled under linux & you can copy the hex digits from your x3270 terminal to your xterm if you are debugging from a linuxbox.

This is quite useful when looking at a parameter passed in as a text string under VM ( unless you are good at decoding ASCII in your head ).

e.g. consider tracing an open syscall

TR SVC 5

We have stopped at a breakpoint

000151B0' SVC 0A05 -> 0001909A' CC 0

D 20.8 to check the SVC old psw in the prefix area & see was it from userspace ( for the layout of the prefix area consult P18 of the s/390 390 Reference

Summary

if you have it available ).

V00000020 070C2000 800151B2

The problem state bit wasn't set & it's also too early in the boot sequence for it to be a userspace SVC if it was we would have to temporarily switch the psw to user space addressing so we could get at the first parameter of the open in gpr2.

Next do a

D G2

GPR 2 = 00014CB4

Now display what gpr2 is pointing to

D 00014CB4.20

V00014CB4 2F646576 2F636F6E 736F6C65 00001BF5

V00014CC4 FC00014C B4001001 E0001000 B8070707

Now copy the text till the first 00 hex ( which is the end of the string to an xterm & do hex2ascii on it.

hex2ascii 2F646576 2F636F6E 736F6C65 00

outputs

Decoded Hex:=/ d e v / c o n s o l e 0x00

We were opening the console device,

You can compile the code below yourself for practice :-),

/\*

\* hex2ascii.c

\* a useful little tool for converting a hexadecimal command line to ascii

\*

\* Author(s): Denis Joseph Barrow (djbarrow@de.ibm.com, barrow\_dj@yahoo.com)

\* (C) 2000 IBM Deutschland Entwicklung GmbH, IBM Corporation.

\*/

#include <stdio.h>

int main(int argc, char \*argv[])

{

int cnt1, cnt2, len, toggle=0;

int startcnt=1;

unsigned char c, hex;

```

if(argc>1&&(strcmp(argv[1], "-a")==0))
 startcnt=2;
printf("Decoded Hex:=");
for(cnt1=startcnt;cnt1<argc;cnt1++)
{
 len=strlen(argv[cnt1]);
 for(cnt2=0;cnt2<len;cnt2++)
 {
 c=argv[cnt1][cnt2];
 if(c>='0' && c<='9')
 c=c-'0';
 if(c>='A' && c<='F')
 c=c-'A'+10;
 if(c>='a' && c<='f')
 c=c-'a'+10;
 switch(toggle)
 {
 case 0:
 hex=c<<4;
 toggle=1;
 break;
 case 1:
 hex+=c;
 if(hex<32||hex>127)
 {
 if(startcnt==1)
 printf("0x%02X ", (int)hex);
 else
 printf(".");
 }
 else
 {
 printf("%c", hex);
 if(startcnt==1)
 printf(" ");
 }
 toggle=0;
 break;
 }
 }
}
printf("\n");
}

```

## Stack tracing under VM

### A basic backtrace

Here are the tricks I use 9 out of 10 times it works pretty well,

When your backchain reaches a dead end

-----  
This can happen when an exception happens in the kernel & the kernel is entered twice

if you reach the NULL pointer at the end of the back chain you should be able to sniff further back if you follow the following tricks.

1) A kernel address should be easy to recognise since it is in primary space & the problem state bit isn't set & also The Hi bit of the address is set.

2) Another backchain should also be easy to recognise since it is an address pointing to another address approximately 100 bytes or 0x70 hex behind the current stackpointer.

Here is some practice.

boot the kernel & hit PA1 at some random time

d g to display the gprs, this should display something like

GPR 0 = 00000001 00156018 0014359C 00000000

GPR 4 = 00000001 001B8888 000003E0 00000000

GPR 8 = 00100080 00100084 00000000 000FE000

GPR 12 = 00010400 8001B2DC 8001B36A 000FFED8

Note that GPR14 is a return address but as we are real men we are going to trace the stack.

display 0x40 bytes after the stack pointer.

V000FFED8 000FFF38 8001B838 80014C8E 000FFF38

V000FFEE8 00000000 00000000 000003E0 00000000

V000FFEF8 00100080 00100084 00000000 000FE000

V000FFF08 00010400 8001B2DC 8001B36A 000FFED8

Ah now look at whats in sp+56 (sp+0x38) this is 8001B36A our saved r14 if you look above at our stackframe & also agrees with GPR14.

now backchain

d 000FFF38.40

we now are taking the contents of SP to get our first backchain.

V000FFF38 000FFFA0 00000000 00014995 00147094

V000FFF48 00147090 001470A0 000003E0 00000000

V000FFF58 00100080 00100084 00000000 001BF1D0

V000FFF68 00010400 800149BA 80014CA6 000FFF38

This displays a 2nd return address of 80014CA6

now do d 000FFFA0.40 for our 3rd backchain

V000FFFA0 04B52002 0001107F 00000000 00000000

V000FFFB0 00000000 00000000 FF000000 0001107F

V000FFFC0 00000000 00000000 00000000 00000000

V000FFFD0 00010400 80010802 8001085A 000FFFA0

our 3rd return address is 8001085A

as the 04B52002 looks suspiciously like rubbish it is fair to assume that the kernel entry routines

for the sake of optimisation don't set up a backchain.

now look at System.map to see if the addresses make any sense.

```
grep -i 0001b3 System.map
outputs among other things
0001b304 T cpu_idle
so 8001B36A
is cpu_idle+0x66 (quiet the cpu is asleep, don't wake it)
```

```
grep -i 00014 System.map
produces among other things
00014a78 T start_kernel
so 0014CA6 is start_kernel+some hex number I can't add in my head.
```

```
grep -i 00108 System.map
this produces
00010800 T _stext
so 8001085A is _stext+0x5a
```

Congrats you've done your first backchain.

#### s/390 & z/Architecture IO Overview

I am not going to give a course in 390 IO architecture as this would take me quite a while & I'm no expert. Instead I'll give a 390 IO architecture summary for Dummies if you have the s/390 principles of operation available read this instead. If nothing else you may find a few useful keywords in here & be able to use them on a web search engine like altavista to find more useful information.

Unlike other bus architectures modern 390 systems do their IO using mostly fibre optics & devices such as tapes & disks can be shared between several mainframes, also S390 can support up to 65536 devices while a high end PC based system might be choking with around 64. Here is some of the common IO terminology

#### Subchannel:

This is the logical number most IO commands use to talk to an IO device there can be up to

0x10000 (65536) of these in a configuration typically there is a few hundred. Under VM

for simplicity they are allocated contiguously, however on the native hardware they are not

they typically stay consistent between boots provided no new hardware is inserted or removed.

Under Linux for 390 we use these as IRQ's & also when issuing an IO command (CLEAR SUBCHANNEL,



### Debugging390.txt

HALT SUBCHANNEL, MODIFY SUBCHANNEL, RESUME SUBCHANNEL, START SUBCHANNEL, STORE SUBCHANNEL & TEST SUBCHANNEL ) we use this as the ID of the device we wish to talk to, the most important of these instructions are START SUBCHANNEL ( to start IO ), TEST SUBCHANNEL ( to check whether the IO completed successfully ), & HALT SUBCHANNEL ( to kill IO ), a subchannel can have up to 8 channel paths to a device this offers redundancy if one is not available.

#### Device Number:

This number remains static & Is closely tied to the hardware, there are 65536 of these also they are made up of a CHPID ( Channel Path ID, the most significant 8 bits ) & another lsb 8 bits. These remain static even if more devices are inserted or removed from the hardware, there is a 1 to 1 mapping between Subchannels & Device Numbers provided devices aren't inserted or removed.

#### Channel Control Words:

CCWS are linked lists of instructions initially pointed to by an operation request block (ORB), which is initially given to Start Subchannel (SSCH) command along with the subchannel number for the IO subsystem to process while the CPU continues executing normal code. These come in two flavours, Format 0 ( 24 bit for backward ) compatibility & Format 1 ( 31 bit ). These are typically used to issue read & write ( & many other instructions ) they consist of a length field & an absolute address field. For each IO typically get 1 or 2 interrupts one for channel end ( primary status ) when the channel is idle & the second for device end ( secondary status ) sometimes you get both concurrently, you check how the IO went on by issuing a TEST SUBCHANNEL at each interrupt, from which you receive an Interruption response block (IRB). If you get channel & device end status in the IRB without channel checks etc. your IO probably went okay. If you didn't you probably need a doctor to examine the IRB & extended status word etc. If an error occurs, more sophisticated control units have a facility known as concurrent sense this means that if an error occurs Extended sense information will be presented in the Extended status word in the IRB if not you have to issue a subsequent SENSE CCW command after the test subchannel.

TPI( Test pending interrupt) can also be used for polled IO but in multitasking multiprocessor systems it isn't recommended except for checking special cases ( i.e. non looping checks for

pending IO etc. ).

Store Subchannel & Modify Subchannel can be used to examine & modify operating characteristics of a subchannel ( e.g. channel paths ).

Other IO related Terms:

Sysplex: S390's Clustering Technology

QDIO: S390's new high speed IO architecture to support devices such as gigabit ethernet,

this architecture is also designed to be forward compatible with up & coming 64 bit machines.

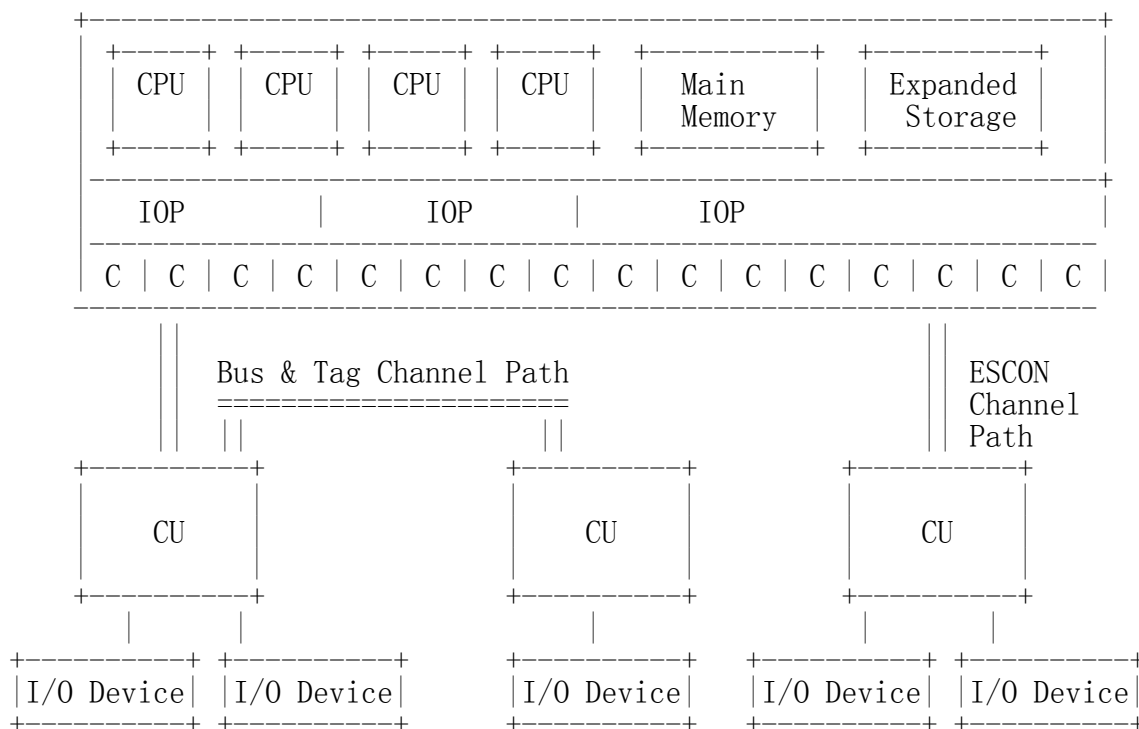
## General Concepts

Input Output Processors (IOP's) are responsible for communicating between the mainframe CPU's & the channel & relieve the mainframe CPU's from the burden of communicating with IO devices directly, this allows the CPU's to concentrate on data processing.

IOP's can use one or more links ( known as channel paths ) to talk to each IO device. It first checks for path availability & chooses an available one, then starts ( & sometimes terminates IO ).

There are two types of channel path: ESCON & the Parallel IO interface.

IO devices are attached to control units, control units provide the logic to interface the channel paths & channel path IO protocols to the IO devices, they can be integrated with the devices or housed separately & often talk to several similar devices ( typical examples would be raid controllers or a control unit which connects to 1000 3270 terminals ).



CPU = Central Processing Unit  
C = Channel  
IOP = IP Processor  
CU = Control Unit

The 390 IO systems come in 2 flavours the current 390 machines support both

The Older 360 & 370 Interface, sometimes called the Parallel I/O interface, sometimes called Bus-and Tag & sometimes Original Equipment Manufacturers Interface (OEMI).

This byte wide Parallel channel path/bus has parity & data on the "Bus" cable & control lines on the "Tag" cable. These can operate in byte multiplex mode for sharing between several slow devices or burst mode & monopolize the channel for the whole burst. Up to 256 devices can be addressed on one of these cables. These cables are about one inch in diameter. The maximum unextended length supported by these cables is 125 Meters but this can be extended up to 2km with a fibre optic channel extended such as a 3044. The maximum burst speed supported is 4.5 megabytes per second however some really old processors support only transfer rates of 3.0, 2.0 & 1.0 MB/sec. One of these paths can be daisy chained to up to 8 control units.

ESCON if fibre optic it is also called FICON  
Was introduced by IBM in 1990. Has 2 fibre optic cables & uses either leds or lasers for communication at a signaling rate of up to 200 megabits/sec. As 10bits are transferred for every 8 bits info this drops to 160 megabits/sec & to 18.6 Megabytes/sec once control info & CRC are added. ESCON only operates in burst mode.

ESCONs typical max cable length is 3km for the led version & 20km for the laser version known as XDF ( extended distance facility ). This can be further extended by using an ESCON director which triples the above mentioned ranges. Unlike Bus & Tag as ESCON is serial it uses a packet switching architecture the standard Bus & Tag control protocol is however present within the packets. Up to 256 devices can be attached to each control unit that uses one of these interfaces.

Common 390 Devices include:

Network adapters typically OSA2, 3172's, 2116's & OSA-E gigabit ethernet adapters, Consoles 3270 & 3215 ( a teletype emulated under linux for a line mode console ).

DASD's direct access storage devices ( otherwise known as hard disks ).

Tape Drives.

CTC ( Channel to Channel Adapters ),

ESCON or Parallel Cables used as a very high speed serial link

between 2 machines. We use 2 cables under linux to do a bi-directional serial link.

### Debugging IO on s/390 & z/Architecture under VM

Now we are ready to go on with IO tracing commands under VM

A few self explanatory queries:

Q OSA

Q CTC

Q DISK ( This command is CMS specific )

Q DASD

Q OSA on my machine returns

OSA 7C08 ON OSA 7C08 SUBCHANNEL = 0000

OSA 7C09 ON OSA 7C09 SUBCHANNEL = 0001

OSA 7C14 ON OSA 7C14 SUBCHANNEL = 0002

OSA 7C15 ON OSA 7C15 SUBCHANNEL = 0003

If you have a guest with certain privileges you may be able to see devices which don't belong to you. To avoid this, add the option V.

e.g.

Q V OSA

Now using the device numbers returned by this command we will

Trace the io starting up on the first device 7c08 & 7c09

In our simplest case we can trace the

start subchannels

like TR SSCH 7C08-7C09

or the halt subchannels

or TR HSCH 7C08-7C09

MSCH's , STSCH's I think you can guess the rest

Ingo's favourite trick is tracing all the IO's & CCWS & spooling them into the reader of another

VM guest so he can ftp the logfile back to his own machine. I'll do a small bit of this & give you

a look at the output.

1) Spool stdout to VM reader

SP PRT TO (another vm guest ) or \* for the local vm guest

2) Fill the reader with the trace

TR IO 7c08-7c09 INST INT CCW PRT RUN

3) Start up linux

i 00c

4) Finish the trace

TR END

5) close the reader

C PRT

6) list reader contents

RDRLIST

7) copy it to linux4's minidisk

RECEIVE / LOG TXT A1 ( replace

8)

file1 & press F11 to look at it

You should see something like:

```
00020942' SSCH B2334000 0048813C CC 0 SCH 0000 DEV 7C08
 CPA 000FFDF0 PARM 00E2C9C4 KEY 0 FPI C0 LPM 80
 CCW 000FFDF0 E4200100 00487FE8 0000 E4240100
 IDAL 43D8AFE8
 IDAL 0FB76000
00020B0A' I/O DEV 7C08 -> 000197BC' SCH 0000 PARM 00E2C9C4
00021628' TSCH B2354000 >> 00488164 CC 0 SCH 0000 DEV 7C08
 CCWA 000FFDF8 DEV STS 0C SCH STS 00 CNT 00EC
 KEY 0 FPI C0 CC 0 CTLS 4007
00022238' STSCH B2344000 >> 00488108 CC 0 SCH 0000 DEV 7C08
```

If you don't like messing up your readed ( because you possibly booted from it )  
you can alternatively spool it to another readers guest.

#### Other common VM device related commands

-----  
These commands are listed only because they have  
been of use to me in the past & may be of use to  
you too. For more complete info on each of the commands  
use type HELP <command> from CMS.

detaching devices

DET <devno range>

ATT <devno range> <guest>

attach a device to guest \* for your own guest

READY <devno> cause VM to issue a fake interrupt.

The VARY command is normally only available to VM administrators.

VARY ON PATH <path> TO <devno range>

VARY OFF PATH <PATH> FROM <devno range>

This is used to switch on or off channel paths to devices.

Q CHPID <channel path ID>

This displays state of devices using this channel path

D SCHIB <subchannel>

This displays the subchannel information SCHIB block for the device.

this I believe is also only available to administrators.

DEFINE CTC <devno>

defines a virtual CTC channel to channel connection

2 need to be defined on each guest for the CTC driver to use.

COUPLE devno userid remote devno

Joins a local virtual device to a remote virtual device

( commonly used for the CTC driver ).

Building a VM ramdisk under CMS which linux can use

def vfb-<blocksize> <subchannel> <number blocks>

blocksize is commonly 4096 for linux.

Formatting it

## Debugging390.txt

format <subchannel> <driver letter e.g. x> (blksize <blocksize>

Sharing a disk between multiple guests

LINK userid devno1 devno2 mode password

## GDB on S390

=====

N.B. if compiling for debugging gdb works better without optimisation  
( see Compiling programs for debugging )

## invocation

=====

`gdb <victim program> <optional corefile>`

## Online help

=====

`help:` gives help on commands

e.g.

`help`

`help display`

Note `gdb`'s online help is very good use it.

## Assembly

=====

`info registers:` displays registers other than floating point.

`info all-registers:` displays floating points as well.

`disassemble:` disassembles

e.g.

`disassemble` without parameters will disassemble the current function

`disassemble $pc $pc+10`

## Viewing & modifying variables

=====

`print` or `p:` displays variable or register

e.g. `p/x $sp` will display the stack pointer

`display:` prints variable or register each time program stops

e.g.

`display/x $pc` will display the program counter

`display argc`

`undisplay :` undo's `display`'s

`info breakpoints:` shows all current breakpoints

`info stack:` shows stack back trace ( if this doesn't work too well, I'll show you the

`stacktrace` by hand below ).

`info locals:` displays local variables.

`info args:` display current procedure arguments.

## Debugging390.txt

set args: will set argc & argv each time the victim program is invoked.

```
set <variable>=value
set argc=100
set $pc=0
```

### Modifying execution

---

```
step: steps n lines of sourcecode
step steps 1 line.
step 100 steps 100 lines of code.
```

next: like step except this will not step into subroutines

```
stepi: steps a single machine code instruction.
e.g. stepi 100
```

nexti: steps a single machine code instruction but will not step into subroutines.

finish: will run until exit of the current routine

run: (re)starts a program

cont: continues a program

quit: exits gdb.

### breakpoints

---

```
break
sets a breakpoint
e.g.
```

```
break main
```

```
break *$pc
```

```
break *0x400618
```

Here's a really useful one for large programs

```
rbr
```

Set a breakpoint for all functions matching REGEXP

```
e.g.
```

```
rbr 390
```

will set a breakpoint with all functions with 390 in their name.

```
info breakpoints
```

lists all breakpoints

delete: delete breakpoint by number or delete them all

```
e.g.
```

## Debugging390.txt

delete 1 will delete the first breakpoint  
delete will delete them all

watch: This will set a watchpoint ( usually hardware assisted ),  
This will watch a variable till it changes

e.g.

watch cnt, will watch the variable cnt till it changes.

As an aside unfortunately gdb's, architecture independent watchpoint code  
is inconsistent & not very good, watchpoints usually work but not always.

info watchpoints: Display currently active watchpoints

condition: ( another useful one )

Specify breakpoint number N to break only if COND is true.

Usage is `condition N COND', where N is an integer and COND is an  
expression to be evaluated whenever breakpoint N is reached.

## User defined functions/macros

---

define: ( Note this is very very useful, simple & powerful )

usage define <name> <list of commands> end

examples which you should consider putting into .gdbinit in your home directory

define d

stepi

disassemble \$pc \$pc+10

end

define e

nexti

disassemble \$pc \$pc+10

end

## Other hard to classify stuff

---

signal n:

sends the victim program a signal.

e.g. signal 3 will send a SIGQUIT.

info signals:

what gdb does when the victim receives certain signals.

list:

e.g.

list lists current function source

list 1,10 list first 10 lines of current file.

list test.c:1,10

directory:

Adds directories to be searched for source if gdb cannot find the source.  
(note it is a bit sensitive about slashes)

e.g. To add the root of the filesystem to the searchpath do



directory //

call <function>

This calls a function in the victim program, this is pretty powerful

e.g.

(gdb) call printf("hello world")

outputs:

\$1 = 11

You might now be thinking that the line above didn't work, something extra had to be done.

(gdb) call fflush(stdout)

hello world\$2 = 0

As an aside the debugger also calls malloc & free under the hood to make space for the "hello world" string.

hints

-----

1) command completion works just like bash  
( if you are a bad typist like me this really helps )  
e.g. hit br <TAB> & cursor up & down :-).

2) if you have a debugging problem that takes a few steps to recreate put the steps into a file called .gdbinit in your current working directory if you have defined a few extra useful user defined commands put these in your home directory & they will be read each time gdb is launched.

A typical .gdbinit file might be.

break main

run

break runtime\_exception

cont

stack chaining in gdb by hand

-----

This is done using a the same trick described for VM  
p/x (\*(\$sp+56))&0x7fffffff get the first backchain.

For z/Architecture

Replace 56 with 112 & ignore the &0x7fffffff

in the macros below & do nasty casts to longs like the following  
as gdb unfortunately deals with printed arguments as ints which  
messes up everything.

i.e. here is a 3rd backchain dereference

p/x \*(long \*) (\*\*\*(long \*\*\*)\$sp+112)

this outputs

\$5 = 0x528f18

on my machine.

Now you can use

info symbol (\*(\$sp+56))&0x7fffffff

## Debugging390.txt

you might see something like.  
rl\_getc + 36 in section .text telling you what is located at address 0x528f18  
Now do.  
p/x ((\*\$sp+56))&0x7fffffff  
This outputs  
\$6 = 0x528ed0  
Now do.  
info symbol ((\*\$sp+56))&0x7fffffff  
rl\_read\_key + 180 in section .text  
now do  
p/x ((\*\$sp+56))&0x7fffffff  
& so on.

## Disassembling instructions without debug info

---

gdb typically complains if there is a lack of debugging symbols in the disassemble command with  
"No function contains specified address." To get around this do  
x/<number lines to disassemble>xi <address>  
e.g.  
x/20xi 0x400730

Note: Remember gdb has history just like bash you don't need to retype the whole line just use the up & down arrows.

## For more info

---

From your linuxbox do  
man gdb or info gdb.

## core dumps

---

What a core dump ?,  
A core dump is a file generated by the kernel ( if allowed ) which contains the registers,  
& all active pages of the program which has crashed.  
From this file gdb will allow you to look at the registers & stack trace & memory of the  
program as if it just crashed on your system, it is usually called core & created in the  
current working directory.  
This is very useful in that a customer can mail a core dump to a technical support department  
& the technical support department can reconstruct what happened.  
Provided they have an identical copy of this program with debugging symbols compiled in &  
the source base of this build is available.  
In short it is far more useful than something like a crash log could ever hope to be.

In theory all that is missing to restart a core dumped program is a kernel patch

## Debugging390.txt

which

will do the following.

- 1) Make a new kernel task structure
- 2) Reload all the dumped pages back into the kernel's memory management structures.
- 3) Do the required clock fixups
- 4) Get all files & network connections for the process back into an identical state ( really difficult ).
- 5) A few more difficult things I haven't thought of.

Why have I never seen one ?.

Probably because you haven't used the command

`ulimit -c unlimited` in bash

to allow core dumps, now do

`ulimit -a`

to verify that the limit was accepted.

A sample core dump

To create this I'm going to do

`ulimit -c unlimited`

`gdb`

to launch gdb (my victim app. ) now be bad & do the following from another telnet/xterm session to the same machine

`ps -aux | grep gdb`

`kill -SIGSEGV <gdb's pid>`

or alternatively use `killall -SIGSEGV gdb` if you have the `killall` command.

Now look at the core dump.

`./gdb core`

Displays the following

GNU gdb 4.18

Copyright 1998 Free Software Foundation, Inc.

GDB is free software, covered by the GNU General Public License, and you are welcome to change it and/or distribute copies of it under certain conditions. Type "show copying" to see the conditions.

There is absolutely no warranty for GDB. Type "show warranty" for details.

This GDB was configured as "s390-ibm-linux"...

Core was generated by `./gdb'.

Program terminated with signal 11, Segmentation fault.

Reading symbols from /usr/lib/libncurses.so.4...done.

Reading symbols from /lib/libm.so.6...done.

Reading symbols from /lib/libc.so.6...done.

Reading symbols from /lib/ld-linux.so.2...done.

#0 0x40126d1a in read () from /lib/libc.so.6

Setting up the environment for debugging gdb.

Breakpoint 1 at 0x4dc6f8: file utils.c, line 471.

Breakpoint 2 at 0x4d87a4: file top.c, line 2609.

(top-gdb) info stack

#0 0x40126d1a in read () from /lib/libc.so.6

#1 0x528f26 in rl\_getc (stream=0x7ffffde8) at input.c:402

#2 0x528ed0 in rl\_read\_key () at input.c:381

#3 0x5167e6 in readline\_internal\_char () at readline.c:454

#4 0x5168ee in readline\_internal\_charloop () at readline.c:507

#5 0x51692c in readline\_internal () at readline.c:521

#6 0x5164fe in readline (prompt=0x7ffff810 "\177ÿøx\177ÿ÷ø\177ÿøxÀ")

## Debugging390.txt

```
at readline.c:349
#7 0x4d7a8a in command_line_input (prompt=0x564420 "(gdb) ", repeat=1,
 annotation_suffix=0x4d6b44 "prompt") at top.c:2091
#8 0x4d6cf0 in command_loop () at top.c:1345
#9 0x4e25bc in main (argc=1, argv=0x7ffffdf4) at main.c:635
```

## LDD

===

This is a program which lists the shared libraries which a library needs, Note you also get the relocations of the shared library text segments which help when using objdump --source.

e.g.

```
ldd ./gdb
```

outputs

```
libncurses.so.4 => /usr/lib/libncurses.so.4 (0x40018000)
libm.so.6 => /lib/libm.so.6 (0x4005e000)
libc.so.6 => /lib/libc.so.6 (0x40084000)
/lib/ld-linux.so.2 => /lib/ld-linux.so.2 (0x40000000)
```

## Debugging shared libraries

=====

Most programs use shared libraries, however it can be very painful when you single step instruction into a function like printf for the first time & you end up in functions like \_dl\_runtime\_resolve this is the ld.so doing lazy binding, lazy binding is a concept in ELF where shared library functions are not loaded into memory unless they are actually used, great for saving memory but a pain to debug.

To get around this either relink the program -static or exit gdb type export LD\_BIND\_NOW=true this will stop lazy binding & restart the gdb'ing the program in question.

## Debugging modules

=====

As modules are dynamically loaded into the kernel their address can be anywhere to get around this use the -m option with insmod to emit a load map which can be piped into a file if required.

## The proc file system

=====

What is it ?.

It is a filesystem created by the kernel with files which are created on demand by the kernel if read, or can be used to modify kernel parameters, it is a powerful concept.

e.g.

```
cat /proc/sys/net/ipv4/ip_forward
```

On my machine outputs

```
0
```

telling me ip\_forwarding is not on to switch it on I can do

```
echo 1 > /proc/sys/net/ipv4/ip_forward
```

```
cat it again
```

## Debugging390.txt

```
cat /proc/sys/net/ipv4/ip_forward
```

On my machine now outputs

```
1
```

IP forwarding is on.

There is a lot of useful info in here best found by going in & having a look around,

so I'll take you through some entries I consider important.

All the processes running on the machine have there own entry defined by  
/proc/<pid>

So lets have a look at the init process

```
cd /proc/1
```

```
cat cmdline
```

```
emits
```

```
init [2]
```

```
cd /proc/1/fd
```

This contains numerical entries of all the open files,  
some of these you can cat e.g. stdout (2)

```
cat /proc/29/maps
```

on my machine emits

```
00400000-00478000 r-xp 00000000 5f:00 4103 /bin/bash
00478000-0047e000 rw-p 00077000 5f:00 4103 /bin/bash
0047e000-00492000 rwxp 00000000 00:00 0
40000000-40015000 r-xp 00000000 5f:00 14382 /lib/ld-2.1.2.so
40015000-40016000 rw-p 00014000 5f:00 14382 /lib/ld-2.1.2.so
40016000-40017000 rwxp 00000000 00:00 0
40017000-40018000 rw-p 00000000 00:00 0
40018000-4001b000 r-xp 00000000 5f:00 14435 /lib/libtermcap.so.2.0.8
4001b000-4001c000 rw-p 00002000 5f:00 14435 /lib/libtermcap.so.2.0.8
4001c000-4010d000 r-xp 00000000 5f:00 14387 /lib/libc-2.1.2.so
4010d000-40111000 rw-p 000f0000 5f:00 14387 /lib/libc-2.1.2.so
40111000-40114000 rw-p 00000000 00:00 0
40114000-4011e000 r-xp 00000000 5f:00 14408 /lib/libnss_files-2.1.2.so
4011e000-4011f000 rw-p 00009000 5f:00 14408 /lib/libnss_files-2.1.2.so
7fffd000-80000000 rwxp fffff000 00:00 0
```

Showing us the shared libraries init uses where they are in memory  
& memory access permissions for each virtual memory area.

/proc/1/cwd is a softlink to the current working directory.

/proc/1/root is the root of the filesystem for this process.

/proc/1/mem is the current running processes memory which you  
can read & write to like a file.

strace uses this sometimes as it is a bit faster than the  
rather inefficient ptrace interface for peeking at DATA.

```
cat status
```

```
Name: init
```

```

State: S (sleeping)
Pid: 1
PPid: 0
Uid: 0 0 0 0
Gid: 0 0 0 0
Groups:
VmSize: 408 kB
VmLck: 0 kB
VmRSS: 208 kB
VmData: 24 kB
VmStk: 8 kB
VmExe: 368 kB
VmLib: 0 kB
SigPnd: 0000000000000000
SigBlk: 0000000000000000
SigIgn: 7fffffff0d8fc
SigCgt: 00000000280b2603
CapInh: 00000000fffffeff
CapPrm: 00000000ffffffff
CapEff: 00000000fffffeff

```

```

User PSW: 070de000 80414146
task: 004b6000 tss: 004b62d8 ksp: 004b7ca8 pt_regs: 004b7f68

```

```

User GPRS:
00000400 00000000 0000000b 7ffffa90
00000000 00000000 00000000 0045d9f4
0045cafc 7ffffa90 7ffffff18 0045cb08
00010400 804039e8 80403af8 7ffff8b0

```

```

User ACRS:
00000000 00000000 00000000 00000000
00000001 00000000 00000000 00000000
00000000 00000000 00000000 00000000
00000000 00000000 00000000 00000000

```

```

Kernel BackChain CallChain BackChain CallChain
 004b7ca8 8002bd0c 004b7d18 8002b92c
 004b7db8 8005cd50 004b7e38 8005d12a
 004b7f08 80019114

```

Showing among other things memory usage & status of some signals & the processes' registers from the kernel task\_structure as well as a backchain which may be useful if a process crashes in the kernel for some unknown reason.

### Some driver debugging techniques

#### =====

#### debug feature

Some of our drivers now support a "debug feature" in /proc/s390dbf see s390dbf.txt in the linux/Documentation directory for more info.

e.g.

to switch on the lcs "debug feature"

```
echo 5 > /proc/s390dbf/lcs/level
```

& then after the error occurred.

```
cat /proc/s390dbf/lcs/sprintf >/logfile
```

the logfile now contains some information which may help

tech support resolve a problem in the field.

## high level debugging network drivers

---

ifconfig is a quite useful command  
it gives the current state of network drivers.

If you suspect your network device driver is dead  
one way to check is type

ifconfig <network device>

e.g. tr0

You should see something like

```
tr0 Link encap:16/4 Mbps Token Ring (New) HWaddr 00:04:AC:20:8E:48
 inet addr:9.164.185.132 Bcast:9.164.191.255 Mask:255.255.224.0
 UP BROADCAST RUNNING MULTICAST MTU:2000 Metric:1
 RX packets:246134 errors:0 dropped:0 overruns:0 frame:0
 TX packets:5 errors:0 dropped:0 overruns:0 carrier:0
 collisions:0 txqueuelen:100
```

if the device doesn't say up

try

/etc/rc.d/init.d/network start

( this starts the network stack & hopefully calls ifconfig tr0 up ).

ifconfig looks at the output of /proc/net/dev & presents it in a more  
presentable form

Now ping the device from a machine in the same subnet.

if the RX packets count & TX packets counts don't increment you probably  
have problems.

next

cat /proc/net/arp

Do you see any hardware addresses in the cache if not you may have problems.

Next try

ping -c 5 <broadcast\_addr> i.e. the Bcast field above in the output of  
ifconfig. Do you see any replies from machines other than the local machine  
if not you may have problems. also if the TX packets count in ifconfig  
hasn't incremented either you have serious problems in your driver  
(e.g. the txbusy field of the network device being stuck on )  
or you may have multiple network devices connected.

## chandev

---

There is a new device layer for channel devices, some  
drivers e.g. lcs are registered with this layer.

If the device uses the channel device layer you'll be  
able to find what interrupts it uses & the current state  
of the device.

See the manpage chandev.8 & type cat /proc/chandev for more info.

Starting points for debugging scripting languages etc.

---

bash/sh

## Debugging390.txt

bash -x <scriptname>  
e.g. bash -x /usr/bin/bashbug  
displays the following lines as it executes them.  
+ MACHINE=i586  
+ OS=linux-gnu  
+ CC=gcc  
+ CFLAGS= -DPROGRAM='bash' -DHOSTTYPE='i586' -DOSTYPE='linux-gnu'  
-DMACHTYPE='i586-pc-linux-gnu' -DSHELL -DHAVE\_CONFIG\_H -I. -I. -I./lib -O2  
-pipe  
+ RELEASE=2.01  
+ PATCHLEVEL=1  
+ RELSTATUS=release  
+ MACHTYPE=i586-pc-linux-gnu

perl -d <scriptname> runs the perlscript in a fully interactive debugger  
<like gdb>.  
Type 'h' in the debugger for help.

for debugging java type  
jdb <filename> another fully interactive gdb style debugger.  
& type ? in the debugger for help.

### Dumptool & Lcrash ( lkcd )

=====

Michael Holzhau & others here at IBM have a fairly mature port of  
SGI's lcrash tool which allows one to look at kernel structures in a  
running kernel.

It also complements a tool called dumptool which dumps all the kernel's  
memory pages & registers to either a tape or a disk.  
This can be used by tech support or an ambitious end user do  
post mortem debugging of a machine like gdb core dumps.

Going into how to use this tool in detail will be explained  
in other documentation supplied by IBM with the patches & the  
lcrash homepage <http://oss.sgi.com/projects/lkcd/> & the lcrash manpage.

### How they work

-----

Lcrash is a perfectly normal program, however, it requires 2  
additional files, Kerntypes which is built using a patch to the  
linux kernel sources in the linux root directory & the System.map.

Kerntypes is an objectfile whose sole purpose in life  
is to provide stabs debug info to lcrash, to do this  
Kerntypes is built from kerntypes.c which just includes the most commonly  
referenced header files used when debugging, lcrash can then read the  
.stabs section of this file.

Debugging a live system it uses /dev/mem  
alternatively for post mortem debugging it uses the data  
collected by dumptool.



## Debugging390.txt

### SysRq

=====

This is now supported by linux for s/390 & z/Architecture.

To enable it do compile the kernel with

Kernel Hacking -> Magic SysRq Key Enabled

echo "1" > /proc/sys/kernel/sysrq

also type

echo "8" > /proc/sys/kernel/printk

To make printk output go to console.

On 390 all commands are prefixed with

^\_

e.g.

^\_t will show tasks.

^\_? or some unknown command will display help.

The sysrq key reading is very picky ( I have to type the keys in an xterm session & paste them into the x3270 console )

& it may be wise to predefine the keys as described in the VM hints above

This is particularly useful for syncing disks unmounting & rebooting if the machine gets partially hung.

Read Documentation/sysrq.txt for more info

### References:

=====

Enterprise Systems Architecture Reference Summary

Enterprise Systems Architecture Principles of Operation

Hartmut Penners s390 stack frame sheet.

IBM Mainframe Channel Attachment a technology brief from a CISCO webpage

Various bits of man & info pages of Linux.

Linux & GDB source.

Various info & man pages.

CMS Help on tracing commands.

Linux for s/390 Elf Application Binary Interface

Linux for z/Series Elf Application Binary Interface ( Both Highly Recommended )

z/Architecture Principles of Operation SA22-7832-00

Enterprise Systems Architecture/390 Reference Summary SA22-7209-01 & the

Enterprise Systems Architecture/390 Principles of Operation SA22-7201-05

### Special Thanks

=====

Special thanks to Neale Ferguson who maintains a much prettier HTML version of this page at

<http://penguinvm.princeton.edu/notes.html#Debug390>

Bob Grainger Stefan Bader & others for reporting bugs