

The OSD Standard

OSD (Object-Based Storage Device) is a T10 SCSI command set that is designed to provide efficient operation of input/output logical units that manage the allocation, placement, and accessing of variable-size data-storage containers, called objects. Objects are intended to contain operating system and application constructs. Each object has associated attributes attached to it, which are integral part of the object and provide metadata about the object. The standard defines some common obligatory attributes, but user attributes can be added as needed.

See: <http://www.t10.org/ftp/t10/drafts/osd2/> for the latest draft for OSD 2 or search the web for "OSD SCSI"

OSD in the Linux Kernel

osd-initiator:

The main component of OSD in Kernel is the osd-initiator library. Its main user is intended to be the pNFS-over-objects layout driver, which uses objects as its back-end data storage. Other clients are the other osd parts listed below.

osd-uld:

This is a SCSI ULD that registers for OSD type devices and provides a testing platform, both for the in-kernel initiator as well as connected targets. It currently has no useful user-mode API, though it could have if need be.

exofs:

Is an OSD based Linux file system. It uses the osd-initiator and osd-uld, to export a usable file system for users.
See Documentation/filesystems/exofs.txt for more details

osd target:

There are no current plans for an OSD target implementation in kernel. For all needs, a user-mode target that is based on the scsi tgt target framework is available from Ohio Supercomputer Center (OSC) at:
<http://www.open-osd.org/bin/view/Main/OscOsdProject>
There are several other target implementations. See <http://open-osd.org> for more links.

Files and Folders

This is the complete list of files included in this work:

```
include/scsi/
    osd_initiator.h    Main API for the initiator library
    osd_types.h        Common OSD types
    osd_sec.h          Security Manager API
    osd_protocol.h     Wire definitions of the OSD standard protocol
    osd_attributes.h   Wire definitions of OSD attributes

drivers/scsi/osd/
    osd_initiator.c    OSD-Initiator library implementation
    osd_uld.c          The OSD scsi ULD
    osd_ktest.{h,c}    In-kernel test suite (called by osd_uld)
    osd_debug.h        Some printk macros
    Makefile           For both in-tree and out-of-tree compilation
```

	osd.txt
Kconfig	Enables inclusion of the different pieces
osd_test.c	User-mode application to call the kernel tests

The OSD-Initiator Library

osd_initiator is a low level implementation of an osd initiator encoder. But even though, it should be intuitive and easy to use. Perhaps over time an higher lever will form that automates some of the more common recipes.

init/fini:

- osd_dev_init() associates a scsi_device with an osd_dev structure and initializes some global pools. This should be done once per scsi_device (OSD LUN). The osd_dev structure is needed for calling osd_start_request().
- osd_dev_fini() cleans up before a osd_dev/scsi_device destruction.

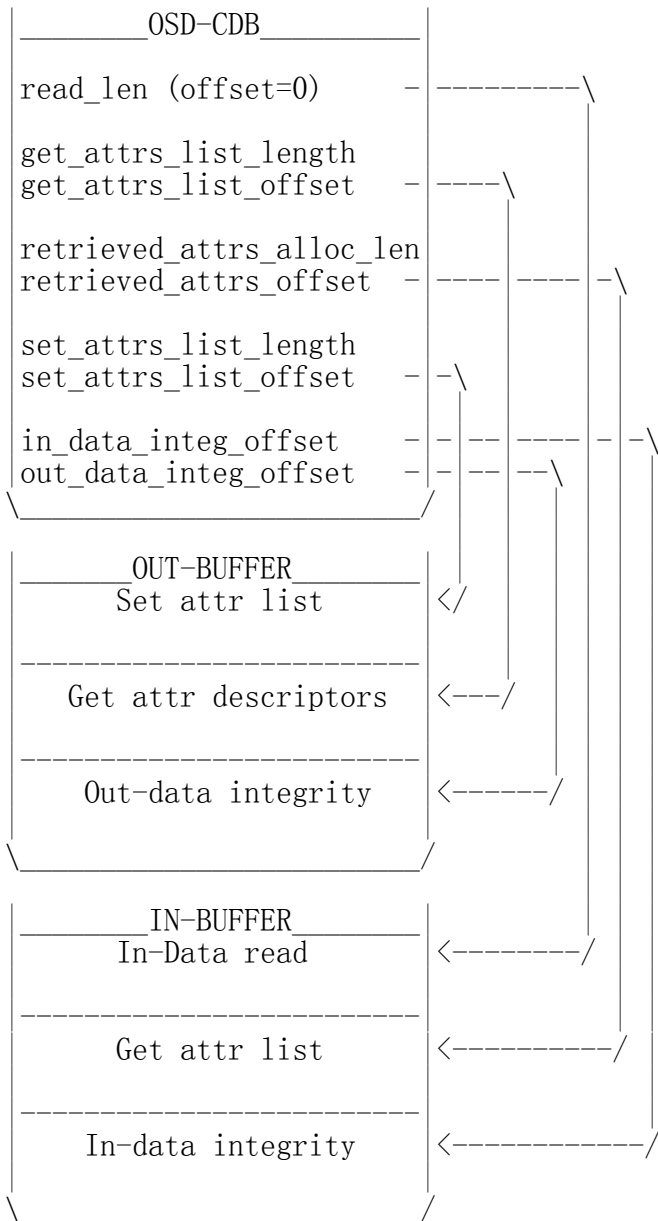
OSD commands encoding, execution, and decoding of results:

struct osd_request's is used to iteratively encode an OSD command and carry its state throughout execution. Each request goes through these stages:

- a. osd_start_request() allocates the request.
- b. Any of the osd_req_* methods is used to encode a request of the specified type.
- c. osd_req_add_{get,set}_attr_* may be called to add get/set attributes to the CDB. "List" or "Page" mode can be used exclusively. The attribute-list API can be called multiple times on the same request. However, only one attribute-page can be read, as mandated by the OSD standard.
- d. osd_finalize_request() computes offsets into the data-in and data-out buffers and signs the request using the provided capability key and integrity-check parameters.
- e. osd_execute_request() may be called to execute the request via the block layer and wait for its completion. The request can be executed asynchronously by calling the block layer API directly.
- f. After execution, osd_req_decode_sense() can be called to decode the request's sense information.
- g. osd_req_decode_get_attr() may be called to retrieve osd_add_get_attr_list() values.
- h. osd_end_request() must be called to deallocate the request and any resource associated with it. Note that osd_end_request cleans up the request at any stage and it must always be called after a successful osd_start_request().

osd_request's structure:

The OSD standard defines a complex structure of IO segments pointed to by members in the CDB. Up to 3 segments can be deployed in the IN-Buffer and up to 4 in the OUT-Buffer. The ASCII illustration below depicts a secure-read with associated get+set of attributes-lists. Other combinations very on the same basic theme. From no-segments-used up to all-segments-used.



A block device request can carry bidirectional payload by means of associating a bidi_read request with a main write-request. Each in/out request is described by a chain of BIOs associated with each request.

The CDB is of a SCSI VARLEN CDB format, as described by OSD standard.

The OSD standard also mandates alignment restrictions at start of each segment.

In the code, in struct `osd_request`, there are two `_osd_io_info` structures to describe the IN/OUT buffers above, two BIOs for the data payload and up to five `_osd_req_data_segment` structures to hold the different segments allocation and information.

Important: We have chosen to disregard the assumption that a BIO-chain (and the resulting sg-list) describes a linear memory buffer. Meaning only first and last scatter chain can be incomplete and all the middle chains are of `PAGE_SIZE`. For us, a scatter-gather-list, as its name implies and as used by the Networking layer, is to describe a vector of buffers that will be transferred to/from the

osd.txt

wire. It works very well with current iSCSI transport. iSCSI is currently the only deployed OSD transport. In the future we anticipate SAS and FC attached OSD devices as well.

The OSD Testing ULD

TODO: More user-mode control on tests.

Authors, Mailing list

Please communicate with us on any deployment of osd, whether using this code or not.

Any problems, questions, bug reports, lonely OSD nights, please email:
OSD Dev List <osd-dev@open-osd.org>

More up-to-date information can be found on:
<http://open-osd.org>

Boaz Harrosh <bharrosh@panasas.com>
Benny Halevy <bhalevy@panasas.com>

References

Weber, R., "SCSI Object-Based Storage Device Commands",
T10/1355-D ANSI/INCITS 400-2004,
<http://www.t10.org/ftp/t10/drafts/osd/osd-r10.pdf>

Weber, R., "SCSI Object-Based Storage Device Commands -2 (OSD-2)"
T10/1729-D, Working Draft, rev. 3
<http://www.t10.org/ftp/t10/drafts/osd2/osd2r03.pdf>