

SCSI EH

=====

This document describes SCSI midlayer error handling infrastructure. Please refer to Documentation/scsi/scsi_mid_low_api.txt for more information regarding SCSI midlayer.

TABLE OF CONTENTS

- [1] How SCSI commands travel through the midlayer and to EH
 - [1-1] struct scsi_cmnd
 - [1-2] How do scmd's get completed?
 - [1-2-1] Completing a scmd w/ scsi_done
 - [1-2-2] Completing a scmd w/ timeout
 - [1-3] How EH takes over
- [2] How SCSI EH works
 - [2-1] EH through fine-grained callbacks
 - [2-1-1] Overview
 - [2-1-2] Flow of scmds through EH
 - [2-1-3] Flow of control
 - [2-2] EH through transportt->eh_strategy_handler()
 - [2-2-1] Pre transportt->eh_strategy_handler() SCSI midlayer conditions
 - [2-2-2] Post transportt->eh_strategy_handler() SCSI midlayer conditions
 - [2-2-3] Things to consider

[1] How SCSI commands travel through the midlayer and to EH

[1-1] struct scsi_cmnd

Each SCSI command is represented with struct scsi_cmnd (== scmd). A scmd has two list_head's to link itself into lists. The two are scmd->list and scmd->eh_entry. The former is used for free list or per-device allocated scmd list and not of much interest to this EH discussion. The latter is used for completion and EH lists and unless otherwise stated scmds are always linked using scmd->eh_entry in this discussion.

[1-2] How do scmd's get completed?

Once LLDD gets hold of a scmd, either the LLDD will complete the command by calling scsi_done callback passed from midlayer when invoking hostt->queuecommand() or SCSI midlayer will time it out.

[1-2-1] Completing a scmd w/ scsi_done

For all non-EH commands, scsi_done() is the completion callback. It does the following.

1. Delete timeout timer. If it fails, it means that timeout timer has expired and is going to finish the command. Just return.
2. Link scmd to per-cpu scsi_done_q using scmd->eh_entry

3. Raise SCSI_SOFTIRQ

SCSI_SOFTIRQ handler `scsi_softirq` calls `scsi_decide_disposition()` to determine what to do with the command. `scsi_decide_disposition()` looks at the `scmd->result` value and sense data to determine what to do with the command.

- SUCCESS
 `scsi_finish_command()` is invoked for the command. The function does some maintenance chores and notify completion by calling `scmd->done()` callback, which, for fs requests, would be HLD completion callback - `sd:sd_rw_intr`, `sr:rw_intr`, `st:st_intr`.
- NEEDS_RETRY
- ADD_TO_MLQUEUE
 `scmd` is requeued to blk queue.
- otherwise
 `scsi_eh_scmd_add(scmd, 0)` is invoked for the command. See [1-3] for details of this function.

[1-2-2] Completing a `scmd` w/ timeout

The timeout handler is `scsi_times_out()`. When a timeout occurs, this function

1. invokes optional `hostt->eh_timed_out()` callback. Return value can be one of
 - EH_HANDLED
 This indicates that `eh_timed_out()` dealt with the timeout. The `scmd` is passed to `__scsi_done()` and thus linked into per-cpu `scsi_done_q`. Normal command completion described in [1-2-1] follows.
 - EH_RESET_TIMER
 This indicates that more time is required to finish the command. Timer is restarted. This action is counted as a retry and only allowed `scmd->allowed + 1(!)` times. Once the limit is reached, action for `EH_NOT_HANDLED` is taken instead.

 NOTE This action is racy as the LLDD could finish the `scmd` after the timeout has expired but before it's added back. In such cases, `scsi_done()` would think that timeout has occurred and return without doing anything. We lose completion and the command will time out again.
 - EH_NOT_HANDLED
 This is the same as when `eh_timed_out()` callback doesn't exist. Step #2 is taken.
2. `scsi_eh_scmd_add(scmd, SCSI_EH_CANCEL_CMD)` is invoked for the command. See [1-3] for more information.

[1-3] How EH takes over

scmds enter EH via `scsi_eh_scmd_add()`, which does the following.

1. Turns on `scmd->eh_eflags` as requested. It's 0 for error completions and `SCSI_EH_CANCEL_CMD` for timeouts.
2. Links `scmd->eh_entry` to `shost->eh_cmd_q`
3. Sets `SHOST_RECOVERY` bit in `shost->shost_state`
4. Increments `shost->host_failed`
5. Wakes up SCSI EH thread if `shost->host_busy == shost->host_failed`

As can be seen above, once any `scmd` is added to `shost->eh_cmd_q`, `SHOST_RECOVERY` `shost_state` bit is turned on. This prevents any new `scmd` to be issued from blk queue to the host; eventually, all `scmds` on the host either complete normally, fail and get added to `eh_cmd_q`, or time out and get added to `shost->eh_cmd_q`.

If all `scmds` either complete or fail, the number of in-flight `scmds` becomes equal to the number of failed `scmds` - i.e. `shost->host_busy == shost->host_failed`. This wakes up SCSI EH thread. So, once woken up, SCSI EH thread can expect that all in-flight commands have failed and are linked on `shost->eh_cmd_q`.

Note that this does not mean lower layers are quiescent. If a LLDD completed a `scmd` with error status, the LLDD and lower layers are assumed to forget about the `scmd` at that point. However, if a `scmd` has timed out, unless `host->eh_timed_out()` made lower layers forget about the `scmd`, which currently no LLDD does, the command is still active as long as lower layers are concerned and completion could occur at any time. Of course, all such completions are ignored as the timer has already expired.

We'll talk about how SCSI EH takes actions to abort - make LLDD forget about - timed out `scmds` later.

[2] How SCSI EH works

LLDD's can implement SCSI EH actions in one of the following two ways.

- Fine-grained EH callbacks
LLDD can implement fine-grained EH callbacks and let SCSI midlayer drive error handling and call appropriate callbacks. This will be discussed further in [2-1].
- `eh_strategy_handler()` callback
This is one big callback which should perform whole error handling. As such, it should do all chores SCSI midlayer performs during recovery. This will be discussed in [2-2].

Once recovery is complete, SCSI EH resumes normal operation by calling `scsi_restart_operations()`, which

1. Checks if door locking is needed and locks door.
2. Clears `SHOST_RECOVERY` `shost_state` bit
3. Wakes up waiters on `shost->host_wait`. This occurs if someone calls `scsi_block_when_processing_errors()` on the host.
(*QUESTION* why is it needed? All operations will be blocked anyway after it reaches blk queue.)
4. Kicks queues in all devices on the host in the asses

[2-1] EH through fine-grained callbacks

[2-1-1] Overview

If `eh_strategy_handler()` is not present, SCSI midlayer takes charge of driving error handling. EH's goals are two - make LLDD, host and device forget about timed out `scmds` and make them ready for new commands. A `scmd` is said to be recovered if the `scmd` is forgotten by lower layers and lower layers are ready to process or fail the `scmd` again.

To achieve these goals, EH performs recovery actions with increasing severity. Some actions are performed by issuing SCSI commands and others are performed by invoking one of the following fine-grained hostt EH callbacks. Callbacks may be omitted and omitted ones are considered to fail always.

```
int (* eh_abort_handler)(struct scsi_cmnd *);
int (* eh_device_reset_handler)(struct scsi_cmnd *);
int (* eh_bus_reset_handler)(struct scsi_cmnd *);
int (* eh_host_reset_handler)(struct scsi_cmnd *);
```

Higher-severity actions are taken only when lower-severity actions cannot recover some of failed `scmds`. Also, note that failure of the highest-severity action means EH failure and results in offlining of all unrecovered devices.

During recovery, the following rules are followed

- Recovery actions are performed on failed `scmds` on the to do list, `eh_work_q`. If a recovery action succeeds for a `scmd`, recovered `scmds` are removed from `eh_work_q`.

Note that single recovery action on a `scmd` can recover multiple `scmds`. e.g. resetting a device recovers all failed `scmds` on the device.

- Higher severity actions are taken iff `eh_work_q` is not empty after lower severity actions are complete.

- EH reuses failed cmds to issue commands for recovery. For timed-out cmds, SCSI EH ensures that LLDD forgets about a cmd before reusing it for EH commands.

When a cmd is recovered, the cmd is moved from eh_work_q to EH local eh_done_q using scsi_eh_finish_cmd(). After all cmds are recovered (eh_work_q is empty), scsi_eh_flush_done_q() is invoked to either retry or error-finish (notify upper layer of failure) recovered cmds.

cmds are retried iff its sdev is still online (not offlined during EH), REQ_FAILFAST is not set and ++cmd->retries is less than cmd->allowed.

[2-1-2] Flow of cmds through EH

1. Error completion / time out
 ACTION: scsi_eh_cmd_add() is invoked for cmd
 - set cmd->eh_eflags
 - add cmd to shost->eh_cmd_q
 - set SHOST_RECOVERY
 - shost->host_failed++
 LOCKING: shost->host_lock
2. EH starts
 ACTION: move all cmds to EH's local eh_work_q. shost->eh_cmd_q is cleared.
 LOCKING: shost->host_lock (not strictly necessary, just for consistency)
3. cmd recovered
 ACTION: scsi_eh_finish_cmd() is invoked to EH-finish cmd
 - shost->host_failed--
 - clear cmd->eh_eflags
 - scsi_setup_cmd_retry()
 - move from local eh_work_q to local eh_done_q
 LOCKING: none
4. EH completes
 ACTION: scsi_eh_flush_done_q() retries cmds or notifies upper layer of failure.
 - cmd is removed from eh_done_q and cmd->eh_entry is cleared
 - if retry is necessary, cmd is requeued using scsi_queue_insert()
 - otherwise, scsi_finish_command() is invoked for cmd
 LOCKING: queue or finish function performs appropriate locking

[2-1-3] Flow of control

EH through fine-grained callbacks start from scsi_unjam_host().

<<scsi_unjam_host>>

1. Lock shost->host_lock, splice_init shost->eh_cmd_q into local

scsi_eh.txt

eh_work_q and unlock host_lock. Note that shost->eh_cmd_q is cleared by this action.

2. Invoke scsi_eh_get_sense.

<<scsi_eh_get_sense>>

This action is taken for each error-completed (!SCSI_EH_CANCEL_CMD) commands without valid sense data. Most SCSI transports/LLDDs automatically acquire sense data on command failures (autosense). Autosense is recommended for performance reasons and as sense information could get out of sync inbetween occurrence of CHECK CONDITION and this action.

Note that if autosense is not supported, scmd->sense_buffer contains invalid sense data when error-completing the scmd with scsi_done(). scsi_decide_disposition() always returns FAILED in such cases thus invoking SCSI EH. When the scmd reaches here, sense data is acquired and scsi_decide_disposition() is called again.

1. Invoke scsi_request_sense() which issues REQUEST_SENSE command. If fails, no action. Note that taking no action causes higher-severity recovery to be taken for the scmd.

2. Invoke scsi_decide_disposition() on the scmd

- SUCCESS
scmd->retries is set to scmd->allowed preventing scsi_eh_flush_done_q() from retrying the scmd and scsi_eh_finish_cmd() is invoked.
- NEEDS_RETRY
scsi_eh_finish_cmd() invoked
- otherwise
No action.

3. If !list_empty(&eh_work_q), invoke scsi_eh_abort_cmds().

<<scsi_eh_abort_cmds>>

This action is taken for each timed out command. hostt->eh_abort_handler() is invoked for each scmd. The handler returns SUCCESS if it has succeeded to make LLDD and all related hardware forget about the scmd.

If a timedout scmd is successfully aborted and the sdev is either offline or ready, scsi_eh_finish_cmd() is invoked for the scmd. Otherwise, the scmd is left in eh_work_q for higher-severity actions.

Note that both offline and ready status mean that the sdev is ready to process new cmds, where processing also implies immediate failing; thus, if a sdev is in one of the two states, no further recovery action is needed.

scsi_eh.txt

Device readiness is tested using `scsi_eh_tur()` which issues `TEST_UNIT_READY` command. Note that the `scmd` must have been aborted successfully before reusing it for `TEST_UNIT_READY`.

4. If `!list_empty(&eh_work_q)`, invoke `scsi_eh_ready_devs()`

<<`scsi_eh_ready_devs`>>

This function takes four increasingly more severe measures to make failed `sdevs` ready for new commands.

1. Invoke `scsi_eh_stu()`

<<`scsi_eh_stu`>>

For each `sdev` which has failed `scmds` with valid sense data of which `scsi_check_sense()`'s verdict is `FAILED`, `START_STOP_UNIT` command is issued w/ `start=1`. Note that as we explicitly choose error-completed `scmds`, it is known that lower layers have forgotten about the `scmd` and we can reuse it for `STU`.

If `STU` succeeds and the `sdev` is either offline or ready, all failed `scmds` on the `sdev` are `EH-finished` with `scsi_eh_finish_cmd()`.

NOTE If `hostt->eh_abort_handler()` isn't implemented or failed, we may still have timed out `scmds` at this point and `STU` doesn't make lower layers forget about those `scmds`. Yet, this function `EH-finish` all `scmds` on the `sdev` if `STU` succeeds leaving lower layers in an inconsistent state. It seems that `STU` action should be taken only when a `sdev` has no timed out `scmd`.

2. If `!list_empty(&eh_work_q)`, invoke `scsi_eh_bus_device_reset()`.

<<`scsi_eh_bus_device_reset`>>

This action is very similar to `scsi_eh_stu()` except that, instead of issuing `STU`, `hostt->eh_device_reset_handler()` is used. Also, as we're not issuing `SCSI` commands and resetting clears all `scmds` on the `sdev`, there is no need to choose error-completed `scmds`.

3. If `!list_empty(&eh_work_q)`, invoke `scsi_eh_bus_reset()`

<<`scsi_eh_bus_reset`>>

`hostt->eh_bus_reset_handler()` is invoked for each channel with failed `scmds`. If bus reset succeeds, all failed `scmds` on all ready or offline `sdevs` on the channel are `EH-finished`.

4. If `!list_empty(&eh_work_q)`, invoke `scsi_eh_host_reset()`

<<scsi_eh_host_reset>>

This is the last resort. `hostt->eh_host_reset_handler()` is invoked. If host reset succeeds, all failed `scmds` on all ready or offline `sdevs` on the host are EH-finished.

5. If `!list_empty(&eh_work_q)`, invoke `scsi_eh_offline_sdevs()`

<<scsi_eh_offline_sdevs>>

Take all `sdevs` which still have unrecovered `scmds` offline and EH-finish the `scmds`.

5. Invoke `scsi_eh_flush_done_q()`.

<<scsi_eh_flush_done_q>>

At this point all `scmds` are recovered (or given up) and put on `eh_done_q` by `scsi_eh_finish_cmd()`. This function flushes `eh_done_q` by either retrying or notifying upper layer of failure of the `scmds`.

[2-2] EH through `transportt->eh_strategy_handler()`

`transportt->eh_strategy_handler()` is invoked in the place of `scsi_unjam_host()` and it is responsible for whole recovery process. On completion, the handler should have made lower layers forget about all failed `scmds` and either ready for new commands or offline. Also, it should perform SCSI EH maintenance chores to maintain integrity of SCSI midlayer. IOW, of the steps described in [2-1-2], all steps except for #1 must be implemented by `eh_strategy_handler()`.

[2-2-1] Pre `transportt->eh_strategy_handler()` SCSI midlayer conditions

The following conditions are true on entry to the handler.

- Each failed `scmd`'s `eh_flags` field is set appropriately.
- Each failed `scmd` is linked on `scmd->eh_cmd_q` by `scmd->eh_entry`.
- `SHOST_RECOVERY` is set.
- `shost->host_failed == shost->host_busy`

[2-2-2] Post `transportt->eh_strategy_handler()` SCSI midlayer conditions

The following conditions must be true on exit from the handler.

- `shost->host_failed` is zero.
- Each `scmd`'s `eh_eflags` field is cleared.
- Each `scmd` is in such a state that `scsi_setup_cmd_retry()` on the

scsi_eh.txt

scmd doesn't make any difference.

- shost->eh_cmd_q is cleared.
- Each scmd->eh_entry is cleared.
- Either scsi_queue_insert() or scsi_finish_command() is called on each scmd. Note that the handler is free to use scmd->retries and ->allowed to limit the number of retries.

[2-2-3] Things to consider

- Know that timed out cmds are still active on lower layers. Make lower layers forget about them before doing anything else with those cmds.
- For consistency, when accessing/modifying shost data structure, grab shost->host_lock.
- On completion, each failed sdev must have forgotten about all active cmds.
- On completion, each failed sdev must be ready for new commands or offline.

--

Tejun Heo
htejun@gmail.com
11th September 2005