

vwsnd - Sound driver for the Silicon Graphics 320 and 540 Visual Workstations' onboard audio.

Copyright 1999 Silicon Graphics, Inc. All rights reserved.

At the time of this writing, March 1999, there are two models of Visual Workstation, the 320 and the 540. This document only describes those models. Future Visual Workstation models may have different sound capabilities, and this driver will probably not work on those boxes.

The Visual Workstation has an Analog Devices AD1843 "SoundComm" audio codec chip. The AD1843 is accessed through the Cobalt I/O ASIC, also known as Lithium. This driver programs both chips.

QUICK CONFIGURATION

```
# insmod soundcore
# insmod vwsnd
```

I/O CONNECTIONS

On the Visual Workstation, only three of the AD1843 inputs are hooked up. The analog line in jacks are connected to the AD1843's AUX1 input. The CD audio lines are connected to the AD1843's AUX2 input. The microphone jack is connected to the AD1843's MIC input. The mic jack is mono, but the signal is delivered to both the left and right MIC inputs. You can record in stereo from the mic input, but you will get the same signal on both channels (within the limits of A/D accuracy). Full scale on the Line input is +/- 2.0 V. Full scale on the MIC input is 20 dB less, or +/- 0.2 V.

The AD1843's LOUT1 outputs are connected to the Line Out jacks. The AD1843's HPOUT outputs are connected to the speaker/headphone jack. LOUT2 is not connected. Line out's maximum level is +/- 2.0 V peak to peak. The speaker/headphone out's maximum is +/- 4.0 V peak to peak.

The AD1843's PCM input channel and one of its output channels (DAC1) are connected to Lithium. The other output channel (DAC2) is not connected.

CAPABILITIES

The AD1843 has PCM input and output (Pulse Code Modulation, also known as wavetable). PCM input and output can be mono or stereo in any of four formats. The formats are 16 bit signed and 8 bit unsigned, u-Law, and A-Law format. Any sample rate from 4 KHz to 49 KHz is available, in 1 Hz increments.

The AD1843 includes an analog mixer that can mix all three input signals (line, mic and CD) into the analog outputs. The mixer has a separate gain control and mute switch for each input.

There are two outputs, line out and speaker/headphone out. They always produce the same signal, and the speaker always has 3 dB more gain than the line out. The speaker/headphone output can be muted, but this driver does not export that function.

The hardware can sync audio to the video clock, but this driver does not have a way to specify syncing to video.

PROGRAMMING

This section explains the API supported by the driver. Also see the Open Sound Programming Guide at <http://www.opensound.com/pguide/>. This section assumes familiarity with that document.

The driver has two interfaces, an I/O interface and a mixer interface. There is no MIDI or sequencer capability.

PROGRAMMING PCM I/O

The I/O interface is usually accessed as `/dev/audio` or `/dev/dsp`. Using the standard Open Sound System (OSS) `ioctl` calls, the sample rate, number of channels, and sample format may be set within the limitations described above. The driver supports triggering. It also supports getting the input and output pointers with one-sample accuracy.

The `SNDCCTL_DSP_GETCAP` `ioctl` returns these capabilities.

`DSP_CAP_DUPLEX` - driver supports full duplex.

`DSP_CAP_TRIGGER` - driver supports triggering.

`DSP_CAP_REALTIME` - values returned by `SNDCCTL_DSP_GETIPTR` and `SNDCCTL_DSP_GETOPTR` are accurate to a few samples.

Memory mapping (`mmap`) is not implemented.

The driver permits subdivided fragment sizes from 64 to 4096 bytes. The number of fragments can be anything from 3 fragments to however many fragments fit into 124 kilobytes. It is up to the user to determine how few/small fragments can be used without introducing glitches with a given workload. Linux is not realtime, so we can't promise anything. (sigh...)

When this driver is switched into or out of mu-Law or A-Law mode on output, it may produce an audible click. This is unavoidable. To prevent clicking, use signed 16-bit mode instead, and convert from mu-Law or A-Law format in software.

PROGRAMMING THE MIXER INTERFACE

The mixer interface is usually accessed as `/dev/mixer`. It is accessed

through ioctls. The mixer allows the application to control gain or mute several audio signal paths, and also allows selection of the recording source.

Each of the constants described here can be read using the `MIXER_READ(SOUND_MIXER_xxx)` ioctl. Those that are not read-only can also be written using the `MIXER_WRITE(SOUND_MIXER_xxx)` ioctl. In most cases, `<sys/soundcard.h>` defines constants `SOUND_MIXER_READ_xxx` and `SOUND_MIXER_WRITE_xxx` which work just as well.

`SOUND_MIXER_CAPS` Read-only

This is a mask of optional driver capabilities that are implemented. This driver's only capability is `SOUND_CAP_EXCL_INPUT`, which means that only one recording source can be active at a time.

`SOUND_MIXER_DEVMASK` Read-only

This is a mask of the sound channels. This driver's channels are PCM, LINE, MIC, CD, and RECLEV.

`SOUND_MIXER_STEREODEVS` Read-only

This is a mask of which sound channels are capable of stereo. All channels are capable of stereo. (But see caveat on MIC input in I/O CONNECTIONS section above).

`SOUND_MIXER_OUTMASK` Read-only

This is a mask of channels that route inputs through to outputs. Those are LINE, MIC, and CD.

`SOUND_MIXER_RECMASK` Read-only

This is a mask of channels that can be recording sources. Those are PCM, LINE, MIC, CD.

`SOUND_MIXER_PCM` Default: 0x5757 (0 dB)

This is the gain control for PCM output. The left and right channel gain are controlled independently. This gain control has 64 levels, which range from -82.5 dB to +12.0 dB in 1.5 dB steps. Those 64 levels are mapped onto 100 levels at the ioctl, see below.

`SOUND_MIXER_LINE` Default: 0x4a4a (0 dB)

This is the gain control for mixing the Line In source into the outputs. The left and right channel gain are controlled independently. This gain control has 32 levels, which range from -34.5 dB to +12.0 dB in 1.5 dB steps. Those 32 levels are mapped onto 100 levels at the ioctl, see below.

`SOUND_MIXER_MIC` Default: 0x4a4a (0 dB)

This is the gain control for mixing the MIC source into the outputs. The left and right channel gain are controlled independently. This

vwsnd..txt

gain control has 32 levels, which range from -34.5 dB to +12.0 dB in 1.5 dB steps. Those 32 levels are mapped onto 100 levels at the ioctl, see below.

SOUND_MIXER_CD Default: 0x4a4a (0 dB)

This is the gain control for mixing the CD audio source into the outputs. The left and right channel gain are controlled independently. This gain control has 32 levels, which range from -34.5 dB to +12.0 dB in 1.5 dB steps. Those 32 levels are mapped onto 100 levels at the ioctl, see below.

SOUND_MIXER_RECLEV Default: 0 (0 dB)

This is the gain control for PCM input (RECORDing LEVel). The left and right channel gain are controlled independently. This gain control has 16 levels, which range from 0 dB to +22.5 dB in 1.5 dB steps. Those 16 levels are mapped onto 100 levels at the ioctl, see below.

SOUND_MIXER_RECSRC Default: SOUND_MASK_LINE

This is a mask of currently selected PCM input sources (RECORDing SouRCes). Because the AD1843 can only have a single recording source at a time, only one bit at a time can be set in this mask. The allowable values are SOUND_MASK_PCM, SOUND_MASK_LINE, SOUND_MASK_MIC, or SOUND_MASK_CD. Selecting SOUND_MASK_PCM sets up internal resampling which is useful for loopback testing and for hardware sample rate conversion. But software sample rate conversion is probably faster, so I don't know how useful that is.

SOUND_MIXER_OUTSRC DEFAULT: SOUND_MASK_LINE|SOUND_MASK_MIC|SOUND_MASK_CD

This is a mask of sources that are currently passed through to the outputs. Those sources whose bits are not set are muted.

GAIN CONTROL

There are five gain controls listed above. Each has 16, 32, or 64 steps. Each control has 1.5 dB of gain per step. Each control is stereo.

The OSS defines the argument to a channel gain ioctl as having two components, left and right, each of which ranges from 0 to 100. The two components are packed into the same word, with the left side gain in the least significant byte, and the right side gain in the second least significant byte. In C, we would say this.

```
#include <assert.h>
```

```
...
```

```
assert(leftgain >= 0 && leftgain <= 100);
assert(rightgain >= 0 && rightgain <= 100);
arg = leftgain | rightgain << 8;
```

So each OSS gain control has 101 steps. But the hardware has 16, 32, or 64 steps. The hardware steps are spread across the 101 OSS steps nearly evenly. The conversion formulas are like this, given N equals 16, 32, or 64.

```
int round = N/2 - 1;
OSS_gain_steps = (hw_gain_steps * 100 + round) / (N - 1);
hw_gain_steps = (OSS_gain_steps * (N - 1) + round) / 100;
```

Here is a snippet of C code that will return the left and right gain of any channel in dB. Pass it one of the predefined gain_desc_t structures to access any of the five channels' gains.

```
typedef struct gain_desc {
    float min_gain;
    float gain_step;
    int nbits;
    int chan;
} gain_desc_t;

const gain_desc_t gain_pcm    = { -82.5, 1.5, 6, SOUND_MIXER_PCM    };
const gain_desc_t gain_line   = { -34.5, 1.5, 5, SOUND_MIXER_LINE   };
const gain_desc_t gain_mic    = { -34.5, 1.5, 5, SOUND_MIXER_MIC    };
const gain_desc_t gain_cd     = { -34.5, 1.5, 5, SOUND_MIXER_CD     };
const gain_desc_t gain_reclev = {  0.0, 1.5, 4, SOUND_MIXER_RECLEV  };

int get_gain_dB(int fd, const gain_desc_t *gp,
                float *left, float *right)
{
    int word;
    int lg, rg;
    int mask = (1 << gp->nbits) - 1;

    if (ioctl(fd, MIXER_READ(gp->chan), &word) != 0)
        return -1; /* fail */
    lg = word & 0xFF;
    rg = word >> 8 & 0xFF;
    lg = (lg * mask + mask / 2) / 100;
    rg = (rg * mask + mask / 2) / 100;
    *left = gp->min_gain + gp->gain_step * lg;
    *right = gp->min_gain + gp->gain_step * rg;
    return 0;
}
```

And here is the corresponding routine to set a channel's gain in dB.

```
int set_gain_dB(int fd, const gain_desc_t *gp, float left, float right)
{
    float max_gain =
        gp->min_gain + (1 << gp->nbits) * gp->gain_step;
    float round = gp->gain_step / 2;
    int mask = (1 << gp->nbits) - 1;
    int word;
    int lg, rg;
```

```

                                vwsnd..txt
if (left < gp->min_gain || right < gp->min_gain)
    return EINVAL;
lg = (left - gp->min_gain + round) / gp->gain_step;
rg = (right - gp->min_gain + round) / gp->gain_step;
if (lg >= (1 << gp->nbits) || rg >= (1 << gp->nbits))
    return EINVAL;
lg = (100 * lg + mask / 2) / mask;
rg = (100 * rg + mask / 2) / mask;
word = lg | rg << 8;

return ioctl(fd, MIXER_WRITE(gp->chan), &word);
}

```