

Started Oct 1999 by Kanoj Sarcar <kanojsarcar@yahoo.com>

The intent of this file is to have an uptodate, running commentary from different people about how locking and synchronization is done in the Linux vm code.

page_table_lock & mmap_sem

Page stealers pick processes out of the process pool and scan for the best process to steal pages from. To guarantee the existence of the victim mm, a mm_count inc and a mmdrop are done in swap_out(). Page stealers hold kernel_lock to protect against a bunch of races. The vma list of the victim mm is also scanned by the stealer, and the page_table_lock is used to preserve list sanity against the process adding/deleting to the list. This also guarantees existence of the vma. Vma existence is not guaranteed once try_to_swap_out() drops the page_table_lock. To guarantee the existence of the underlying file structure, a get_file is done before the swapout() method is invoked. The page passed into swapout() is guaranteed not to be reused for a different purpose because the page reference count due to being present in the user's pte is not released till after swapout() returns.

Any code that modifies the vmlist, or the vm_start/vm_end/vm_flags:VM_LOCKED/vm_next of any vma *in the list* must prevent kswapd from looking at the chain.

The rules are:

1. To scan the vmlist (look but don't touch) you must hold the mmap_sem with read bias, i.e. down_read(&mm->mmap_sem)
2. To modify the vmlist you need to hold the mmap_sem with read&write bias, i.e. down_write(&mm->mmap_sem) *AND* you need to take the page_table_lock.
3. The swapper takes _just_ the page_table_lock, this is done because the mmap_sem can be an extremely long lived lock and the swapper just cannot sleep on that.
4. The exception to this rule is expand_stack, which just takes the read lock and the page_table_lock, this is ok because it doesn't really modify fields anybody relies on.
5. You must be able to guarantee that while holding page_table_lock or page_table_lock of mm A, you will not try to get either lock for mm B.

The caveats are:

1. find_vma() makes use of, and updates, the mmap_cache pointer hint. The update of mmap_cache is racy (page stealer can race with other code that invokes find_vma with mmap_sem held), but that is okay, since it is a hint. This can be fixed, if desired, by having find_vma grab the page_table_lock.

Code that add/delete elements from the vmlist chain are

1. callers of insert_vm_struct
2. callers of merge_segments
3. callers of avl_remove

locking..txt

Code that changes `vm_start/vm_end/vm_flags:VM_LOCKED` of `vma`'s on the list:

1. `expand_stack`
2. `mprotect`
3. `mlock`
4. `mremap`

It is advisable that changes to `vm_start/vm_end` be protected, although in some cases it is not really needed. Eg, `vm_start` is modified by `expand_stack()`, it is hard to come up with a destructive scenario without having the `vmlist` protection in this case.

The `page_table_lock` nests with the inode `i_mmap_lock` and the `kmem` cache `c_spinlock` spinlocks. This is okay, since the `kmem` code asks for pages after dropping `c_spinlock`. The `page_table_lock` also nests with `pagecache_lock` and `pagemap_lru_lock` spinlocks, and no code asks for memory with these locks held.

The `page_table_lock` is grabbed while holding the `kernel_lock` spinning monitor.

The `page_table_lock` is a spin lock.

Note: PTL can also be used to guarantee that no new clones using the `mm` start up ... this is a loose form of stability on `mm_users`. For example, it is used in `copy_mm` to protect against a racing `tlb_gather_mmu` single address space optimization, so that the `zap_page_range` (from `truncate`) does not lose sending `ipi`'s to cloned threads that might be spawned underneath it and go to user mode to drag in `pte`'s into `tlbs`.

swap_lock

The swap devices are chained in priority order from the "swap_list" header. The "swap_list" is used for the round-robin swaphandle allocation strategy. The #free swaphandles is maintained in "nr_swap_pages". These two together are protected by the `swap_lock`.

The `swap_lock` also protects all the device reference counts on the corresponding swaphandles, maintained in the "swap_map" array, and the "highest_bit" and "lowest_bit" fields.

The `swap_lock` is a spinlock, and is never acquired from `intr` level.

To prevent races between swap space deletion or `async` readahead swpins deciding whether a swap handle is being used, ie worthy of being read in from disk, and an `unmap` -> `swap_free` making the handle unused, the swap delete and readahead code grabs a temp reference on the swaphandle to prevent warning messages from `swap_duplicate` <- `read_swap_cache_async`.

Swap cache locking

Pages are added into the swap cache with `kernel_lock` held, to make sure that multiple pages are not being added (and hence lost) by associating all of them with the same swaphandle.

Pages are guaranteed not to be removed from the scache if the page is "shared": ie, other processes hold reference on the page or the associated

locking..txt

swap handle. The only code that does not follow this rule is `shrink_mmap`, which deletes pages from the swap cache if no process has a reference on the page (multiple processes might have references on the corresponding swap handle though). `lookup_swap_cache()` races with `shrink_mmap`, when establishing a reference on a scache page, so, it must check whether the page it located is still in the swapcache, or `shrink_mmap` deleted it. (This race is due to the fact that `shrink_mmap` looks at the page ref count with `pagecache_lock`, but then drops `pagecache_lock` before deleting the page from the scache).

`do_wp_page` and `do_swap_page` have MP races in them while trying to figure out whether a page is "shared", by looking at the `page_count` + `swap_count`. To preserve the sum of the counts, the page lock `_must_` be acquired before calling `is_page_shared` (else processes might switch their `swap_count` refs to the page count refs, after the page count ref has been snapshotted).

Swap device deletion code currently breaks all the scache assumptions, since it grabs neither `mmap_sem` nor `page_table_lock`.