

Last reviewed: 06/02/2009

HP iLO2 NMI Watchdog Driver  
NMI sourcing for iLO2 based ProLiant Servers  
Documentation and Driver by  
Thomas Mingarelli <thomas.mingarelli@hp.com>

The HP iLO2 NMI Watchdog driver is a kernel module that provides basic watchdog functionality and the added benefit of NMI sourcing. Both the watchdog functionality and the NMI sourcing capability need to be enabled by the user. Remember that the two modes are not dependant on one another. A user can have the NMI sourcing without the watchdog timer and vice-versa.

Watchdog functionality is enabled like any other common watchdog driver. That is, an application needs to be started that kicks off the watchdog timer. A basic application exists in the Documentation/watchdog/src directory called watchdog-test.c. Simply compile the C file and kick it off. If the system gets into a bad state and hangs, the HP ProLiant iLO 2 timer register will not be updated in a timely fashion and a hardware system reset (also known as an Automatic Server Recovery (ASR)) event will occur.

The hpwdt driver also has four (4) module parameters. They are the following:

soft\_margin - allows the user to set the watchdog timer value  
allow\_kdump - allows the user to save off a kernel dump image after an NMI  
nowayout - basic watchdog parameter that does not allow the timer to be restarted or an impending ASR to be escaped.  
priority - determines whether or not the hpwdt driver is first on the die\_notify list to handle NMIs or last. The default value for this module parameter is 0 or LAST. If the user wants to enable NMI sourcing then reload the hpwdt driver with priority=1 (and boot with nmi\_watchdog=0).

NOTE: More information about watchdog drivers in general, including the ioctl interface to /dev/watchdog can be found in Documentation/watchdog/watchdog-api.txt and Documentation/IPMI.txt.

The priority parameter was introduced due to other kernel software that relied on handling NMIs (like oprofile). Keeping hpwdt's priority at 0 (or LAST) enables the users of NMIs for non critical events to be work as expected.

The NMI sourcing capability is disabled by default due to the inability to distinguish between "NMI Watchdog Ticks" and "HW generated NMI events" in the Linux kernel. What this means is that the hpwdt nmi handler code is called each time the NMI signal fires off. This could amount to several thousands of NMIs in a matter of seconds. If a user sees the Linux kernel's "dazed and confused" message in the logs or if the system gets into a hung state, then the hpwdt driver can be reloaded with the "priority" module parameter set (priority=1).

1. If the kernel has not been booted with nmi\_watchdog turned off then edit /boot/grub/menu.lst and place the nmi\_watchdog=0 at the end of the currently booting kernel line.
2. reboot the sever
3. Once the system comes up perform a `rmmod hpwdt`
4. `insmod /lib/modules/`uname -r`/kernel/drivers/char/watchdog/hpwdt.ko`

hpwdt.txt

priority=1

Now, the hpwdt can successfully receive and source the NMI and provide a log message that details the reason for the NMI (as determined by the HP BIOS).

Below is a list of NMIs the HP BIOS understands along with the associated code (reason):

No source found	00h
Uncorrectable Memory Error	01h
ASR NMI	1Bh
PCI Parity Error	20h
NMI Button Press	27h
SB_BUS_NMI	28h
ILO Doorbell NMI	29h
ILO IOP NMI	2Ah
ILO Watchdog NMI	2Bh
Proc Throt NMI	2Ch
Front Side Bus NMI	2Dh
PCI Express Error	2Fh
DMA controller NMI	30h
Hypertransport/CSI Error	31h

-- Tom Mingarelli  
(thomas.mingarelli@hp.com)