

PAT (Page Attribute Table)

x86 Page Attribute Table (PAT) allows for setting the memory attribute at the page level granularity. PAT is complementary to the MTRR settings which allows for setting of memory types over physical address ranges. However, PAT is more flexible than MTRR due to its capability to set attributes at page level and also due to the fact that there are no hardware limitations on number of such attribute settings allowed. Added flexibility comes with guidelines for not having memory type aliasing for the same physical memory with multiple virtual addresses.

PAT allows for different types of memory attributes. The most commonly used ones that will be supported at this time are Write-back, Uncached, Write-combined and Uncached Minus.

PAT APIs

There are many different APIs in the kernel that allows setting of memory attributes at the page level. In order to avoid aliasing, these interfaces should be used thoughtfully. Below is a table of interfaces available, their intended usage and their memory attribute relationships. Internally, these APIs use a `reserve_memtype()/free_memtype()` interface on the physical address range to avoid any aliasing.

API	RAM	ACPI,...	Reserved/Holes
<code>ioremap</code>	---	UC-	UC-
<code>ioremap_cache</code>	---	WB	WB
<code>ioremap_nocache</code>	---	UC-	UC-
<code>ioremap_wc</code>	---	---	WC
<code>set_memory_uc</code> <code>set_memory_wb</code>	UC-	---	---
<code>set_memory_wc</code> <code>set_memory_wb</code>	WC	---	---
<code>pci sysfs resource</code>	---	---	UC-
<code>pci sysfs resource_wc</code> <code>is IORESOURCE_PREFETCH</code>	---	---	WC
<code>pci proc</code> <code>!PCIIOC_WRITE_COMBINE</code>	---	---	UC-
<code>pci proc</code> <code>PCIIOC_WRITE_COMBINE</code>	---	---	WC

		pat. txt	
/dev/mem read-write	--	WB/WC/UC-	WB/WC/UC-
/dev/mem mmap SYNC flag	--	UC-	UC-
/dev/mem mmap !SYNC flag and any alias to this area	--	WB/WC/UC- (from exist- ing alias)	WB/WC/UC- (from exist- ing alias)
/dev/mem mmap !SYNC flag no alias to this area and MTRR says WB	--	WB	WB
/dev/mem mmap !SYNC flag no alias to this area and MTRR says !WB	--	--	UC-

Advanced APIs for drivers

A. Exporting pages to users with `remap_pfn_range`, `io_remap_pfn_range`, `vm_insert_pfn`

Drivers wanting to export some pages to userspace do it by using `mmap` interface and a combination of

- 1) `pgprot_noncached()`
- 2) `io_remap_pfn_range()` or `remap_pfn_range()` or `vm_insert_pfn()`

With PAT support, a new API `pgprot_writecombine` is being added. So, drivers can continue to use the above sequence, with either `pgprot_noncached()` or `pgprot_writecombine()` in step 1, followed by step 2.

In addition, step 2 internally tracks the region as UC or WC in memtype list in order to ensure no conflicting mapping.

Note that this set of APIs only works with IO (non RAM) regions. If driver wants to export a RAM region, it has to do `set_memory_uc()` or `set_memory_wc()` as step 0 above and also track the usage of those pages and use `set_memory_wb()` before the page is freed to free pool.

Notes:

-- in the above table mean "Not suggested usage for the API". Some of the --'s are strictly enforced by the kernel. Some others are not really enforced today, but may be enforced in future.

pat.txt

For ioremap and pci access through /sys or /proc - The actual type returned can be more restrictive, in case of any existing aliasing for that address. For example: If there is an existing uncached mapping, a new ioremap_wc can return uncached mapping in place of write-combine requested.

set_memory_[uc|wc] and set_memory_wb should be used in pairs, where driver will first make a region uc or wc and switch it back to wb after use.

Over time writes to /proc/mtrr will be deprecated in favor of using PAT based interfaces. Users writing to /proc/mtrr are suggested to use above interfaces.

Drivers should use ioremap_[uc|wc] to access PCI BARs with [uc|wc] access types.

Drivers should use set_memory_[uc|wc] to set access type for RAM ranges.

PAT debugging

With CONFIG_DEBUG_FS enabled, PAT memtype list can be examined by

```
# mount -t debugfs debugfs /sys/kernel/debug
# cat /sys/kernel/debug/x86/pat_memtype_list
PAT memtype list:
uncached-minus @ 0x7fadf000-0x7fae0000
uncached-minus @ 0x7fb19000-0x7fb1a000
uncached-minus @ 0x7fb1a000-0x7fb1b000
uncached-minus @ 0x7fb1b000-0x7fb1c000
uncached-minus @ 0x7fb1c000-0x7fb1d000
uncached-minus @ 0x7fb1d000-0x7fb1e000
uncached-minus @ 0x7fb1e000-0x7fb25000
uncached-minus @ 0x7fb25000-0x7fb26000
uncached-minus @ 0x7fb26000-0x7fb27000
uncached-minus @ 0x7fb27000-0x7fb28000
uncached-minus @ 0x7fb28000-0x7fb2e000
uncached-minus @ 0x7fb2e000-0x7fb2f000
uncached-minus @ 0x7fb2f000-0x7fb30000
uncached-minus @ 0x7fb31000-0x7fb32000
uncached-minus @ 0x80000000-0x90000000
```

This list shows physical address ranges and various PAT settings used to access those physical address ranges.

Another, more verbose way of getting PAT related debug messages is with "debugpat" boot parameter. With this parameter, various debug messages are printed to dmesg log.